

Question #1

Topic 1

You execute this command:

```
CREATE BIGFILE TABLESPACE big_tbs
    DATAFILE '/u01/oracle/data/big_f1.dbf'
    SIZE 20G;
```


Sufficient storage is available in filesystem /u01.

Which two statements are true about the BIG_TBS tablespace? (Choose two.)

- A. AUTOEXTEND is possible for the datafile
- B. It must be bigger than the largest SMALLFILE tablespace
- C. Additional data files may not be added
- D. It will be a dictionary-managed tablespace by default
- E. It will always have a 32K blocksize

Correct Answer: AB

Community vote distribution



danito Highly Voted 3 years, 7 months ago
A C i think the correct answers, becuase B has not sense and i did not references about that
upvoted 18 times

HassanShami Most Recent 1 month, 1 week ago
Selected Answer: AC
B is not correct. The size of file doesn't related to other normal tablespaces
upvoted 1 times

nautil2 4 months, 2 weeks ago
Selected Answer: AC
Quotations originate in official Oracle Database 19c documentation.
A – true; "With a bigfile tablespace, you can use the AUTOEXTEND clause outside of the ADD DATAFILE clause."
B – false; there is no restriction for minimal size of a bigfile tablespace mentioned in Oracle documentation
C – true; "A bigfile tablespace is a tablespace with a single, but potentially very large (up to 4G blocks) data file. Traditional smallfile tablespaces, in contrast, can contain multiple data files, but the files cannot be as large."
D – false; "Bigfile tablespaces are supported only for locally managed tablespaces with automatic segment space management, with three exceptions: locally managed undo tablespaces, temporary tablespaces, and the SYSTEM tablespace."
E – false; "A.2 Physical Database Limits ... Database Block Size ... Maximum ...
Operating system dependent; never more than 32 KB"
upvoted 1 times

auwia 6 months, 3 weeks ago
Selected Answer: AC
A and C are corrects.
upvoted 1 times

Methio 7 months, 3 weeks ago
A c is answer
upvoted 1 times

0z0n3 8 months, 1 week ago
Selected Answer: AC
AC is correct
upvoted 1 times

AberTroth 8 months, 2 weeks ago
A and C
upvoted 1 times

RaNik69 9 months, 1 week ago
Selected Answer: AC

A= Correct - Tested
C= Correct - Tested
upvoted 1 times

  **FelipeC** 12 months ago

Selected Answer: AC

A= Correct - Tested
B= False - Tested
C= Correct - Tested (ORA-32771: cannot add file to bigfile tablespace)
D= False - Tested (Getting DDL)
E= False -Tested
upvoted 2 times

  **alexbraga93** 1 year, 1 month ago

Selected Answer: AC



I've tested in my lab and:
B - is not correct 'cause having a 30GB smallfile tablespace, I can create a 10GB bigfile tablespace
upvoted 2 times

  **amayorcas** 1 year, 4 months ago

Selected Answer: AC

Correct answer is A and C
A = Because autoextend is a feature that is usually set by default in TBS BIGFILE.
C = Because the tablespace of type BIGFILE can only have one datafile.

B is not correct: Because it is possible to create a 10GB TBS BIGFILE, but in a year it will grow to 100GB while we have another 25GB NORMAL TBS.
upvoted 1 times

  **Blob44** 1 year, 4 months ago

Selected Answer: AC

A = Because autoextend allowed in TBS BIGFILE.
C = Because the tablespace of type BIGFILE can only have one datafile.
upvoted 1 times

  **Albertopology** 1 year, 5 months ago

Selected Answer: AC

100% A and C
upvoted 2 times

  **LeandroHPN** 1 year, 6 months ago

Selected Answer: AC

A and C
upvoted 2 times

  **amayorcas** 1 year, 6 months ago

Correct answer is A and C
A = Because autoextend is a feature that is usually set by default in TBS BIGFILE.
C = Because the tablespace of type BIGFILE can only have one datafile.

B is not correct: Because it is possible to create a 10GB TBS BIGFILE, but in a year it will grow to 100GB while we have another 25GB NORMAL TBS.
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

AC, Correct Answer!
upvoted 1 times

  **Aramazd** 1 year, 9 months ago

oracle@localhost ~]\$ cat date_format.sql
alter session set NLS_DATE_FORMAT='DD-MON-YYYY HH24:MI:SS' ;

select current_timestamp from dual ;

select current_date from dual ;

[oracle@localhost ~]\$ sqlplus system/oracle @date_format.sql

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Apr 6 04:41:16 2022
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.

Last Successful login time: Wed Apr 06 2022 04:41:01 -04:00

Connected to:

Session altered.

CURRENT_TIMESTAMP

06-APR-22 04.41.16.562674 AM -04:00

CURRENT_DATE

06-APR-2022 04:41:16
upvoted 1 times

Which statement is true about the INTERSECT operator used in compound queries?

- A. Multiple INTERSECT operators are not possible in the same SQL statement
- B. It processes NULLs in the selected columns
- C. INTERSECT is of lower precedence than UNION or UNION ALL
- D. It ignores NULLs

Correct Answer: B

Community vote distribution



elvegaa Highly Voted 3 years, 2 months ago

B.
SET operators all have the same priority in Oracle Database.
NULL values are processed with set operators
upvoted 5 times

Ekos 3 years, 1 month ago

i agree
upvoted 1 times

nautil2 Most Recent 4 months, 2 weeks ago

Selected Answer: B

Quotations originate in official Oracle Database 19c documentation.
A – false; following query is executed successfully: SELECT empno FROM history WHERE beginyear <= 1995
INTERSECT
SELECT empno FROM employees WHERE job='TRAINER'
INTERSECT
SELECT trainer FROM offerings WHERE location <> 'SEATTLE';
B – true; “The INTERSECT operator with the keyword ALL returns the result of two or more SELECT statements in which rows appear in all result sets. Null values that are common across the component queries of INTERSECT ALL are returned at the end of the result set.”
C – false; “You can combine multiple queries using the set operators UNION, UNION ALL, INTERSECT, INTERSECT ALL, EXCEPT, EXCEPT ALL, MINUS, and MINUS ALL. All set operators have equal precedence.”
D – false; see answer B
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: B

Provided answer is correct.
upvoted 1 times

duongntsav 7 months, 2 weeks ago

In IZO-082 exam topic, aren't there SQL fudatmental? right
upvoted 1 times

AberTroth 8 months, 2 weeks ago

Selected Answer: B

B 100%
upvoted 1 times

j_tw 9 months ago

SET operators process NULL values
upvoted 1 times

FirstJM 9 months, 3 weeks ago

B is correct
upvoted 1 times

Albertopology 1 year, 5 months ago

Selected Answer: B

agree, answer should be B
upvoted 2 times

Cobra 1 year, 8 months ago



B is correct

upvoted 1 times

  **ryuah** 2 years ago



B is correct

upvoted 1 times

  **Fan** 2 years, 10 months ago

B is right.

upvoted 2 times

  **ora404** 3 years, 4 months ago

SELECT 1, 2 FROM DUAL
UNION
SELECT 2, 3 FROM DUAL
INTERSECT
SELECT NULL, NULL FROM DUAL
INTERSECT
SELECT NULL, NULL FROM DUAL;
returns "no data found".

upvoted 3 times

  **SimoneF** 3 years, 2 months ago

That is expected: the second two selects have no values in common with the corresponding columns in the first two, and the first two have nothing in common between them.
For any row to be returned, with NULL values or not, it should be present in all of the intersected tables.

upvoted 1 times

Which three statements are true about advanced connection options supported by Oracle Net for connection to Oracle Database instances?

(Choose three.)

- A. Connect Time Failover requires the use of Transparent Application Failover (TAF)
- B. Source Routing requires the use of a name server
- C. Source Routing enables the use of Connection Manager (CMAN) which enables network traffic to be routed through a firewall
- D. Load Balancing can balance the number of connections to dispatchers when using a Shared Server configuration
- E. Load Balancing requires the use of a name server
- F. Connect Time Failover requires the connect string to have two or more listener addresses configured

Correct Answer: BCF

Reference:

https://docs.oracle.com/cd/E11882_01/network.112/e41945/advcfg.htm#NETAG013

Community vote distribution



Rivaldo11 Highly Voted 3 years, 6 months ago
I think C, D and F are correct.
upvoted 13 times

ama 3 years, 6 months ago
D seems a legit answer..

<https://docs.oracle.com/en/database/oracle/oracle-database/12.2/rilin/shared-server-configuration-for-an-oracle-rac-database.html#GUID-2EFBA08D-FEEE-407F-BC42-E548DA946DF9>
upvoted 2 times

auwia Most Recent 6 months, 3 weeks ago
Selected Answer: CDF
A. Connect Time Failover requires the use of Transparent Application Failover (TAF) --> False
B. Source Routing requires the use of a name server --> False, IP address are also accepted.
C. Source Routing enables the use of Connection Manager (CMAN) which enables network traffic to be routed through a firewall --> True
D. Load Balancing can balance the number of connections to dispatchers when using a Shared Server configuration --> True
E. Load Balancing requires the use of a name server --> False, IP address are also accepted.
F. Connect Time Failover requires the connect string to have two or more listener addresses configured. --> True
upvoted 2 times

Ashikraja 11 months, 2 weeks ago
BCF is right
upvoted 1 times



yarsalan 2 years ago
D is incorrect. Standalone Oracle databases perform load balancing by distributing connections among the shared server dispatcher processes. Check out <https://docs.oracle.com/en/database/oracle/oracle-database/19/rilin/shared-server-configuration-for-an-oracle-rac-database.html#GUID-2EFBA08D-FEEE-407F-BC42-E548DA946DF9>
upvoted 2 times



Kira_Sat 2 years ago
Selected Answer: CDF
Correct Answer : C,D and F
upvoted 3 times



adozoo 2 years, 1 month ago
CDF is right!!!
upvoted 1 times



elvegaa 3 years, 2 months ago
C, D & F
upvoted 3 times

you1234 3 years, 6 months ago
ama, Please suggest suitable answer
upvoted 2 times

  **ama** 3 years, 6 months ago
C, D, F
upvoted 6 times

  **you1234** 3 years, 6 months ago
A & C & F is correct answer
upvoted 2 times

  **elvegaa** 3 years, 2 months ago
A is wrong.
Connect-Time Failover involves not yet connected sessions to failover a different address in case the first attempt is unreachble.
Transparent Application Failover involves sessions already connected in the database, to failover to the next address when the session is terminated.
Pretty similar, but not dependent. You can have a configuration with ConnectTime Failover but not Transparent Application Failover
upvoted 2 times

  **ama** 3 years, 6 months ago
A is wrong...

Connect-Time Failover is a feature of local naming method, and can be easily set up by adding a tns entry into tnsnames. ora file on the client machine, then the client can connect to the second database (listener) if the first database (listener) cannot be connected
upvoted 6 times


Which two statements are true about date/time functions in a session where NLS_DATE_FORMAT is set to DD-MON-YYYY HH24:MI:SS? (Choose two.)


- A. CURRENT_TIMESTAMP returns the same date and time as SYSDATE with additional details of fractional seconds
- B. SYSDATE can be queried only from the DUAL table
- C. CURRENT_DATE returns the current date and time as per the session time zone
- D. SYSDATE can be used in expressions only if the default date format is DD-MON-RR
- E. SYSDATE and CURRENT_DATE return the current date and time set for the operating system of the database server
- F. CURRENT_TIMESTAMP returns the same date as CURRENT_DATE

Correct Answer: CF

Community vote distribution




 **you1234** Highly Voted 3 years, 6 months ago
C & F is correct
upvoted 15 times


 **nautil2** Most Recent 4 months, 2 weeks ago
Selected Answer: CF
A - false; it is the same only in case database server operating system uses the same time zone as the operating system of the database client (session time zone)
B - false; can be queried in any table, e.g. SELECT emp_no,ename,salary,mgr_no,sysdate FROM emp;
C - true;
SQL> select sys_extract_utc(systimestamp) from dual;
SYS_EXTRACT_UTC(SYSTIMESTAMP)

11.09.23 13:58:24,316896
SQL> SELECT SESSIONTIMEZONE FROM DUAL;
SESSIONTIMEZONE

+02:00
SQL> SELECT current_date FROM dual;
CURRENT_DATE

11-SEP-2023 15:58:41
D - false; sysdate can be used independently of the setting of the NLS_DATE_FORMAT
E - false; sydate returns date and time of the operating system of the database server, current_date returns date and time in the session time zone (i.e. in the time zone of the client)
F - true; both variables returns current time and date set in the session time zone
upvoted 1 times

 **feixiang** 5 months, 2 weeks ago
I think C and E are correct,
CURRENT_TIMESTAMP shows more content than current_date
upvoted 2 times

 **chandanchoudhary** 7 months, 4 weeks ago
C & F are correct
SQL>
SQL> SELECT sessiontimezone FROM DUAL;

SESSIONTIMEZONE

+10:00

SQL> select current_date from dual;

CURRENT_DATE

31-MAY-2023 02:24:46

SQL> select CURRENT_TIMESTAMP from dual;

CURRENT_TIMESTAMP

SQL> select SYSDATE from dual;

SYSDATE

30-MAY-2023 22:25:23

SQL> select CURRENT_DATE from dual;

CURRENT_DATE

31-MAY-2023 02:25:41

SQL> select CURRENT_TIMESTAMP from dual;

CURRENT_TIMESTAMP



31-MAY-23 02.25.58.563965 AM +10:00

SQL> select sysdate from v\$database;

SYSDATE

30-MAY-2023 22:26:58

SQL>
upvoted 3 times

  **Blob44** 1 year, 4 months ago

Selected Answer: CF

CURRENT_DATE and CURRENT_TIMESTAMP are tied to session time zone
SYSDATE is tied to database server date and time. It can be used querying any table, not just dual, no matter the session nls_date_format chosen
upvoted 1 times

  **Aramazd** 1 year, 9 months ago

F cannot be correct

oracle@localhost ~]\$ cat date_format.sql
alter session set NLS_DATE_FORMAT='DD-MON-YYYY HH24:MI:SS' ;

select current_timestamp from dual ;

select current_date from dual ;

[oracle@localhost ~]\$ sqlplus system/oracle @date_format.sql

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Apr 6 04:41:16 2022
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.

Last Successful login time: Wed Apr 06 2022 04:41:01 -04:00

Connected to:
Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0

Session altered.

CURRENT_TIMESTAMP

06-APR-22 04.41.16.562674 AM -04:00

CURRENT_DATE

06-APR-2022 04:41:16
upvoted 1 times

  **Aramazd** 1 year, 9 months ago

F cannot be correct:

[oracle@localhost ~]\$ cat date_format.sql
alter session set NLS_DATE_FORMAT is set to 'DD-MON-YYYY HH24:MI:SS' ;

select current_timestamp from dual ;

店长微信 : hj feng128

select current_date from dual ;

[oracle@localhost ~]\$ sqlplus system/oracle @date_format.sql

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Apr 6 04:39:17 2022
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.

Last Successful login time: Wed Apr 06 2022 04:38:02 -04:00

Connected to:
Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0

alter session set NLS_DATE_FORMAT is set to 'DD-MON-YYYY HH24:MI:SS'
*



ERROR at line 1:
ORA-00927: missing equal sign



CURRENT_TIMESTAMP



06-APR-22 04.39.18.153607 AM -04:00



CURRENT_D

06-APR-22
upvoted 1 times

  **ryuah** 2 years ago
C,F is correct
upvoted 2 times

  **alimdyahoo** 2 years ago
Current_date is session specific, Current_timestamp is also session specific with zone and sysdate is operating system date . So only left with Valid choice B and C
upvoted 1 times

  **alimdyahoo** 2 years ago
oops my mistake . C and F are only valid choices. So C and F are correct
upvoted 2 times



  **eyildiz** 2 years, 3 months ago
I think A and C are correct.



SQL> ALTER SESSION SET NLS_DATE_FORMAT = 'DD-MON-YYYY HH24:MI:SS';
SQL> select sysdate,current_date,current_timestamp from dual;



SYSDATE CURRENT_DATE CURRENT_TIMESTAMP

28-SEP-2021 14:32:48 28-SEP-2021 14:32:48 28-SEP-21 02.32.48.201457 PM +03:00

https://docs.oracle.com/cd/B19306_01/server.102/b14200/functions172.htm
https://docs.oracle.com/cd/B19306_01/server.102/b14200/functions036.htm
https://docs.oracle.com/cd/B19306_01/server.102/b14200/functions037.htm
upvoted 3 times

  **j_tw** 9 months, 2 weeks ago
A cannot be correct. Becasue your DB server and your session are in same region, current_date and sysdate are same. If you connect to your DB at another timezone, that command return another result. So....it can be same, but not always do.
upvoted 1 times

  **DarrenChenSHCMB** 2 years, 2 months ago
AC is correct
upvoted 1 times

  **kawsar** 2 years, 5 months ago
SQL> set serveroutput on
SQL> declare
2 ld date;
3 begin
4 ld:=sysdate;
5 dbms_output.put_line(ld);
6 end;

🗨️ 👤 **kawsar** 2 years, 5 months ago
ALTER SESSION SET NLS_DATE_FORMAT = 'DD-MON-YYYY HH24:MI:SS';
select current_timestamp from dual;
select current_date from dual;
select sysdate from dual;
08-AUG-21 02.00.13.303335 AM US/PACIFIC
08-AUG-2021 02:00:15
08-AUG-2021 09:00:17
Now check the answers!!!!
upvoted 2 times

🗨️ 👤 **mianjee** 2 years, 6 months ago
B and C are correct!
upvoted 1 times

🗨️ 👤 **Ekos** 3 years, 3 months ago
the correct answer is should be A and F
upvoted 1 times

🗨️ 👤 **Ekos** 3 years, 2 months ago
sorry CF is the correct one
upvoted 7 times

🗨️ 👤 **ama** 3 years, 6 months ago
so corrects are C , F
upvoted 3 times

🗨️ 👤 **ama** 3 years, 6 months ago
E is also wrong becuae sysdate Displays only date
upvoted 1 times

🗨️ 👤 **danito** 3 years, 5 months ago
if you read the question "where nls_date_format is set to"
I have done this and this is the result:
SQL> ALTER SESSION SET NLS_DATE_FORMAT = 'DD-MON-YYYY HH24:MI:SS';

Sesion modificada.

SQL> select sysdate from dual;

SYSDATE

17-AGO-2020 17:59:34

SQL> select current_date from dual;

CURRENT_DATE

17-AGO-2020 17:59:53

maybe E is correct
upvoted 1 times

🗨️ 👤 **ama** 3 years, 6 months ago
A is wrong becuae sysdate Displays only date but not the time.
upvoted 1 times

🗨️ 👤 **elvegaa** 3 years, 2 months ago
SYSDATE also shows time if you change the date display format to include it.
However, A is wrong because SYSDATE shows database resident OS date&time while CURRENT_TIMESTAMP shows session date&time. Different time zones can affect the result.
upvoted 2 times

A database is configured to use automatic undo management with temporary undo enabled.

An UPDATE is executed on a temporary table.

Where is the UNDO stored?

- A. in the undo tablespace
- B. in the SYSAUX tablespace
- C. in the SGA
- D. in the PGA
- E. in the temporary tablespace

Correct Answer: A

Reference:

<https://docs.oracle.com/database/121/ADMIN/undo.htm#ADMIN11478>

(16.7)

Community vote distribution

E (100%)

  **Rivaldo11**  3 years, 6 months ago

E is correct answer.

upvoted 19 times

  **ama** 3 years, 6 months ago

Good hint, i think you have Right !

Oracle database 12c Release 1 (12.1) introduced the concept of temporary undo, allowing the undo segments for global temporary tables to be stored in the temporary tablespace. This allows global temporary tables to be used in physical standby databases and read-only databases, as well as removing the need to create redo.

upvoted 6 times

  **NowOrNever** 3 years, 6 months ago

I would agree

http://www.dba-oracle.com/t_temp_undo_enabled.htm

upvoted 1 times

  **adoptc94**  3 years, 5 months ago

Funny how the reference they put up supports answer E and not answer A:

16.7 Managing Temporary Undo

By default, undo records for temporary tables are stored in the undo tablespace and are logged in the redo, which is the same way undo is managed for persistent tables. However, you can use the TEMP_UNDO_ENABLED initialization parameter to separate undo for temporary tables from undo for persistent tables. When this parameter is set to TRUE, the undo for temporary tables is called temporary undo.

16.7.1 About Managing Temporary Undo

Temporary undo records are stored in the database's temporary tablespaces and thus are not logged in the redo log. When temporary undo is enabled, some of the segments used by the temporary tablespaces store the temporary undo, and these segments are called temporary undo segments.

Reference: <https://docs.oracle.com/database/121/ADMIN/undo.htm#ADMIN11479>

The correct answer is: E

upvoted 10 times

  **Abdullejr**  3 weeks ago

Selected Answer: E

E is the correct answer. The reference attached here states that.

16.7 Managing Temporary Undo

By default, undo records for temporary tables are stored in the undo tablespace and are logged in the redo, which is the same way undo is managed for persistent tables. However, you can use the TEMP_UNDO_ENABLED initialization parameter to separate undo for temporary tables from undo for persistent tables. When this parameter is set to TRUE, the undo for temporary tables is called temporary undo.

16.7.1 About Managing Temporary Undo

Temporary undo records are stored in the database's temporary tablespaces and thus are not logged in the redo log. When temporary undo is enabled, some of the segments used by the temporary tablespaces store the temporary undo, and these segments are called temporary undo segments.

upvoted 1 times

  **doyinbare** 4 months, 1 week ago

By default, undo records for temporary tables are stored in the undo tablespace . This means that If the temporary undo is set to false, then undo records generated by temporary tables are logged to undo tablespace. However, when temporary undo is enabled, the undo for temporary tables is called temporary undo which are then logged in to a separate tablespace called temporary tablespace.

upvoted 1 times

  **nachomago** 6 months, 2 weeks ago

Selected Answer: E

E is correct answer.


upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: E

As many user commented, the correct answer is E, this is a particular case.


upvoted 1 times

  **0z0n3** 8 months, 1 week ago

Selected Answer: E


undo of temp table would be stored in temp tablespace if temp undo is enabled

upvoted 1 times

  **cristycool** 9 months, 1 week ago

E is correct answer.



upvoted 1 times

  **j_tw** 9 months, 2 weeks ago

Selected Answer: E

E. If there is no 'temporary undo enabled', then it will be stored in the undo tablespace.



upvoted 1 times

  **Blob44** 1 year ago

Selected Answer: E

E is correct



upvoted 1 times

  **Blob44** 1 year, 4 months ago

Selected Answer: E

Enabling temporary undo allows temporary undo segments to be created in temp tablespaces for the corresponding temporary tables. This can improve performance a lot

upvoted 1 times

  **Soiram** 1 year, 11 months ago

Selected Answer: E

...with temporary undo enabled.

upvoted 1 times

  **emburria** 2 years ago

Selected Answer: E

E-->Temporary undo records are stored in the database's temporary tablespaces and thus are not logged in the redo log. When temporary undo is enabled, some of the segments used by the temporary tablespaces store the temporary undo, and these segments are called temporary undo segments.

upvoted 1 times

  **ryuah** 2 years ago

E is correct

upvoted 1 times

  **ryuah** 2 years ago



E is correct

upvoted 1 times

  **elvegaa** 3 years, 2 months ago

Who's marking the right answers? They're almost all wrong!
A is wrong, E is the right one. The same reference supports the E answer

upvoted 6 times

  **Ekos** 3 years, 3 months ago

E is the correct answer

Temporary undo records are stored in the database's temporary tablespaces and thus are not logged in the redo log

upvoted 2 times

You have been tasked to create a table for a banking application.
One of the columns must meet three requirements:

1. Be stored in a format supporting date arithmetic without using conversion functions
2. Store a loan period of up to 10 years
3. Be used for calculating interest for the number of days the loan remains unpaid

Which data type should you use?

- A. INTERVAL YEAR TO MONTH
- B. INTERVAL DAY TO SECOND
- C. TIMESTAMP WITH LOCAL TIMEZONE
- D. TIMESTAMP
- E. TIMESTAMP WITH TIMEZONE

Correct Answer: B

Community vote distribution



nautil2 4 months, 2 weeks ago

Selected Answer: B

A - false; does not meet condition 3, because this interval does not store days
B - true; this interval can store days
C - false; cannot store period
D - false; cannot store period
E - false; cannot store period
upvoted 1 times

trgbighero 1 year, 3 months ago

Selected Answer: B

B is correct
upvoted 1 times

LeandroHPN 1 year, 6 months ago

Selected Answer: B

B is correct
upvoted 2 times

thuto 1 year, 8 months ago

It is B... :)
upvoted 1 times

ryuah 2 years ago

B is correct
upvoted 1 times

ty_1 2 years, 9 months ago



yes B is correct one
upvoted 1 times

khizar782 2 years, 9 months ago

Correct Answer is B.
Since the third requirement for banking application is
'Be used for calculating interest for the number of days the loan remains unpaid'
So, Data Type 'INTERVAL DAY TO SECOND' must be used to calculate duration in number of days.
upvoted 2 times

pr0glnx 3 years, 2 months ago

Yes: B'
upvoted 2 times

  **Ekos** 3 years, 3 months ago
agree, b is the correct answer
upvoted 4 times

店长微信：hj feng128

In the spfile of a single instance database, LOCAL_LISTENER is set to LISTENER_1.

The TNSNAMES.ORA file in \$ORACLE_HOME/network/admin in the database home contains:

```
LISTENER_1 =  
  (ADDRESS =  
    (PROTOCOL = TCP)  
    (HOST = host1.abc.com)  
    (PORT = 1521)  
  )
```

Which statement is true?

- A. Dynamic service registration cannot be used for this database instance
- B. The LREG process registers services dynamically with the LISTENER_1 listener
- C. LISTENER_1 must also be defined in the LISTENER.ORA file to enable dynamic service registration
- D. There are two listeners named LISTENER and LISTENER_1 running simultaneously using port 1521 on the same host as the database instances
- E. The definition for LISTENER_1 requires a CONNECT_DATA section to enable dynamic service registration

Correct Answer: C

Reference:

<https://docs.oracle.com/database/121/NETAG/listenercfg.htm#NETAG292>

Community vote distribution



Rivaldo11 Highly Voted 3 years, 6 months ago

Answer is B.

upvoted 19 times

NowOrNever 3 years, 6 months ago

I would agree

Dynamic service registration does not require any manual configuration in the listener.ora file.

upvoted 2 times

SimoneF 3 years, 2 months ago

I disagree: the manual configuration to which the documentation refers is intended to be "configuration [relative to dynamic registration]". But the listener has to be defined in the listener.ora nonetheless, as stated in the same document:

"Synchronization occurs when the protocol address of the listener is specified in the listener.ora file and the location of the listener is specified in the initialization parameter file."

upvoted 3 times

Muhab Highly Voted 3 years, 5 months ago

Answer is B.

The listener forwards client requests to supported services. These services are dynamically registered with the listener. This dynamic registration feature is called service registration. The registration is performed by the Listener Registration (LREG) process. Dynamic service registration does not require any manual configuration in the listener.ora file.

Reference:

<https://docs.oracle.com/database/121/NETAG/listenercfg.htm#NETAG292>

upvoted 9 times

elvegaa 3 years, 2 months ago

You DO need to define the listener in the listener.ora file if you're using an alias, however you don't define the instances registered in the _SID_LIST argument.

upvoted 4 times

SimoneF 3 years, 2 months ago

Exactly, if the listener specified as local_listener is not configured, the service cannot be registered there.

upvoted 2 times

Abdullejr Most Recent 2 weeks, 6 days ago

Selected Answer: C


C is the correct answer.

upvoted 1 times

  **SpyderSh** 4 months, 4 weeks ago

B is correct answer C is wrong due to not required to configure in listener.ora file to use DSR

upvoted 1 times

  **mamadu** 6 months, 2 weeks ago



Selected Answer: C

Try creating a new PDB without you not specifying the LISTENER_1 on the listener.ora, you will not see any service registered for that PDB. Its a Rule:

To get Dynamic Service Registration on user created listeners, you must:

- configure the LISTENER_1 on tnsname.ora
- configure the LISTENER_1 on listener.ora
- set or append the LISTENER_1 to the LOCAL_LISTENER parameter

upvoted 2 times

  **Darkseid1231** 7 months ago

D is correct

upvoted 2 times

  **RaNik69** 9 months, 1 week ago

Selected Answer: B

<https://docs.oracle.com/en/database/oracle/oracle-database/19/netag/configuring-and-administering-oracle-net-listener.html#GUID-A3263EB4-C3F2-4529-ABC2-ADE749114D33>

The listener forwards client requests to supported services. These services are dynamically registered with the listener. This dynamic registration feature is called service registration.

The registration is performed by the Listener Registration (LREG) process. Dynamic service registration does not require any manual configuration in the listener.ora file.

upvoted 1 times

  **canijho** 11 months ago

B is correct.

Dynamic service registration does not use the listener.ora file.

upvoted 1 times

  **jonsnoow** 1 year ago

B is the correct answer. C is wrong - dynamic service registration does not use the listener.ora file.

upvoted 1 times

  **Blob44** 1 year ago

Selected Answer: B

To my understanding of the question, LISTENER_1 is just an alias for the DEFAULT listener listening on host1 and port 1521. So it's already on listener.ora file with the default name LISTENER or the name given when database was created with DBCA

upvoted 1 times

  **nautil2** 1 year, 1 month ago

Selected Answer: B

My vote is for answer B, because it's the closest to my findings. The question text does not imply the listener name is LISTENER_1. The string "LISTENER_1" is just an alias, that does not need a record in listener.ora, just a setting of LOCAL_LISTENER parameter in the database and alias explanation in tnsnames.ora as described at <https://docs.oracle.com/en/database/oracle/oracle-database/19/netag/configuring-and-administering-oracle-net-listener.html#GUID-0E7C39E3-4627-403A-AE69-E9AA2C7E4C57>

When reproduced a scenario on a testing instance and typing \$ lsnrctl status LISTENER_1, I can see status of default LISTENER with services registered dynamically after database was bounced. In listener.ora there is a definition of standard LISTENER, but no mention about LISTENER_1, however dynamic service registration works fine.

upvoted 3 times

  **alan_resende** 1 year, 1 month ago

Hi, I think the correct is C.

I did the test with configuration about this question:

If you want the LREG do the dynamic service registration with LISTENER_1, you have to set both tnsnames.ora to resolve the address/port and listener.ora with LISTENER_1 address/port. In another case the database do the dynamic register with default LISTENER on default port 1521 regardless local_listener parameter configuration.

Take account the question ask about dynamic service register with LISTENER_1.

upvoted 1 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: D

My teacher said D

upvoted 1 times

  **saad3577** 1 year, 4 months ago

Answer is C.

The listener MUST be defined in the listener.ora file in order for a listener to work, not in the tnsnames file, since tnsnames its only used for aliases. You can check multiple sources, for example: <https://docs.oracle.com/database/121/NETAG/listenercfg.htm#NETAG292> The default listener name is LISTENER. If you want to use LISTENER_1, it must be defined in LISTENER.ORA Source: <https://docs.oracle.com/en/database/oracle/oracle-database/19/netag/configuring-and-administering-oracle-net-listener.html#GUID-A3263EB4-C3F2-4529-ABC2-ADE749114D33> The local_listener service name, address or address list is must be specified in the tnsnames.ora file or you will get an ORA-00119 error.

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 5 months ago

The answer is C. I validated this practically, this also came in exam and have passed the exam with high percentage. There were no mention of this being a problem.

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 5 months ago

sorry not C. D. There are two listeners named LISTENER and LISTENER_1 running simultaneously using port 1521 on the same host as the database instances

upvoted 1 times

🗨️ 👤 **majla** 1 year, 3 months ago

It's not possible to have 2 processes listening on the same port - just the last one can listen. In this particular scenario you only have one listener LISTENER and the alias LISTENER_1 for the same process.

upvoted 1 times

🗨️ 👤 **joseank** 1 year, 9 months ago

B is a true statement, if you dont define a service name in the listener then it listens to every request (that's what a dynamic service registration means)

HOWEVER!!!

C is the correct answer, the listener MUST be defined in the listener.ora file in order for a listener to work, not in the tnsnames file, since tnsnames its only used for aliases.

You can check multiple sources, for example: <https://docs.oracle.com/database/121/NETAG/listenercfg.htm#NETAG292>

upvoted 2 times

🗨️ 👤 **Franky_T** 1 year, 9 months ago

Selected Answer: B

Dynamic service registration is configured in the database initialization file. It does not require any configuration in the listener.ora file. However, listener configuration must be set to listen on the ports named in the database initialization file, and must not have parameters set that prevent automatic registration, such as COST parameters.

By default, the LREG process registers service information with its local listener on the default local address of TCP/IP, port 1521. If the listener configuration is synchronized with the database configuration, then LREG can register service information with a nondefault local listener or a remote listener on another node.

Ref: <https://docs.oracle.com/en/database/oracle/oracle-database/19/netag/configuring-and-administering-oracle-net-listener.html#GUID-A7F49477-1F13-49E7-9B23-9A50A4627862>

upvoted 2 times

Which three statements are true concerning logical and physical database structures? (Choose three.)

- A. All tablespaces may have one or more data files
- B. The extents of a segment must always reside in the same datafile
- C. A smallfile tablespace might be bigger than a bigfile tablespace
- D. A segment can span multiple data files in some tablespaces
- E. A segment's blocks can be of different sizes
- F. A segment might have only one extent
- G. Segments can span multiple tablespaces

Correct Answer: CDF

Community vote distribution



- Rivaldo11**

Highly Voted

3 years, 6 months ago

C, D and F. Answer A is wrong, it says - "All tablespaces may have one or more data files", but bigfile tablespaces can have only one datafile.
upvoted 13 times
- ama**

3 years, 5 months ago

what you think About this Information :

Smallfile tablespaces can contain up to 1024 files, but bigfile tablespaces contain only one file that can be 1024 times larger than a smallfile tablespace.

<https://web.stanford.edu/dept/itss/docs/oracle/10gR2/server.102/b14220/physical.htm>
upvoted 1 times
- adoptc94**

3 years, 5 months ago

It's just like Rivaldo11 said, not all tablespaces may have one or more data files! Bigfile tablespaces can only have ONE datafile (which can be and in most cases are much bigger in size than smallfile datafiles) - as stated in your statement.

A is definitely incorrect!

My picks would also be C, D and F
upvoted 5 times
- nyale**

2 years, 4 months ago

They said one or more!!! so a bigfile can have one and this makes A correct because one more is ≥ 1
upvoted 3 times
- algerianphoenix**

1 year, 8 months ago

Nope, "Not All tablespaces might have one or more data files", this is one of question when they play with your logic and common sense, I fell into the same logic as yours, before I negate the question.
upvoted 1 times
- Franky_T**

Highly Voted

1 year, 9 months ago

Selected Answer: CDF

A is wrong. Because bigfile tablespaces may have only one datafile
B is wrong. Because if a tablespace has more than one datafile then the extents of the segment can be in any one or all of the data files. Remember that tablespaces are LOGICAL objects and the segment lives in the logical object, regardless of which file the extent exists on.
C is correct. Smallfile tablespaces may have many datafiles that collectively exceed the the size of the datafile that is in the bigfile tablespace
D is correct. The opposite of option B. A tablespace is a logical object made up of one or more datafiles. With locally managed tablespaces the server will decide where the next extent allocation should be done for a segment
E is wrong. A segment must have all it's extents in the same tablespace, and block size is defined at tablespace level
F is correct. If a segment has row pieces that collectively are less than or equal to the total storage requirement of one extent then the segment will have only one extent allocated to it
G is wrong. A segment can reside in one tablespace only
upvoted 7 times
- SpyderSh**

Most Recent

4 months, 3 weeks ago

A B D is correct answers
upvoted 1 times
- SpyderSh**

4 months, 3 weeks ago

A C D are correct answers
upvoted 1 times

  **Hell3n** 5 months ago

Selected Answer: ACF

A C F, checked.

A is true, it say ---> one OR more

D is false.

D. A segment can span multiple data files in some tablespaces

'A segment can span one or more data files, but it cannot span multiple tablespaces'.

So in SOME tablespaces is not correct

upvoted 1 times

  **Hell3n** 5 months ago

A C F, checked.

A is true, it say ---> one OR more

D is false.

D. A segment can span multiple data files in some tablespaces

'A segment can span one or more data files, but it cannot span multiple tablespaces'.

So in SOME tablespaces is not correct

upvoted 1 times

  **linhao** 6 months, 1 week ago

A C F

A is correct. Bigfile tablespace have only one datafile and smallfile tablespace have several datafiles.

B is wrong. The extents of a segment could reside in the some others datafiles in the same tablespace

C is correct. The difference between smallfile tablespace and bigfile tablespace is the number of datafile. The max size of datafile is same.



D is wrong. A segment can span multiple data files in same tablespace, but not in some tablespaces

E is wrong. The block size can be defined in tablespace in creation.

F is correct.

G is wrong. Refer to D

upvoted 1 times

  **AliKhan345** 9 months ago

B is wrong .

If the initial extent become full, and if more space is required, then the database automatically allocates an incremental extent for this segment. An incremental extent is a subsequent extent created for the segment.

The allocation algorithm depends on whether the tablespace is locally managed or dictionary-managed. In the locally managed case, the database searches the bitmap of a data file for adjacent free blocks. If the data file has insufficient space, then the database looks in another data file. Extents for a segment are always in the same tablespace but may be in different data files.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html#GUID-C205981E-182A-466D-B151-090455F9ECE7>

upvoted 1 times

  **j_tw** 9 months, 2 weeks ago

Selected Answer: BCF

One tablespace may have multiple datafile.

A segment can be span multile datafile which reside in the same tablespace.

A extent, however, should be reside in same datafile.

So...I think the answer is B,C,AorF

A&F are ambiguous, but D&E&G are wrong.

D - A segment cannot span multiple data files in 'some' tablespaces. They should be reside in one tablespace.

reference : https://docs.oracle.com/cd/A57673_01/DOC/server/doc/SCN73/ch3.htm#extents

"... The segments also can span files, but the individual extents cannot."

upvoted 1 times

  **Minoro** 1 year, 1 month ago

Some tests said G is right ... G. Segments can span multiple tablespaces

A Segment can't span multiple tablespaces ... but the Segments, (table, Indexes,cluster,...) may be span multiple tablespaces it is on plural ... so may be possible correct

upvoted 1 times



  **alan_resende** 1 year, 1 month ago



Hi,

B is correct.



<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html#GUID-E3D8E030-A056-40AC-9B7E-8C957E28EE75>



"Allocation of Extents
By default, the database allocates an initial extent for a data segment when the segment is created. An extent is always contained in one data file."
upvoted 1 times

  **Darkseid1231** 1 year, 7 months ago
AMA, what is the correct answer please ?
D & F i'm sure
upvoted 1 times

  **Aramazd** 1 year, 9 months ago
A segment belongs only to one tablespace, but a segment's extents might span on different Datafiles of a tablespace.



That's why I vote for B C D
upvoted 1 times

  **Aramazd** 1 year, 9 months ago
G is wrong because
Each segment belongs to one and only one tablespace. Thus, all extents for a segment are stored in the same tablespace. Within a tablespace, a segment can include extents from multiple data files, as shown in Figure 12-2. For example, one extent for a segment may be stored in users01.dbf, while another is stored in users02.dbf. A single extent can never span data files.
upvoted 1 times

  **Aramazd** 1 year, 9 months ago
A segment can have several extents:

Segments

The level of logical database storage above an extent is called a segment. A segment is a set of extents that have been allocated for a specific type of data structure, and that all are stored in the same tablespace. For example, each table's data is stored in its own data segment, while each index's data is stored in its own index segment.
upvoted 1 times

  **ryuah** 2 years ago
C,D,F is correct
but I think A can also be the answer.
upvoted 2 times

  **adozoo** 2 years, 1 month ago
Answer is C D F
upvoted 1 times



Which two statements are true regarding a SAVEPOINT? (Choose two.)



- A. Rolling back to a SAVEPOINT can undo a CREATE INDEX statement
- B. Rolling back to a SAVEPOINT can undo a TRUNCATE statement
- C. Only one SAVEPOINT may be issued in a transaction
- D. A SAVEPOINT does not issue a COMMIT
- E. Rolling back to a SAVEPOINT can undo a DELETE statement



Correct Answer: DE


Community vote distribution



DE (100%)



  **Sharif1** Highly Voted 3 years, 5 months ago
DE provided answer's are correct.
upvoted 14 times



  **Ekos** 3 years, 1 month ago
agree, DE are correct
upvoted 3 times



  **Abdullejr** Most Recent 1 week, 4 days ago
Selected Answer: DE
D and E are correct answers because DDL statements such as CREATE, ALTER, TRUNCATE, etc commit trunsactions. and you can also have multiple savepoints in the same transaction.
upvoted 1 times

  **nautil2** 4 months, 1 week ago
Selected Answer: DE
A - FALSE; DDL changes cannot be undone
B - FALSE; see A
C - FALSE; there can be multiple savepoints, see <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/SAVEPOINT.html#GUID-78EEA746-0021-42E8-9971-3BA6DFFEE794> where text "Savepoint names must be distinct within a given transaction. If you create a second savepoint with the same identifier as an earlier savepoint, then the earlier savepoint is erased." is mentioned
D - TRUE; savepoint helps to rollback some changes, thus it cannot issue a COMMIT
E - TRUE; DELETE is a DML change, and thus it can be undone with ROLLBACK
upvoted 1 times

  **Hell3n** 5 months ago
Selected Answer: DE
I agree
upvoted 1 times

  **auwia** 6 months, 4 weeks ago
Selected Answer: DE
It is still correct :-)
upvoted 1 times

  **Aramazd** 1 year, 9 months ago
DE I agree too
upvoted 1 times

  **ryuah** 2 years ago
D,E is correct
upvoted 1 times


Which three functions are performed by dispatchers in a shared server configuration? (Choose three.)

- A. writing inbound request to the common request queue from all shared server connections
- B. checking for outbound shared server responses on the common outbound response queue
- C. receiving inbound requests from processes using shared server connections
- D. sending each connection input request to the appropriate shared server input queue
- E. broadcasting shared server session responses back to requesters on all connections
- F. sending shared server session responses back to requesters on the appropriate connection

Correct Answer: ACD

Community vote distribution

ACD (100%)

 **Abhirup**  3 years, 5 months ago

A,C,F correct
upvoted 9 times

 **NiciMilo**  3 years, 5 months ago

i think A, C and F

D is wrong

Ref: https://docs.oracle.com/cd/B28359_01/server.111/b28310/manproc001.htm#ADMIN11168

An ***idle shared server process picks up*** the virtual circuit from the common queue

upvoted 8 times

 **nautil2**  4 months, 1 week ago

Selected Answer: ACF

see <https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-processes.html#GUID-A522BFBF-5452-4038-92DB-C5116A0B1B50>

A - TRUE; All dispatchers are required to write incoming requests to a common response queue, which the (idle) shared server processes will monitor and pick up.

B - FALSE; Each dispatcher has its own response queue.

C - TRUE; This is the main purpose of dispatchers in a shared server architecture.

D - FALSE; Connection input requests are handled by the listener and, shared server processes do not have their own input queues, they monitor the common input queue.

E - FALSE; Shared server responses cannot be broadcasted to all connections.

F - TRUE; Shared server responses must be sent back to the requester on its own connection.

upvoted 2 times


 **Guhborges** 10 months ago

C,D,F

In a shared server configuration, dispatchers are responsible for managing incoming requests and outgoing responses. When a user connects to the database using a shared server configuration, the dispatcher receives the user's request and places it in the common request queue. The dispatcher then selects an available shared server and assigns the request to it by placing it in the appropriate shared server input queue.

The shared server then processes the request and generates a response, which is placed in the appropriate outbound response queue. The dispatcher checks for outbound shared server responses on the common outbound response queue and sends the response back to the requester on the appropriate connection.

upvoted 1 times

 **Franky_T** 1 year, 9 months ago

Selected Answer: ACF

A is correct. All dispatchers are required to write incoming requests to a common response queue, which the (idle) shared server processes will monitor and pick up.

B is wrong. Each dispatcher has it's own response queue.

C is correct. This is the main purpose of dispatchers in a shared server architecture.

D is wrong. Connection input requests are handled by the listener and, shared server processes do not have their own input queues, they monitor the common input queue.


E is wrong. Shared server responses cannot be broadcasted to all connections.


F is correct. Shared server responses must be sent back to the requester on it's own connection.

Reference:

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/application-and-networking-architecture.html#GUID-EC8F8778-02B7-4CF9-9E6F-A5D88C777235>


upvoted 6 times

- 


 **Aramazd** 1 year, 9 months ago

So correct answers can only be C D F

upvoted 1 times
- 


 **Aramazd** 1 year, 9 months ago


A is wrong too because the dispatcher does not writes inbounds requests from shared servers.

upvoted 1 times
- 


 **Aramazd** 1 year, 9 months ago


E is wrong because the concept of a broadcast does not exist in this mechanism.

upvoted 1 times
- 

 **Aramazd** 1 year, 9 months ago

B is wrong because there is no common response queue. Instead each dispatcher has it's own queue.

upvoted 1 times
- 

 **brolem** 2 years, 5 months ago

I think C, D, F, only because of the wording. "A" states " common request queue FROM all shared server connections". The dispatches send requests TO shared server connections. The send responses FROM the shared server connections.


upvoted 1 times

Which two statements are true about the SET VERIFY ON command? (Choose two.)

- A. It can be used only in SQL*Plus
- B. It displays values for variables used only in the WHERE clause of a query
- C. It can be used in SQL Developer and SQL*Plus
- D. It displays values for variables created by the DEFINE command
- E. It displays values for variables prefixed with &&

Correct Answer: AD

Community vote distribution



danito Highly Voted 3 years, 7 months ago
C D correct answers
upvoted 10 times

elbelgounetos 2 years, 10 months ago
If guess D is a trap. If you use DEFINE and you're not using the variable, it'll never be displayed with SET VERIFY ON.

Example :
set verify on
DEFINE dummy_char = 'X'
SELECT * FROM dual
upvoted 5 times

RaNik69 Most Recent 9 months, 1 week ago
Selected Answer: CE
both tested
upvoted 1 times

J4vi 9 months, 1 week ago
Selected Answer: CE
I think C D E, but you can use DEFINE without & or && and VERIFY ON won't work, so C E seems to be the most correct then
upvoted 1 times

ivanadj 10 months, 2 weeks ago
Selected Answer: CD
I think C,D is correct.
upvoted 1 times

jonsnoow 1 year ago
C, D, E
upvoted 1 times

Franky_T 1 year, 9 months ago
Selected Answer: CE
The SET VERIFY ON command will enable the display of any substitution variable used in our last executed SQL statement. When you use the DEFINE clause to create a variable, and the variable is not used in the next SQL statement then that variable will not be "shown" by the SET VERIFY command. Oh, and the command is available in all apps that allow the execution of SQL language against your database.
upvoted 4 times

ryuah 2 years ago
C,E is correct
upvoted 1 times

yukclam9 2 years, 9 months ago
C,E. D is not correct because 'define' as a individual command will not trigger the prompt. Only when defined variable is used with & or && the verify on thing will take effect.
upvoted 4 times

SimoneF 3 years, 2 months ago
This is impossible, I just tested the behaviour to be sure and, unsurprisingly, I confirmed that C-D-E seems to be true, thus invalidating A:
upvoted 2 times

  **SimoneF** 3 years, 2 months ago

店长微信：hj feng128

C: It works on SqlDevelope 3.22, so it's not even a new functionality.
upvoted 1 times

  **SimoneF** 3 years, 2 months ago

D - It works for variables created via DEFINE:

```
SQL> set verify off
SQL> define emp = 200;
SQL> select employee_id from hr.employees where employee_id = &emp;

EMPLOYEE_ID
-----
200

SQL> set verify on
SQL> select employee_id from hr.employees where employee_id = &emp;
vecchio 1: select employee_id from hr.employees where employee_id = &emp
nuovo 1: select employee_id from hr.employees where employee_id = 200
```

```
EMPLOYEE_ID
-----
200
upvoted 1 times
```

  **SimoneF** 3 years, 2 months ago

E - It works on && prefixed variables:

```
SQL> set verify off
SQL> select employee_id from hr.employees where employee_id =&&emp;
Immettere un valore per emp: 200
```

```
EMPLOYEE_ID
-----
200
```

```
SQL> undefine emp
SQL> set verify on
SQL> select employee_id from hr.employees where employee_id =&&emp;
Immettere un valore per emp: 200
vecchio 1: select employee_id from hr.employees where employee_id =&&emp
nuovo 1: select employee_id from hr.employees where employee_id =200
```

```
EMPLOYEE_ID
-----
200
```

So, either there are 3 answers instead of (and different from) the 2 suggested, or I'm seriously missing something here.
upvoted 2 times

  **Guhborges** 2 years, 2 months ago

Simone, you are missing something here! SET VERIFY does not display values for variables created by the DEFINE command, it shows the corresponding values of the variables we are currently using. If you don't use the create by define variable, you won't be able to see the value of DEFINE.
upvoted 3 times

  **esarregui** 3 years, 2 months ago

Answer: C + ¿D or E?



```
SQL> set verify on
SQL> select &&x from dual;
Enter value for x: dummy
old 1: select &&x from dual
new 1: select dummy from dual
```

```
D
-
X
```

```
SQL> set verify on
SQL> define y='dummy';
SQL> select &y from dual;
old 1: select &y from dual
new 1: select dummy from dual
```

```
D
-
X
```

upvoted 1 times

  **Chansi** 3 years, 5 months ago

Did u check if it works on sql developer? It doesn't work on toad. I tested
upvoted 1 times

  **adoptc94** 3 years, 4 months ago

If works in SQL Developer, there are also some YouTube Tutorials that show the command using SQL Developer.
upvoted 1 times

  **you1234** 3 years, 6 months ago

C & E is correct
upvoted 2 times

  **ama** 3 years, 5 months ago

i agree, E is working:



```
SQL> set verify off
SQL> select &&var from dual;
Geben Sie einen Wert fur var ein: sysdate
```

```
SYSDATE
-----
04.08.20
```



```
SQL> set verify on
SQL> select &&var from dual;
alt 1: select &&var from dual
neu 1: select sysdate from dual
```

```
SYSDATE
-----
04.08.20
```

upvoted 1 times

  **NowOrNever** 3 years, 6 months ago

why not E?
SET VERIFY ON
UNDEFINE dummy_number;
SELECT employee_id, job_id FROM employees where EMPLOYEE_ID = &&dummy_number;
-- working, he just asks for an input (like 206)
i agree with C, but unsure between D and E
upvoted 2 times

  **ama** 3 years, 6 months ago

E is wrong... look at below example, the variable was disaplyed even without using set verify on!

```
SQL> SELECT owner , TABLE_NAME FROM DBA_TABLES where table_name = '&&dummy_number';
Geben Sie einen Wert fur dummy_number ein: DUAL
alt 1: SELECT owner , TABLE_NAME FROM DBA_TABLES where table_name = '&&dummy_number'
neu 1: SELECT owner , TABLE_NAME FROM DBA_TABLES where table_name = 'DUAL'
```

```
OWNER
-----
TABLE_NAME
-----
SYS
DUAL
```

upvoted 1 times

  **adoptc94** 3 years, 4 months ago

that's because SET VERIFY ON is the default. If you don't specify SET VERIFY OFF it will use SET VERIFY ON
upvoted 3 times


Which three statements are true about a self join? (Choose three.)



- A. The ON clause must be used
- B. The query must use two different aliases for the table
- C. It must be an equijoin
- D. It must be an inner join
- E. The ON clause can be used
- F. It can be an outer join




Correct Answer: BDE




Community vote distribution



BEF (100%)



  **Rivaldo11**  3 years, 6 months ago
B, E and F
upvoted 18 times



  **ama** 3 years, 6 months ago
Good Hint
upvoted 1 times

  **escoletsgo1**  3 years, 4 months ago
B, E, F
upvoted 6 times

  **nautil2**  4 months, 1 week ago
Selected Answer: BEF
A - FALSE; use OLD syntax, e.g. SELECT ... FROM employees e, employees m WHERE e.mgr=m.empno ...
B - TRUE; when omitting, the error "ORA-00918: column ambiguously defined" appears
C - FALSE; e.g. SELECT ... FROM employees e1, employees e2 WHERE e1.msal > e2.msal
D - FALSE; valid query is e.g. SELECT e1.empno,e1.ename,e2.ename ... FROM employees e1 RIGHT OUTER JOIN employees e2 ON e1.mgr = e2.empno;
E - TRUE; valid query is e.g. SELECT e.ename employee,e.bdate birth, m.ename manager FROM employees e JOIN employees m ON e.mgr = m.empno;
F - TRUE; see answer B
upvoted 1 times



  **Oracle2020** 5 months ago
The correct answers are B,E,F
C and D are incorrect because it must not be an equijoin or an inner join. That is, it can be an equijoin , non-equijoin, inner join, outer join.
A is incorrect, because it is not mandatory for an inner join to use the on clause, if you use old inner join syntax. That is, you can join two tables in a select, without using the word inner join and on clause.
upvoted 1 times

  **wchoi189** 8 months, 4 weeks ago
I meant to say BCE
upvoted 2 times

  **wchoi189** 9 months, 1 week ago
B,C,F for me.
Joining a table to itself on a condition is called a self join. A table is joined to itself to compare rows in the same table. The join condition would have to be an equijoin. If it is not there is another name for the join and it's called non equijoins. I guess a self non equijoin would be possible but this goes against the point I'm making. A self join is joining the same table to itself ON a condition and that is why it has to be an equijoin.

For example the left outer joined table would show all the rows from the left and unmatched rows on the right filled with nulls where there was no matches. I don't think in self joins there would be any unmatched rows due to a missing column on the right.

An outer self join is not possible because it's the same table your joining. Again, you can force a outer self join but this is more of an outer join than a self join.
upvoted 1 times

  **Franky_T** 1 year, 9 months ago
Selected Answer: BEF
When you are doing a self-join, keep in mind that you are joining two "individual" data sets. There exists no restriction within Oracle that forces you to use an ON clause join rather than a WHERE clause join, or an INNER join versus an OUTER join.

upvoted 1 times

  **Datajimm** 1 year, 10 months ago

Selected Answer: BEF

- a) false: you can use WHERE clause instead
- b) true: otherwise it reports ORA-00918
- c) false: can be any condition
- d) false: can be any join
- e) true: can be used, but not necessary
- f) true: can be any join

upvoted 2 times

  **ryuah** 2 years ago

B,E,F is correct
upvoted 1 times

  **flaviogcmelo** 2 years, 9 months ago

I would like to understand why the ON clause "can" instead of "must" be used? My answer was A, B and F.
upvoted 2 times

  **nobody347** 2 years, 8 months ago



u can make self join using WHERE
upvoted 1 times

  **saif_alrwiliy** 3 years, 5 months ago

Actually the Self Join CAN be the A, C and D. The word MUST makes the question wrong.
upvoted 2 times

  **avanand** 3 years, 5 months ago

B, D , E looks correct.
I don't think self join can be outer.
upvoted 3 times

  **elvegaa** 3 years, 2 months ago

Self outer joins are also possible, and plausible.
Figure a table with employees like:
id name boss
-- -----
1 John NULL
2 Paul 1
3 Scott 1
4 Kate 3

You can build a self outer join to retrieve a list of employees names and their bosses names, if any.
upvoted 2 times

  **adoptc94** 3 years, 4 months ago

The correct answers should be B, E and F a self join can also be an outer join.

There is no difference for self joins. A self join is just two tables being joined. They just happen to be two copies of the same data
upvoted 3 times

You want to write a query that prompts for two column names and the WHERE condition each time it is executed in a session but only prompts for the table name the first time it is executed.

The variables used in your query are never undefined in your session.


Which query can be used?

- A. SELECT &&col1, &&col2 FROM &table WHERE &&condition = &&cond;
- B. SELECT &col1, &col2 FROM &&table WHERE &condition;
- C. SELECT &col1, &col2 FROM λ€&tableλ€ WHERE &condition;
- D. SELECT '&&col1', '&&col2' FROM &table WHERE '&&condition' = '&cond';
- E. SELECT &&col1, &&col2 FROM &table WHERE &&condition;

Correct Answer: B




Community vote distribution


B (100%)

  **avanand**  3 years, 5 months ago

B is correct



upvoted 8 times


  **nautil2**  4 months, 1 week ago

 **Selected Answer: B**

B – TRUE; value of variables prefixed with && is stored and not requested repeatedly

upvoted 1 times

  **zouve** 5 months, 2 weeks ago



 **Selected Answer: B**


the correct queries are :

- A. SELECT &&col1, &&col2 FROM &table WHERE &&condition = &&cond;
- B. SELECT &col1, &col2 FROM &&table WHERE &condition;
- C. SELECT &col1, &col2 FROM "&&table" WHERE &condition;
- D. SELECT '&&col1', '&&col2' FROM &table WHERE '&&condition' = '&cond';
- E. SELECT &&col1, &&col2 FROM &table WHERE &&condition;

these are copy-paste from 1Z0-071



upvoted 1 times


  **LeandroHPN** 1 year, 6 months ago

 **Selected Answer: B**

B is correct

upvoted 1 times



  **Datajimm** 1 year, 10 months ago

 **Selected Answer: B**

single & prompts for input for every run



double && prompts for input only on the first run

upvoted 1 times

  **ryuah** 2 years ago



B is correct

upvoted 1 times

  **baeji** 3 years, 1 month ago

B is correct

upvoted 2 times

  **adoptc94** 3 years, 4 months ago

Correct answer: B

Tested in SQL-Developer, col1, col2 and condition will be prompted everytime, while table will only be prompted the first time you execute this statement

upvoted 4 times

Examine the description of the CUSTOMERS table:

Name	Null?	Type
CUST_ID	NOT NULL	VARCHAR2(6)
FIRST_NAME		VARCHAR2(50)
LAST_NAME	NOT NULL	VARCHAR2(50)
ADDRESS		VARCHAR2(50)
CITY		VARCHAR2(25)



You want to display details of all customers who reside in cities starting with the letter D followed by at least two characters.
Which query can be used?



- A. SELECT * FROM customers WHERE city LIKE 'D_%';
- B. SELECT * FROM customers WHERE city = '%D_';
- C. SELECT * FROM customers WHERE city LIKE 'D_';
- D. SELECT * FROM customers WHERE city = 'D_%';



Correct Answer: A

Community vote distribution



A (100%)



  **NowOrNever** Highly Voted 3 years, 2 months ago
there is no valid answer here
correct would be (followed by at least two characters):
SELECT * FROM EMPLOYEES WHERE FIRST_NAME LIKE 'D__%';
but A is the most correct
upvoted 13 times

  **elvegaa** 3 years, 2 months ago
You're right. Also, there's no valid answer because double quotes are used in the LIKE clause
upvoted 3 times

  **SimoneF** 3 years, 2 months ago
I assume there's been some issue with copy-pasting the text, and the first ' as been replaced with "~"

I think it's supposed to be A anyway, but you're certainly correct that they're all wrong written as they are.
upvoted 2 times

  **Ekos** Highly Voted 3 years, 3 months ago
yes, A is the correct answer
upvoted 7 times

  **nautil2** Most Recent 4 months, 1 week ago
Selected Answer: A
A is the most close to be valid answer.
LIKE must be used when comparing text strings (B, D are incorrect), metacharacter % must be used to accept strings of variable length (C is incorrect), metacharacter _ stands for any character.
Underscore _ should be doubled to fulfill the request of two additional characters that follows initial letter D, so the query should end with LIKE 'D__%'
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

店长微信：hj feng128

No correct answer here, the correct SQL query would be:
SELECT * FROM EMPLOYEES WHERE FIRST_NAME LIKE 'D__%';

Notice there 2 underscores;
upvoted 1 times

  **saad3577** 1 year, 3 months ago

the table name CUSTOMERS table, from where you came with EMPLOYEES ?
ex: WHERE CustomerName LIKE 'a_%' --> Finds any values that start with "a" and are at least 2 characters in length
https://www.w3schools.com/sql/sql_like.asp
upvoted 1 times

Examine this command:

```
ALTER DATABASE
    MOVE DATAFILE '/u01/sales1.dbf'
    TO          '/u01/sales01.dbf'
    REUSE;
```


Which two statements are true? (Choose two.)

- A. DML may be performed on tables with one or more extents in this data file during the execution of this command.
- B. The tablespace containing SALES1.DBF must be altered READ ONLY before executing the command.
- C. The tablespace containing SALES1.DBF must be altered OFFLINE before executing the command.
- D. If Oracle Managed Files (OMF) is used, then the file is renamed but moved to DB_CREATE_FILE_DEST.
- E. The file is renamed and stored in the same location

Correct Answer: AB

Community vote distribution

AE (100%)

 **Rivaldo11** Highly Voted 3 years, 6 months ago

I would say A and E.
upvoted 26 times

 **NowOrNever** 3 years, 6 months ago

Yes, BCD is wrong
D - like in the posted link described
it will be moved to DB_CREATE_FILE_DEST location only, when TO is omitted
So AE is correct
upvoted 5 times

 **khalilshahin01** Highly Voted 3 years, 6 months ago

answer is A & E
Queries and DML and DDL operations can be performed while the data file is being moved, for example:
1-SELECT statements against tables and partitions
2-Creation of tables and indexes
3- Rebuilding of indexes
Other notes:
1- If objects are compressed while the data file is moved, the compression remains the same.
2- You do not have to shut down the database or take the data file offline while you move a data file to another location, disk, or storage system.
3- You can omit the TO clause only when an Oracle-managed file is used. In this case, the DB_CREATE_FILE_DEST initialization parameter should be set to indicate the new location.
4-If the REUSE option is specified, the existing file is overwritten.
note: The REUSE keyword indicates the new file should be created even if it already exists.
5-If the KEEP clause is specified, the old file will be kept after the move operation. The KEEP clause is not allowed if the source file is an Oracle-managed file.
upvoted 15 times

 **nautil2** Most Recent 4 months, 1 week ago

Selected Answer: AE

See documentation "14.5.1 Renaming and Relocating Online Data Files" in Release 19 Database Administrator's Guide
See documentation "ALTER DATABASE" in Release 19 SQL Language Reference
A - TRUE; ALTER DATABASE MOVE DATAFILE ... This statement enables you to rename or relocate a data file while the database is open and users are accessing the data file."
B - FALSE; There is no condition to alter moved datafile READ ONLY before executing ALTER ... MOVE command.
C - FALSE; The datafile is online during the movement, see answer A.
D - FALSE; the file will be moved to specified folder only if TO is omitted. If TO is not omitted, Oracle will create specified file. "If you are using Oracle Managed Files, then you can omit the TO clause. In this case, Oracle Database creates a unique name for the data file and saves it in the directory specified by the DB_CREATE_FILE_DEST initialization parameter."
E - TRUE; its obvious from the command listed in question
upvoted 1 times

 **doyinbare** 4 months, 1 week ago

A and E are correct answers. The Alter tablespace MOVE command works just like the linux mv command which either renames the file or moves the file to the specified location if it doesnt exist there.
Note, while this movement is ongoing, DML operations are allowed.
Also, the tablespace does not have to be in OFFLINE mode or READ ONLY mode for this operation to be performed.

upvoted 1 times

  **Oracle2020** 5 months ago



The correct answers are C,E, for Oracle 12c I understand
C is correct because to rename a datafile you have to put it offline. Due to the fact that it is a physical Windows file and if it is in use it does not allow you to rename it. Same as any other file you want to rename in windows
E is correct , because it is not just moving the datafile and renaming it
For Oracle 19 the correct answer is A,E. Could be?

upvoted 1 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: AE

AE is correct
upvoted 2 times

  **cratostt** 1 year, 10 months ago



The is no rename in move datafile. The only options are A,B
upvoted 1 times

  **emburria** 2 years ago



A, E. It's an online operation
upvoted 1 times

  **ryuah** 2 years ago

A,E is correct
upvoted 1 times

  **gabriel3600** 2 years, 7 months ago

A and E
upvoted 2 times

  **Ekos** 3 years, 1 month ago

the answer is A and E
upvoted 4 times

  **SimoneF** 3 years, 2 months ago

It's funny how A and B are actually contradicting one another... if the datafile is Read Only, how can you execute any DML on tables residing there?

Considering the ability of renaming datafiles online from 12c, I would say the more correct answers are A and E.

D is wrong because, evn with OMF, you can still force the creation of a datafile with a custom name on filesystem if you give the full path. If you only gave the file name, on the contrary, it would in fact create the new datafile on the default location.

upvoted 3 times

  **elvegaa** 3 years, 2 months ago

For old folks, note that "ALTER DATABASE MOVE..." syntax replaces "ALTER DATABASE RENAME..." (pre-12c era), so is no longer to set the tablespace offline and DML operations can be performed.

upvoted 3 times

  **avanand** 3 years, 5 months ago

A E

ALTER DATABASE MOVE DATAFILE ('filename' | 'ASM_filename' | file_number)
[TO ('filename' | 'ASM_filename')]
[REUSE] [KEEP]

The source file can be specified using the file number or name, while the destination file must be specified by the file name.
The REUSE keyword indicates the new file should be created even if it already exists.
The KEEP keyword indicates the original copy of the datafile should be retained.



When the source file is an OMF file the KEEP option can not be used.
If the destination file is an OMF file, the TO clause can be omitted and the file will be created with an OMF name in the DB_CREATE_FILE_DEST location.

upvoted 8 times

  **you1234** 3 years, 7 months ago

ama,

which one is correct? Plase suggest
upvoted 1 times

  **ama** 3 years, 6 months ago

I would take A, D



See below <https://oracle-base.com/articles/12c/online-move-datafile-12cr1>
upvoted 2 times

  **NiciMilo** 3 years, 6 months ago

Why D but not E?

店长微信：hj feng128

if it were OMF, i think we dont need the TO clause
upvoted 2 times

  **elvegaa** 3 years, 2 months ago

But in the question they're defining a destination in the TO clause, hence the datafile would be created in that destination, not in the OMF destination.
upvoted 3 times



  **SimoneF** 3 years, 2 months ago

I agree, and so I also think it's A and E
upvoted 3 times



  **ama** 3 years, 7 months ago

Why B is correct ? according to the Oracle documentation, Move Option is new and only possible in 12c and can be executed ONLINE without taking anything offline



<https://oracle-base.com/articles/12c/online-move-datafile-12cr1>
upvoted 2 times

  **danito** 3 years, 7 months ago

B D correct options
upvoted 1 times

  **danito** 3 years, 7 months ago

<https://oracle-base.com/articles/12c/online-move-datafile-12cr1>
upvoted 1 times

  **danito** 3 years, 5 months ago

I agree with you A E
upvoted 2 times

Which three statements are true about dropping and unused columns in an Oracle database? (Choose three.)

- A. A primary key column referenced by another column as a foreign key can be dropped if using the CASCADE option.
- B. An UNUSED column's space is reclaimed automatically when the block containing that column is next queried.
- C. An UNUSED column's space is reclaimed automatically when the row containing that column is next queried.
- D. Partition key columns cannot be dropped.
- E. A DROP COLUMN command can be rolled back
- F. A column that is set to UNUSED still counts towards the limit of 1000 columns per table

Correct Answer: ABF

Community vote distribution

ADF (100%)

- dotruonghan

Highly Voted

3 years, 6 months ago

- Answer A is RIGHT: Orders is Parent table with PRIMARY KEY ord_no, Order_items is child table which ord_no is REFERENCE KEY that reference ord_no of Parent table, now drop PRIMARY KEY on Orders by command: ALTER TABLE orders DROP COLUMN ORD_NO CASCADE CONSTRAINTS;
 - Answer D is RIGHT: ORA-12984: cannot drop partitioning column
 - Answer F is RIGHT: Unused Columns Count against 1000-column Table Limit Causing ORA-01792 on Compressed Table (Doc ID 2259600.1) and "ORA-01792: Maximum Number Of Columns In A Table Or View Is 1000" AND HIDDEN_COLUMN name is recreated with date and timestamp (Doc ID 2624150.1)

upvoted 19 times
- Ekos

3 years, 1 month ago

i agree

upvoted 1 times
- Rivaldo11

Highly Voted

3 years, 6 months ago

It should be A, D and F.

upvoted 8 times
- NowOrNever

3 years, 6 months ago

i think so too

F is true - i tested it

BC are false

A is true, but you have to add it to

ALTER TABLE <table_mame> SET unused (<column_name>) CASCADE CONSTRAINTS;

before you drop unused constraints

E should be wrong, could not find a way to undo this with roll back, this seems already the case if it is set to unused

upvoted 2 times
- nautil2

Most Recent

4 months, 1 week ago

Selected Answer: ADF

A - TRUE; e.g. ALTER TABLE departments DROP COLUMN deptno CASCADE CONSTRAINTS;

B - FALSE; In Oracle Database 19c, the space occupied by an unused column is automatically reclaimed when the row containing that column is next updated. The space is released and becomes available for reuse within the block.

C - FALSE; see B

D - TRUE; partition key column cannot be dropped. When an attempt to drop the column is made, an error "ORA-12984: cannot drop partitioning column" occurs.

E - TRUE; ALTER TABLE documentation says "Until you actually drop these columns, they continue to count toward the absolute limit of 1000 columns in a single table."

upvoted 1 times
- nautil2

4 months, 1 week ago

I am sorry, last two options are

E - FALSE; DDL commands cannot be rolled back

F - TRUE; ALTER TABLE documentation says "Until you actually drop these columns, they continue to count toward the absolute limit of 1000 columns in a single table."

upvoted 1 times
- Oz0n3

8 months, 1 week ago

Selected Answer: ADF

D is correct, partition table will not work otherwise

Same with A and F

upvoted 1 times

- cristycool** 9 months, 1 week ago

A is correct
B is correct
F is correct : Unused columns count against the 1000-column limit for tables in an Oracle database. The total count of columns that count against this limit can be obtained by getting a count from DBA_TAB_COLS

upvoted 1 times
- alexbraga93** 1 year, 2 months ago

Selected Answer: ADF

B is not correct, once, the data remains in the table, but not available for querying.

upvoted 1 times
- trgbighero** 1 year, 3 months ago

Selected Answer: ADF

ADF is correct

upvoted 1 times
- Phineas** 1 year, 6 months ago

Selected Answer: ADF

Its ADF

upvoted 1 times
- mitosenoriko** 1 year, 7 months ago

Selected Answer: ADF

A and D and F

upvoted 1 times
- emburria** 2 years ago

A, D and F

upvoted 1 times
- ryuah** 2 years ago

A,D,F is correct

upvoted 1 times
- gabriel3600** 2 years, 7 months ago

A,D e F

upvoted 1 times
- dotruonghan** 3 years, 6 months ago

F is true
Unused Columns Count against 1000-column Table Limit Causing ORA-01792 on Compressed Table (Doc ID 2259600.1)

upvoted 2 times
- NiciMilo** 3 years, 6 months ago

Is B really correct?

According to:
<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-tables.html#GUID-29CA8FDC-55FB-4BEF-AE4A-5D1B4D2BE1E1>

"The ALTER TABLE...DROP UNUSED COLUMNS statement is the only action allowed on unused columns. It physically removes unused columns from the table and reclaims disk space. "

upvoted 2 times

Which two statements are true regarding Oracle database space management within blocks managed by Automatic Segment Space Management (ASSM)?
(Choose two.)

- A. PCTFREE defaults to 10% for all blocks in all segments for all compression methods
- B. ASSM assigns blocks to one of four fullness categories based on what percentage of the block is allocated for rows
- C. Update operations always attempt to find blocks with free space appropriate to the length of the row being updated
- D. Insert operations always attempt to find blocks with free space appropriate to the length of the row being inserted
- E. A block will always be eligible for inserts if the row is short enough to fit into the block

Correct Answer: CE

Community vote distribution



ioio Highly Voted 2 years, 9 months ago
the answer is B and D

Characteristics of Bitmap Segment Management

Bitmap space management uses four bits inside each data block header to indicate the amount of available space in the data block. Unlike traditional space management with a fixed relink and unlink threshold, bitmap space managements allow Oracle to compare the actual row space for an INSERT with the actual available space on the data block. This enables better reuse of the available free space especially for objects with rows of highly varying size. Here are the values inside the four-bit space:

Value	Meaning
0000	Unformatted Block
0001	Block is logically full
0010	<25% free space
0011	>25% but <50% free space
0100	> 50% but <75% free space
0101	>75% free space

Table 1: Bitmap value meanings.
upvoted 6 times

nautil2 Most Recent 4 months, 1 week ago
Selected Answer: BD
A - FALSE; default 0 for compressed tables
B - TRUE; each block is divided into four sections, named FS1 (between 0 and 25% of free space), FS2 (25% to 50% free), FS3 (50% to 75% free), and FS4 (75% to 100% free)
C - FALSE; only before INSERT operations free space in block is tested and when there is a lack of space, new blocks are allocated. In case of UPDATE operations, when there is a lack of space, row migration occurs.
D - TRUE; see C
E - FALSE; if ASSM is enabled and free space in a block is less than PCTFREE, insert to the block will be rejected
upvoted 1 times

Oracle2020 5 months ago
Exam Topic is not supposed to have already responded to the answer as correct. CY E. Why do you argue?
upvoted 1 times

auwia 6 months, 2 weeks ago
Selected Answer: BE
The magic word is BLOCK.
upvoted 1 times

auwia 6 months, 2 weeks ago
I believe C and D are wrong because I've seen many cases where Oracle "recycle" blocks deleted that are still occupying space, but they are in the freelist so can be used to insert or updated data. So the clause ALWAYS in question makes both False in my opinion.
A is False because there is not fixed default for PCTFREE.
By exclusion correct answer are BE.
upvoted 2 times

jonsnoow 1 year ago
B and D
upvoted 1 times

🗨️ 👤 **trgbighero** 1 year, 3 months ago

店长微信：hj feng128

Selected Answer: AD

my teacher said AD
upvoted 1 times

🗨️ 👤 **Darkseid1231** 1 year, 7 months ago

AMA, what is the correct answer please ?
upvoted 1 times

🗨️ 👤 **Franky_T** 1 year, 9 months ago

Selected Answer: BD

A is wrong for obvious reasons. All segments include (tables, indexes, undo and temp). On top of that "all" compression is mentioned in the answer.
B is correct. This is a well known feature of ASSM and can be queried via the dbms_space.SPACE_USAGE procedure.
C is wrong. The key to this answer is "always". If a block has sufficient free space to accomodate the row update then Oracle does not need to search for a different block with enough free space.
D is correct. Oracle does this by scanning the BMB in the segment to determine a block that can store the entire row. The main purpose here is to avoid row chanining.
E is wrong. It all depends what the PCTFREE value is for the segment. A block with a PCTFREE of 10% could have rows that uses 90% of the available space, that block is deemed full and not eligible for insert operations.
upvoted 3 times

🗨️ 👤 **emburria** 2 years ago

Selected Answer: BC

B,C for sure
upvoted 1 times

🗨️ 👤 **emburria** 2 years ago

Selected Answer: BD

B,D are correct
upvoted 1 times

🗨️ 👤 **ryuah** 2 years ago

B,D is correct
upvoted 1 times

🗨️ 👤 **Fan** 2 years, 10 months ago

It's A , D.
upvoted 1 times

🗨️ 👤 **Eric_F** 1 year, 10 months ago

PCTFREE dufault is 0 when used COMPRESS BASIC, so A is incurrect.
upvoted 1 times

🗨️ 👤 **ogdru** 2 years, 10 months ago

A...for all compression methods is not correct!!
create table T1
(col1 varchar2(100));

create table T2
(col1 varchar2(100))
row store compress basic;

create table T3
(col1 varchar2(100))
row store compress advanced;

SQL> select TABLE_NAME,COMPRESSION,COMPRESS_FOR,PCT_FREE from dba_tables where OWNER='SCOTT' order by TABLE_NAME;

TABLE_NAME COMPRESS COMPRESS_FOR PCT_FREE

T1 DISABLED 10
T2 ENABLED BASIC 0
T3 ENABLED ADVANCED 10

This make me very confused!!
upvoted 1 times

🗨️ 👤 **adoptc94** 3 years, 5 months ago

Someone has any suggestion?



A could be true as the default value for pctfree is 10, but I am not sure for all compression methods



B should be true with ASSM each block is divided into 4 sections named FS1 (between 0 and 25% free space), FS2 (25% to 50% free space), FS3 (50% to 75% free space) and FS4 (75% to 100% free space)

E is not true, a block is only eligible for inserts if it's not full! A block is marked as full when the pctfree threshold is reached, at this point the block is only available for update or delete operations. The block will be available for inserts again only after the block gets under the pctfree treshold



I this case I'd go with answers: B, C



But I am not completely sure, to me they seem like the right answers
upvoted 3 times



  **adoptc94** 3 years, 4 months ago
Update to answer A - this answer is false. PCTFREE is implicitly zero on compressed tables
upvoted 2 times

  **NowOrNever** 3 years, 4 months ago
there are these 4 fullness "categories", but 6 values to represent the fullness
Value | Meaning | SPACE_USAGE Procedure
0000 | Unformatted Block | unformatted_blocks
0001 | Block is logically full | full_blocks
0010 | <25% free space | fs1_blocks
0011 | >=25% but <50% free space | fs2_blocks
0100 | >=50% but <75% free space | fs3_blocks
0101 | >=75% free space | fs4_blocks

i still would agree to B
upvoted 2 times

  **SimoneF** 3 years, 1 month ago
I agree with B, but I would go on D rather than C. Update operations should first try to fit the new data in the same block and only when it's impossible they will look for other blocks where to fit the complete row.
A and E have already been explained to be false, which upon which I agree.
upvoted 6 times

  **Ekos** 3 years, 1 month ago
agree, the answer is B and D
upvoted 6 times

  **you1234** 3 years, 6 months ago
A. PCTFREE defaults to 10% for all blocks in all segments for all compression methods
B. ASSM assigns blocks to one of four fullness categories based on what percentage of the block is allocated for rows
C. Update operations always attempt to find blocks with free space appropriate to the length of the row being updated
D. Insert operations always attempt to find blocks with free space appropriate to the length of the row being inserted
E) The first block with enough free space to accommodate a row being inserted will always be used for that row.
upvoted 1 times

Evaluate these commands which execute successfully:

```
CREATE SEQUENCE ord_seq
  INCREMENT BY 1
  START WITH 1
  MAXVALUE 100000
  CYCLE
  CACHE 5000;

CREATE TABLE ord_items (
  ord_no      NUMBER(4) DEFAULT ord_seq.NEXTVAL NOT NULL,
  item_no     NUMBER(3),
  qty         NUMBER(3),
  expiry_date DATE,
  CONSTRAINT it_pk PRIMARY KEY (ord_no, item_no),
  CONSTRAINT ord_fk FOREIGN KEY (ord_no) REFERENCES orders (ord_no));
```

Which two statements are true about the ORD_ITEMS table and the ORD_SEQ sequence? (Choose two.)

- A. If sequence ORD_SEQ is dropped then the default value for column ORD_NO will be NULL for rows inserted into ORD_ITEMS
- B. Any user inserting rows into table ORD_ITEMS must have been granted access to sequence ORD_SEQ
- C. Column ORD_NO gets the next number from sequence ORD_SEQ whenever a row is inserted into ORD_ITEMS and no explicit value is given for ORD_NO
- D. Sequence ORD_SEQ cycles back to 1 after every 5000 numbers and can cycle 20 times
- E. Sequence ORD_SEQ is guaranteed not to generate duplicate numbers

Correct Answer: BE

Community vote distribution



ama Highly Voted 3 years, 6 months ago
i double checked this Question again.

i am sure B,C are correct
upvoted 25 times

Ekos 3 years, 1 month ago
yes, agree
upvoted 3 times

Sharif1 Highly Voted 3 years, 5 months ago
BC are correct.
upvoted 9 times


Abdullejr Most Recent 1 week, 3 days ago
Selected Answer: C
There is only one correct answer which is C.
B is incorrect because you can insert into the table without using the default value of ord_no i.e.
INSERT INTO ord_items (ord_no,item_no, qty, expiry_date)
VALUES (3,25,11, sysdate +10) ;
upvoted 1 times

mamadu 6 months, 2 weeks ago
Selected Answer: BC
B&C are true.
upvoted 1 times

auwia 6 months, 4 weeks ago
Selected Answer: BC

- A. If sequence ORD_SEQ is dropped then the default value for column ORD_NO will be NULL for rows inserted into ORD_ITEMS --> False, the column is NOT NULL.
- B. Any user inserting rows into table ORD_ITEMS must have been granted access to sequence ORD_SEQ --> True
- C. Column ORD_NO gets the next number from sequence ORD_SEQ whenever a row is inserted into ORD_ITEMS and no explicit value is given for ORD_NO --> True
- D. Sequence ORD_SEQ cycles back to 1 after every 5000 numbers and can cycle 20 times --> False, it goes back to 1 after 1.000.000...
- E. Sequence ORD_SEQ is guaranteed not to generate duplicate numbers --> False, having CYCLE in the create statement, duplicates will arrive after the first cycle.


upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: BC

B,C checked

upvoted 1 times

  **477267** 10 months, 3 weeks ago

BE

I tired to run the sequence again starting from one and giving max value 10 to check if duplicate value is inserted ,

But it complains as below

insert into ord_items (ITEM_NO,QTY,EXPIRY_DATE,ORD_NO)values (1,12,'03-MAR-23',ord_seq.NEXTVAL)
*

ERROR at line 1:

ORA-00001: unique constraint (SYS.IT_PK) violated

So bC OF THE CONSTRAINITY ,iT WILL NEVER BE DUPLICATED

upvoted 1 times

  **Dbi** 11 months, 2 weeks ago

For me B is false. You can insert into the table without having been granted access by not using the column's default value

upvoted 2 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BC

BC is corrent

upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

BC, Correct answers.

upvoted 2 times

  **BM2000** 1 year, 8 months ago

Selected Answer: BE

B: GRANT SELECT ON my_user.my_seq TO another_user; —to give access to others

For E create sequence and table with smaller values:

```
CREATE sequence ord_seq  
INCREMENT BY 1  
START WITH 1  
MAXVALUE 1000  
CYCLE  
CACHE 500;
```

```
CREATE TABLE ord_items (ord_no number(1) DEFAULT ord_seq.NEXTVAL NOT NULL, item_no number(3), PRIMARY KEY (ord_no, item_no));
```

```
select * from ord_items;
```

insert into ord_items values(default,1); — keep entering until it will not take any more, keep going after few errors

After getting values from 1 to 9 it will give this error

“value larger than specified precision allowed for this column”

Sequence never get the chance to CYCLE so E is correct in this scenario only.

upvoted 1 times

  **ryuah** 2 years ago

B,C is correct

upvoted 1 times

  **Guhborges** 2 years, 3 months ago



B, C for sure.

upvoted 1 times

  **gabriel3600** 2 years, 7 months ago



B and C



upvoted 1 times



  **kawsar** 2 years, 9 months ago



Is there any possibility D is right...Hence C and D are correct

upvoted 1 times

  **escoletsgo1** 3 years, 4 months ago
BC 100%
upvoted 7 times

  **NowOrNever** 3 years, 6 months ago
BC are correct
ADE are wrong
upvoted 8 times

  **ama** 3 years, 6 months ago
Wrong, A C are correct..
B is wrong, i tested it
upvoted 2 times

  **elbelgounetos** 2 years, 10 months ago
A is wrong dropping the sequence lead to ORA-02289 : sequence does not exists
upvoted 1 times

Which three instance situations are possible with the Oracle Database server without multi-tenant? (Choose three.)



- A. two or more instances on separate servers all associated with one database
- B. one instance on one server associated with one database
- C. one instance on one server associated with two or more databases on the same server
- D. one instance on one server not associated with any database
- E. one instance on one server associated with two or more databases on separate servers



Correct Answer: CDE

Community vote distribution



ABD (100%)



  **Rivaldo11** Highly Voted 3 years, 6 months ago
A, B and D.
upvoted 21 times



  **Ekos** 3 years, 1 month ago
yes, i agree
upvoted 2 times



  **ama** Highly Voted 3 years, 6 months ago
A: RAC 2,3 or 4 nodes
B: Single instance (standalone) with database
C: one instance running multiple INDEPENDENT databases!!
D: Single instance (standalone) without database
E: absurd configuration!



corrects are A, B and D
upvoted 10 times



  **elbelgounetos** 2 years, 10 months ago
It helped me : https://asktom.oracle.com/pls/apex/f?p=100:11:0::::P11_QUESTION_ID:1631683800346891854
upvoted 1 times

  **Abdullejr** Most Recent 1 week, 3 days ago
Selected Answer: ABD
A, B and D are the correct answers.
upvoted 1 times

  **nautil2** 4 months, 1 week ago
Selected Answer: ABD
A - TRUE; Oracle RAC
B - TRUE; typical configuration for standalone databases
C - FALSE; valid relation between database and instance is 1:N
D - TRUE; standalone database in NOMOUNT/MOUNTED state
E - FALSE; see A
upvoted 1 times



  **nautil2** 4 months, 1 week ago
correction: E - FALSE; see C
upvoted 1 times



  **auwia** 6 months, 2 weeks ago
Selected Answer: ABD
0 or 1 DB! :-)
upvoted 1 times

  **trgbighero** 1 year, 3 months ago
Selected Answer: ABD
ama is right.
A: RAC 2,3 or 4 nodes
B: Single instance (standalone) with database
C: one instance running multiple INDEPENDENT databases!!

D: Single instance (standalone) without database
E: absurd configuration!
upvoted 1 times



店长微信：hj feng128

  **algerianphoenix** 1 year, 4 months ago
A,B,D without a doubt !
upvoted 1 times



  **LeandroHPN** 1 year, 6 months ago

Selected Answer: ABD



A, B and D correct.
upvoted 1 times

  **hadiwuu** 1 year, 8 months ago
AT THIS POINT I REALLY THINK THIS WEBSITE CREATED BY ORACLE TO MAKE PEOPLE FAIL THEIR EXAMS AND REPEAT OVER AND OVER FOR PROFIT!

HOW COME THEY CAN MESSED UP THIS SIMPLE QUESTION WITH WRONG ANSWER??
upvoted 1 times



  **zouve** 5 months, 1 week ago
Yes, it's true st least they provide the questions for the preparation. However, it would be beneficial for us if they hired an experienced DBA to answer all of them before posting them.



In my understanding, all the sites are doing the same thing. They post the same material with the same answers. At least here, we have a forum for each answer where we can support each other with the correct ones.
upvoted 1 times

  **emburria** 2 years ago

Selected Answer: ABD

A, B, D
upvoted 1 times

  **ryuah** 2 years ago
A,B,D is correct
upvoted 1 times


  **ioio** 2 years, 9 months ago
A, B and D.
https://docs.oracle.com/cd/E11882_01/server.112/e40540/startup.htm#CNCPT955
upvoted 1 times

Which two statements are true about the ORDER BY clause when used with a SQL statement containing a SET operator such as UNION? (Choose two.)

- A. Column positions must be used in the ORDER BY clause
- B. Only column names from the first SELECT statement in the compound query are recognized
- C. The first column in the first SELECT of the compound query with the UNION operator is used by default to sort output in the absence of an ORDER BY clause
- D. Each SELECT statement in the compound query must have its own ORDER BY clause
- E. Each SELECT statement in the compound query can have its own ORDER BY clause

Correct Answer: BE

Community vote distribution



danito Highly Voted 3 years, 7 months ago
B C are the corrects
upvoted 23 times

ama 3 years, 7 months ago
evidence ?
upvoted 1 times

Ekos 3 years, 1 month ago
i agree
upvoted 2 times

you1234 Highly Voted 3 years, 6 months ago
B & C is correct answer
upvoted 5 times

ama 3 years, 6 months ago
Yes!!!
B and C are correct
upvoted 4 times

Abdullejr Most Recent 1 week, 3 days ago
Selected Answer: BC
B and C without doubt.
upvoted 1 times

jagz3000 1 month, 2 weeks ago
Selected Answer: BC
i agree
upvoted 1 times

nautil2 4 months, 1 week ago
Selected Answer: BC

A - FALSE; column ord_no is part of the primary key, hence it cannot be NULL. There is no default value after sequence is dropped. When an attempt to insert a row without ord_no specified is made, an error ORA-02289 occurs.
B - TRUE; if GRANT SELECT is missing, an error "ORA-00942: table or view does not exists" occurs after INSERT attempt
C - TRUE; tested
D - FALSE; it cycles after value 100000 is reached, so after value 100000, next value is 1
E - FALSE; sequence cycles, so when MAXVALUE is reached, duplicates can occur
upvoted 1 times



nautil2 4 months, 1 week ago
I am sorry, answers above belong to another question. Right answers are:
A - FALSE; e.g. SELECT dept_no,ename FROM emp UNION SELECT deptno,ename FROM employees ORDER BY ename;
B - TRUE; tested; not documented in SQL Reference, but works
C - TRUE; tested; NULLs are last in output
D - FALSE; Release 19 SQL Reference: "Also, the ORDER BY clause can appear only in the last component query. The ORDER BY clause orders all rows returned by the entire compound query."
E - FALSE; see D

upvoted 1 times

  **Oracle2020** 5 months ago

The E is not correct because in the Union statement composition you add an ORDER BY clause to the last SELECT statement
A is incorrect because the column composition has no reason to be used in the order by, you can use only one of the columns to sort only the first select list , if it has more than one column.
B is correct because indeed the names of the columns of the first select list are those that are shown in the output of the query and are recognized
C is correct
So by default the correct answers are B and C

upvoted 1 times

  **zouve** 5 months, 3 weeks ago



Selected Answer: BC
BC for me

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Tip: to answer this question you should read the answers from bottom and try to find out a possible reason why E is false! :)

upvoted 1 times

  **mamadu** 6 months, 2 weeks ago



Selected Answer: BC
correct statements.

upvoted 1 times

  **auwia** 6 months, 4 weeks ago

Selected Answer: BE
Provided answer are correct:
A. Column positions must be used in the ORDER BY clause --> False, you can use column name too.
B. Only column names from the first SELECT statement in the compound query are recognized --> True
C. The first column in the first SELECT of the compound query with the UNION operator is used by default to sort output in the absence of an ORDER BY clause --> False, to have an order, you must use ORDER BY clause.
D. Each SELECT statement in the compound query must have its own ORDER BY clause. False, ORDER BY clause is optional.
E. Each SELECT statement in the compound query can have its own ORDER BY clause --> True

upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

C = true , try querying a table without an order by, you will still get the records ordered by the first column of the table.

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

it's just a case because probably you used few records, try with milion records maybe splitted in partitions ... maybe it's still working, but Oracle doesn't guarantee you any order if you don't use the ORDER BY clause. At least you can try to find an official Oracle link saying that the sentence is True, but I don't think you'll be able to find one! :)

upvoted 1 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BC
BC is correct

upvoted 1 times

  **algerianphoenix** 1 year, 4 months ago

B,C Correct answers.

upvoted 1 times

  **emburria** 2 years ago

B, C are correct. It's incredible how (nearly) all the answers they give are wrong

upvoted 2 times

  **ryuah** 2 years ago

B,C is correct

upvoted 1 times

  **mianjee** 2 years, 5 months ago

```
select employee_id,last_name "Last Name"
FROM employees
where employee_id=101
union
select employee_id EMP_NO, last_name
from employees
where employee_id=113;
C is correct.
```

upvoted 1 times

  **you1234** 3 years, 6 months ago

Please ignore this - I want to type for other question

店长微信：hj feng128

B is wrong - you can insert if you have insert privilege on table from other schema. not required separate privileges on sequence.
upvoted 1 times

  **you1234** 3 years, 6 months ago

B is wrong - you can insert if you have insert privilege on table from other schema. not required separate privileges on sequence.
upvoted 1 times

Which four account management capabilities can be configured using Oracle profiles? (Choose four.)

- A. the number of hours for which an account is locked after the configured number of login attempts has been reached
- B. the number of days for which an account may be inactive before it is locked
- C. the maximum amount of CPU time allowed for a user's sessions before their account is locked
- D. the ability to prevent a password from ever being reused
- E. the number of password changes required within a period of time before a password can be reused
- F. the number of days for which an account is locked after the configured number of login attempts has been reached
- G. the maximum number of sessions permitted for a user before the account is locked

Correct Answer: ADEF

Community vote distribution



NowOrNever Highly Voted 3 years, 4 months ago

Correct Answer: ABEF
Wrong Answer: CDG
A: 1/24 is one hour, 1/96 is 15 minutes
B: INACTIVE_ACCOUNT_TIME
C: the account will not be locked, it can not exceed this limit
D: not possible so far
E:
PASSWORD_REUSE_MAX - defines the number of changes required before a password can be reused
PASSWORD_REUSE_TIME - Sets the number of days before which a password cannot be reused.
F: PASSWORD_LOCK_TIME
G: the account will not be locked, it can not exceed this limit
upvoted 13 times

Ekos 3 years, 1 month ago
i agree
upvoted 2 times

avanand Highly Voted 3 years, 5 months ago
A D E F

A. the number of hours for which an account is locked after the configured number of login attempts has been reached -- [it accepts in days,you can enter fraction of days to achieve hours].
B. the number of days for which an account may be inactive before it is locked [Inactive is not a state for an account]
C. the maximum amount of CPU time allowed for a user's sessions before their account is locked [account never get locked in such settings]
D. the ability to prevent a password from ever being reused [PASSWORD_REUSE_MAX]
E. the number of password changes required within a period of time before a password can be reused [PASSWORD_REUSE_TIME]
F. the number of days for which an account is locked after the configured number of login attempts has been reached [PASSWORD_LOCK_TIME]
G. the maximum number of sessions permitted for a user before the account is locked [account never get locked in such settings]
upvoted 12 times

auwia 6 months, 2 weeks ago
B, it's true:

INACTIVE_ACCOUNT_TIME
Specify the permitted number of consecutive days of no logins to the user account, after which the account will be locked. The minimum value is 15 days. The maximum value is 24855. If you omit this clause, then the default is UNLIMITED.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-PROFILE.html#GUID-ABC7AE4D-64A8-4EA9-857D-BEF7300B64C3>
upvoted 2 times

ama 3 years, 5 months ago
This time i have to give you right ! LOL

Thx
upvoted 3 times

Darkseid1231 7 months, 1 week ago
B D E F because to lock an account after a configured number of login attempts is reached, the value is set in days, not hours. So A is wrong
upvoted 2 times

  **elvegaa** 3 years, 2 months ago

Not D, but B.
D - PASSWORD_REUSE_MAX and PASSWORD_REUSE_TIME cannot avoid to EVER reuse a password. You can set a max number of passwords and days before you can reuse a password, but not forever.
B - Since 12.2 theres a new parameter The INACTIVE_ACCOUNT_TIME profile parameter locks a user account that has not logged in to the database instance in a specified number of days.

Users are considered active users if they log in periodically. The INACTIVE_ACCOUNT_TIME timing is based on the number of days after the last time a user successfully logs in.

Ref <https://docs.oracle.com/en/database/oracle/oracle-database/12.2/dbseg/configuring-authentication.html#GUID-ED98E6DA-A30C-4052-A343-B516CD641737>
upvoted 3 times

  **SimoneF** 3 years, 1 month ago

But for D, there is this, speaking about PASSWORD_REUSE_MAX and PASSWORD_REUSE_TIME:



"If you specify a value for either of these parameters and specify UNLIMITED for the other, then the user can never reuse a password."

source: https://docs.oracle.com/database/121/SQLRF/statements_6012.htm#SQLRF01310
upvoted 5 times

  **SimoneF** 3 years, 1 month ago

I would say instead that E is wrong, since it doesn't define "the number of password changes required within a period of time" but the time that has to pass before reusing the same password.



For this and what you all said before, I would go with A-B-D-F
upvoted 8 times

  **Emperor11** 10 months ago

I agree if you set the PASSWORD_REUSE_TIME and PASSWORD_REUSE_MAX then you cannot reuse your password even if you change the password all times that is required, you'll need to wait the time to use it again, A-B-D-F and hope for luck
upvoted 1 times

  **elbelgounetos** 2 years, 10 months ago

I agree with your analyse? E cannot be true. And D is in the documentation so A B D F
upvoted 1 times



  **Chebas** 1 year, 6 months ago

If you specify a value for either of these parameters and specify UNLIMITED for the other, then the user can never reuse a password.
<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-PROFILE.html#GUID-ABC7AE4D-64A8-4EA9-857D-BEF7300B64C3>
upvoted 1 times

  **nautil2** Most Recent 4 months, 1 week ago



Selected Answer: ABDE

A - TRUE; PASSWORD_LOCK_TIME 3/24 = 3 hours
B - TRUE; INACTIVE_ACCOUNT_TIME
C - FALSE; RESOURCE=KERNEL, RESOURCE_NAME=CPU_PER_SESSION or CPU_PER_CALL allows to set limit of CPU usage, but user account is not locked when limit is reached
D - TRUE; PASSWORD_REUSE_TIME and PASSWORD_REUSE_MAX, If you specify a value for either of these parameters and specify UNLIMITED for the other, then the user can never reuse a password.
E - TRUE; same parameters as those mentioned in D
F - FALSE; PASSWORD_LOCK_TIME ... see A ... this parameter is related to FAILED login attempts, not to all login attempts
G - FALSE; SESSIONS_PER_USER, if limit is reached by a user, the user is not locked, but prevented to establish a new connection.
upvoted 2 times

  **zouve** 5 months, 1 week ago

Selected Answer: BDEF

A is wrong cause PASSWORD_LOCK_TIME specifies the number of days NOT the number of hours for which the account is locked after the specified number of failed login attempts
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

Selected Answer: BDEF

There is a confusion between A & F but if you look at the documentation the value for PASSWORD_LOCK_TIME is in days not hours.
upvoted 1 times



  **Methio** 8 months ago

the answer is BDEF
upvoted 2 times

  **j_tw** 9 months ago

A B D F
A : PASSWORD_LOCK_TIME
B : INACTIVE_ACCOUNT_TIME

C : false
D : PASSWORD_REUSE_MAX = UNLIMITED
E : ? 'within a period of time'?
F : PASSWORD_LOCK_TIME
G : false
upvoted 3 times

  **piontk** 10 months, 2 weeks ago

Selected Answer: ABEF
ABEF is the correct
upvoted 1 times

  **Tunglinfo9** 1 year, 1 month ago

Selected Answer: ABEF
A: password_lock_time: number of days after account is locked
you can use"alter profile xxx limit password_lock_time 5/24" it works
B: inactive_lock_time
E: password_reuse_time & password_reuse_max
F: password_lock_time, like the option A
upvoted 1 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BDEF
I choose BDEF
upvoted 2 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BDEF
i choose BDEF
upvoted 2 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BDEF
i choose BDEF
upvoted 1 times

  **daniel33** 1 year, 5 months ago

Selected Answer: BDEF
I believe BDEF to be the correct answer. People are debating "B", however B is correct according to:
https://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_6010.htm
"If you specify an integer for either of these parameters and specify UNLIMITED for the other, then the user can never reuse a password."

as the above pertains to parameters: PASSWORD_REUSE_TIME and PASSWORD_REUSE_MAX
upvoted 1 times



  **OCP19c** 1 year, 6 months ago

Correct Answers are: A, B, D, E
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: ABDE
Firstly, I have an issue with the answers, there are 5 correct ones and not 4.



A is correct. The time interval is not important here, but you can lock an account for a specified time when this threshold has been reached (failed_login_attempts & password_lock_time).
B is correct. Can be set with the incative_account_time of the profile.
C is wrong. When the CPU limit has been reached the account is not locked, but rather the session receives either an ORA-02392 or an ORA-02393 errors.
D is correct. This can be done by setting a value for either password_reuse_time or password_reuse_max and unlimited for the other, then a user can never reuse a password.
E is correct. Again we use the password_reuse_time and password_reuse_max parameters (must be used together to have any effect).
F is correct. It's essentially the same answer as option A, the only difference being the time interval.
G is wrong. Accounts are not locked when this threshold is reached, additional sessions are simply denied with error ORA-00018.
upvoted 2 times

  **nautil2** 4 months, 1 week ago

I agree with your answers, except of F. Answer F is IMHO wrong, because it mentions all login attempts, not failed login attempts. "number of login attempts"
upvoted 1 times

  **ryuah** 2 years ago

A,D,E,F is correct
upvoted 1 times

  **krontar** 2 years, 4 months ago

I think that we have 5 correct answers.
My guess is that this question is older than the "INACTIVE_ACCOUNT_TIME " feature.
Correct answers are:
A B D E F
upvoted 2 times




Which three statements are true about single-row functions? (Choose three.)



- A. They can be used only in the WHERE clause of a SELECT statement
- B. The argument can be a column name, variable, literal or an expression
- C. The data type returned can be different from the data type of the argument
- D. They can be nested to any level
- E. They can accept only one argument
- F. They return a single result row per table



Correct Answer: BCD



Community vote distribution





BCD (100%)




  **avanand**  3 years, 8 months ago
Correct choice is B C D.
single row function can be nested to any level.
upvoted 14 times



  **danito** 3 years, 7 months ago
D is not correct, more than three nested functions return ora0935
upvoted 2 times



  **elvegaa** 3 years, 2 months ago
I had no such error in 11.2
select DECODE(COALESCE(NVL(ROUND(TRUNC(5)),4),3),2)1) from dual;
upvoted 2 times

  **saidov1981** 2 years, 9 months ago
statement correction:
select decode(coalesce(nvl(round(trunc(5)),4),3),5,5,1) from dual; -- or
select decode(coalesce(nvl(round(trunc(5)),4),3),2,1) from dual;
upvoted 1 times



  **nautil2**  4 months, 1 week ago
 **Selected Answer: BCD**
A - FALSE; can be used after SELECT, in WHERE clause, in START WITH clause, in CONNECT BY clause and in HAVING clause
B - TRUE
C - TRUE; typically conversion functions, e.g. TO_CHAR, TO_DATE
D - TRUE
E - FALSE; they can accept more arguments, e.g. TO_DATE('2023-09-19 16:15','YYYY-MM-DD HH24:MI')
F - FALSE; group functions return a single result per group, while single-row functions returns a result for each input value
upvoted 1 times

  **emburria** 2 years ago
 **Selected Answer: BCD**
B, C, D for sure
upvoted 1 times

  **ryuah** 2 years ago
B,C,D is correct
upvoted 1 times



  **Guhborges** 2 years, 3 months ago
BCDE is right for me.



E- We can use upper for only one and we can use CONCAT for more than two, isn't that correct?
upvoted 1 times

  **Guhborges** 2 years, 3 months ago
Why not E? They can only accept one argument, they can also accept more. Is different when we use the word MUST.... I'm not a native English speaker...
upvoted 1 times



  **dbaclaudio** 3 years ago



B,C,D
A - Incorrect -> Documentation: "can appear in select lists, WHERE clauses, START WITH and CONNECT BY clauses, and HAVING clauses."
B - Correct -> EXAMPLE: select lower(&data), lower(DUMMY), lower('TABLE'),LOWER(TO_DATE('12-JAN-08','DD-MONTH-RR')) from dual;
RESULT: 12-jan-21 x table 12-jan-08
C - Correct -> Like in length or months_between
D - Correct -> there's no limit in documentation. EXAMPLE: select instr(substr(decode(mod(6,4),4,'abcbcabcb',1,'bcabcbca','cabcbcab'),4),'abc')
from dual;
RESULT: 2
E - Incorrect -> Functions like substr, accepts more than one parameter
F - Incorrect -> Here's a catch, documentation: "Single-row functions return a single result row for every row of a queried table or view" NOT
"...return a single result row per table"
upvoted 4 times



  **thuto** 1 year, 5 months ago
Thank you for Clarifying F...
upvoted 1 times



  **elvegaa** 3 years, 2 months ago
B C D & E seems right to me :\



I guess.. D is the most close to wrong, since i'm sure there's some kind of nesting limit, but this limit is not specified in oracle documentation.
upvoted 2 times



  **SimoneF** 3 years, 1 month ago
There is no such limit for single-row functions, and E is certainly false: for instance, CONCAT is a single-row function and it accepts two arguments.
upvoted 2 times

  **Ekos** 3 years, 3 months ago
C, D, and E
upvoted 1 times



  **Ekos** 3 years, 1 month ago
Sorry, B,C,D is the correct one
upvoted 2 times

  **dkmraju** 3 years, 4 months ago
Correct one B, C, F
upvoted 2 times

  **you1234** 3 years, 6 months ago
B & C & D is correct answer
upvoted 4 times

  **ama** 3 years, 6 months ago
yes, D is also correct ...

see Question 30 > https://www.tutorialspoint.com/sql_certificate/using_single_row_functions_questions.htm
upvoted 1 times

  **GraceYu** 3 years, 7 months ago
This is note regards single row function and aggregate function

<https://oracledbars.blogspot.com/2018/01/using-conversion-functions-and.html>.
Nesting Functions
Single-row functions can be nested to any level. • Nested functions are evaluated from the deepest level to the least deep level.

<https://mahtodeepak05.wordpress.com/2014/12/17/aggregate-function-nesting-in-oracle/>
Hence Two level Deep is the max you can go with nested aggregate functions in Oracle.
Logically that the way it should be, Third aggregation functions wont have any group to act on as previous aggregations functions “SUM” had already produce singular output.
upvoted 2 times

You want to use table compression suitable for OLTP that will:

- 1. Compress rows for all DML statements on that table
- 2. Minimize the overheads associated with compression

Which compression option is best suited for this?

- A. COLUMN STORE COMPRESS FOR QUERY LOW
- B. ROW STORE COMPRESS BASIC
- C. COLUMN STORE COMPRESS FOR ARCHIVE LOW
- D. COLUMN STORE COMPRESS FOR ARCHIVE HIGH
- E. ROW STORE COMPRESS ADVANCED

Correct Answer: E

Reference:

<https://www.oracle.com/technetwork/database/options/compression/advanced-compression-wp-12c-1896128.pdf>

Community vote distribution



E (100%)

  **avanand** Highly Voted 3 years, 5 months ago

E is correct

<https://docs.oracle.com/database/121/ADMIN/tables.htm#ADMIN-GUID-34D15DD1-0925-4C9A-BE8A-3EE91671E526>

upvoted 9 times

  **Ekos** 3 years, 1 month ago

i agree

upvoted 1 times

  **nautil2** Most Recent 4 months, 1 week ago

Selected Answer: E

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-TABLE.html#GUID-F9CE0CC3-13AE-4744-A43C-EAC7A71AAAB6>

When you enable table compression by specifying ROW STORE COMPRESS ADVANCED, you enable Advanced Row Compression. Oracle Database compresses data during all DML operations on the table. This form of compression is recommended for OLTP environments.

upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: E

Advanced row compression.

upvoted 1 times

  **ryuah** 2 years ago

E is correct


upvoted 1 times

Which two statements are true about space-saving features in an Oracle Database? (Choose two.)

- A. Private Temporary Tables (PTTS) store metadata in memory only
- B. An index created with the UNUSABLE attribute has no segment
- C. If they exist for a session, Private Temporary Tables (PTTs) are always dropped at the next COMMIT OR ROLLBACK statement
- D. An index that is altered to be UNUSABLE will retain its segment
- E. A table that is truncated will always have its segment removed

Correct Answer: *BC*

Community vote distribution



avanand Highly Voted 3 years, 8 months ago

Correct should be A and B

C is incorrect as on the session level setting, it preserves the definition.

upvoted 13 times

NowOrNever 3 years, 4 months ago

i agree with AB

A - <https://oracle-base.com/articles/18c/private-temporary-tables-18c>

Oracle 18c introduced the concept of a private temporary table, a memory-based temporary table that is dropped at the end of the session or transaction depending on the setup.

B - https://docs.oracle.com/html/E25494_01/indexes003.htm#CIHFIGDG

Beginning with Oracle Database 11g Release 2, the database does not create an index segment when creating an unusable index.

C - not always, depends on the setup of the table creation

upvoted 5 times

nautil2 Most Recent 4 months, 1 week ago

Selected Answer: AB

A - TRUE; PTT stores metadata and data in memory

B - TRUE; Release 19, SQL Reference, CREATE INDEX: When an index, or some partitions or subpartitions of an index, are created UNUSABLE, no segment is allocated for the unusable object. The unusable index or index partition consumes no space in the database.

C - FALSE; Release 19, SQL Reference, CREATE TABLE: The ON COMMIT clause used with keywords PRESERVE DEFINITION creates a session-specific table whose definition is preserved when the transaction commits.

D - FALSE; Release 19, SQL Reference, ALTER INDEX: The space allocated for an index or index partition or subpartition is freed immediately when the object is marked UNUSABLE.

E - FALSE; Release 19, SQL Reference, TRUNCATE TABLE: Deallocates all space used by the removed rows except that specified by the MINEXTENTS storage parameter - space specified by the MINEXTENTS parameter is not removed

[A]: <https://gavinsoorma.com.au/knowledge-base/oracle-18c-new-feature-private-temporary-tables/>

upvoted 1 times

piontk 8 months, 2 weeks ago

Selected Answer: AD

Well, D is correct. You can try:

```
create table teste(
id number,
id2 number);
```

```
create index idx_teste01 on teste(id);
create index idx_teste02 on teste(id2) unusable;
```

```
col segment_name for a15
```

```
select segment_name, segment_type, extents, blocks, bytes from dba_segments where segment_name in ('IDX_TESTE01', 'IDX_TESTE02');
```

You can see that the segments exist in dba_segments for the valid and the unusable index.

Another example is to switch the valid index to unusable and repeat the query. Also the segments still there.

```
alter index IDX_TESTE01 unusable;
```

```
col segment_name for a15
```

```
select segment_name, segment_type, extents, blocks, bytes from dba_segments where segment_name in ('IDX_TESTE01', 'IDX_TESTE02');
```

upvoted 1 times

Franky_T 1 year, 9 months ago

Selected Answer: AB

A is correct. PTT's exist in memory only.


B is correct. Indexes created this way will not have a segment. Query user_segments.

C is wrong. PTT's will exist for the session if they are created with the ON COMMIT PRESERVE DEFINITION clause.

D is wrong. When an existing index is made unusable then the index segment is dropped.

E is wrong. Truncated tables will release the space that was used by the removed rows EXCEPT that which is specified by the MINEXTENTS storage parameter.

upvoted 2 times

  **ryuah** 2 years ago

A,B is correct

upvoted 1 times

  **ioio** 2 years, 9 months ago


A and B are correct.

E is incorrect:

TRUNCATE TABLE statement:

Deallocates all space used by the removed rows except that specified by the MINEXTENTS storage parameter

upvoted 1 times

  **khalilshahin01** 3 years, 5 months ago

my friend passed the exam and he Choose B&C and get wrong answer about space-saving

i will go for

AE if database 18c

or

BE if database 12c

upvoted 1 times

  **adoptc94** 3 years, 4 months ago



I am not sure about E, the wording is important here.

When truncating a table you can use the option REUSE STORAGE, f.e. TRUNCATE TABLE table1 REUSE STORAGE;

REUSE STORAGE specifies that all space currently allocated for the table or cluster remains allocated to it. Which means a table that is being truncated has not always all it's segments removed!

I'd go with answers: A,B

upvoted 2 times

  **Ekos** 3 years, 1 month ago



i agree

upvoted 1 times

  **you1234** 3 years, 6 months ago



ama, any suggestion on correct answer please

upvoted 1 times

  **ama** 3 years, 6 months ago

I pick B,C

upvoted 1 times

  **FayK** 3 years, 6 months ago

Correct B &C.

Oracle 18c introduced the concept of a private temporary table, a memory-based temporary table that is dropped at the end of the session or transaction depending on the setup.

Private temporary tables are dropped at the end of the session, Global temporary tables have their data definition preserved.

upvoted 1 times

  **NiciMilo** 3 years, 6 months ago

so C is incorrect

Private Temporary Tables (PTTs) are ***NOT*** always dropped at the next COMMIT OR ROLLBACK statement because you can create PTTs with ON COMMIT PRESERVE DEFINITION;

upvoted 3 times

  **ama** 3 years, 6 months ago

what?

The ON COMMIT DROP DEFINITION clause, the default, indicates the table should be dropped at the end of the transaction, or the end of the session.

see here <https://oracle-base.com/articles/18c/private-temporary-tables-18c>

upvoted 1 times

  **avanand** 3 years, 5 months ago

Exactly, as per that link also, you may choose "ON COMMIT PRESERVE DEFINITION" also. So You have a choice you may preserve or you may drop.

So "C" is not true, as its not always dropped.
upvoted 4 times

Which two statements are true about the PMON background process? (Choose two.)

- A. It registers database services with all local and remote listeners known to the database instance
- B. It frees resources held by abnormally terminated processes
- C. It records checkpoint information in the control file
- D. It frees unused temporary segments
- E. It kills sessions that exceed idle time

Correct Answer: AB

Reference:



https://docs.oracle.com/cd/B19306_01/server.102/b14220/process.htm

Community vote distribution

BE (100%)

  **Rivaldo11** Highly Voted 3 years, 6 months ago

Should be B and E. A is not correct - "...LREG registers database services with all local and remote listeners...", in 12c and higher versions.
upvoted 13 times

  **Ekos** 3 years, 1 month ago

yes, agree
upvoted 1 times

  **nautil2** Most Recent 4 months ago

Selected Answer: BE

A - FALSE; LREG
B - TRUE; Release 19, Database Concepts, 15 Process Architecture, Process Monitor Process (PMON) Group: The process monitor (PMON) detects the termination of other background processes. If a server or dispatcher process terminates abnormally, then the PMON group is responsible for performing process recovery.
C - FALSE; CKPT
D - FALSE; SMON
E - TRUE; Release 19, Database Concepts, 15 Process Architecture, Process Monitor Process (PMON) Group: CLMN periodically performs cleanup of terminated processes, ... and detached network connections that have exceeded their idle timeout.
upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago

Selected Answer: BE

B and E is correct.
upvoted 1 times

  **Datajimm** 1 year, 10 months ago

Selected Answer: BE

a) is done by LREG
b) is done by PMON process group
c) is done by CKPT
d) is done by SMON
e) is done by PMON process group
upvoted 4 times

  **ryuah** 2 years ago

B,E is correct
upvoted 1 times

  **abebew** 2 years, 7 months ago

AB
https://docs.oracle.com/cd/B19306_01/server.102/b14220/process.htm
upvoted 1 times

  **LrnsTgh** 2 years, 5 months ago

1z0-082 for Oracle Database 12cR2 or higher. agree with answering B & E.
answering A (register database) for Oracle 11g or lower version.
upvoted 1 times

  **leoanon94** 2 years, 5 months ago



We have to consider that the question is on PMON process and not PMON group (with PMON process, CLMN process and CLnn)

<https://docs.oracle.com/en/database/oracle/oracle-database/21/cncpt/process-architecture.html#GUID-B5CA9579-53DB-442C-A85F-F21FD334833A>

So i think E is wrong for this reason.
The answer B is right.
Answer A could be right or wrong depending by Oracle version.

For me correct answers are: A, B and I think that this is an old question referenced to an old Oracle version (<= 12cR1). Idk.
Maybe during the exam the answers will be different.



upvoted 1 times

  **Muhab** 3 years, 5 months ago

AE is Answer



- Performs process recovery when a user process fails
- Cleans up the database buffer cache
- Frees resources that are used by the user process
- Monitors sessions for idle session timeout

upvoted 1 times

  **Muhab** 3 years, 5 months ago

sorry i mean BE

upvoted 4 times

  **avanand** 3 years, 5 months ago

Correct, Answer is BE.

upvoted 3 times

In which three situations does a new transaction always start? (Choose three.)

- A. when issuing a SELECT FOR UPDATE statement after a CREATE TABLE AS SELECT statement was issued in the same session
- B. when issuing a TRUNCATE statement after a SELECT statement was issued in the same session
- C. when issuing a CREATE TABLE statement after a SELECT statement was issued in the same session
- D. when issuing the first Data Manipulation Language (DML) statement after a COMMIT OR ROLLBACK statement was issued in the same session
- E. when issuing a CREATE INDEX statement after a CREATE TABLE statement completed successfully in the same session
- F. when issuing a DML statement after a DML statement failed in the same session

Correct Answer: ADF

Community vote distribution



dotruonghan Highly Voted 3 years, 6 months ago

ADE should be correct

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/transactions.html#GUID-7690181F-8293-47B9-ADC1-905BD9C3AF57>

- A transaction begins when the first executable SQL statement is encountered. An executable SQL statement is a SQL statement that generates calls to a database instance, including DML and DDL statements and the SET TRANSACTION statement.
- A transaction ends when any of the following actions occurs:
 - o A user issues a COMMIT or ROLLBACK statement without a SAVEPOINT clause. □ Answer D
 - o A user runs a DDL command such as CREATE, DROP, RENAME, or ALTER. The database issues an implicit COMMIT statement before and after every DDL statement. If the current transaction contains DML statements, then Oracle Database first commits the transaction and then runs and commits the DDL statement as a new, single-statement transaction.
 - o SELECT do not have a transaction start/end. Check by command: SELECT XID AS "txn id", XIDUSN AS "undo seg", XIDSLOT AS "slot", XIDSQN AS "seq", STATUS AS "txn status" FROM V\$TRANSACTION;

upvoted 16 times

SimoneF 3 years, 1 month ago

Spot on, I agree

upvoted 1 times

you1234 Highly Voted 3 years, 6 months ago

A & D & F is correct answer

upvoted 12 times

ama 3 years, 6 months ago

why not A, D; E?

upvoted 5 times

Marcello86CT 1 year, 7 months ago

DDL should just close the transaction with an implicit COMMIT. In this case the question is asking when a Transaction START , no when is COMPLETED or CLOSED. So my opinion is that:

A Is correct , DML after a DDL

B not correct, TRUNACATE is a DDL and is executed after a select that is not a DML.

C not correct for the same reason of B

D is Correct because we have a DML after a commit so after the end of a transaction.

E is not Correct : DDL after DDL

F is correct : <https://oracle-base.com/articles/10g/dml-error-logging-10gr2>

By default, when a DML statement fails the whole statement is rolled back, regardless of how many rows were processed successfully before the error was detected.

upvoted 2 times

nautil2 Most Recent 4 months ago

Selected Answer: BCE

A - FALSE; SELECT FOR UPDATE locks table rows, but do not start a transaction

B - TRUE; TRUNCATE is a single DDL command, so it starts a transaction

C - TRUE; CREATE TABLE is a single DDL command, so it starts a transaction

D - FALSE; typically, first DML command issued after COMMIT/ROLLBACK starts a transaction but it is not generally true. Only DML command that makes an atomic change in the database starts a transaction.

E - TRUE; CREATE INDEX is a single DDL command, so it starts a transaction

F - FALSE; DML statement does not start a new transaction when not preceded by COMMIT, ROLLBACK, DDL statement or new session

upvoted 2 times

🗄️ 👤 **zouve** 5 months, 1 week ago

Selected Answer: ADE

the CREATE INDEX statement is now listed in the table, which means that it does start a new transaction in Oracle 19c and later
upvoted 1 times

🗄️ 👤 **musafir** 6 months ago

Selected Answer: ADF

A D F

E are two DDL statements after each other, each statement is a completed transaction, there is no active transaction after each statement. Since the question is asking when a new transaction starts E is incorrect since there is no active transaction after the two statements. Tested in Oracle.
upvoted 1 times

🗄️ 👤 **auwia** 6 months, 3 weeks ago

Selected Answer: ADF

Provided answer are corrects!

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/transactions.html#GUID-7690181F-8293-47B9-ADC1-905BD9C3AF57>
upvoted 1 times

🗄️ 👤 **auwia** 6 months, 3 weeks ago

A new transaction is always created after a DML statement is issued! Remember this! ;-)
upvoted 1 times

🗄️ 👤 **RaNik69** 9 months ago

Selected Answer: ACE

A - tested

C - DDL (always start new transaction)

E - DDL (always start new transaction)

upvoted 1 times

🗄️ 👤 **ivanadj** 10 months, 2 weeks ago

Selected Answer: ADE

<https://www.examttopics.com/discussions/oracle/view/10274-exam-1z0-071-topic-1-question-256-discussion/>
upvoted 1 times

🗄️ 👤 **nautil2** 1 year ago

Selected Answer: ADF

Oracle 19c documentation states a DDL command is performed as a single-statement transaction. See paragraph "End of a Transaction" at <https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/transactions.html#GUID-510DF997-BA9C-4431-95EF-A0622881EC9C>
However if I am requested to choose three answers, I choose those options that causes transaction ID is assigned. It can be printed with `exec dbms_output.put_line(dbms_transaction.local_transaction_id);`
I tried all options and only A, D, F returned non empty transaction ID.
upvoted 1 times

🗄️ 👤 **Marcello86CT** 1 year, 7 months ago

Sorry Guys but i think that E is not Correct. Oracle shouldn't support Transactional DDL. DDL should just close the transaction with an implicit COMMIT. In this case the question is asking when a Transaction START , no when is COMPLETED or CLOSED. So my opinion is that:

A Is correct , DML after a DDL

B not correct, TRUNCATE is a DDL and is executed after a select that is not a DML.

C not correct for the same reason of B

D is Correct because we have a DML after a commit so after the end of a transaction.

E is not Correct : DDL after DDL

F is correct : <https://oracle-base.com/articles/10g/dml-error-logging-10gr2>

By default, when a DML statement fails the whole statement is rolled back, regardless of how many rows were processed successfully before the error was detected.

Then i would say : A , D , F

upvoted 2 times

🗄️ 👤 **Franky_T** 1 year, 9 months ago

Selected Answer: AD

Issue again, only two possible answers here.

A transaction implicitly begins with any operation that obtains a TX lock:

- When a statement that modifies data is issued (DML ONLY)

- When a SELECT ... FOR UPDATE statement is issued

- When a transaction is explicitly started with a SET TRANSACTION statement or the DBMS_TRANSACTION package

A is correct. A select for update clause obtains a transaction lock TX, which is one of the three conditions under which transactions will implicitly begin.

B is wrong. A truncate statement is DDL, and DDL statements perform implicit commits after they are done. Commits ends transactions, they don't start them

C is wrong. A create table statement is DDL, same issue as option B.

D is correct. DML statements start transactions.

E is wrong. A create index statement is DDL, same issue as option B and C.



F is wrong. When a DML statement fails it does not end a transaction, only commit or rollback (implicit or explicit) can do this. New DML statements will simply continue in the context of the current transaction. This option is easily tested.

upvoted 3 times

  **linhao** 6 months, 1 week ago

F may be correct. First DML failed and the new transaction not start. The second DML start new transaction.

upvoted 1 times

  **SantiBZ_07032022_1744** 1 year ago



Right with you

upvoted 1 times

  **ryuah** 2 years ago

A,C,E is correct

upvoted 2 times

  **ozan_a** 2 years, 5 months ago

A. when issuing a SELECT FOR UPDATE statement after a CREATE TABLE AS SELECT statement was issued in the same session --> DDL auto commit + new transaction. true.

B. when issuing a TRUNCATE statement after a SELECT statement was issued in the same session --> truncate doesn't start a transaction. wrong.

C. when issuing a CREATE TABLE statement after a SELECT statement was issued in the same session --> whether have a transaction or not, a DDL always create a transaction. true.



D. when issuing the first Data Manipulation Language (DML) statement after a COMMIT OR ROLLBACK statement was issued in the same session --> a select doesn't create a transaction. wrong.

E. when issuing a CREATE INDEX statement after a CREATE TABLE statement completed successfully in the same session --> a DDL always create a transaction. true.

F. when issuing a DML statement after a DML statement failed in the same session --> we need commit or rollback for a new transaction. wrong.

Answer : A,C,E

upvoted 2 times

  **ozan_a** 2 years, 5 months ago

A. when issuing a SELECT FOR UPDATE statement after a CREATE TABLE AS SELECT statement was issued in the same session --> DDL auto commit + new transaction. true.

B. when issuing a TRUNCATE statement after a SELECT statement was issued in the same session --> truncate doesn't start a transaction. wrong.

C. when issuing a CREATE TABLE statement after a SELECT statement was issued in the same session --> whether have a transaction or not, a DDL always create a transaction. true.

D. when issuing the first Data Manipulation Language (DML) statement after a COMMIT OR ROLLBACK statement was issued in the same session --> a select doesn't create a transaction. wrong.

E. when issuing a CREATE INDEX statement after a CREATE TABLE statement completed successfully in the same session --> a DDL always create a transaction. true.

F. when issuing a DML statement after a DML statement failed in the same session --> we need commit or rollback for a new transaction. wrong.

Answer : A,C,E

upvoted 2 times

  **ioio** 2 years, 9 months ago

C D E are correct

C and E: Create Table and Create Index are DDL and a new single transaction is created anyway.



D: when a DML issued when no transaction is still active, a new transaction will started

Select for update is not a DML

Truncate is not starting a transaction

if a DML fails, the transaction still remains active

upvoted 2 times

  **Fan** 2 years, 10 months ago

E should be DDL after DDL , I don't konw what will happened? a new trans or not.

upvoted 1 times

  **saif_alrwiliy** 3 years, 5 months ago

A,D,F is correct answer

upvoted 3 times

Examine the description of the SALES1 table:

Name	Null	Type
-----	-----	-----
SALES_ID	NOT NULL	NUMBER
STORE_ID	NOT NULL	NUMBER
ITEMS_ID		NUMBER
QUANTITY		NUMBER
SALES_DATE		DATE



SALES2 is a table with the same description as SALES1.
Some sales data is duplicated in both tables.
You want to display the rows from the SALES1 table which are not present in the SALES2 table.
Which set operator generates the required output?



- A. INTERSECT
- B. UNION ALL
- C. UNION
- D. SUBTRACT
- E. MINUS



Correct Answer: E



Community vote distribution

E (100%)

  **Ekos** Highly Voted 3 years, 1 month ago
E. MINUS
upvoted 7 times

  **nautil2** Most Recent 4 months ago
Selected Answer: E
E is correct
upvoted 1 times

  **nautil2** 1 year ago
Selected Answer: E
<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/The-UNION-ALL-INTERSECT-MINUS-Operators.html#GUID-B64FE747-586E-4513-945F-80CB197125EE>
upvoted 1 times

  **ryuah** 2 years ago
E is correct
upvoted 2 times

Your database instance is started with a PFILE.

Examine these parameters:

NAME	TYPE	VALUE
-----	-----	-----
memory_max_target	big integer	0
memory_target	big integer	0
sga_max_size	big integer	2G
sga_target	big integer	2G

You want to increase the size of the buffer cache.

Free memory is available to increase the size of the buffer cache.

You execute the command:

SQL> ALTER SYSTEM SET DB_CACHE_SIZE=1024M;

What is the outcome?

- A. The value is changed only in the PFILE and takes effect at the next instance startup
- B. The value is changed for the current instance and in the PFILE
- C. It fails because the SCOPE clause is missing
- D. Change is applied to the current instance, but does not persist after instance restart

Correct Answer: D

Reference:

https://docs.oracle.com/database/121/SQLRF/statements_2017.htm#SQLRF00902

Community vote distribution



adoptc94 Highly Voted 3 years, 4 months ago

Correct Answer: D, the change is applied in memory only, the change does not persist as the database can't write to the pfile

The default SCOPE option, when you start the instance using a PFILE, is MEMORY (as well as the only scope option you can use with a pfile). If scope is omitted the default scope option is used instead. The SCOPE clause is optional and not mandatory when changing the value of a parameter!

A is wrong, as the pfile is a read only file for the database so it can't write to it. Changes to the pfile have to be done manually by changing it directly using f.e. vi on unix systems.

B is wrong, same reason as for why answer A is wrong

C is wrong, as mentioned above, when omitted the scope clause defaults to MEMORY when using a pfile.

Not related to that question but the more you know: When starting the database using a spfile you have three options for the scope (MEMORY, SPFILE and BOTH). Default value is BOTH.

upvoted 18 times

Ekos 3 years, 1 month ago

yes, agree. D is the correct answer

upvoted 1 times

nautil2 Most Recent 4 months ago

Selected Answer: D

- A - FALSE; DB_CACHE_SIZE is a dynamic parameter, its change is applied immediately to the instance
 - B - FALSE; when starting db instance with PFILE, no parameter changes can be written to PFILE
 - C - FALSE; for dynamic parameters and starting db instance with PFILE, SCOPE=MEMORY is a default /implicit option, so it can be omitted. Change is applied to current instance.
 - D - TRUE; see C
- upvoted 1 times

musafir 6 months ago

Selected Answer: D

D is correct since it's PFILE
If a server parameter file (spfile) was used to start up the database, then BOTH is the default. If a parameter file (pfile) was used to start up the database, then MEMORY is the default, as well as the only scope you can specify.
<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/ALTER-SYSTEM.html#GUID-2C638517-D73A-41CA-9D8E-A62D1A0B7ADB>

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: B

When you execute the "ALTER SYSTEM" command to modify a parameter, the default behavior is to change the value for the current instance and also update the value in the server parameter file (in this case: PFILE) to make it persistent across instance restarts.

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

I was wrong, correct answer is D.

upvoted 1 times

  **canijho** 11 months ago

All answer are wrong. Tested:

a) The change is applied immediately. Database cannot write into the PFILE (init.ora).

b) Database cannot write into the PFILE (init.ora).

c) It does not fail:

SQL> ALTER SYSTEM SET DB_CACHE_SIZE=1024M;

System altered.

Mon Feb 27 13:01:09 2023

ALTER SYSTEM SET db_cache_size=1024M SCOPE=BOTH;

SQL> show parameter db_cache_size

NAME TYPE VALUE

db_cache_size big integer 1G

d) Default SCOPE option is BOTH (as shown in the alert.log). The change persist after instance restart (validated).

upvoted 1 times

  **canijho** 11 months ago

Correct answer should be B regarding SPFILE, but not PFILE.

upvoted 1 times

  **raferen10** 1 year, 1 month ago



Selected Answer: D

Spfile (server parameter file)-> default scope both

Pfile(texto parameter file) -> default scope memory (and only available).

Database can't write un text parameter file.



upvoted 1 times

  **Juv07** 1 year, 1 month ago

Selected Answer: D

right answer : D

upvoted 1 times


  **Xhost** 1 year, 7 months ago

Selected Answer: B

The right answer is B!

Default is SCOPE=BOTH

upvoted 2 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: D

A is wrong. PFILES are static and any changes made to the system does not update the text (p)file.

B is wrong. Issue as with point A.

C is wrong. Statement without the scope parameter does not fail. Easy to test.



D is correct. PFILE based instances change parameters by default happen in memory and as with points A and B, will not modify the PFILE.

upvoted 2 times

  **ryuah** 2 years ago


D is correct

upvoted 1 times

  **brolem** 2 years, 5 months ago

Per adoptc94's final notes, in 19c "When starting the database using a spfile you have three options for the scope (MEMORY, SPFILE and BOTH). Default value is BOTH.". As the default is both, when SCOPE is not used, the change will be applied to both MEMORY and the SPFILE.

upvoted 3 times

  **dirk1** 2 years, 2 months ago

Yes, the default is SCOPE=BOTH if a server parameter file was used to start up the instance, and the default is MEMORY if a text initialization parameter file was used to start up the instance.

upvoted 1 times




Which three Oracle database space management features will work with both Dictionary and Locally managed tablespaces? (Choose three.)



- A. Oracle Managed Files (OMF)
- B. Online table segment shrink
- C. Online index segment shrink
- D. Automatic data file extension (AUTOEXTEND)
- E. Capacity planning growth reports based on historical data in the Automatic Workload Repository (AWR)

Correct Answer: ADE

Community vote distribution

ADE (100%)

  **Ekos**  3 years, 1 month ago
correct, A,D, and E
upvoted 8 times

  **baeji** 3 years, 1 month ago
i agree
upvoted 1 times

  **Franky_T**  1 year, 9 months ago

Selected Answer: ADE

A. OMF implements an automatic file name and location option. This is done at file level and has nothing to do with how the tablespace is "managed".
B is wrong. Segment shrink operations are only available with ASSM, which in turn requires LMTBS.
C is wrong. Same feature as option B.
D is correct. AUTOEXTEND is done at file level.
E is correct. Reports are always available through the AWR and has nothing to do with tablespace management type.
upvoted 5 times

  **nautil2**  3 months, 3 weeks ago

Selected Answer: ADE

A - TRUE; Release 19, Database Administrator’s Guide, 17.3.3.6 CREATE DATABASE Statement Using Oracle Managed Files: Examples
B - FALSE; Release 19, Database Administrator’s Guide, 19.3.3 Shrinking Database Segments Online; Shrink operations can be performed only on segments in locally managed tablespaces with automatic segment space management (ASSM). Within an ASSM tablespace, all segment types are eligible for online segment shrink except these: IOT mapping tables, Tables with rowid based materialized views,...
C - FALSE; see B
D - TRUE; no restriction for autoextension mentioned with DMT in Oracle documentation
E - TRUE; see D
upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: ADE

Provided answer are corrects.
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: ADE

Online segment shrink only works whit local managed tablespaces and ASMM is activated.
upvoted 1 times

You execute this command:

```
[oracle@host01 ~]$ expdp system/oracle \
> FULL=Y \
> DUMPFILE=exp_db_full.dmp \
> PARALLEL=4 \
> LOGFILE=exp_db_full.log \
> JOB_NAME=exp_db_full
```

During the export operation, you detach from the job by using CTRL+C and then execute this command:

Export> STOP_JOB=immediate -

Are you sure you wish to stop the job ([yes]/no): yes

Which two statements are true about the job? (Choose two.)


- A. You can no longer monitor it
- B. You can reattach to it and monitor it
- C. It is paused and can be resumed
- D. It continues to run in the background
- E. It terminates

Correct Answer: AE

Reference:

<https://blog.oracle48.nl/killing-and-resuming-datapump-expdp-and-impdp-jobs/>

Community vote distribution



you1234 Highly Voted 3 years, 6 months ago
typo. B & C is correct answer.
upvoted 13 times

you1234 Highly Voted 3 years, 6 months ago
B & C is correct. see the example which is done on 12c version

B. You can reattach to it and monitor it
C. It is paused and can be resumed
upvoted 9 times

auwia 6 months, 3 weeks ago
It's related to a normal stop (without argument), immediate: destroy all :-) probably the log will continue to write for X time, but the process is cancelled you can be sure! :-)
upvoted 1 times

nautil2 Most Recent 3 months, 3 weeks ago
Selected Answer: BC
A - FALSE; the job can be attached and restarted at a later time with the START_JOB command. A KILL_JOB command performs non-reversible termination of the job.
B - TRUE; see A
C - TRUE; see A
D - FALSE; it is completely stopped
E - FALSE; it is stopped, but not terminated. A KILL_JOB command terminates the job.

Release 19, Utilities; Part I Oracle Data Pump; 2.5.11 STOP_JOB; If the Data Pump control job table and dump file set are not disturbed when or after the STOP_JOB command is issued, then the job can be attached to and restarted at a later time with the START_JOB command. To perform an immediate shutdown, specify STOP_JOB=IMMEDIATE. A warning requiring confirmation will be issued. All attached clients, including the one issuing the STOP_JOB command, receive a warning that the job is being stopped by the current user and they will be detached. After all clients are detached, the process structure of the job is immediately run down. That is, the Data Pump control job process will not wait for the child processes to finish their current tasks.
upvoted 1 times

musafir 6 months ago
Selected Answer: BC
A. You can no longer monitor it -- Semi-wrong, can be monitored if reattached.
B. You can reattach to it and monitor it -- Correct, STOP_JOB can be reattached and monitored

- 店长微信: hifeng128
- C. It is paused and can be resumed -- Correct, As per doc it can be resumed but some tasks may have to restarted.
 - D. It continues to run in the background -- Wrong, The job is stopped
 - E. It terminates -- Wrong, Terminates with KILL_JOB not STOP_JOB

B and C seems to be the most correct

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sutil/oracle-data-pump-export-utility.html#GUID-1A2BF459-680D-4D2B-B35F-9A3F4EE72530>

upvoted 2 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: AE

Provided answer are corrects:

<https://docs.oracle.com/database/121/SUTIL/GUID-2BF78573-239C-4D56-8C40-B1F30F3E8609.htm#SUTIL953>

To perform an immediate shutdown, specify STOP_JOB=IMMEDIATE. A warning requiring confirmation will be issued. All attached clients, including the one issuing the STOP_JOB command, receive a warning that the job is being stopped by the current user and they will be detached. After all clients are detached, the process structure of the job is immediately run down. That is, the master process will not wait for the worker processes to finish their current tasks. There is no risk of corruption or data loss when you specify STOP_JOB=IMMEDIATE. However, some tasks that were incomplete at the time of shutdown may have to be redone at restart time.



upvoted 1 times

  **trgbighero** 1 year, 3 months ago

Selected Answer: BC

I choose BC

upvoted 1 times

  **ESK** 1 year, 3 months ago

I think it is AE, the STOP_JOB will effectively stop the export. It is resumable but the question asks what happens when the STOP_JOB=IMMEDIATE option.

upvoted 2 times

  **auwia** 6 months, 3 weeks ago

You have to run back from scratch.

upvoted 1 times

  **ryuah** 2 years ago



B,C is correct

upvoted 1 times

  **van2020** 2 years, 4 months ago

"If the master table and dump file set are not disturbed when or after the STOP_JOB command is issued, then the job can be attached to and restarted at a later time with the START_JOB command." <https://docs.oracle.com/database/121/SUTIL/GUID-2BF78573-239C-4D56-8C40-B1F30F3E8609.htm#SUTIL953>

upvoted 2 times

  **Ekos** 3 years, 1 month ago

i choose B and C for the correct answer

upvoted 2 times

  **dotruonghan** 3 years, 6 months ago

Export> start_job



upvoted 1 times

  **you1234** 3 years, 6 months ago

A is wrong because you can still be monitor & see the status of Job via attach=

correct is B & C.

upvoted 3 times

  **NowOrNever** 3 years, 6 months ago


AC

B - it can be reattached, but it is no longer running, so it can't be monitored in that meaning

D - wrong

E - KILL_JOB does that not STOP_JOB

upvoted 1 times

  **you1234** 3 years, 6 months ago

A & B is correct answer

On 12c

Export>stop_job=immediate;

Export> start_job

店长微信：hjfeng128

[illegible]

 NowOrNever 3 years, 6 months ago

.. ORA-39012: Client detached before the job started. ...
if the export startet and you perform the command "STOP_JOB=immediate" it also exits EXPORT> and you cant execute "start_job"
upvoted 1 times

 NowOrNever 3 years, 6 months ago

i have to correct myself
if you reconnect "expdp system/oracle"
it continues ... thats new to me
upvoted 1 times

In one of your databases, you create a user, HR, and then execute this command:

```
GRANT CREATE SESSION TO hr WITH ADMIN OPTION;
```

Which four actions can HR perform? (Choose four.)

- A. Revoke the CREATE SESSION privilege from other users
- B. Revoke the CREATE SESSION privilege from user HR
- C. Log in to the database instance
- D. Grant the CREATE SESSION privilege with ADMIN OPTION to other users
- E. Execute DDL statements in the HR schema
- F. Execute DML statements in the HR schema

Correct Answer: ABCD

Community vote distribution

ABCD (100%)

  **Guhborges** 10 months ago

- A.TRUE: SQL> revoke create session from g;
Revoke succeeded.
- B. TRUE: SQL> revoke create session from hr;
Revoke succeeded.
- C. TRUE: SQL> conn hr/hr@pdb3;
Connected.
- D. TRUE: SQL> grant create session to hr with admin option;
Grant succeeded.
- E. No make sense, you don't need this option to execute DDL
- F. Same of letter E>

upvoted 1 times

  **M_ORACLE_RU** 11 months, 2 weeks ago

- A. Revoke the CREATE SESSION privilege from other users - ORA-01031: insufficient privileges - false.
 - B. Revoke the CREATE SESSION privilege from user HR – it works, true.
 - C. Log in to the database instance – true.
 - D. Grant the CREATE SESSION privilege with ADMIN OPTION to other users – true.
 - E. Execute DDL statements in the HR schema – false.
 - F. Execute DML statements in the HR schema – it works if somebody with appropriate rights creates some objects in hr's schema. – so, it is also true.
- So, BCDF.

upvoted 3 times

  **raferen10** 1 year, 1 month ago

Selected Answer: ABCD

ABCD is right

upvoted 2 times

  **crucinal1** 1 year, 4 months ago

Selected Answer: ABCD

ABCD right

upvoted 3 times

Which two statements are true about the WHERE and HAVING clauses in a SELECT statement? (Choose two.)



- A. Aggregating functions and columns used in HAVING clauses must be specified in the SELECT list of a query
- B. WHERE and HAVING clauses can be used in the same statement only if applied to different table columns
- C. The HAVING clause can be used with aggregating functions in subqueries
- D. The WHERE clause can be used to exclude rows before dividing them into groups
- E. The WHERE clause can be used to exclude rows after dividing them into groups

Correct Answer: CD

Community vote distribution

CD (100%)

  **NiciMilo** Highly Voted 3 years, 6 months ago
C & D are correct
upvoted 27 times

  **Ekos** 3 years, 1 month ago
i agree, CD is the correct answer
upvoted 2 times

  **nautil2** Most Recent 3 months, 3 weeks ago

Selected Answer: CD

32
A - FALSE; there can be different aggregation function after SELECT and in HAVING clause; e.g. SELECT deptno, MIN(msal) AS "Lowest salary" FROM employees GROUP BY deptno HAVING MAX(msal) > 4000;
B - FALSE; same columns can be used in WHERE and HAVING clause; e.g. SELECT deptno, MIN(msal) AS "Lowest salary" FROM employees WHERE msal<1000 GROUP BY deptno HAVING MIN(msal) < 1400;
C - TRUE; HAVING can be used in subqueries, e.g. SELECT empno, deptno, msal FROM employees WHERE msal < (SELECT MIN(msal) AS "Lowest salary" FROM employees GROUP BY deptno HAVING MIN(msal) > 1200);
D - TRUE; first rows are filtered, then group functions are applied to remaining rows
E - FALSE; see D
upvoted 1 times

  **zouve** 5 months, 3 weeks ago


Selected Answer: CD

CD for me
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: CD

C and D are correct
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: CD

A is wrong. The issue in this answer is the "and columns" portion. Having clauses may only contain a group function that was used in the select list.
B is wrong. This answer is silly. The where clause "looks" at table (column) data and the having clause "looks" at group function results. Remember that the WHERE clause pre-filters table data and the HAVING clause post-filters group function results.
C is correct. Having clauses can be used with aggregate functions in a subquery, as long as the result is properly used in the main query.
D is correct. The WHERE clause pre-filters data from the table before the GROUP BY clause divides them into groups.
E is wrong. The HAVING clause should be used as a post-filter in aggregate queries.
upvoted 1 times

  **Aramazd** 1 year, 9 months ago

HR user can connect to the DB
can also give create session to other users but he cannot do DML

QL> create table students (ID number(10), name varchar(20)) ;
create table students (ID number(10), name varchar(20))
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>


upvoted 1 times

  **ryuah** 2 years ago

B,C,D is correct
upvoted 1 times

  **ryuah** 2 years ago

C,D is correct
upvoted 1 times

  **NowOrNever** 3 years, 4 months ago

B is wrong
<https://www.techonthenet.com/oracle/having.php>
SELECT expression1, expression2, ... expression_n,
aggregate_function (aggregate_expression)
FROM tables
[WHERE conditions]
GROUP BY expression1, expression2, ... expression_n
HAVING having_condition;
upvoted 2 times

  **you1234** 3 years, 6 months ago

B & D is correct one
upvoted 1 times

  **you1234** 3 years, 6 months ago

A & D is correct answer
upvoted 1 times

  **ama** 3 years, 6 months ago

A is wrong ... look at this example :

select file#,avg(rfile#) from v\$datafile Group by file# having Count(TS#) >0

file# avg(rfile#)

1 1
6 6
5 5
8 1024
3 3
7 10
upvoted 6 times

Which two statements are true about UNDO and REDO? (Choose two.)

- A. The generation of UNDO generates REDO
- B. DML modifies Oracle database objects and only generates UNDO
- C. The generation of REDO generates UNDO
- D. DML modifies Oracle database objects and only generates REDO
- E. DML modifies Oracle database objects and generates UNDO and REDO

Correct Answer: AC

Community vote distribution



Rivaldo11 Highly Voted 3 years, 6 months ago

I would say A is correct - "Whenever we do insert in the database, it does following things:
- Insert the row in the buffer
- generate the undo in the buffer
- Write the redo changes i.e insert rows in redo logs
- Write the redo changes for undo segments in the redo logs.
So A and E should be correct answers.
upvoted 21 times

NowOrNever Highly Voted 3 years, 6 months ago

AE
<https://www.experts-exchange.com/articles/13880/UNDO-AND-REDO-IN-ORACLE.html>
... undo is stored in redo logs also. During roll forward phase of instance/media recovery, as changes (redo in redo logs) are applied to datafiles, undo stored in redo logs is used to generate undo segments. These segments are subsequently used to rollback uncommitted changes during rollback phase of instance/media recovery.
upvoted 11 times

nautil2 Most Recent 3 months, 3 weeks ago

Selected Answer: AE
A - TRUE; The redo is the capture of *everything* that ever happens in your database. Writing information to undo is a change to the database, hence it needs to go into redo., https://asktom.oracle.com/pls/apex/f?p=100:11:::P11_QUESTION_ID:9545686800346340839
B - FALSE; every DML modifies/generates both UNDO and REDO
C - FALSE; If the generation of REDO caused the generation of UNDO and at the same time the answer A was valid, the database would loop in the generation of UNDO/REDO changes.
D - FALSE; see B
E - TRUE; see B
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: AC
By exclusion provided answer are corrects:

DDL (Alter, create, drop, rename, truncate, comment) modify DB objects and not DML (select, insert, update, delete, and merge).
upvoted 1 times

nautil2 1 year ago

Selected Answer: AE
A) correct: REDO is generated at every change, even UNDO change. This can be read in Oracle doc "Redo entries record data that you can use to reconstruct all changes made to the database, including the undo segments." <https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-the-redo-log.html#GUID-4625A35C-EF8A-4A9E-8D19-829C1A665A34>
B) wrong: that would cause Oracle does not follow ACID
C) wrong: if A is correct, that would generates neverending generation of both UNDO and REDO
D) wrong: without generating UNDO a ROLLBACK of uncommitted changes cannot be performed
E) correct: with UNDO a ROLLBACK can be performed and with REDO an instance recovery can be performed
upvoted 1 times

trgbighero 1 year, 3 months ago

Selected Answer: AE
I choose AE
upvoted 2 times

Chebas 1 year, 6 months ago

Selected Answer: AC

DML does not modify Oracle database objects hence this exclude B,D and E. Correct answer is AC
upvoted 2 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: CE

A is wrong. Not all UNDO generates REDO. When temp_undo_enabled parameter is set to true then UNDO is stored in the temporary tablespace and does not generate REDO.



B is wrong. DML by default generates both UNDO and REDO.

C is correct. All REDO entries must have associated UNDO segments, which both are needed for consistent instance recovery.

D is wrong. Same as answer B.

E is correct. Same as answer B.

upvoted 4 times

  **josue1** 1 year, 6 months ago

"The database records undo changes in the redo log. But not temp writes.

So if you use temporary tables, setting this to true reduces your redo"

<https://asktom.oracle.com/pls/apex/asktom.search?tag=temp-undo-enabled-parameter>

upvoted 1 times

  **Aramazd** 1 year, 9 months ago

AE sounds correct

upvoted 1 times

  **Datajimm** 1 year, 10 months ago

Selected Answer: AE

a) true: changes in undo has to be stored in the redo, so during instance recovery the undo can be reconstructed in order to rollback transactions

b) false: see above

c) false

d) false: dml has to generate undo, so transaction can be rolled back



e) true: dml generates undo, which in turn generates redo

upvoted 3 times

  **ryuah** 2 years ago

A,E is correct

upvoted 3 times

  **danito** 3 years, 7 months ago

I think C and E

upvoted 1 times

The SCOTT/TIGER user exists in two databases, BOSTON_DB and DALLAS_DB, in two different locations. Each database has a tnsnames.ora file defining DALLAS_DB as a service name. Examine this command:
CREATE DATABASE LINK dblink1 CONNECT TO scott IDENTIFIED BY tiger USING 'dallas_db';
How do you execute the command so that only SCOTT in BOSTON_DB can access the SCOTT schema in DALLAS_DB?

- A. as SCOTT in DALLAS_DB
- B. as SCOTT in BOSTON_DB
- C. as SCOTT in BOSTON_DB and SYS in DALLAS_DB
- D. as SYS in both the databases
- E. as SCOTT in both the databases

Correct Answer: D

Community vote distribution

B (100%)

- NowOrNever**

Highly Voted

3 years, 6 months ago

B

https://docs.oracle.com/cd/B28359_01/server.111/b28310/ds_concepts002.htm#ADMIN12085

Database links are either private or public. If they are private, then only the user who created the link has access; if they are public, then all database users have access.

https://docs.oracle.com/database/121/SQLRF/statements_5006.htm#SQLRF01205

upvoted 12 times
- avanand**

Highly Voted

3 years, 5 months ago

B is correct.

upvoted 6 times
- Phat**

3 years ago

agree. if Scott create the link, only Scott can use it. to access dallas_db from boston_db, Scoot has to create the private link from boston_db.

upvoted 1 times
- nautil2**

Most Recent

3 months, 3 weeks ago

Selected Answer: B

A - FALSE; source database is BOSTON_DB, in that db the database link must be created

B - TRUE; source database is BOSTON_DB and owner is SCOTT => private database link is created

C - FALSE; correct private dblink is created, but also SYS in DALLAS_DB can access to SCOTT schema

D - FALSE; both db links are accessible only to SYS users in both databases

E - FALSE; both SCOTT in BOSTON_DB and SCOTT in DALLAS_DB can access SCOTT schema in DALLAS_DB

upvoted 1 times
- Hell3n**

4 months, 2 weeks ago

Selected Answer: B

https://docs.oracle.com/cd/B28359_01/server.111/b28310/ds_concepts002.htm#ADMIN12085

upvoted 1 times
- alan_resende**

1 year, 1 month ago

I think B is correct.

You must execute the statement connected to scott in boston_db to create a private db link. so, to access scott schema in dallas_db.

upvoted 1 times
- Minoro**

1 year, 1 month ago

Db link was created at SCOTT/BOSTON_DB, So scott/Boston_db can access Scott/Dallas_db, using dblink @db_link ... and scott/Boston_db can access the objects due to it is the owner ...

Letter E ... Scott in both databases

upvoted 1 times
- Minoro**

1 year, 1 month ago

Db link was created at SCOTT/BOSTON_DB, So scott/Boston_db can access Scott/Dallas_db, using dblink @db_link ... and scott/Boston_db can access the objects due to it is the owner ... So for me letter C

can access the SCOTT schema in DALLAS_DB?

upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago

Selected Answer: B

B is correct

upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: B

Firstly, there is no reference to Scott's permissions, so we have to assume that Scott has the ability to create the DB Link.

A is wrong. DALLAS_DB is the target, so the link has to be created from the source which is BOSTON_DB.
B is correct. Source and target is correct and this has to be a private link only accessible to Scott.
C is wrong. This link is one-directional and only needs to be done from BOSTON_DB.
D is wrong. DB Links created by SYS becomes publicly available, we only want Scott to access this link.
E is wrong. Just silly, we don't need a DB Link from a database to itself.

upvoted 2 times

  **ryuah** 2 years ago

B is correct

upvoted 1 times

  **Rivaldo11** 3 years, 6 months ago

How do you know that user SCOTT lacks "create database link" privilege? I would say B.

upvoted 2 times

  **ama** 3 years, 6 months ago

this is the script for scott Schema...

<https://github.com/oracle/dotnet-db-samples/blob/master/schemas/scott.sql>

i didn't find any create database link priveleges!

upvoted 2 times

  **SimoneF** 3 years, 1 month ago

But if you create the DBLINK as SYS schema, SCOTT cannot access it. I think the question is not about the privileges of the schema, which is more in a hypothetical situation where the privilege is there, while rather on the functioning of dblinks.

upvoted 2 times

  **you1234** 3 years, 6 months ago


ama, D is correct?

upvoted 2 times

  **you1234** 3 years, 7 months ago



B is correct? any suggestion

upvoted 4 times

  **ama** 3 years, 7 months ago



easy, SCOT User lacks the "create database link" privilege
so B is wrong!

upvoted 1 times

  **ESK** 1 year, 3 months ago

They don't indicate the rights SCOTT has, i therefore don't think they have any bearing on the question. I think it is about how the the link is created though.

upvoted 1 times

  **Ekos** 3 years, 1 month ago

agree, B is correct

upvoted 2 times

Which three statements are true about the DESCRIBE command? (Choose three.)

- A. It displays the PRIMARY KEY constraint for any column or columns that have that constraint
- B. It can be used from SQL Developer
- C. It displays the NOT NULL constraint for any columns that have that constraint
- D. It can be used to display the structure of an existing view
- E. It displays all constraints that are defined for each column
- F. It can be used only from SQL*Plus



Correct Answer: CDF

Community vote distribution

BCD (100%)

  **husseinyounes** Highly Voted 3 years, 7 months ago

BCD is the correct ans
upvoted 19 times

  **Ekos** 3 years, 1 month ago

i agree
upvoted 1 times

  **SimoneF** 3 years, 1 month ago

But F is also definitely true.
upvoted 1 times

  **SimoneF** 3 years, 1 month ago

Sorry my mistake, I didn't read "only". BCD it is :)
upvoted 1 times

  **NowOrNever** Highly Voted 3 years, 6 months ago

BCD are correct
upvoted 6 times

  **nautil2** Most Recent 3 months, 2 weeks ago

Selected Answer: BCD

A - FALSE; describe displays only NOT NULL constraint
B - TRUE;
C - TRUE; see A
D - TRUE; e.g. DESC user_cons_columns;
E - FALSE; dictionary views USER_CONS_COLUMNS, ALL_CONS_COLUMNS and DBA_CONS_COLUMNS shows constraints, not DESCRIBE
F - FALSE; it can be used from SQL Developer
upvoted 1 times

  **Tsie** 8 months ago

BCD is correct
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: BCD

B C D are correct
upvoted 1 times



  **LeandroHPN** 1 year, 4 months ago

Selected Answer: BCD

BCD are correct
upvoted 1 times

  **ryuah** 2 years ago



B,C,D is correct
upvoted 2 times

  **tko86** 3 years, 6 months ago

It can be used in SQL developer as well as SQL Plus. I tried it as well.

upvoted 4 times

店长微信：hj feng128

  **danito** 3 years, 7 months ago

I use describe comand in sqldeveloper, i guess the answers are not totaly right

upvoted 4 times

Examine the description of the CUSTOMERS table:

Name	Null?	Type
-----	-----	-----
CUST_ID	NOT NULL	NUMBER
CUST_FIRST_NAME	NOT NULL	VARCHAR2 (20)
CUST_LAST_NAME	NOT NULL	VARCHAR2 (30)
CUST_INCOME_LEVEL		VARCHAR2 (30)
CUST_CREDIT_LIMIT		NUMBER


For customers whose income level has a value, you want to display the first name and due amount as 5% of their credit limit. Customers whose due amount is null should not be displayed.

Which query should be used?

- A. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level IS NOT NULL AND due_amount IS NOT NULL;
- B. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level != NULL AND cust_credit_level !=NULL;
- C. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level <> NULL AND due_amount <> NULL;
- D. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level != NULL AND due_amount != NULL;
- E. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level IS NOT NULL AND cust_credit_limit IS NOT NULL;

Correct Answer: E

Community vote distribution



Borislone Highly Voted 3 years, 3 months ago
E is correct
upvoted 13 times

Ekos 3 years, 1 month ago
i agree
upvoted 2 times

Guhborges 2 years, 1 month ago
I agree
upvoted 1 times

nautil2 Most Recent 3 months, 2 weeks ago
Selected Answer: E

A - FALSE; returns error 'ORA-00904: "DUE_AMMOUNT": invalid identifier'. Conditions in WHERE clause are evaluated before column names between SELECT and FROM, so expression alias DUE_AMMOUNT is not known and error is returned.
B - FALSE; != is not a valid operator for handling with NULL values. No rows are returned when this operator is used.
C - FALSE; <> is not a valid operator for handling with NULL values. No rows are returned when this operator is used.
D - FALSE; same asnwer as B
E - TRUE; IS NOT NULL is a valid operator, it is applied on both cust_income_level column and cust_credit_limit column. The other column is a source for computeate DUE_AMMOUNT value. When a value of cust_credit_limit is NULL, then also DUE_AMMOUNT is null and therefore the appropriate row is not displayed.
upvoted 2 times

zouve 5 months, 3 weeks ago
Selected Answer: A

A. SELECT cust_first_name, cust_credit_limit * .05 AS DUE_AMOUNT FROM customers WHERE cust_income_level IS NOT NULL AND due_amount IS NOT NULL;

This query selects the cust_first_name and calculates the DUE_AMOUNT as 5% of the cust_credit_limit from the customers table. The WHERE clause filters the rows to include only those where the cust_income_level is not null and the due_amount is not null. This means that only customers whose income level has a value and whose due amount is not null will be displayed in the result.
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: E

E is correct , IS NULL/IS NOT NULL is used to compare nulls (discard B,C,D),and due_amount is a alias columns cannot be used in where (discard A) , only posible option is E

upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

E, You can only filter by NULL using 'IS NULL' or 'IS NOT NULL'.

upvoted 2 times

  **Aramazd** 1 year, 9 months ago



there is not due_amount in the table. So E is correct

upvoted 1 times

  **ryuah** 2 years ago

E is correct

upvoted 1 times

  **brolem** 2 years, 5 months ago

A is the correct answer. E does not address the fact that customer's with a due_amount of null should not be displayed

upvoted 2 times

  **LrnsTgh** 2 years, 5 months ago

it is impossible WHERE clause used alias column. A is wrong.

upvoted 4 times

The INVOICE table has a QTY_SOLD column of data type NUMBER and an INVOICE_DATE column of data type DATE. NLS_DATE_FORMAT is set to DD-MON-RR.

Which two are true about data type conversions involving these columns in query expressions? (Choose two.)

- A. CONCAT (qty_sold, invoice_date) : requires explicit conversion
- B. invoice_date = '15-march-2019' : uses implicit conversion
- C. invoie_date > '01-02-2019' : uses implicit conversion
- D. qty_sold BETWEEN '101' AND '110' : uses implicit conversion
- E. qty_sold = '0554982' uses implicit conversion

Correct Answer: BD

Community vote distribution



NowOrNever Highly Voted 3 years, 4 months ago

- A. CONCAT (qty_sold, invoice_date) : requires explicit conversion
- B. invoice_date = '15-march-2019' : uses implicit conversion
- C. invoie_date > '01-02-2019' : uses implicit conversion
- D. qty_sold BETWEEN '101' AND '110' : uses implicit conversion
- E. qty_sold = '0554982' uses implicit conversion

ABDE - producing no errors
BDE - are the correct ones
upvoted 8 times

Ekos 3 years, 1 month ago

i agree
upvoted 1 times

Guhborges 2 years, 1 month ago

B is wrong, it is in the wrong format, the NLS_DATE_FORMAT is set to DD-MON-RR, and the answer B the date is set to DD-MONTH-YYYY
upvoted 1 times

ioio Highly Voted 2 years, 9 months ago

sorry but you are wrong.

Oracle always converts string to number and string to date.

VARCHAR → NUMBER
VARCHAR → DATE
only D and E are correct.

as the NLS_DATE_FORMAT is set to DD-MON-RR then B and C can not use implicit conversion because the showed dates are not in defined NLS format.

A is also incorrect
upvoted 8 times

leozanon94 2 years, 5 months ago

Right, i agree.

A is false cause CONCAT doesn't require explicit conversion.
B is false cause it gives to me error ORA-00933 is not the correct format (should be 15-MAR-2019)
C is false cause the month is not valid
D and E are true!
upvoted 1 times

jonsnoow Most Recent 4 months, 2 weeks ago

B,D,E are correct
upvoted 1 times

sheilawu 3 months, 3 weeks ago

HEY it's only 2 answers but you have 3
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: BD

Provided answer are corrects, but in the exam the E answer is formulated like this:

E. qty_sold = "'0554982' : requires explicit conversion

As it is written now, it is also a plausible case. So be careful in the exam which are answer are provided!

upvoted 4 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: DE

with NLS_DATE_FORMAT = DD-MON-RR you must use '15-march-19' not '15-march-2019', so B is incorrect.

upvoted 1 times

  **ivanadj** 11 months ago

Selected Answer: BD

<https://www.examtopycs.com/discussions/oracle/view/21058-exam-1z0-071-topic-2-question-39-discussion/>

upvoted 1 times

  **ivanadj** 11 months ago

Selected Answer: BD

<https://www.examtopycs.com/discussions/oracle/view/21058-exam-1z0-071-topic-2-question-39-discussion/>

upvoted 1 times

  **FelipeC** 11 months ago

Selected Answer: BD

BDE , All tested on 19c

upvoted 1 times

  **Tunglinfo9** 1 year, 1 month ago

Selected Answer: DE

A: select concat(2,'30-NOV-22') from dual; -- 230-NOV-22 implicit conversion


B: select to_char(sysdate,'DD-month-RR') from dual; -- need explicit conversion

C: select to_char(sysdate,'DD-mm-RR') from dual; -- need explicit conversion

D: select * from employees where employee_id between '101' and '110'; -- employee_id is NUMBER datatype, it may use implicit conversion

E: like D option

upvoted 1 times

  **saimmuz** 1 year, 3 months ago

BDE are correct

select salary from employees where salary between '2000' and '3000';

select salary from employees where salary = '012008';

select start_date from job_history where start_date = '24-march-2006';

upvoted 2 times

  **Franky_T** 1 year, 9 months ago

Again a question with 3 correct answers and not two. Assuming that the values supplied in answers D and E should read '101' AND '110', and '0554982' respectively (we have seen "mistakes" like this on the site before), the we can easily test the options against the sample HR schema. I tested all the options against Oracle 12.2.0.1.

A is wrong. Oracle performs implicit conversion here.

B is correct. Implicit conversion does work here. The supplied date format is DD-month-YYYY, it is close enough to the default DD-MON-RR, which allows Oracle to do the conversion.

C is wrong. Oracle cannot perform implicit conversion here. The supplied format is DD-MM-YYYY, the MONTH portion of the format causes the issue here, it is supplied in a number format where the default MON format requires a text based value. Error ORA-01843: not a valid month is returned.

D is correct. Oracle will convert the supplied values of 101 and 110 into number format to match the data type of the column.

E is correct. Oracle will convert the supplied value into a number format to match the data type of the column. The preceding 0 will also be dropped after conversion.

upvoted 3 times

  **ryuah** 2 years ago



D,E is correct

upvoted 1 times

  **escoletsgo1** 3 years, 3 months ago

BD is the answer

upvoted 5 times

  **Chansi** 3 years, 5 months ago

Is that a type in D and E as they are not quoted properly

upvoted 2 times

Which three are types of segments in an Oracle Database? (Choose three.)

- A. undo
- B. index
- C. stored procedures
- D. sequences
- E. tables
- F. clusters

Correct Answer: BEF

Reference:

<http://www.adp-gmbh.ch/ora/concepts/segments.html>

Community vote distribution



- ogdru** Highly Voted 2 years, 10 months ago

1.User Segments
-Table, table partition, or table cluster
-LOB or LOB partition
-Index or index partition
2.Temporary Segments
3.Undo Segments
Ref: <https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html#GUID-7DA83E64-9FF1-45A7-A9AC-D4997DDE0866>
A B E F should be correct??
upvoted 11 times
- Guhborges** 2 years, 3 months ago

I agree with A B E F, so this question is impossible, because they need only three
upvoted 1 times
- NowOrNever** Highly Voted 3 years, 4 months ago

BEF
https://docs.oracle.com/en/database/oracle/oracle-database/18/refrn/DBA_SEGMENTS.html#GUID-7BF4F8C7-1917-4532-B048-C044E17A1754
upvoted 6 times
- Ekos** 3 years, 1 month ago

i agree
upvoted 1 times
- lahiruwic** Most Recent 3 weeks, 6 days ago

Selected Answer: BEF

SQL> select distinct SEGMENT_TYPE from dba_segments;
SEGMENT_TYPE

INDEX
CLUSTER
TABLE PARTITION
LOBINDEX
TABLE SUBPARTITION
SYSTEM STATISTICS
LOBSEGMENT
INDEX PARTITION
ROLLBACK
TABLE
LOB PARTITION
NESTED TABLE
TYPE2 UNDO
upvoted 1 times
- zouve** 5 months, 2 weeks ago

Selected Answer: BEF

Undo are called Rollback segments, so correct is BEF
upvoted 2 times

🗨️ 👤 **auwia** 6 months, 3 weeks ago

Selected Answer: ABE

A, B, and E the correct answers.
upvoted 1 times

🗨️ 👤 **auwia** 6 months, 2 weeks ago

B, E, and F are correct.
upvoted 1 times

🗨️ 👤 **auwia** 6 months, 3 weeks ago

I'm sorry I was wrong, provided answer are corrects.
upvoted 1 times

🗨️ 👤 **FelipeC** 11 months ago

Selected Answer: ABE

ABEF - All tested on 19c
upvoted 1 times

🗨️ 👤 **Franky_T** 1 year, 9 months ago

Again a question with four possible answers and not three.

select the distinct segment_type from dba_segments to get a list of all the segment types.

A is correct.
B is correct.
C is wrong.
D is wrong.
E is correct.
F is correct.

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/DBA_SEGMENTS.html#GUID-7BF4F8C7-1917-4532-B048-C044E17A1754
upvoted 4 times

🗨️ 👤 **Aramazd** 1 year, 9 months ago

Segment types
There are 11 types of segments in Oracle 10g:
table
table partition
index
index partition
cluster
rollback
deferred rollback
temporary
cache
lobsegment
lobindex
These types can be grouped into four segment classes:
data segments: table and cluster
index
rollback
temporary data
upvoted 1 times

🗨️ 👤 **ryuah** 2 years ago

B,E,F is correct
upvoted 1 times

🗨️ 👤 **brolem** 2 years, 5 months ago

From the Oracle 19c documentation:
User Segments

A single data segment in a database stores the data for one user object.

There are different types of segments. Examples of user segments include:

Table, table partition, or table cluster



LOB or LOB partition



Index or index partition

Each nonpartitioned object and object partition is stored in its own segment. For example, if an index has five partitions, then five segments contain the index data.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html#GUID-EFB292CB-87EA-42AA-808C-BD85E540BACC>



upvoted 1 times



  **Borislone** 2 years, 6 months ago
AEF is correct
upvoted 1 times

  **adoptc94** 3 years, 4 months ago
A is also correct

Taken from the Oracle Administrators Workshop Study Guide:

The different types of segments include:
- Table and cluster
- Index
- Undo
- Temporary
upvoted 3 times

  **Phat** 2 years, 10 months ago
From Oracle 9i, undo is managed by undo tablespace, and not with rollback segment anymore. Hence we have rollback segment, but not undo segment.
upvoted 4 times

  **jonsnoow** 1 year ago
Undo segment: Oracle Database 19c maintains information to reverse changes made to the database. This information consists of records of the actions of transactions, collectively known as undo. Undo is stored in undo segments in an undo tablespace.
upvoted 1 times



Which two statements are true about the results of using the INTERSECT operator in compound queries? (Choose two.)



- A. Column names in each SELECT in the compound query can be different
- B. The number of columns in each SELECT in the compound query can be different
- C. Reversing the order of the intersected tables can sometimes affect the output
- D. INTERSECT returns rows common to both sides of the compound query
- E. INTERSECT ignores NULLs

Correct Answer: AE



Community vote distribution



AD (100%)



  **danito** Highly Voted 3 years, 7 months ago
I think A D
upvoted 19 times

  **ama** 3 years, 7 months ago
D yes but A no!

The reason why B is wrong is Illustrated here: in case of two compound Statements : The number and the order of columns must be the same in the two queries.
<https://www.oracletutorial.com/oracle-basics/oracle-intersect/>
upvoted 1 times



  **SimoneF** 3 years, 1 month ago
The number has to be the same but the names can be different. In that page you linked, in fact, I see no reference to the fact that the names must be the same, even though they are in the example.
upvoted 3 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times

  **you1234** Highly Voted 3 years, 6 months ago
A & D is correct one
upvoted 9 times

  **nautil2** Most Recent 3 months, 2 weeks ago

Selected Answer: AD
A - TRUE; Release 19, SQL Language Reference, 9 SQL Queries and Subqueries, The UNION [ALL], INTERSECT, MINUS Operators; The corresponding expressions in the select lists of the component queries of a compound query must match in number and must be in the same data type group (such as numeric or character). => column names can be different, important is a datatype
B - FALSE; see A, there's a mention: The corresponding expressions in the select lists of the component queries of a compound query must match in number ...
C - FALSE; Release 19, SQL Language Reference, 4 Operators, Set Operators; All distinct rows selected by both queries => order is not important
D - TRUE; see C
E - FALSE; <https://blog.jooq.org/how-to-use-sql-intersect-to-work-around-sqls-null-logic/> Like UNION or EXCEPT (MINUS) in Oracle, as well as SELECT DISTINCT, these set operations handle two NULL values as NOT DISTINCT. Yes, they're not equal but also not distinct.
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Selected Answer: AD
If you read fast the A looks like False, but if you get the particular "column name" then it's easy to say True. My final answers: A, D.
upvoted 1 times

  **josue1** 1 year, 7 months ago
Selected Answer: AD
-> example tables t1 and t2 in 19.3.
SQL> select * from t1 /*rows: 1,2,3,null*/;

N

1
2
3

4 rows selected.

SQL> select * from t2 /*rows: 2,null*/;

NN

2

2 rows selected.

--> Test A, D and E.
SQL> select n from t1
2 intersect
3 select nn from t2;

N

2

2 rows selected.

--> Test B
SQL> select n, n from t1
2 intersect
3 select nn from t2;
select n, n from t1
*
ERROR at line 1:
ORA-01789: query block has incorrect number of result columns

--> Test C
SQL> select nn from t2
2 intersect
3 select n from t1;

NN



2



2 rows selected.



Summary:
A. Column names in each SELECT in the compound query can be different --> True
B. The number of columns in each SELECT in the compound query can be different --> False
C. Reversing the order of the intersected tables can sometimes affect the output --> False
D. INTERSECT returns rows common to both sides of the compound query --> True
E. INTERSECT ignores NULLs --> False
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: AD
A is correct. Column names in a compound query is not important, column data types are as well as the number of columns.
B is wrong. Compound queries must have the same amount of columns in each of the query blocks.
C is wrong. Intersect shows rows that are common to both queries so the order of the query blocks do not matter.
D is correct. This is the purpose of the INTERSECT operator.
E is wrong. NULLS are processed in SET queries.
upvoted 1 times

  **ryuah** 2 years ago
A,D is correct
upvoted 1 times

  **Sharif1** 3 years, 6 months ago
A, D are correct.
upvoted 6 times

  **NowOrNever** 3 years, 6 months ago
AD are correct
<https://www.examttopics.com/discussions/oracle/view/8224-exam-1z0-071-topic-1-question-239-discussion/>
upvoted 6 times

Which two statements are true about single row functions? (Choose two.)

- A. MOD : returns the quotient of a division operation
- B. FLOOR : returns the smallest integer greater than or equal to a specified number
- C. TRUNC : can be used with NUMBER and DATE values
- D. CONCAT : can be used to combine any number of values
- E. CEIL : can be used for positive and negative numbers



Correct Answer: AC


Reference:



<https://docs.oracle.com/en/database/oracle/oracle-database/18/sqlrf/Single-Row-Functions.html#GUID-B93F789D-B486-49FF-B0CD-0C6181C5D85C>



Community vote distribution



CE (100%)



  **you1234** Highly Voted 3 years, 6 months ago
C & E is correct one
upvoted 22 times

  **ama** 3 years, 6 months ago
I agree
upvoted 3 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 4 times

  **nautil2** Most Recent 3 months, 2 weeks ago
Selected Answer: CE
A - FALSE; MOD returns the remainder of a division operation
B - FALSE; FLOOR returns the greatest integer smaller or equal to a specified number
C - TRUE; TRUNC(2.152) and TRUNC(sysdate) are valid expressions
D - FALSE; CONCAT joins only two strings (or other datatype expressions that are implicitly converted to a string)
E - TRUE; CEIL as well as FLOOR can be used for both possitive and negative numbers
upvoted 1 times



  **auwia** 6 months, 3 weeks ago
Selected Answer: CE
A is clearly wrong, correct answer C, E.
upvoted 1 times



  **Minoro** 1 year, 1 month ago
Concat may be used with number datatype... for me C,D,E are right
MANAGER_ID NUMBER(6)
DEPARTMENT_ID NUMBER(4)

SQL> select concat(manager_id,department_id) from employees;

CONCAT(MANAGER_ID,DEPARTMENT_ID)

90
10090
10090
upvoted 1 times

  **FelipeC** 11 months, 1 week ago
D is wrong, it says concat can be used with any numbers of ARGUMENTS ,and its false because concat just accept 2 arguments
upvoted 2 times

  **raferen10** 1 year, 1 month ago
Selected Answer: CE
C and E are correct
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

A is wrong. A quotient is the result of a division operation. MOD returns the remainder of n2 divided by n1.
B is wrong. FLOOR returns the largest integer equal to or less than n.
C is correct. The TRUNC (number) function returns n1 truncated to n2 decimal places. The TRUNC (date) function returns date with the time portion of the day truncated to the unit specified by the format model fmt.
D is wrong. CONCAT returns char1 concatenated with char2. Both char1 and char2 can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.
E is correct. Not clearly documented but implied with: "This function takes as an argument any numeric data type". Negative numbers are still numbers.

upvoted 2 times

  **Aldo_7** 1 year, 1 month ago

an example of ceil with negative numbers is correct.

```
SQL> SELECT id, value,CEIL(value) FROM myTable
2 /
```

ID	VALUE	CEIL(VALUE)

1	9	9
2	2.11	3
3	3.44	4
4	-4.21	-4
5	10	10
6	3	3
7	-5.88	-5
8	123.45	124
9	98.23	99

9 rows selected.

upvoted 1 times

  **ryuah** 2 years ago

C,E is correct
upvoted 1 times

  **SimoneF** 3 years, 1 month ago

It's amusing how the link in the alleged solution is correct, but the given solution is wrong. MOD returns the remainder of course.

I agree with C & E
upvoted 2 times

The EMPLOYEES table contains columns EMP_ID of data type NUMBER and HIRE_DATE of data type DATE.
You want to display the date of the first Monday after the completion of six months since hiring.
The NLS_TERRITORY parameter is set to AMERICA in the session and, therefore, Sunday is the first day on the week.
Which query can be used?

- A. SELECT emp_id, ADD_MONTHS(hire_date, 6), NEXT_DAY('MONDAY') FROM employees;
- B. SELECT emp_id, NEXT_DAY(ADD_MONTHS(hire_date, 6), 'MONDAY') FROM employees;
- C. SELECT emp_id, NEXT_DAY(MONTHS_BETWEEN(hire_date, SYSDATE), 6) FROM employees;
- D. SELECT emp_id, NEXT_DAY(ADD_MONTHS(hire_date, 6), 1) FROM employees;

Correct Answer: B



Community vote distribution

(B:110%)

  **ama** Highly Voted 3 years, 6 months ago

I double checked this Question, B is correct !

the Setting of NLS_TERRITORY changed the behavior
upvoted 15 times

  **Ekos** 3 years, 1 month ago

i agree
upvoted 1 times

  **saif_alrwiliy** Highly Voted 3 years, 5 months ago



"The NLS_TERRITORY parameter is set to AMERICA in the session and, therefore, Sunday is the first day of the week."
1 is Sunday in that case.
Correct answer is B.
upvoted 10 times

  **nautil2** Most Recent 3 months, 1 week ago

Selected Answer: B
A - false; query returns emp_id and two dates where any of those dates does not meet required conditions
B - true; 6 months are added to hire_date properly with ADD_MONTHS, then first monday following the date is found with NEXT_DAY function
C - false; function MONTHS_BETWEEN returns the number of months, which is a NUMBER value. This number is then first argument of NEXT_DAY which accepts DATE value as its first parameter. Therefore an error "ORA-00932: inconsistent datatypes: expected DATE got NUMBER" occurs.
D - false; it's very similar to B and almost correct/true. Key information is NLS_TERRITORY=AMERICA : number 1 as a week day number links to Sunday. If NLS_TERRITORY was set to some European country, number 1 would link to Monday and both queries (B, D) would return the same results.
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: B
B is correct
upvoted 1 times

  **Tony1122** 1 year, 6 months ago
what is "'MONDAY' ? but not 'MONDAY'
is it my browser problem?
upvoted 1 times

  **AnnWang** 1 year, 6 months ago

B is correct
upvoted 2 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: B
A is wrong. The NEXT_DAY function needs 2 arguments to work. NEXT_DAY returns the date of the first weekday named by char that is later than the date date.
B is correct. First we add 6 months to the hire date using the ADD_MONTHS function, then passing that date value to the NEXT_DAY function.
C is wrong. Issues with both the MONTHS_BETWEEN function (wrong usage and argument order) and the NEXT_DAY function (second argument should be a day name).
D is wrong. Not because of the usage but because of the Day returned. Using a "1" as the second argument in the NEXT_DAY function it returns a Sunday date and not the first Monday. Looking at the official Oracle documentation the second argument should be character (day name) only.



店长微信: hifeng128

"The argument char must be a day of the week in the date language of your session, either the full name or the abbreviation." Just goes to show that 1, Oracle documentation is never 100% complete and 2, test test test if you are not sure.

upvoted 2 times

  **ryuah** 2 years ago

B is correct
upvoted 1 times

  **mianjee** 2 years, 5 months ago


ALTER SESSION SET NLS_TERRITORY = "AMERICA";
SELECT employee_id, NEXT_DAY(ADD_MONTHS(hire_date, 6), 'MONDAY') FROM employees;
Result = 22-DEC-03 (Monday)

SELECT employee_id, NEXT_DAY(ADD_MONTHS(hire_date, 6), 1) FROM employees;
Result = 21-DEC-03 (Sunday)

Correct answer is B!
upvoted 1 times

  **you1234** 3 years, 6 months ago

D is correct
upvoted 1 times

  **ama** 3 years, 7 months ago

Today is 24.06.2020...

SQL> alter session set nls_date_format='DD.MM.YYYY';

Session wurde geandert.

SQL>
SQL>
SQL> select sysdate, NEXT_DAY(ADD_MONTHS(sysdate, 6), 1) from dual;

SYSDATE NEXT_DAY(A

24.06.2020 28.12.2020

The first Monday after six months is 28.12.2020 > so correct answer must be D
upvoted 2 times

  **SimoneF** 3 years, 1 month ago

But if your NLS_TERRITORY is 'AMERICA' and your week starts with sunday, it will give you the wrong result.

For me, this works until i leave the territory as ITALY but, when I set it to AMERICA like the question states, it's wrong while 'MONDAY' works.
upvoted 1 times

The ORCL database has RESUMABLE__TIMEOUT = 7200 and DEFERRED_SEGMENT_CREATION = FALSE

User U1 has a 1 MB quota in tablespace DATA.

U1 executes this command:

SQL> CREATE TABLE t1 AS -

(SELECT object_name, sharing, created

FROM dba_objects);

U1 complains that the command is taking too long to execute.

In the alert log, the database administrator (DBA) finds this:

2017-03-06T12:15:17.183438+05:30


statement in resumable session 'User U1(136), Session 1, Instance 1' was suspended due to ORA-01536: space quota exceeded for tablespace 'DATA'

Which are three actions any one of which the DBA could take to resume the session? (Choose three.)

- A. Add a data file to DATA
- B. Drop other U1 objects in DATA
- C. Increase U1's quota sufficiently in DATA
- D. Set DEFERRED_SEGMENT_CREATION to TRUE
- E. Grant UNLIMITED TABLESPACE to U1
- F. Set AUTOEXTEND ON for data files in DATA

Correct Answer: BCE

Community vote distribution



elcaso34 Highly Voted 3 years, 7 months ago
I think BCE

Set DEFERRED_SEGMENT_CREATION to TRUE, when creating the table you are already inserting rows as select, therefore the segment has already been created the D is not.

Configure Set AUTOEXTEND ON for data files in DATA, if you have 1M quota it will not work.
upvoted 18 times

ama 3 years, 6 months ago
DEFERRED_SEGMENT_CREATION is Default TRUE.
it doesnt make Change at all. so this Option is wrong and wont solve the problem
upvoted 1 times

SimoneF 3 years, 1 month ago
He didn't said that DEFERRED_SEGMENT_CREATION would help. And it wouldn't but not for the reason you're saying: in the question is written that the parameter is set at FALSE in this case. The reason for which it wouldn't help is that, with a CREATE AS SELECT, the first row is inserted immediately, so the segment is also created. That could have helped in the case of a normal create table.
upvoted 3 times

avanand Highly Voted 3 years, 5 months ago
BCE,

- A. Add a data file to DATA -- Won't help as problem is quota to user not the space
 - B. Drop other U1 objects in DATA -- Will help as it will reduce the used space from quota for that user
 - C. Increase U1's quota sufficiently in DATA -- Will help.
 - D. Set DEFERRED_SEGMENT_CREATION to TRUE -- Won't help, segments has to be created as its CTAS
 - E. Grant UNLIMITED TABLESPACE to U1 -- Will help.
 - F. Set AUTOEXTEND ON for data files in DATA -- Won't help as problem is quota to user not the space
- upvoted 11 times

nautil2 Most Recent 3 months, 1 week ago

Selected Answer: BCE

A - false; the problem (user quota exceeded) persists when any number of datafiles is added because user quota is not related to free space in the tablespace DATA
B - true; dropping user objects decreases occupied space

C - true; increasing a user quota allows session to continue in segment allocation for table t1
D - false; the table t1 is created with CTAS, therefore is filled with data at the time of creation, table segments are allocated at the time of creating the table and segment allocation deferment is not used
E - true; when UNLIMITED TABLESPACE is granted to a user, any quotas are ignored, therefore the restriction "space quota exceeded" ceases to apply
F - false; user quota is not related to free space in tablespace DATA. Even if the datafiles were huge and almost empty, the user quota would still prevent the data of table t1 from being written.

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: BCE

Provided answer are corrects.

upvoted 1 times

  **j_tw** 9 months, 1 week ago

Selected Answer: BCE

BCE

No matter how big the DATA datafile will be, U1 Only have 1MB quota at DATA datafile. So, adding a new datafile or setting autoextend on is not an important thing to solve this problem. We should do something regards to U1's quota.



upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: BCE

B C E correct

upvoted 1 times

  **Xhost** 1 year, 7 months ago

Selected Answer: BCF

I would say B, C, F!

upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: BCE

The DEFERRED_SEGMENT_CREATION = FALSE was added to the question to throw you off. DEFERRED_SEGMENT_CREATION set to TRUE will only delay the allocation of the first extent (thus creating the segment) until the first row is inserted into the table.

The issue here is simply space required for table creation vs quota allocated to the user.

Answers can only be B, C and E.

upvoted 1 times

  **Aramazd** 1 year, 9 months ago

I guess B C E too.

upvoted 1 times

  **ryuah** 2 years ago



B,C,E is correct

upvoted 1 times

  **Borislone** 2 years, 6 months ago



CDE is correct.

upvoted 1 times

  **Ekos** 3 years, 2 months ago

i think the answer is BCE



upvoted 1 times

  **Chansi** 3 years, 5 months ago

B CE . D may help you get the table created but may fail while inserting.

Also the question says " Any of the options should resolve the issue." B is not ideal but given the options that we have it is one of the solutions that work.

upvoted 4 times



  **ama** 3 years, 7 months ago

ORA-01536 space quota exceeded for tablespace 'string'

Cause: The space quota for the segment owner in the tablespace has been exhausted and the operation attempted the creation of a new segment extent in the tablespace.

Action: Either drop unnecessary objects in the tablespace to reclaim space or have a privileged user increase the quota on this tablespace for the segment owner

upvoted 3 times

  **Sharif1** 3 years, 5 months ago

so BCE seems to be the right answers.

upvoted 8 times



  **adoptc94** 3 years, 5 months ago

I'd also go with BCE
as elcaso34 mentioned, your're not only creating the table but you're also inserting rows at the same time, so setting deferred_segment_creation to true won't help as rows are inserted therefore segments will be allocated which use up quota


While B isn't really something I'd recommend it still would resume the session if enough quota is available again.
upvoted 2 times

  **SimoneF** 3 years, 1 month ago



Yes, I agree on BCE
upvoted 1 times

  **ama** 3 years, 7 months ago

i think CDE ... dropping existing objects is never a choice to solve a Problem! so B is false
upvoted 2 times

  **SimoneF** 3 years, 1 month ago

It's probably not something you would do without consulting with the owner of the objects, or at least without knowing that it's disposable data. But it is indeed a course of action that would resume the session. Hence, it's a correct answer.
upvoted 1 times

  **danito** 3 years, 7 months ago

I think BCD
upvoted 1 times

Which three statements are true about the Oracle Data Dictionary? (Choose three.)

- A. Data dictionary views are created by joins of dictionary base tables and DBA-defined tables
- B. The data dictionary is created and maintained by the database administrator
- C. Views with the same name but different prefixes, such as CDB, DBA, ALL and USER, reference the same base tables from the data dictionary
- D. Base tables can be queried directly
- E. It is owned by the SYSTEM user
- F. Usernames of all users including database administrators are stored in the data dictionary



Correct Answer: CDF

Community vote distribution

CDF (100%)

  **NowOrNever** Highly Voted 3 years, 4 months ago

CDF - correct
https://docs.oracle.com/cd/B28359_01/server.111/b28318/datadict.htm#CNCPT1210
SYS, Owner of the Data Dictionary
upvoted 11 times

  **Ekos** 3 years, 1 month ago

i agree
upvoted 3 times

  **nautil2** Most Recent 3 months, 1 week ago

Selected Answer: CDF
A - false; data dictionary views are based only on dictionary base tables, not on DBA-defined tables
B - false; data dictionary is created and maintained by the Oracle database server
C - true; Oracle Database Workshop: "Generally, each view is a subset of the higher-privileged view set, row-wise and column wise."
D - true; SELECT * FROM dual; or SELECT * FROM ts\$;
E - false; owner is SYS, tablespace is SYSTEM
F - true; usernames are stored in USER\$ base table, they are accessible typically with DBA_USERS dictionary view
upvoted 1 times

  **raferen10** 1 year, 1 month ago

Selected Answer: CDF
C D F are correct
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: CDF
A is wrong. "DBA_defined tables" does not fit here. Dictionary views are based on system base tables only.
B is wrong. Internally created and managed by the Oracle server.
C is correct. They all have to reference the same base tables, they only display filtered rowsets from these tables.
D is correct. You can query these tables, however not recommended as data is stored in a cryptic format. Consider logging in as the SYS user, since you "own" the tables you can select from them. And as DBA you can grant select on any table in any schema to any user. This is a permission (ownership) issue and not a data format issue.
E is wrong. Owned by the SYS user, stored in the SYSTEM tablespace.
F is correct. All user's usernames are stored in the dictionary and are queries via DBA_USERS.
upvoted 3 times

  **Guhborges** 2 years, 3 months ago

GUYS, the correctly answer are ACDF

A- User-Accessible Views
The views that summarize and display the information stored in the base tables of the data dictionary. These views decode the base table data into useful information, such as user or table names, using joins and WHERE clauses to simplify the information. Most users are given access to the views rather than the base tables.

https://docs.oracle.com/cd/B19306_01/server.102/b14220/datadict.htm
upvoted 1 times

  **Guhborges** 2 years, 3 months ago

Sorry,

店长微信：hj feng128

A-C-F

We are not able to display all base table, most of the data is stored in a cryptic format

upvoted 1 times

  **Borislone** 2 years, 6 months ago

ACF is correct

upvoted 2 times

  **Datajimm** 1 year, 10 months ago

A is wrong: they are certainly not joins of tables defined by the database administrator (DBA), they are joins of system tables

upvoted 1 times

  **Borislone** 2 years, 6 months ago

D is wrong. Base Tables

The underlying tables that store information about the associated database. Only Oracle should write to and read these tables. Users rarely access them directly because they are normalized, and most of the data is stored in a cryptic format.

upvoted 1 times

  **ioio** 2 years, 9 months ago

CDF - correct

<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/data-dictionary-and-dynamic-performance-views.html#GUID-6F1EA52F-C3AF-407C-B4FA-AE8C8651055E>

upvoted 1 times

  **adoptc94** 3 years, 4 months ago

Correct answers: C,E,F

upvoted 2 times

In one of your databases, the user HR has the password HRMGR.

You want to connect to a database instance whose listener listens on port 1531 by using this statement:

```
CONNECT HR/HRMGR@orcl -
```

No name server is used.

Which statement is true about ORCL?

- A. It must be the value of the SERVICE_NAMES parameter on the client side
- B. It must resolve to a valid connect descriptor in the server's tnsnames.ora file
- C. It must resolve to a valid connect descriptor in the client's tnsnames.ora file
- D. It must be the name of the database to whose instance HR wishes to connect
- E. It must be the name of the server running the database to whose instance HR wishes to connect

Correct Answer: E

Community vote distribution

C (100%)

🗲️ 👤 **Rivaldo11** Highly Voted 3 years, 6 months ago

It should be C.

upvoted 23 times

🗲️ 👤 **adoptc94** 3 years, 4 months ago

I agree correct answer should be C

upvoted 1 times

🗲️ 👤 **Ekos** 3 years, 1 month ago

i agree

upvoted 1 times

🗲️ 👤 **raferen10** Most Recent 1 year, 1 month ago

Selected Answer: C

C is correct

upvoted 2 times

🗲️ 👤 **OCP19c** 1 year, 4 months ago

Selected Answer: C

It's Answer C

upvoted 2 times

🗲️ 👤 **Franky_T** 1 year, 9 months ago

Selected Answer: C

A is wrong. SERVICE_NAMES is an instance parameter on server side.

B is wrong. The connect descriptor must be accessible to the client and therefore cannot be on the server side.

C is correct. Service name resolution must happen at client side so we need a local TNSNAMES.ORA file with the appropriate connect descriptor.

D is wrong. You never connect to a database, but rather to a database instance via a service.

E is wrong. The server name has nothing to do with the service.

upvoted 3 times

🗲️ 👤 **ryuah** 2 years ago

C is correct

upvoted 1 times

🗲️ 👤 **NowOrNever** 3 years, 4 months ago

E - correct

https://docs.oracle.com/cd/E17781_01/server.112/e18804/connecting.htm#ADMQS142

username/password@[//]host[:port][/]service_name]

CONNECT system/mypassword@mydbserver.example.com/XE

upvoted 3 times

🗲️ 👤 **SimoneF** 3 years, 1 month ago

"port (optional) is the TCP port number on which the Oracle Net listener is listening. If not specified, the default port number 1521 is assumed."

So you cannot use that syntax in this case, since it would try to connect on port 1521 while the listener, as per question, is configured to listen

on port 1531.

店长微信：hj feng128

C is correct because "Name server" is referred to the record in a DNS but it doesn't change the fact that the tns must point to the correct server, either using an IP or direct resolution of the address via hosts file.

upvoted 5 times

  **NowOrNever** 3 years, 4 months ago

C is correct

upvoted 1 times

Which two statements are true about views used for viewing tablespace and datafile information? (Choose two.)

A. Tablespace free space can be viewed in V\$TABLESPACE

B. V\$TABLESPACE displays information that is contained in the controlfile about tablespaces

C. V\$TABLESPACE displays information about tablespaces contained in the data dictionary


D. Tablespace free space can be viewed in DBA_TABLESPACES

E. A datafile can be renamed when the database is in MOUNT state and the new file name is displayed when querying DBA_DATA_FILES after the database is opened

Correct Answer: BE

Community vote distribution

BE (100%)

 **nautil2** 3 months, 1 week ago

Selected Answer: BE

A - false; v\$tablespace is a dynamic performance view and doesn't contain any information referring to space usage


B - true; see Oracle documentation, Release 19, Database Reference, Part III Dynamic Performance Views, 9 Dynamic Performance (V\$) Views: V\$SCHEDULER_RUNNING_JOBS to V\$ZONEMAP_USAGE_STATS, 9.105 V\$TABLESPACE, V\$TABLESPACE displays tablespace information from the control file.

C - false; dynamic performance views (like e.g. V\$TABLESPACE) gets its content from instance, not from a data dictionary; these views are independent of data dictionary

D - false; free space in a tablespace cannot be get from the view DBA_TABLESPACES. It can be get using views DBA_SEGMENTS and DBA_DATA_FILES

E - true; see <https://oracle-base.com/articles/misc/renaming-or-moving-oracle-files>

upvoted 2 times

 **auwia** 6 months, 2 weeks ago

Selected Answer: BE

Provided answer are corrects:

https://www.oracle.com/webfolder/technetwork/tutorials/obe/db/9i/r2/paa/obe-dbf/html/Move_DF.htm

E is the old way to rename datafile and it's correct as link above.

upvoted 1 times

 **Franky_T** 1 year, 9 months ago

Selected Answer: BE

A is wrong. This view displays basic TS information with no reference to block consumption.

B is correct.

C is wrong. V\$-views query the instance and not the dictionary.

D is wrong. You have to query dba_tablespaces, dba_data_files and dba_free_space to obtain this information.


E is correct, but could be wrong also. If we rename the datafile with the SQL command "alter database move datafile" then yes. If we assume that the renaming is done on the OS then the "ALTER DATABASE RENAME FILE" command should be issued when the database is in mount state so that the dictionary can be updated.

upvoted 1 times

 **maiccrosas** 4 months, 4 weeks ago

you mean with the database in open mode to use the rename clause, right?


upvoted 1 times

 **casfdsaf** 1 year, 10 months ago

Selected Answer: BE

b e is right

upvoted 1 times

 **zorel** 2 years, 5 months ago

SQL> shutdown immediate;

Database closed.

E correct.

Database dismounted.


ORACLE instance shut down.

SQL> startup mount;

ORACLE instance started.


Total System Global Area 1375731600 bytes
Fixed Size 8896400 bytes
Variable Size 838860800 bytes
Database Buffers 520093696 bytes
Redo Buffers 7880704 bytes
Database mounted.
SQL> alter database move datafile '/oradata/ORA19/testing1.dbf' to '/oradata/ORA19/testing01.dbf';

Database altered.
upvoted 1 times

 **saidov1981** 2 years, 9 months ago

B & E are the right answers.
B: v\$tablespace gets information from control file as per article shared by SimoneF
https://docs.oracle.com/cd/B19306_01/server.102/b14237/dynviews_2156.htm#REFRN30277
E: since all DBA_* dictionaries are not accessible until DB is opened


upvoted 4 times

 **batavo27** 3 years, 2 months ago

Answer C.
SQL> select * from v\$tablespace;

TS# NAME INC BIG FLA ENC CON_ID


1 SYSAUX YES NO YES 1
0 SYSTEM YES NO YES 1
2 UNDOTBS1 YES NO YES 1
4 USERS YES NO YES 1
3 TEMP NO NO YES 1
0 SYSTEM YES NO YES 2
1 SYSAUX YES NO YES 2
2 UNDOTBS1 YES NO YES 2
3 TEMP NO NO YES 2
0 SYSTEM YES NO YES 3
1 SYSAUX YES NO YES 3
Have some informations about the data files
upvoted 2 times

 **SimoneF** 3 years, 1 month ago

Answer C doesn't speak about data files thought, but states that v\$tablespace view would get data from the data dictionary. That's not true and the view actually gets data from the controlfile, as stated in answer B which is correct:

https://docs.oracle.com/cd/B19306_01/server.102/b14237/dynviews_2156.htm#REFRN30277


upvoted 1 times

 **batavo27** 3 years, 2 months ago

B and C is correct. TO rename the datafile, database needs is open state.
SQL> alter tablespace teste10
2 rename datafile '/u02/oradata/ORCL/teste10.dbf'
3 to '/u02/oradata/ORCL/teste11.dbf';
alter tablespace teste10
*


ERROR at line 1:
ORA-01109: database not open

upvoted 3 times

 **Ekos** 3 years, 1 month ago


i agree

upvoted 2 times

 **Ekos** 3 years, 1 month ago


sorry, it should be B E the correct answer

upvoted 1 times

 **SimoneF** 3 years, 1 month ago

Was the database in mount? Because E is true and C is definitely wrong:
https://docs.oracle.com/cd/B19306_01/server.102/b14237/dynviews_2156.htm#REFRN30277

upvoted 2 times

 **avanand** 3 years, 5 months ago

B, E should be correct

upvoted 2 times

Examine the description of the PROMOTIONS table:

Name	Null?	Type

PROMO_ID	NOT NULL	NUMBER (6)
PROMO_NAME	NOT NULL	VARCHAR2 (30)
PROMO_CATEGORY	NOT NULL	VARCHAR2 (30)
PROMO_COST	NOT NULL	NUMBER (10, 2)

You want to display the unique promotion costs in each promotion category.



Which two queries can be used? (Choose two.)



- A. SELECT promo_cost, promo_category FROM promotions ORDER BY by 1;
- B. SELECT DISTINCT promo_cost || ' in ' || DISTINCT promo_category FROM promotions ORDER BY 1;
- C. SELECT DISTINCT promo_category || ' has ' || promo_cost AS COSTS FROM promotions ORDER BY 1;
- D. SELECT promo_category, DISTINCT promo_cost FROM promotions ORDER BY 2;
- E. SELECT DISTINCT promo_category, promo_cost FROM promotions ORDER BY 1;



Correct Answer: DE



Community vote distribution



CE (100%)



  **you1234** Highly Voted 3 years, 6 months ago
C & E is correct answer
upvoted 17 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 2 times

  **Sharif1** Highly Voted 3 years, 6 months ago
yes, C and E are correct.
upvoted 8 times

  **nautil2** Most Recent 3 months, 1 week ago
One could understand the task specification as asking to list only those categories for which the value of promo_cost was unique. That is, if several rows of promo_id, promo_category had the same promo_cost, this row was not listed. Such a task specification is noticeably more complex, it is not in the offered answers, and the corresponding SELECT would look like this:
SELECT promo_category, promo_cost
FROM (
SELECT promo_id, promo_name, promo_category, promo_cost,
COUNT(*) OVER (PARTITION BY promo_category, promo_cost) as cost_count
FROM promotions
)
WHERE cost_count = 1;
upvoted 1 times

  **nautil2** 3 months, 1 week ago
Or simplier:
SELECT promo_category, promo_cost
FROM promotions
GROUP BY promo_category, promo_cost
HAVING COUNT(*) = 1;
upvoted 2 times

  **nautil2** 3 months, 1 week ago
Selected Answer: CE
A - false; the query returns all rows
B - false; using of second DISTINCT returns an error ORA-00936 missing expression
C - true; returns category name and promo_cost, where each promo_cost for given cathegory is returned only once
D - false; using of DISTINCT returns an error ORA-00936 missing expression
E - true; same answer as C
upvoted 1 times

  **Oracle2020** 5 months ago

The correct answers are D and E
A is incorrect because it does not show the unique promotion costs for each promotion category
B is incorrect because it is concatenating varchar and numeric data types without doing data conversion
C is incorrect because it is concatenating varchar and numeric data types without doing data conversion
D provides the correct data
And by discarding the previous ones that are wrong, you can give the solution
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: CE

The base! :-)
upvoted 1 times

  **acr23cd** 10 months, 3 weeks ago

Selected Answer: CE

A. SELECT promo_cost, promo_category FROM promotions ORDER BY by 1; -- FALSE: No distinct
B. SELECT DISTINCT promo_cost || ' in ' || DISTINCT promo_category FROM promotions ORDER BY 1; -- FALSE: Distinct just once, next to SELECT
C. SELECT DISTINCT promo_category || ' has ' || promo_cost AS COSTS FROM promotions ORDER BY 1; -- TRUE
D. SELECT promo_category, DISTINCT promo_cost FROM promotions ORDER BY 2; -- FALSE: Distinct just once, next to SELECT
E. SELECT DISTINCT promo_category, promo_cost FROM promotions ORDER BY 1; -- TRUE
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: CE

A is wrong. No DISTINCT clause used.
B is wrong. DISTINCT can be used only once in a select clause at the beginning.
C is correct.
D is wrong. DISTINCT can be used only once in a select clause at the beginning.
E is correct.
upvoted 1 times

  **casfdsaf** 1 year, 10 months ago

Selected Answer: CE

ce right
upvoted 1 times

  **Guhborges** 2 years ago

Selected Answer: CE

c & e for sure
upvoted 1 times

  **ryuah** 2 years ago

C,E is correct
upvoted 1 times

In the ORCL database, UNDOTBS1 is the active undo tablespace with these properties:

- 1. A size of 100 MB
- 2. AUTOEXTEND is off
- 3. UNDO_RETENTION is set to 15 minutes
- 4. It has RETENTION GUARANTEE

UNDOTBS1 fills with uncommitted undo 10 minutes after the database opens.


What will happen when the next update is attempted by any transaction?



- A. It succeeds and the generated undo is stored in SYSTEM.
- B. It fails and returns the error message `ORA-30036: unable to extend segment by 8 in undo tablespace 'UNDOTBS1'`.
- C. It succeeds and the least recently written undo block of UNDOTBS1 is overwritten by the generated undo.
- D. It succeeds and the generated undo is stored in SYSAUX.
- E. It succeeds and the least recently read undo block of UNDOTBS1 is overwritten by the generated undo.



Correct Answer: C



Community vote distribution



B (100%)

  **Rivaldo11** Highly Voted 3 years, 6 months ago
It should be B.
upvoted 26 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 2 times



  **fgdgdg** Highly Voted 3 years, 5 months ago
B. To guarantee the success of long-running queries or Oracle Flashback operations, you can enable retention guarantee. If retention guarantee is enabled, the specified minimum undo retention is guaranteed; the database never overwrites unexpired undo data even if it means that transactions fail due to lack of space in the undo tablespace. If retention guarantee is not enabled, the database can overwrite unexpired undo when space is low, thus lowering the undo retention for the system. This option is disabled by default.
upvoted 7 times

  **nautil2** Most Recent 3 months, 1 week ago
Selected Answer: B
A - false; if UNDO tablespace UNDOTBS1 is active, undo records are stored in it, not in SYSTEM
B - true; database cannot extend UNDO tablespace due to AUTOEXTEND OFF and cannot rewrite old UNDO records due to RETENTION GUARANTEE
C - false; undo records cannot be overwritten due to UNDO_RETENTION and RETENTION GUARANTEE settings
D - false; see A
E - false; see C
upvoted 1 times

  **Franky_T** 1 year, 9 months ago
Selected Answer: B
Tricky question indeed. The documentation clearly shows that: "The UNDO_RETENTION parameter is ignored for a fixed size undo tablespace. The database always tunes the undo retention period for the best possible retention, based on system activity and undo tablespace size." This shows that Unexpired segments will be overwritten by new transactional data. However, the issue here is with the part of the question that states "UNDOTBS1 fills with uncommitted undo 10 minutes after the database opens". This implies that all segments have an ACTIVE state. ASKTOM states: "We will not reuse space (wrap the rollback segment back around over itself) while there is an active transaction there."

With this in mind the answer has to be B.
upvoted 2 times

  **casfdsaf** 1 year, 10 months ago
Selected Answer: B
b right
upvoted 1 times

  **Datajimm** 1 year, 10 months ago
Selected Answer: B
correct answer is B: The transaction is not committed, therefore the undo space is active. It cannot be overwritten by the next update, so it throws an error.

upvoted 1 times

🗨️ **emburria** 2 years ago

B. "UNDOTBS1 fills with uncommitted undo 10 minutes after the database opens."...data is uncommitted, so oracle is not going to discard uncommitted data.....

upvoted 1 times

🗨️ **ryuah** 2 years ago

B is correct

upvoted 1 times

🗨️ **iamjavacompiler** 2 years, 8 months ago

I did not understand why the answer is B?

upvoted 2 times

🗨️ **ioio** 2 years, 9 months ago

in fact it is an interesting question!
it is C:

The UNDO_RETENTION parameter can only be honored if the current undo tablespace has enough space. If an active transaction requires undo space and the undo tablespace does not have available space, then the system starts reusing unexpired undo space. This action can potentially cause some queries to fail with a "snapshot too old" message.

For AUTOEXTEND undo tablespaces, the system retains undo for at least the time specified in this parameter, and automatically tunes the undo retention period to satisfy the undo requirements of the queries. For fixed- size undo tablespaces, the system automatically tunes for the maximum possible undo retention period, based on undo tablespace size and usage history, and ignores UNDO_RETENTION unless retention guarantee is enabled.

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/UNDO_RETENTION.html#GUID-D238E854-A2C5-463E-A146-31F61300124E

upvoted 1 times

🗨️ **Datajimm** 1 year, 10 months ago

Your quote only applies to unexpired undo space. In our case the undo is active and that cannot be reused by another transaction. C is the wrong answer, B is the correct.

upvoted 1 times

🗨️ **avanand** 3 years, 5 months ago

Its a tricky question. To me answer should be C.

The undo_retention parameter is ignored for a fixed size undo tablespace.

The database may overwrite unexpired undo information when tablespace space becomes low.

Go to this link and search this: "The following points explain the exact impact of the UNDO_RETENTION parameter on undo retention"

https://docs.oracle.com/cd/B19306_01/server.102/b14231/undo.htm

upvoted 6 times

🗨️ **adoptc94** 3 years, 5 months ago

I think you might be right!

Retention Guarantee enforces unexpired undo to be retained for the specified undo retention, but as undo retention is ignored for fixed size undo tablespaces, unexpired undo can be overwritten to support new transactions

But my issue with this question is, that it states that the undo tablespace is filled with uncommitted undo, which means that it is active undo which will never be overwritten right? Only unexpired & expired undo may be overwritten to support new transations

upvoted 1 times

🗨️ **adoptc94** 3 years, 4 months ago

Please ignore my previous comment, I was just confused at that time

Answer C is 100% wrong!

While it is true, that the undo retention will be ignored for fixed size undo tablespaces, it won't be ignored if retention guarantee is enabled!

Source: https://docs.oracle.com/cd/B19306_01/server.102/b14237/initparams222.htm#REFRN10225

For fixed- size undo tablespaces, the system automatically tunes for the maximum possible undo retention period, based on undo tablespace size and usage history, and ignores UNDO_RETENTION !!!unless retention guarantee is enabled.!!!

So the correct answer for this question is B!

upvoted 4 times

Which two statements are true about the rules of precedence for operators? (Choose two.)

- A. The concatenation operator || is always evaluated before addition and subtraction in an expression
- B. NULLS influence the precedence of operators in an expression
- C. The + binary operator has the highest precedence in an expression in a SQL statement
- D. Arithmetic operators with equal precedence are evaluated from left to right within an expression
- E. Multiple parentheses can be used to override the default precedence of operators in an expression


Correct Answer: CE



Reference:



https://docs.oracle.com/cd/B19306_01/server.102/b14200/operators001.htm



Community vote distribution

DE (100%)



  **husseinyounes** Highly Voted 3 years, 7 months ago
DE is correct
upvoted 20 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times



  **you1234** Highly Voted 3 years, 6 months ago
D & E is correct answer
upvoted 6 times



  **nautil2** Most Recent 3 months, 1 week ago
Selected Answer: DE
A - false; concatenation, addition and subtraction has the same precedence, in that case it is evaluated from left to right within an expression.
B - false; precedence does not change, result is NULL except of concatenation
C - false; only unary +, unary -, PRIOR, CONNECT_BY_ROOT and COLLATE have the highest precedence
D - true; see A
E - true; You can use parentheses in an expression to override operator precedence. Oracle evaluates expressions inside parentheses before evaluating those outside.

Source: Oracle documentation: Release 19, SQL Language Reference, 4 Operators, About SQL Operators
upvoted 1 times

  **Oracle2020** 5 months ago
DE is correct
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Selected Answer: DE
The base! :-) part II
upvoted 1 times

  **auwia** 6 months, 2 weeks ago
Rules of precedence:
1. arithmetic operators
2. concatenation operator
3. comparison conditions
4. is (NOT) NULL, LIKE (NOT) IN
5. (NOT) Between
6. NOT equal to
7. NOT logical operator
8. AND logical operator
9. OR logical operator
upvoted 1 times

  **Kuraudio** 1 year, 1 month ago
Selected Answer: DE
DE is correct
upvoted 1 times

  **casfdsaf** 1 year, 10 months ago

Selected Answer: DE

de is right
upvoted 1 times

  **ryuah** 2 years ago

D,E is correct
upvoted 1 times

  **Guhborges** 2 years, 4 months ago

C & E is correct.

+ has the higher precedence.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/About-SQL-Operators.html#GUID-FEF44762-F45C-41D9-B380-F6A61AD25338>

upvoted 1 times

  **Guhborges** 2 years ago

D & E :)
upvoted 1 times

  **emburria** 2 years ago

No, in this case is + binary (not unary operator)
upvoted 1 times

  **SimoneF** 3 years, 1 month ago

D and E are obviously correct BUT, since this was too strange of a mistake to be true, I went and review the precedence rules, finding out something I honestly forgot or ignored altogether:



"+" IS, actually, the operator with the highest precedence, together with "-", yet only when they're used a unary operators "identity" and "negation".

For example, in a simple operation +X*Y, the "+" is the first operator evaluated.

As a binary operator it still comes, obviously, after multiplication and division, I've just been thinking that in a similar question where they could state "unary", I could have easily been tricked overlooking that simple detail.

https://docs.oracle.com/cd/A87860_01/doc/server.817/a85397/operator.htm#997667

upvoted 4 times

  **Joejoel** 2 years, 2 months ago


please check on the same link.

Precedence Example

In the following expression, multiplication has a higher precedence than addition, so Oracle first multiplies 2 by 3 and then adds the result to 1.

1+2*3



upvoted 1 times

  **Datajimm** 1 year, 10 months ago

SimoneF is talking about the + unary operator, and of course he is right about it. You on the other hand talking about the + binary operator.
upvoted 1 times

  **avanand** 3 years, 5 months ago

DE is obviously right answers
upvoted 3 times

  **Trowa** 3 years, 5 months ago



DE are correct.

Precedence

Precedence is the order in which Oracle evaluates different operators in the same expression. When evaluating an expression containing multiple operators, Oracle evaluates operators with higher precedence before evaluating those with lower precedence. Oracle evaluates operators with equal precedence from left to right within an expression.

https://docs.oracle.com/cd/A87860_01/doc/server.817/a85397/operator.htm

upvoted 6 times

  **danito** 3 years, 7 months ago

D E i think
upvoted 5 times

In one of your databases, user KING is:

- 1. Not a DBA user
- 2. An operating system (OS) user

Examine this command and its output:

SHOW PARAMETER OS_AUTHENT_PREFIX		
NAME	TYPE	VALUE
-----	-----	-----
os_authen_prefix	string	



What must you do so that KING is authenticated by the OS when connecting to the database instance?



- A. Set OS_AUTHENT_PREFIX to OPS\$
- B. Have the OS administrator add KING to the OSDBA group
- C. Grant DBA to KING
- D. Unset REMOTE_LOGIN_PASSWORDFILE
- E. Alter user KING to be IDENTIFIED EXTERNALLY



Correct Answer: D

Community vote distribution

E (100%)



  **Sharif1** Highly Voted 3 years, 5 months ago
Correct answer is E, os user must be identified externally.
upvoted 17 times



  **ama** 3 years, 5 months ago
oh yea E is correct!
upvoted 5 times



  **Muhab** Highly Voted 3 years, 5 months ago
The Answer should be A .



Ref : <https://oracle-base.com/articles/misc/os-authentication>
upvoted 5 times

  **Schmeili** 3 years, 5 months ago
when you read <https://oracle-base.com/articles/misc/os-authentication> E is correct...
upvoted 4 times

  **SimoneF** 3 years, 1 month ago
Agreed
upvoted 1 times

  **nautil2** Most Recent 3 months, 1 week ago
Selected Answer: E
A - false; in that case both db user and OS user must have name OPS\$KING
B - false; it would grant DB user DBA role, but login would not succeeded because the user is not identified externally
C - false; user KING is not in DBA group
D - false; REMOTE_LOGIN_PASSWORDFILE must be used to enable user login
E - true; when user is IDENTIFIED EXTERNALLY, it can login with OS authentication
upvoted 1 times



  **Oracle2020** 5 months, 1 week ago
Correct Answer is D, REMOTE_LOGIN_PASSWORDFILE=NONE Oracle ignores any password file. Therefore, privileged users must be authenticated by the operating system.
upvoted 1 times



  **auwia** 6 months, 3 weeks ago
Selected Answer: E
The base! :-) part III
upvoted 1 times

  **Guhborges** 2 years ago

Selected Answer: E

E is correct.
upvoted 2 times

  **ryuah** 2 years ago
E is correct
upvoted 1 times

  **saadkhan45** 2 years, 2 months ago
REMOTE_LOGIN_PASSWORDFILE: NONE



NONE: Instructs the Oracle system not to use password files, and the login of privileged users is authenticated by the operating system.



NONE Oracle ignores any password file. Therefore, privileged users must be authenticated by the operating system.



NONE

Oracle ignores any password file. Therefore, privileged users must be authenticated by the operating system.

https://docs.oracle.com/cd/B14117_01/server.101/b10755/initparams176.htm
<https://titanwolf.org/Network/Articles/Article?AID=cf58b9e3-cc76-4fe0-b7c2-3073595bc869>
upvoted 1 times



  **saadkhan45** 2 years, 2 months ago
So, Correct Answer: D
upvoted 1 times



  **ofranco** 2 years, 4 months ago
Agree with A
upvoted 1 times

  **eddkzew** 3 years ago
E is the correct one:
https://docs.oracle.com/cd/B19306_01/network.102/b14266/admnauth.htm#i1006459
Creating a User Who Is Authenticated Externally
The following statement creates a user who is identified by Oracle and authenticated by the operating system or a network service. This example assumes that OS_AUTHENT_PREFIX = "".



CREATE USER scott IDENTIFIED EXTERNALLY;
Using CREATE USER ... IDENTIFIED EXTERNALLY, you create database accounts that must be authenticated by the operating system or network service. Oracle will then rely on this external login authentication when it provides that specific operating system user with access to the database resources of a specific user.



upvoted 2 times

  **Ekos** 3 years, 2 months ago
B is the correct one.
Members of OSDBA group are granted the SYSDBA system privileges to administer the database.
upvoted 1 times



  **SimoneF** 3 years, 1 month ago
But King is NOT a DBA user, as per stated in the question, and I assume is a requirement to leave that unchanged. So you don't allow a user the DBA role (even worse actually, it would SYSDBA) just for it to be able to connect to the database.

E is the most correct if you want to allow the user to be authenticated by OS without changing its role.
upvoted 3 times

  **mberrios** 3 years, 3 months ago
Correct is E
upvoted 5 times

  **avanand** 3 years, 5 months ago
I think answer should be A.
Because unless you set OS_AUTHENT_PREFIX you can't have this user created on database which can identified externally.

And when you have OS_AUTHENT_PREFIX set to OPS\$ (eg) , you need to alter OPS\$KING to identified externally not to KING.
upvoted 4 times

  **NowOrNever** 3 years, 4 months ago
A - ist wrong
a prefix can be empty, it is not necessary to set it back to default OPS\$
upvoted 5 times

View the Exhibits and examine the structure of the COSTS and PROMOTIONS tables.

You want to display PROD_IDS whose promotion cost is less than the highest cost PROD_ID in a promotion time interval.

Examine this SQL statement:

```
SELECT prod_id
FROM costs
WHERE promo_id IN
  (SELECT promo_id
   FROM promotions
   WHERE promo_cost < ALL
     (SELECT MAX(promo_cost)
      FROM promotions
      GROUP BY (promo_end_date - promo_begin_date)));
```

Exhibit 1.

Table COSTS		
Name	Null?	Type
PROD_ID	NOT NULL	NUMBER
TIME_ID	NOT NULL	DATE
PROMO_ID	NOT NULL	NUMBER
CHANNEL_ID	NOT NULL	NUMBER
UNIT_COST	NOT NULL	NUMBER(10,2)
UNIT_PRICE	NOT NULL	NUMBER(10,2)

Exhibit 2.

Table PROMOTIONS		
Name	Null?	Type
PROMO_ID	NOT NULL	NUMBER(6)
PROMO_NAME	NOT NULL	VARCHAR2(30)
PROMO_SUBCATEGORY	NOT NULL	VARCHAR2(30)
PROMO_SUBCATEGORY_ID	NOT NULL	NUMBER
PROMO_CATEGORY	NOT NULL	VARCHAR2(30)
PROMO_CATEGORY_ID	NOT NULL	NUMBER
PROMO_COST	NOT NULL	NUMBER(10,2)
PROMO_BEGIN_DATE	NOT NULL	DATE
PROMO_END_DATE	NOT NULL	DATE

What will be the result?

- A. It gives an error because the ALL keyword is not valid
- B. It gives an error because the GROUP BY clause is not valid
- C. It executes successfully but does not give the required result
- D. It executes successfully and gives the required result

Correct Answer: D

Community vote distribution

0% 100%

  **ioio** Highly Voted 2 years, 9 months ago

the correct answer is C but so far nobody explained the correct reason!

the issue here is not the ALL operator but the last line of the query:

(promo_end_date - promo_begin_date) is returning the number of days which the promotion was running, but no interval, which is requested by the question. so logically the query is not generating desired output.

upvoted 7 times

  **yukclam9** 2 years, 9 months ago

dear ioio, you are right with the returned data type. however I see the promotion time interval as time difference between promotion, so it is more of a logical concept rather than requirement on data type.

upvoted 2 times


  **NowOrNever**  3 years, 4 months ago

C- correct answer
display PROD_IDS whose promotion cost is less than the highest cost PROD_ID

```
WHERE promo_cost < ALL
(SELECT MAX(promo_cost)
FROM promotions
GROUP BY (promo_end_date - promo_begin_date));
```



"< ALL ..." here is the mistake, if the next selection returns not only one max value, it returns the wrong, maybe even an empty result

upvoted 5 times

  **kaz40**  2 months, 2 weeks ago

C must be correct. D would be correct with WHERE promo_cost < ANY

upvoted 1 times

  **nautil2** 3 months, 1 week ago


Selected Answer: D

A - false; keyword ALL is valid, it is used in Group Comparison Conditions, see Oracle doc: Oracle Database, Release 19, SQL Language Reference, 6 Conditions, Comparison Conditions

B - false; GROUP BY clause is valid, it is used with Aggregate Functions, see Oracle doc: Oracle Database, Release 19, SQL Language Reference, 7 Functions, Aggregate Functions: C - false; The query gives required results. The most inner query returns maximum promotional costs within groups formed with all promotions with same duration in days. The inner query which starts SELECT promo_id returns those ID of promotions, which costs is smaller than any maximum costs returned with the most inner query. The outer query returns those product IDs, which have the promo_id returned with the previous query. There must be some promo_cost which is smaller than any maximum promo_cost within groups. Only in that case the query returns some rows.

D - true; see C

upvoted 1 times

  **Lydia1054** 6 months ago

I think both C and D can be right, and it depends on the data:

1. Get the result: Two promotion data that have the same period but with different cost, and one cost has the minimum cost of all promotions.

>> if set the data like this:

1-\$5000-period 10 days/ 2- \$3000-period 15 days/ 3- \$2500- period 30 days/ 4-\$3500-period 30 days

>> We can get the ALL list of (\$5000, \$3000, \$3500), and \$2500 will be left to the group but this cost is meet the < ALL situation and the related product ID will show up in the end.

2. No result in the end:

If the example above change the third data to \$3000, then we can still get the same ALL list (\$5000, \$3000, \$3500) but there are no other cost that is less than the cost in the list. Therefore, we get no results.

upvoted 1 times

  **Darkseid1231** 6 months, 4 weeks ago

D D D IS CORRECT

upvoted 1 times

  **fthusa** 8 months ago

CORRECT C , executing success but not giving expected answer

upvoted 1 times

  **J4vi** 9 months ago

Selected Answer: D



I see no errors here

upvoted 1 times

  **hadiwuu** 1 year ago

who give the suggested anwer? why are all of it are wrong?

upvoted 1 times

  **Zairlam** 1 year, 10 months ago

I have recreated the tables and inserted some random data on them and tested the query.



The correct answer is C:



```
When you run
select prod_id from costs where prod_id in
(select prod_id from promotions where promo_cost < all
(select max(promo_cost) from promotions group by (promo_end-promo_begin)));
```



The query does go through but returns no results.



When you run the same query without the group by part you get the expected results (prod_id 2 had the highest promo_cost)



PROD_ID
1
3
4
upvoted 2 times



  **ryuah** 2 years ago
C is correct
upvoted 1 times



  **ryuah** 2 years ago
sry D is correct
upvoted 1 times



  **EIYA** 2 years ago
D is correct
upvoted 2 times



  **ogdru** 2 years, 10 months ago
<ALL: less than the lowest value returned by subquery.
<ANY: less than the highest value returned by subquery.
upvoted 5 times



  **Shakhzod1999** 11 months, 1 week ago
<ALL: less than the highest value returned by subquery.
<ANY: less than the lowest value returned by subquery.
upvoted 1 times

  **Fan** 2 years, 10 months ago
thank you.
upvoted 1 times

  **Phat** 3 years ago
so what is the correct answer ???
upvoted 2 times

  **sseet40** 3 years ago
The correct answer is D.
<ALL means that The value must be smaller than the smallest value in the list to evaluate to TRUE
upvoted 4 times

  **mberrios** 3 years, 3 months ago
c - correct
upvoted 2 times

  **adoptc94** 3 years, 4 months ago
I'd suggest answer A, but I am not completely sure!
upvoted 3 times


Which three statements are true about GLOBAL TEMPORARY TABLES? (Choose three.)


- A. A TRUNCATE command issued in a session causes all rows in a GLOBAL TEMPORARY TABLE for the issuing session to be deleted.
- B. GLOBAL TEMPORARY TABLE rows inserted by a session are available to any other session whose user has been granted select on the table.
- C. GLOBAL TEMPORARY TABLE space allocation occurs at session start.
- D. Any GLOBAL TEMPORARY TABLE rows existing at session termination will be deleted.
- E. A GLOBAL TEMPORARY TABLE'S definition is available to multiple sessions.
- F. A DELETE command on a GLOBAL TEMPORARY TABLE cannot be rolled back.


Correct Answer: ADE

Community vote distribution

ADE (100%)


 **danito** Highly Voted 3 years, 8 months ago
A D E i think the correct answers
upvoted 21 times

 **nautil2** Most Recent 3 months ago
Selected Answer: ADE
A - true; TRUNCATE command works in the same way for GTT as for standard tables, it removes all rows. However it is not recommended, because GTT are cleaned at the end of the transaction or the session.
B - false; from CREATE TABLE documentation: The data in a temporary table is visible only to the session that inserts the data into the table.
C - false; from documentation: Space is allocated for the table segment at the time of the first DML operation on the table.
D - true; from documentation: A session becomes unbound to a temporary table with a TRUNCATE statement or at session termination, or, for a transaction-specific temporary table, by issuing a COMMIT or ROLLBACK statement.
E - true; from documentation: Specify GLOBAL TEMPORARY to create a temporary table, whose definition is visible to all sessions with appropriate privileges.
F - false; no mention in documentation about non-standard consequence of DELETE command for GTT. Documentation mentions ROLLBACK.
upvoted 1 times


 **Xjackfbo** 10 months, 2 weeks ago
Only D & E are correct


B tested and couldn't see the data

A & F are not valid in case I create the GTT with a clause ON COMMIT PRESERVE in the end of the creation
C is wrong as Franky_T said
upvoted 2 times

 **Franky_T** 1 year, 9 months ago
Once again a question with four correct answers and not three.

A is correct. Can be easily tested. The default behavior for the ON COMMIT action is to delete rows if no hold-able cursor is open on the table. A TRUNCATE (DDL) performs an implicit COMMIT which closes any open cursor in the session.
B is wrong. You can grant other users SELECT privilege on your own GTT, but even if you do so the other user cannot see it. Can be easily tested.
C is wrong. GTT space allocation (TEMP tablespace) occurs when it's created and not when the session starts.
D is correct. The GTT exists at session level.
E is correct. GTT's metadata is stored on disk and visible to all sessions, GTT's data however is session specific.
F is correct. Can be easily tested. ON ROLLBACK DELETE ROWS is the default behavior for the NOT LOGGED option, thus removing all rows from a GTT when a ROLLBACK command is issued in the session.
upvoted 1 times

 **Aramazd** 1 year, 9 months ago
D E F are correct answers.
upvoted 2 times

 **Aramazd** 1 year, 9 months ago
At the termination of the session the rows are deleted but the table still exists:

SQL> INSERT INTO temp1(id,description)
VALUES(1,'Transaction specific global temp table');
2
1 row created.

```
SQL>
SQL> select * from TEMP1;

ID
-----
DESCRIPTION
-----
1
Transaction specific global temp table

SQL> exit
Disconnected from Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0
[oracle@localhost ~]$ sqlplus system/oracle

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Apr 6 00:41:30 2022
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.



Last Successful login time: Wed Apr 06 2022 00:40:47 -04:00

Connected to:
Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0
```

```
SQL> select * from TEMP1;

no rows selected
```

```
SQL>
upvoted 1 times
```

  **Aramazd** 1 year, 9 months ago
A delete from temp table cannot be rolled back:
SQL> select * from TEMP1;

```
ID
-----
DESCRIPTION
-----
1
Transaction specific global temp table
```

```
SQL> delete from temp1;

1 row deleted.
```

```
SQL> select * from TEMP1;

no rows selected
```

```
SQL> rollback;

Rollback complete.
```

```
SQL> select * from TEMP1;

no rows selected
```

```
SQL>
upvoted 1 times
```



  **josue1** 1 year, 7 months ago
SQL> create global temporary table tmp_t (n integer) on commit preserve rows;

Table created.

```
SQL>
SQL> insert into tmp_t values(1);

1 row created.
```

```
SQL> commit;

Commit complete.
```

```
SQL> select * from tmp_t;
```

N

1

SQL> delete from tmp_t;

1 row deleted.

SQL> rollback;


Rollback complete.

SQL> select * from tmp_t;

N

1

upvoted 2 times



  **casfdsaf** 1 year, 10 months ago

Selected Answer: ADE

ADE correct
upvoted 1 times

  **ryuah** 2 years ago

A,D,E is correct
upvoted 1 times

  **ogdru** 2 years, 10 months ago

F is true
SQL> select * from gtt1;

ID

1

1 row selected.

SQL> delete from gtt1;

1 row deleted.

SQL> select * from gtt1;



no rows selected



SQL> rollback;

Rollback complete.

SQL> select * from gtt1;

no rows selected
upvoted 1 times

  **_gio_** 2 years, 3 months ago
but if you write a savepoint it works....
upvoted 1 times

  **NowOrNever** 3 years, 4 months ago

BCF - are definitely WRONG
upvoted 1 times

  **saif_alrwiliy** 3 years, 5 months ago

- A. GLOBAL TEMPORARY TABLE rows inserted by a session are available to any other session whose user has been granted select on the table -> FALSE, you cant view data from other session
- B. GLOBAL TEMPORARY TABLE space allocation occurs at session start. -> AT First dml sentence, then space is allocated ¿false?
- C. A DELETE command on a GLOBAL TEMPORARY TABLE cannot be rolled back. -> Depends. If you create with on commit preserve rows, you can rollback and commit. -> FALSE
- D. A GLOBAL TEMPORARY TABLE's definition is available to multiple sessions. -> Yes, with the appropriate permissions ¿TRUE?
- E. Any GLOBAL TEMPORARY TABLE rows existing at session termination will be deleted. -> TRUE
- F. A TRUNCATE command issued in a session causes all rows in a GLOBAL TEMPORARY TABLE for the issuing session to be deleted. -> TRUE, truncate delete rows from table

¿D,E,F?
Very well presented explanations from Primisser
<https://www.examttopics.com/discussions/oracle/view/8895-exam-1z0-071-topic-1-question-298-discussion/>
upvoted 3 times



  **NiciMilo** 3 years, 6 months ago

店长微信 : hj feng128

why is E correct?
Reference:

"Other users do not need the same table structure."
"Multiple connections can define declared global temporary tables with the same name"

<https://docs.oracle.com/javadb/10.8.3.0/ref/rrefdeclaretemptable.html>
upvoted 2 times

  **ama** 3 years, 5 months ago

Global temporary Tables are permanent tables created and available to users.
only the data in the table is private to the user who inserted into it.

so the Definition (DDL) is available to multiple sessions
upvoted 1 times

Which two statements are true about trace files produced by the Oracle Database server? (Choose two.)

- A. They can be written by server processes
- B. Trace files are written to the Fast Recovery Area (FRA)
- C. They can be written by background processes
- D. All trace files contain error information that require contacting Oracle Support
- E. Trace file names are based on the database name concatenated with a sequential number

Correct Answer: AC

Reference:

https://gerardnico.com/db/oracle/trace_file

Community vote distribution

AC (100%)



  **GraceYu** Highly Voted 3 years, 7 months ago

AC

https://docs.oracle.com/html/E25494_01/monitoring001.htm

Each server and background process can write to an associated trace file. When an internal error is detected by a process, it dumps information about the error to its trace file. Some of the information written to a trace file is intended for the database administrator, and other information is for Oracle Support Services. Trace file information is also used to tune applications and instances.

upvoted 13 times

  **Ekos** 3 years, 1 month ago

i agree

upvoted 2 times

  **nautil2** Most Recent 3 months ago

Selected Answer: AC

A - true; Oracle documentation: Each server and background process can periodically write to an associated trace file.
B - false; Oracle documentation: The ADR base can contain multiple ADR homes, where each ADR home is the root directory for all diagnostic data—traces, dumps, the alert log, and so on—for an instance of an Oracle product or component.
C - true; see answer A
D - false; Oracle documentation: Each server and background process can periodically write to an associated trace file. The files contain information on the process environment, status, activities, and errors.
E - false; database name does not need to be same as SID; Oracle documentation: Typically, database background process trace file names contain the Oracle SID, the background process name, and the operating system process number. An example of a trace file for the RECO process is mytest_reco_10355.trc.

upvoted 1 times

  **casfdsaf** 1 year, 10 months ago

Selected Answer: AC

AC correct

upvoted 1 times

  **ryuah** 2 years ago

A,C is correct

upvoted 1 times

st

You need to calculate the number of days from 1 January 2019 until today.

Dates are stored in the default format of DD-MON-RR.

Which two queries give the required output? (Choose two.)

- A. SELECT TO_CHAR(SYSDATE, 'DD-MON-YYYY') 01' "€λ-JAN-2019' FROM DUAL;
- B. SELECT ROUND(SYSDATE 01' "€λ-JAN-2019') FROM DUAL;
- C. SELECT ROUND(SYSDATE λ€" TO_DATE('01/JANUARY/2019')) FROM DUAL;
- D. SELECT TO_DATE(SYSDATE, 'DD/MONTH/YYYY') 01' "€λ/JANUARY/2019' FROM DUAL;
- E. SELECT SYSDATE λ€" TO_DATE('01-JANUARY-2019') FROM DUAL;

Correct Answer: CE

Community vote distribution

CE (100%)



  **nautil2** 3 months ago

Selected Answer: CE

A - false; query <SELECT TO_CHAR(sysdate,'DD-MON-YYYY') - '01-JAN-2019' FROM dual;> returns ORA-01722: invalid number
B - false; query <SELECT ROUND(sysdate - '01-JAN-2019') FROM dual;> returns ORA-01722: invalid number
C - true; query <SELECT ROUND(sysdate - TO_DATE('01/JANUARY/2019')) FROM dual;> returns number of days as natural number
D - false; query <SELECT TO_DATE(sysdate,'DD/MONTN/YYYY') - '01/JANUARY/2019' FROM dual;> returns ORA-01821: date format not recognized
E - true; query <SELECT sysdate - TO_DATE('01/JANUARY/2019') FROM dual;> returns number of days as floating number
upvoted 1 times

  **Oracle2020** 5 months ago

The Answer: C, E
The dates in the select '01-jan-2019' have to be converted with To_Date
upvoted 1 times

  **musafir** 6 months, 1 week ago

The Options should be:

A. SELECT TO_CHAR (SYSDATE, 'DD-MON-YYYY') – '01-JAN-2019' FROM DUAL;
B. SELECT ROUND (SYSDATE – '01-JAN-2019') FROM DUAL;
C. SELECT ROUND (SYSDATE – TO_DATE('01/JANUARY/2019')) FROM DUAL;
D. SELECT TO_DATE (SYSDATE, 'DD/MONTH/YYYY') – '01/JANUARY/2019' FROM DUAL;
E. SELECT SYSDATE – TO_DATE ('01-JANUARY-2019') FROM DUAL;

C and E are correct
upvoted 1 times

  **RaNik69** 9 months ago

Selected Answer: CE

SELECT TO_CHAR(SYSDATE, 'DD-MON-YYYY') - '01-JAN-2019' FROM DUAL; [Error] Execution (16: 44): ORA-01722: numero non valido;
SELECT ROUND(SYSDATE - '01-gen-2019') FROM DUAL; [Error] Execution (16: 44): ORA-01722: numero non valido;
SELECT ROUND(SYSDATE - TO_DATE('01/gen/2019')) FROM DUAL; 1576;
SELECT TO_DATE(SYSDATE, 'DD/MONTH/YYYY') - '01/gennaio/2019' FROM DUAL; [Error] Execution (16: 44): ORA-01722: numero non valido;
SELECT SYSDATE - TO_DATE('01-gennaio-2019') FROM DUAL; 1575,681400462962962962962962962963
upvoted 1 times

  **jonsnoow** 1 year, 4 months ago

C and E
upvoted 2 times

Which two statements are true about the DUAL table? (Choose two.)

- A. It can be accessed only by the SYS user
- B. It consists of a single row and single column of VARCHAR2 data type
- C. It can display multiple rows but only a single column
- D. It can be used to display only constants or pseudo columns
- E. It can be accessed by any user who has the SELECT privilege in any schema
- F. It can display multiple rows and columns

Correct Answer: AB

Reference:

https://en.wikipedia.org/wiki/DUAL_table

Community vote distribution



Sharif1 Highly Voted 3 years, 5 months ago

B and F are correct answers.
F is correct :
SQL> SELECT level, sysdate
2 FROM dual
3 CONNECT BY
4 level <= 4;

```
LEVEL SYSDATE
-----
1 06-AUG-20
2 06-AUG-20
3 06-AUG-20
4 06-AUG-20
```

It can return multiple rows and columns.

D is wrong, you just create a brand new user and do not grant any privilege but still it can select dual table.
upvoted 19 times

AnimeshOracle 1 year, 8 months ago

F is incorrect because Dual has only one column. Here in your example what you are doing is a projection of multiple column through select statement which can be done in any table even though the table has limited number of columns.
upvoted 1 times

AnimeshOracle 1 year, 8 months ago
B,E seems to be more correct options.
upvoted 1 times

SimoneF 3 years, 1 month ago

Thank you, I got to F by exclusion but I was having some trouble finding a good example of a case in which it returned multiple rows.
upvoted 4 times

danito Highly Voted 3 years, 8 months ago

B E https://docs.oracle.com/cd/B19306_01/server.102/b14200/queries009.htm
upvoted 11 times



ama 3 years, 6 months ago
E seems wrong, i would choose A and D
upvoted 2 times



ama 3 years, 6 months ago
sorry i mean, B & D
upvoted 5 times

NiciMilo 3 years, 6 months ago
D is incorrect

"Alternatively, you can select a constant, pseudocolumn, or ***expression*** from any table"

upvoted 2 times

  **ama** 3 years, 5 months ago
well, then B, F seems fine
upvoted 1 times

  **Trowa** 3 years, 5 months ago
BD are the correct answers.

<https://stackoverflow.com/questions/50212079/is-it-possible-to-have-multiple-rows-with-select-from-dual>
upvoted 1 times

  **nautil2** Most Recent 3 months ago



Selected Answer: BF



A - false; DUAL is a table automatically created by Oracle Database along with the data dictionary. DUAL is in the schema of the user SYS but is accessible by the name DUAL to all users. Source: Oracle documentation - Oracle Database, Release 19, SQL Language Reference, 9 SQL Queries and Subqueries, Selecting from the DUAL Table
B - true; single column DUMMY with datatype VARCHAR2(1) and value 'X'
C - false; query SELECT sysdate, sysdate FROM dual;
D - false; Alternatively, you can select a constant, pseudocolumn, or expression from any table, but the value will be returned as many times as there are rows in the table. Source: same as mentioned in answer A
E - false; GRANT SELECT any table is not necessary to query DUAL table. Newly created user is not granted with any SELECT privilege, however he/she can query any own object and table DUAL.
F - true; example SELECT with multiple row and column output is SELECT sysdate,sysdate+ 1 FROM dual CONNECT BY LEVEL <= 5;
upvoted 1 times

  **ismoil** 3 months ago

If I run
desc dual;
Name Null Type

DUMMY VARCHAR2(1)
upvoted 1 times

  **ismoil** 3 months ago
So answer is B E
upvoted 1 times



  **Oracle2020** 5 months, 1 week ago


The correct answer is A,B.
the dual table has only one row, it is accessed with the user sys and the column is of type varchar2 .
upvoted 1 times

  **zouve** 5 months, 2 weeks ago

Selected Answer: BE

BE for me
upvoted 1 times

  **zouve** 5 months ago
DUAL is a table automatically created by Oracle Database along with the data dictionary.
DUAL is in the schema of the user SYS but is accessible by the name DUAL to all users. It has one column, DUMMY, defined to be VARCHAR2(1), and contains one row with a value X.
Selecting from the DUAL table is useful for computing a constant expression with the SELECT statement. Because DUAL has only one row, the constant is returned only once.
Alternatively, you can select a constant, pseudocolumn, or expression from any table, but the value will be returned as many times as there are rows in the table.
upvoted 1 times


  **auwia** 6 months, 3 weeks ago

Selected Answer: BE

A. It can be accessed only by the SYS user --> False, anybody can select from dual.
B. It consists of a single row and single column of VARCHAR2 data type. --> True
C. It can display multiple rows but only a single column. --> False, only 1 single row and column (by default: select * from dual)
D. It can be used to display only constants or pseudo columns. --> False, Expression too.
E. It can be accessed by any user who has the SELECT privilege in any schema. --> True
F. It can display multiple rows and columns. False, 1 row/column only.
upvoted 1 times

  **fthusa** 8 months ago



Answers should be BE
upvoted 1 times

  **Emperor11** 11 months ago

Selected Answer: BF

I correct myself B and F are correct because when you give the create session to any user, he is capable to select the dual table without give that exactly permission.

upvoted 1 times



  **Emperor11** 11 months ago

Selected Answer: BE

B and E are 100% correct.

"DUAL is a table automatically created by Oracle Database along with the data dictionary. DUAL is in the schema of the user SYS but is accessible by the name DUAL to all users. It has one column, DUMMY, defined to be VARCHAR2(1), and contains one row with a value X."
<https://docs.oracle.com/database/121/SQLRF/queries009.htm#SQLRF20036>

upvoted 1 times

  **Emperor11** 11 months ago

B and E are 100% correct.

"DUAL is a table automatically created by Oracle Database along with the data dictionary. DUAL is in the schema of the user SYS but is accessible by the name DUAL to all users. It has one column, DUMMY, defined to be VARCHAR2(1), and contains one row with a value X."
<https://docs.oracle.com/database/121/SQLRF/queries009.htm#SQLRF20036>

upvoted 1 times

  **chenwen** 1 year, 8 months ago

B,E IS CORRECT

upvoted 2 times

  **casfdsaf** 1 year, 10 months ago

Selected Answer: EF

I THINK EF

upvoted 2 times

  **yarsalan** 1 year, 12 months ago

Selected Answer: BF

B , F are correct.

upvoted 2 times

  **ryuah** 2 years ago

B,F is correct

upvoted 2 times

  **dummydba** 2 years, 8 months ago



B is correct as you guys explained here. And the other one should be E.
E. It can be accessed by any user who has the select privilege in any schema.
Any schema owner can select data from its own tables. so that argument is valid here.

upvoted 2 times

  **hp9000** 3 years ago

We agree on B, what's the second right answer?

upvoted 1 times

  **Phat** 2 years, 10 months ago

D as Sharif1 explained.

upvoted 1 times

Which four statements are true regarding primary and foreign key constraints and the effect they can have on table data? (Choose four.)

- A. A table can have only one primary key but multiple foreign keys
- B. A table can have only one primary key and one foreign key
- C. The foreign key columns and parent table primary key columns must have the same names
- D. It is possible for child rows that have a foreign key to remain in the child table at the time the parent row is deleted
- E. It is possible for child rows that have a foreign key to be deleted automatically from the child table at the time the parent row is deleted
- F. Only the primary key can be defined at the column and table level
- G. Primary key and foreign key constraints can be defined at both the column and table level

Correct Answer: ACEG

Community vote distribution



- danito**

Highly Voted

3 years, 7 months ago

ADEG i think
upvoted 27 times
- NiciMilo**

3 years, 5 months ago

How is D possible?
upvoted 1 times
- nhsang**

3 years, 5 months ago

In the relational 1-0, you can use ON DELETE SET NULL to keep child-records, but child-keys will be updated to NULL.
upvoted 5 times
- ama**

3 years, 5 months ago

this means that D is wrong
upvoted 1 times
- SimoneF**

3 years, 1 month ago

No because the child row stays, just some of its values are nullified.
upvoted 3 times
- renzoku**

2 years, 10 months ago

But you can be achieved by disabling the FK constraint and re enabling with no validate clause after performing delete on parent table.
upvoted 2 times
- Ekos**

3 years, 1 month ago

i agree
upvoted 1 times
- avanand**

Highly Voted

3 years, 5 months ago

ADEG in my opinion too...

A. A table can have only one primary key but multiple foreign keys --> For sure
B. A table can have only one primary key and one foreign key --> There is no such restrictions
C. The foreign key columns and parent table primary key columns must have the same names --> There is no such restriction.
D. It is possible for child rows that have a foreign key to remain in the child table at the time the parent row is deleted --> Can be achieved by disabling the FK constraint and re enabling with no validate clause after performing delete on parent table.
E. It is possible for child rows that have a foreign key to be deleted automatically from the child table at the time the parent row is deleted --> CASCADE option
F. Only the primary key can be defined at the column and table level --> No such restrictions
G. Primary key and foreign key constraints can be defined at both the column and table level --> True.
upvoted 16 times
- LrnsTgh**

2 years, 5 months ago

might correct A, B (can have), D, E, G
wrong C ("Must"), F (PK,FK, UQ, CK defined at column and table level. only not null in column level).

I vote ADEG
upvoted 1 times

Selected Answer: ADEG

A - true; quote1: A table or view can have only one primary key., quote2: You can define multiple foreign keys in a table or view.
B - false; see quote2 in answer A
C - false; example, when names of FK columns and parent table columns are different: ALTER TABLE employees2 ADD CONSTRAINT fk_dept2 FOREIGN KEY (dept_id) REFERENCES departments2(id);
D - true; example: ALTER TABLE employees ADD CONSTRAINT fk_dept FOREIGN KEY (dept_id) REFERENCES departments (id) ON DELETE SET NULL; perform INSERT(s) into employees with dept_id=17; DELETE FROM departments WHERE dept_id=17;
E - true; example: ALTER TABLE employees ADD CONSTRAINT fk_dept FOREIGN KEY (dept_id) REFERENCES departments (id) ON DELETE CASCADE;
F - false; Either primary key, unique constraint, foreign key or check constraints can be defined on both table and column levels.
G - true; see answer F
upvoted 1 times

🗨️ **nautil2** 2 months, 3 weeks ago

quotes originate in following sources:
A - Oracle documentation: Database > Oracle > Oracle Database > Release 19 > SQL Language Reference > 8 Common SQL DDL Clauses > constraint
F,G - Oracle documentation: Database > Oracle > Oracle Database > Release 19 > SQL Language Reference > 12 SQL Statements: ALTER SYNONYM to COMMENT > ALTER TABLE
upvoted 1 times

🗨️ **nautil2** 2 months, 4 weeks ago

Selected Answer: ADEG

A - true; You can define multiple foreign keys in a table or view.
B - false; see A
C - false; a foreign key needn't be based on a primary key, it can be based on unique key, therefore its columns can have different names.
D - true; parent row can be deleted, but parent column cannot be deleted (ORA-12992: cannot drop parent key column)
E - true; foreign key must be defined with ON DELETE CASCADE clause
F - false; not only primary key, but also other constraints can be defined both on table level and column level, e.g. CREATE TABLE a (b NUMBER UNIQUE) or CREATE TABLE a (b NUMBER, UNIQUE(b))
G - true; e.g. CREATE TABLE child (a number, b number REFERENCES parent(c)) or CREATE TABLE CHILD (a number, b number, FOREIGN KEY (parent_id) REFERENCES parent(c))

Source:
A,B,C Oracle documentation: Oracle Database, Release 19, SQL Language Reference, 8 Common SQL DDL Clauses, constraint
upvoted 1 times

🗨️ **Oracle2020** 5 months, 1 week ago

The correct answers are A,C,E,G
The D is not correct because for referential integrity it does not allow you to delete rows from the parent table and leave related rows in the child tables.
The E is correct because effectively when you delete rows from the parent table and you specify on delete cascade in the delete statement, it will delete the related rows in the child tables. If it is not specified, it does not allow deleting records from the parent table if you have related records.
upvoted 2 times

🗨️ **zouve** 5 months, 2 weeks ago

Selected Answer: ADEG

AEG are for sure right
BFC are definitely wrong so i believe ADEG is the right combination
upvoted 1 times

🗨️ **fthusa** 8 months ago

ADEG are true ,
upvoted 1 times

🗨️ **ryuah** 1 year, 7 months ago

Selected Answer: ADEG

good choice
upvoted 1 times


🗨️ **algerianphoenix** 1 year, 8 months ago


ADEG, without a shadow of doubt, just common DB knowledge.
upvoted 1 times

🗨️ **Franky_T** 1 year, 9 months ago

Selected Answer: ADEG


A is correct. The only restriction is on the PK, a table can have only one. No restrictions exist on the number of FK's a table may have.
B is wrong. Tables may have many foreign keys in them, it all depends on the db's relational model.
C is wrong. The DB does not care about column names, only data type, size and contents of the PK and FK.
D is correct. The FK must be created with the ON DELETE SET NULL option. Without this option ORA-02292 will be raised when attempting to delete a row from the parent table that has child record.
E is correct. The FK must be created with ON DELETE CASCADE option. Without this option ORA-02292 will be raised when attempting to delete a row from the parent table that has child record.
F is wrong. The only constraint that has to be defined inline (column level) is NOT NULL.
G is correct. The only constraint that has to be defined inline (column level) is NOT NULL.
upvoted 1 times


- 

 **casfdsaf** 1 year, 10 months ago


Selected Answer: ADEF


adeg is correct

upvoted 1 times
- 


 **ryuah** 2 years ago


A,D,E,G is correct

upvoted 2 times
- 

 **LongHi** 2 years, 1 month ago


ADEG correct


upvoted 1 times
- 

 **LongHi** 2 years, 1 month ago


D: using ON DELETE cascade


F: using on delete set null

upvoted 1 times
- 

 **LongHi** 2 years, 1 month ago

sorry error typo: E: using on delete set null

upvoted 1 times
- 

 **you1234** 3 years, 6 months ago

ACEG is correct answer

upvoted 4 times

Examine the description of the EMPLOYEES table:

Name	Null?	Type
-----	-----	-----
EMP_ID	NOT NULL	NUMBER
EMP_NAME		VARCHAR2 (40)
DEPT_ID		NUMBER (2)
SALARY		NUMBER (8, 2)
JOIN_DATE		DATE

Which query is valid?

- A. SELECT dept_id, AVG(MAX(salary)) FROM employees GROUP BY dept_id;
- B. SELECT dept_id, join_date, SUM(salary) FROM employees GROUP BY dept_id;
- C. SELECT dept_id, join_date, SUM(salary) FROM employees GROUP BY dept_id, join_date;
- D. SELECT dept_id, MAX(AVG(salary)) FROM employees GROUP BY dept_id;

Correct Answer: B

Community vote distribution



you1234 Highly Voted 3 years, 6 months ago
C is correct answer
upvoted 23 times

danito Highly Voted 3 years, 7 months ago
I think C
upvoted 9 times

Ekos 3 years, 1 month ago
i agree
upvoted 2 times

nautil2 Most Recent 2 months, 3 weeks ago
Selected Answer: C
A - false; returns error ORA-00937: not a single-group group function; first occurrence of dept_id is marked as the cause of the error
B - false; returns error ORA-00979: not a GROUP BY expression; join_date is marked as the cause of the error
C - true; columns mentioned between SELECT and FROM keywords are also mentioned after GROUP BY clause
D - false; same as answer A
upvoted 1 times

auwia 6 months, 3 weeks ago
Selected Answer: B
I see only B as invalid query, probably the question is wrong and it should be:

Which query is INvalid?
upvoted 1 times

fthusa 8 months ago
C is correct
upvoted 1 times

FelipeC 11 months ago
Selected Answer: C
A (FALSE). SELECT dept_id, AVG(MAX(salary)) FROM employees GROUP BY dept_id;
select DEPARTMENT_ID, AVG(MAX(SALARY)) from EMPLOYEES group by DEPARTMENT_ID;
ERROR at line 1:
ORA-00937: not a single-group group function

B (FALSE). SELECT dept_id, join_date, SUM(salary) FROM employees GROUP BY dept_id;
select DEPARTMENT_ID, HIRE_DATE,sum(salary) from EMPLOYEES group by DEPARTMENT_ID
*
ERROR at line 1:
ORA-00979: not a GROUP BY expression

店长微信: hifeng128

5150128

C (TRUE). SELECT dept_id, join_date, SUM(salary) FROM employees GROUP BY dept_id, join_date;
SQL> select DEPARTMENT_ID, HIRE_DATE,SUM(SALARY) from EMPLOYEES group by DEPARTMENT_ID , HIRE_DATE;

DEPARTMENT_ID HIRE_DATE SUM(SALARY)



60 25-JUN-05 4800

D (FALSE). SELECT dept_id, MAX(AVG(salary)) FROM employees GROUP BY dept_id;
SELECT DEPARTMENT_ID , MAX(AVG(salary)) FROM employees GROUP BY DEPARTMENT_ID
*

ERROR at line 1:
ORA-00937: not a single-group group function
upvoted 2 times

  **OCP19c** 1 year, 4 months ago

C is correct answer
B is wrong answer, because not working!
But why not A not ?
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

A and D are wrong simply because we cannot use nested group functions while grouping by any specific column, we should use the whole table as one group (by not using a group by clause) otherwise the query would not make sense.
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

C, by testing on OE DB sample.
upvoted 1 times

  **Franky_T** 1 year, 9 months ago

Selected Answer: C

A is wrong. When using nested aggregate functions you are not allowed to display any additional column(s) other than the nested aggregate function.
B is wrong. Missing the join_date column in the GROUP BY clause. The rule states: "Any column in the select list that's not part of an aggregate function MUST BE in the group by clause".
C is correct.
D is wrong. Same issue as with option A.
upvoted 2 times

  **Sankar1406** 1 year, 9 months ago

Selected Answer: C

C for sure
upvoted 1 times

  **sagartake** 1 year, 10 months ago

Selected Answer: B



B is correct
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

C, in B there is a missing column in the group by clause.
upvoted 2 times

  **ryuah** 2 years ago

C is correct
upvoted 1 times

  **Taush** 2 years, 4 months ago

I created the same table with same data type: Option A and B gave errors while C and D gave the same output.
upvoted 1 times

  **Guhborges** 2 years, 4 months ago

C is the correct answer:

SELECT dept_id, join_date, SUM(salary) FROM employees GROUP BY dept_id, join_date;

All options with non-groupable functions need to be added in GROUP BY.
upvoted 1 times

What is true about non-equijoin statement performance? (Choose two.)

- A. Table aliases can improve performance
- B. The BETWEEN condition always performs better than using the >= and <= conditions
- C. The join syntax used makes no difference to performance
- D. The BETWEEN condition always performs less well than using the >= and <= conditions
- E. The Oracle join syntax performs better than the SQL:1999 compliant ANSI join syntax

Correct Answer: AE

Community vote distribution

AC (100%)

  **khalilshahin01** Highly Voted 3 years, 5 months ago

Dear all,
AC is correct answer for sure
i faced this question in 1z0-071 exam and i select AC and got correct answers :)
Regards.
upvoted 18 times


  **avanand** Highly Voted 3 years, 5 months ago

A and C.

There's no performance benefit or hit by using ANSI joins rather than traditional joins, but by using ANSI joins, your queries are more portable between DBMS platforms, and they're a bit easier to read.

http://www.dba-oracle.com/oracle_news/2004_2_19_rittman.htm

upvoted 5 times

  **Ekos** 3 years, 1 month ago

i agree, AC is the correct answer
upvoted 1 times

  **adoptc94** 3 years, 5 months ago


So now we have to choose between A and C or A and E

After reading through following explanation it seems to me that there are some specific cases where traditional joins perform better
Source (taken from ama's post up there): <https://stackoverflow.com/questions/58476906/1z0-071-oracle-exam-question-which-answers-is-correct>

But generally speaking there shouldn't be any big differences in performance between traditional and ansi join syntax - I am pretty sure that there are also some specific cases where ansi joins perform better than traditional joins.

So in that case I would go with A and C and just hope for the best!

upvoted 4 times

  **Fan** 2 years, 10 months ago

ORACLE will never issue a software worse than others. so A E are correct.
upvoted 1 times

  **yukclam9** 2 years, 9 months ago

your point is valid and they make all joining syntax equal in performance. there is no winner
upvoted 1 times

  **nautil2** Most Recent 2 months, 3 weeks ago

Selected Answer: AC

- A - true; there can be a performance benefit when using table aliases
- B - false; there is no performance benefit difference when using keyword BETWEEN or operators <= =>
- C - true; there is no difference in performance when using SQL 1999 ANSI join syntax or Oracle syntax
- D - false; see answer B
- E - false; see answer C

Sources

A - <https://orawiss.wordpress.com/table-alias-speed-up-database-access/>

According to the Oracle Sql fundamentals written by Nancy Greenberg and reviewed by a few technical contributors in the page 151, the Table aliases speed up database access and the Table aliases can help to keep SQL code smaller by conserving memory.

upvoted 1 times

- auwia** 6 months, 3 weeks ago

Selected Answer: AC

Provided answer are wrong.

upvoted 1 times
- RaNik69** 8 months, 1 week ago

Selected Answer: AC

AC is correct

upvoted 1 times
- ryuah** 2 years ago

A,C is correct

upvoted 1 times
- Chansi** 3 years, 5 months ago

Thin its A and C. any evidence for E?

upvoted 2 times
- ama** 3 years, 5 months ago

E is correct... check this> <https://stackoverflow.com/questions/58476906/1z0-071-oracle-exam-question-which-answers-is-correct>

so i bet on A, E

upvoted 1 times
- you1234** 3 years, 6 months ago

B & C looks close answers

upvoted 1 times
- ama** 3 years, 5 months ago

A, E !!!

upvoted 1 times
- dotruonghan** 3 years, 6 months ago

A is correct

<https://orawiss.wordpress.com/table-alias-speed-up-database-access/>

upvoted 1 times
- you1234** 3 years, 6 months ago

B & E is correct one

upvoted 1 times
- ama** 3 years, 6 months ago

can you Elaborate more ? is this a guess or you have evidence?

upvoted 1 times

In your data center, Oracle Managed Files (OMF) is used for all databases.

All tablespaces are smallfile tablespaces.

SALES_Q1 is a permanent user-defined tablespace in the SALES database.

Examine this command which is about to be issued by a DBA logged in to the SALES database:

ALTER TABLESPACE sales_q1 ADD DATAFILE;

Which are two actions, either one of which you could take to ensure that the command executes successfully? (Choose two.)

- A. Ensure that DB_RECOVERY_FILE_DEST and DB_CREATE_FILE_DEST each specify with at least 50 Mb of available space.
- B. Specify a path in the DATAFILE clause of the command specifying a location with at least 100M of available space.
- C. Ensure that DB_CREATE_FILE_DEST specifies a location with at least 100 Mb of available space.
- D. Add the AUTOEXTEND ON clause with NEXT set to 100M.
- E. Ensure that DB_RECOVERY_FILE_DEST and DB_CREATE_FILE_DEST each specify locations with at least 50 Mb of available space.

Correct Answer: BD

Community vote distribution

BC (100%)

 **NowOrNever** Highly Voted 3 years, 5 months ago

BC are the correct one

<https://oracle-base.com/articles/9i/oracle-managed-files>

- Files typically have a default size of 100M

- if OMF is used you do not need to specify a path for the new created datafile

BUT

we have to "ensure that the command executes successfully"

so we need a location with at least 100 MB free space and we can ensure that with B or C

upvoted 12 times

 **ama** 3 years, 5 months ago

really? if my database using OMF i should not add a path for the datafile.

The Database will do that for me !

upvoted 4 times

 **Phat** 2 years, 10 months ago


yes, db does that for you, but if you want to specify the path, it also fine.

upvoted 2 times

 **Chansi** Highly Voted 3 years, 5 months ago

This B,C because the question says either one of the solution should work. In a situation where there isn't enough space in the "db_create_file_dest" location, you either make sure to have enough space in "db_create_file_dest" or you can override by providing the path where there is enough space

upvoted 7 times

 **Ekos** 3 years, 1 month ago

i agree

upvoted 1 times

 **nautil2** Most Recent 2 months, 2 weeks ago

Selected Answer: BC

A - false; without specifying SIZE, a new OMF datafile is created with size of 100 MB in DB_CREATE_FILE_DEST, therefore 50 MB space is not enough

B - true; when using ALTER TABLESPACE sales_q1 ADD DATAFILE '<filename>'; command, <filename> can specify path of location. When the location contains at least 100 MB of available space, a new OMF datafile can be created.

C - true; DB_CREATE_FILE dest is a default directory where new OMF datafiles are created with default size 100 MB

D - false; AUTOEXTEND ON does not cause there is enough space for a new datafile

E - false; it is the same as answer A - same words

upvoted 1 times

 **cksim** 9 months, 3 weeks ago

it is c and d.

c is correct because of omf and make sure we have 100m space coz it is the default size



d is correct because even though autoextend and next have default values, it is no harm to statement execution.

upvoted 1 times

 **Kuraudio** 1 year, 1 month ago

Selected Answer: BC

BC is correct
upvoted 1 times

  **ryuah** 1 year, 7 months ago

Selected Answer: BC

good choice
upvoted 2 times

  **Borislone** 2 years, 6 months ago

BD is correct because in this case the database has both OMF and unmanaged files. with add data file clause, the path must be added
upvoted 1 times

  **adoptc94** 3 years, 4 months ago



Correct answers are B,C

Answers A and E are false as OMF defaults to a file size of at least 100MB

As for answer D, you can omit the AUTOEXTEND clause, the database will use the default value of AUTOEXTEND UNLIMITED instead. Using AUTOEXTEND ON won't ensure that the command executes successfully, as the file will be autoextensible either way.

So the only possible answers left are B and C.

upvoted 5 times

  **ama** 3 years, 5 months ago

Here the evidence that D is correct



Very unfortunately, the default AUTOEXTEND ON NEXT size is 1 Database block (based on the blocksize of the tablespace). But if you create your Tablespace using OMF (i.e. where "db_create_file_dest" is configured), then Oracle defaults the initial size to 100MB and *also* defaults the AUTOEXTEND to ON with a of 100MB !

upvoted 2 times

  **SimoneF** 3 years ago

I think what you just quoted actually disproves D: if the autoextend is already automatically set ON with a next extent of 100M, then giving the exact same option wouldn't change anything and especially wouldn't help ensuring that the statement completes correctly.

upvoted 3 times

  **ama** 3 years, 6 months ago



Well,

first of all OMF is in use so no Need to specify a path for the new created datafile.

Correct answers must be C, D

<https://oracle-base.com/articles/9i/oracle-managed-files>

upvoted 3 times

  **Fan** 2 years, 10 months ago

Oracle Managed Files (OMF) is used, so no need to specify a datafile location.

upvoted 2 times

Examine this command and some partial output:

```
LSNRCTL> start LISTENER_1

Listening Endpoints Summary...
  (DESCRIPTION=(ADDRESS=PROTOCOL=tcp) (HOST=host01.abc.com) (PORT=1562)))
Services Summary...
Service "DB01.abc.com " has 1 instance(s).
  Instance "DB01", status UNKNOWN, has 1 handler(s) for this service...
The command completed successfully
```

Why does the DB01.abc.com service show unknown status?

- A. The service DB01.abc.com is dynamically registered
- B. The LOCAL_LISTENER database parameter is not set to a service name that refers to LISTENER_1
- C. The service DB01.abc.com is statically registered
- D. The listener is not listening on the default port 1521
- E. The SID_LIST_LISTENER section is not contained in the LISTENER.ORA file

Correct Answer: B

Community vote distribution

6 (100%)



  **Rivaldo11** Highly Voted 3 years, 6 months ago

C is correct answer.
upvoted 23 times

  **ama** 3 years, 6 months ago

Good hint,



An UNKNOWN status means that the instance is registered statically (with a SID_LIST) in the listener. ora file rather than dynamically with service registration. Therefore, the database status is "not known" to the listener.
upvoted 4 times

  **Ekos** 3 years, 1 month ago


i agree
upvoted 1 times

  **nautil2** Most Recent 1 month, 4 weeks ago



Selected Answer: C
59
A - false; if a service is registered dynamically, its status handler is either READY, RESTRICTED or BLOCKED. UNKNOWN handler refers to static service registration
B - false; LOCAL_LISTENER parameter can be set regardless static or dynamic registration is used
C - true; service is registered statically using SID_LIST keyword in the config file listener.ora
D - false; listener port does not affect status handler
E - false; when SID_LIST_LISTENER is contained in listener.ora file, statical service registration is configured
upvoted 1 times

  **lps0530** 1 year ago

Selected Answer: C
C is correct answer.
upvoted 1 times

  **XhostSI** 1 year, 11 months ago

C is the answer!
upvoted 1 times

  **ofranco** 2 years, 4 months ago

Agree, it is C
upvoted 1 times



  **avanand** 3 years, 5 months ago

Its C without doubt.

upvoted 1 times

店长微信：hj feng128

  **danito** 3 years, 5 months ago
and D? the picture shows port 1562?
upvoted 1 times

  **danito** 3 years, 5 months ago
please forget my comment. totally agree the correct is C
upvoted 1 times

Which three statements are true about the tools used to configure Oracle Net Services? (Choose three.)

- A. The Oracle Net Configuration Assistant is only used when running the Oracle installer
- B. Oracle Net Manager can be used to centrally configure listeners on any database server target
- C. The lsnrctl utility requires a listener.ora file to exist before it is started
- D. Oracle Net Manager can be used to locally configure naming methods on a database server
- E. Enterprise Manager Cloud Control can be used to centrally configure listeners on any managed database server
- F. Enterprise Manager Cloud Control can be used to centrally configure net service names for any database server target

Correct Answer: BDE

Community vote distribution

DEF (100%)

  **GraceYu** Highly Voted 3 years, 7 months ago

DEF
Oracle Net Manager to locally control and Enterprise Manager Cloud Control to centrally control.
upvoted 15 times

  **dotruonghan** Highly Voted 3 years, 6 months ago

Answer: D,E,F

Explanation:
<https://docs.oracle.com/en/database/oracle/oracle-database/19/netag/introducing-oracle-net-services.html#GUID-6B1F0584-887D-4D2A-B982-40ACF15DF985>
• Oracle Enterprise Manager Cloud Control combines configuration functionality across multiple file systems, along with listener administrative control to provide an integrated environment for configuring and managing Oracle Net Services.
• Oracle Net Manager provides configuration functionality for an Oracle home on a local client or server host.
upvoted 8 times



  **nautil2** Most Recent 1 month, 4 weeks ago

Selected Answer: DEF
A - false; netca can be run independently on the Oracle Installer
B - false; Oracle Net Manager can be used to configure listeners on the local server or server host
C - false; there need not exist the listener.ora file before lsnrctl utility is started
D - true; see answer B
E - true; EM Cloud Control: All the systems and services including enterprise application systems, databases, hosts, middleware application servers, listeners, and so on, are easily managed from one central location.
F - true; see answer E

Source:
B - Oracle database documentation: "Oracle Net Manager supports configuration of connect descriptors in a local tnsnames.ora file or directory server."
E - Oracle Enterprise Manager Cloud Control 13.5 documentation
upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: DEF
Provided answer are wrong.
ref.
<https://docs.oracle.com/en/database/oracle/oracle-database/21/netag/managing-oracle-net-services.html#GUID-1F145906-24CC-4A63-86FD-31776C022209>
upvoted 1 times

  **mikimiki01** 8 months, 4 weeks ago

Answer: DEF
upvoted 1 times

  **Kuraudio** 1 year, 1 month ago

Selected Answer: DEF
As comments below
upvoted 1 times

  **Borislone** 2 years, 6 months ago

A is correct

upvoted 1 times

店长微信：hj feng128

  **Marcello86CT** 1 year, 7 months ago

Oracle Net Configuration Assistant is a post installation tool so i think that A is Wrong
upvoted 1 times



Which two statements are true regarding the UNION and UNION ALL operators? (Choose two.)



- A. Duplicates are eliminated automatically by the UNION ALL operator
- B. The number of columns selected in each SELECT statement must be identical
- C. The names of columns selected in each SELECT statement must be identical
- D. The output is sorted by the UNION ALL operator
- E. NULLS are not ignored during duplicate checking



Correct Answer: BE



Community vote distribution

BE (100%)



  **you1234** Highly Voted 3 years, 6 months ago
B & E is correct answer
upvoted 20 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times



  **dotruonghan** Highly Voted 3 years, 6 months ago
BE is correct answer
upvoted 7 times



  **nautil2** Most Recent 1 month, 4 weeks ago
Selected Answer: BE
A - false
B - true
C - false
D - false
E - true



Sources:
Oracle Documentation: Database > Oracle > Oracle Database > Release 19 > SQL Language Reference > 9 SQL Queries and Subqueries > a) The UNION [ALL], INTERSECT, MINUS Operators, b) Sorting Query Results
A The UNION ALL operator does not eliminate duplicate selected rows
B,C The corresponding expressions in the select lists of the component queries of a compound query must match in number and must be in the same data type group (such as numeric or character).
D For compound queries containing set operators UNION, INTERSECT, MINUS, or UNION ALL, the ORDER BY clause must specify positions or aliases rather than explicit expressions. Also, the ORDER BY clause can appear only in the last component query. The ORDER BY clause orders all rows returned by the entire compound query.
upvoted 1 times

  **babyjaan** 10 months, 2 weeks ago
Selected Answer: BE
B & E CORRECT ANSWER
upvoted 1 times

  **ryuah** 2 years ago
B,E is correct
upvoted 2 times

  **Guhborges** 2 years, 1 month ago
B - E is correct!
upvoted 1 times

  **Oye** 3 years, 2 months ago
B and E
upvoted 1 times

  **avanand** 3 years, 5 months ago
B and E guys.
upvoted 1 times

Which two statements are true about the Automatic Diagnostic Repository (ADR)? (Choose two.)

- A. The ADR base defaults to \$ORACLE_HOME/rdbms/admin if neither DIAGNOSTIC_DEST nor ORACLE_BASE is set
- B. The ADR base defaults to \$ORACLE_HOME/dbs if the DIAGNOSTIC_DEST parameter and the ORACLE_BASE environment variable are not set
- C. It supports diagnostics for Automatic Storage Management (ASM)
- D. It supports diagnostics for Oracle Clusterware
- E. It is held inside an Oracle database schema

Correct Answer: AE

Reference:

https://docs.oracle.com/cd/E11882_01/install.112/e27508/admin.htm



Community vote distribution



CD (100%)

  **elcaso34** Highly Voted 3 years, 7 months ago
I think C D.

The database, Oracle Automatic Storage Management (Oracle ASM), the listener, Oracle Clusterware, and other Oracle products or components store all diagnostic data in the ADR.

<https://docs.oracle.com/database/121/ADMIN/diag.htm#ADMIN11260>
upvoted 20 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times



  **Sharif1** Highly Voted 3 years, 6 months ago
yes elcaso34 is correct, options C and D are correct.
upvoted 11 times

  **nautil2** Most Recent 1 month, 4 weeks ago

Selected Answer: CD

A - false; default value for DIAGNOSTIC_DEST: If \$ORACLE_BASE is not set, then derived from ORACLE_BASE as set by the Oracle Universal Installer. If ORACLE_BASE is not set, then \$ORACLE_HOME/rdbms/log is used.
B - false; see answer A
C - true; Beginning with Oracle Database 11g and later releases, Oracle Automatic Storage Management (Oracle ASM), and other Oracle Database products or components store all diagnostic data in the ADR.
D - true; Review this content to understand clusterware-specific aspects of how Oracle Clusterware uses ADR.
E - false; data are not held inside an Oracle database schema

Sources:
A, B - Oracle documentation: 1.108 DIAGNOSTIC_DEST
C - Oracle documentation: 21.2 Definitions for Oracle Database ADRC
D - Oracle documentation: J Troubleshooting Oracle Clusterware
upvoted 1 times

  **Oracle2020** 4 months, 4 weeks ago
The correct answers are: A, E

If this parameter is omitted or left null, then the database sets DIAGNOSTIC_DEST upon startup as follows:

If environment variable ORACLE_BASE is set, then DIAGNOSTIC_DEST is set to the directory designated by ORACLE_BASE.

If environment variable ORACLE_BASE is not set, then DIAGNOSTIC_DEST is set to ORACLE_HOME/log

Following the url indicated by the correct answer of exam topic. We understand this is the same for Oracle 19
upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: CD

By exclusion:
A and B are wrong because the default path is ORACLE_HOME/log for this case:
The ADR root directory is known as ADR base. Its location is set by the DIAGNOSTIC_DEST initialization parameter. If this parameter is omitted or



left null, the database sets DIAGNOSTIC_DEST upon startup as follows.

- If environment variable ORACLE_BASE is set, DIAGNOSTIC_DEST is set to the directory designated by ORACLE_BASE.
- If environment variable ORACLE_BASE is not set, DIAGNOSTIC_DEST is set to ORACLE_HOME/log.

E is wrong because:
The Automatic Diagnostic Repository (ADR) is a directory structure that is stored outside of the database. It is therefore available for problem diagnosis when the database is down.
ref.



<https://docs.oracle.com/en/database/oracle/oracle-database/12.2/admin/managing-diagnostic-data.html#GUID-C78CCAD4-732B-4B14-9C6B-D47AC2A99E30>

upvoted 1 times



  **RaNik69** 8 months, 1 week ago

Selected Answer: CD

answer is CD
upvoted 1 times

  **mikimiki01** 8 months, 4 weeks ago

answer is C and D
upvoted 1 times

  **brolem** 2 years, 4 months ago

A, B - Wrong: The ADR root directory is known as the ADR base. Its location is set by the DIAGNOSTIC_DEST initialization parameter. If this parameter is omitted or left null, the database sets DIAGNOSTIC_DEST upon startup as follows: If environment variable ORACLE_BASE is set, DIAGNOSTIC_DEST is set to SORACLE_BASE. If environment variable ORACLE_BASE is not set, DIAGNOSTIC_DEST is set to SORACLE_HOME/log.

C - Correct: For example, in an Oracle Real Application Clusters environment with Oracle ASM, each database instance, Oracle ASM instance, and listener has an ADR home.

D - Correct: Oracle Clusterware uses ADR and has its own Oracle home and Oracle base.

E - Wrong: The Automatic Diagnostic Repository (ADR) is a directory structure that is stored outside of the database. It is therefore available for problem diagnosis when the database is down.

From <<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/diagnosing-and-resolving-problems.html#GUID-951A06EE-DDF7-4C2A-B0BB-B24418BB2E33>>



upvoted 5 times

  **elbelgounetos** 2 years, 10 months ago

The ADR root directory is known as the ADR base. Its location is set by the DIAGNOSTIC_DEST initialization parameter. If this parameter is omitted or left null, the database sets DIAGNOSTIC_DEST upon startup as follows:

If environment variable ORACLE_BASE is set, DIAGNOSTIC_DEST is set to SORACLE_BASE.
If environment variable ORACLE_BASE is not set, DIAGNOSTIC_DEST is set to SORACLE_HOME/log.

upvoted 3 times

  **danito** 3 years, 7 months ago

I think A D
upvoted 1 times

You want to apply the principle of Least Privilege in all your live databases.

One of your requirements is to revoke unnecessary privileges from all users who have them using Privilege Analysis.

Which three types of analyses can be done using the DBMS_PRIVILEGE_CAPTURE package? (Choose three.)


- A. analysis of all privileges used by all users including administrative users in the database
- B. analysis of all privileges used by all users but excluding administrative users in the database
- C. analysis of privileges that a user has on their own schema objects that they did not use
- D. analysis of privileges that a user has on their own schema objects that they did use
- E. analysis of privileges granted directly to a role that are then used by a user who has been granted that role
- F. analysis of privileges granted indirectly to a role that are then used by a user who has been granted that role

Correct Answer: ACF

Reference:

https://docs.oracle.com/database/121/ARPLS/d_priv_prof.htm#ARPLS74328

Community vote distribution



ama **Highly Voted** 3 years, 5 months ago

i would say B, E, F are correct!

upvoted 19 times

Ekos 3 years, 1 month ago

i agree

upvoted 1 times

Yema 2 years, 11 months ago

I agree also

<http://dbaparadise.com/2019/02/7-steps-to-setup-privilege-analysis/>

upvoted 1 times

Fw1z 2 years, 7 months ago

b is wrong only SYS!!!!!!

upvoted 2 times

guretto **Most Recent** 3 months, 1 week ago

SYS, SYSTEM, SYSBACKUP, SYSDG, SYSKM, SYSRAC, SYSMAN, and DBSNMP are administrative users as reported in 19c Database Administrator's Guide, Par. 7.5 (Predefined User Accounts):

Furthermore (19c Database PL/SQL Packages and Types Reference)

given that DBMS_PRIVILEGE_CAPTURE.CREATE_CAPTURE "Analyzes all privilege use, except privileges used by the SYS user.", it follows that A cannot be right since SYS is included in "administrative users"

The DBMS_PRIVILEGE_CAPTURE works on database, role and context level. Therefore B is right and C, D are not.

E, F are right as "When using role-based analysis for the CREATE_CAPTURE procedure, privilege use is analyzed even if the privilege is indirectly granted to the specified role."

upvoted 1 times

zouve 5 months, 1 week ago

Selected Answer: BEF

<https://www.databasejournal.com/oracle/oracle-database-12c-new-privilege-management-features/>

upvoted 1 times

auwia 6 months, 2 weeks ago

Selected Answer: CDF

Provided answer are wrong.

C, D, and F the correct answer, everything described here:

https://docs.oracle.com/en/database/oracle/oracle-database/19/arpls/DBMS_PRIVILEGE_CAPTURE.html

upvoted 1 times

auwia 6 months, 2 weeks ago

A. analysis of all privileges used by all users including administrative users in the database --> False, SYS is excluded

B. analysis of all privileges used by all users but excluding administrative users in the database --> False, SYSTEM is included.

C. analysis of privileges that a user has on their own schema objects that they did not use



D. analysis of privileges that a user has on their own schema objects that they did use
E. analysis of privileges granted directly to a role that are then used by a user who has been granted that role
F. analysis of privileges granted indirectly to a role that are then used by a user who has been granted that role

upvoted 1 times

  **auwia** 6 months, 2 weeks ago


F a TRUE

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

I would say C, D, and F ; but D makes no sense, but it's possible to do (as the question is asking).

upvoted 1 times

  **auwia** 6 months, 2 weeks ago


D makes sense! :-)

https://docs.oracle.com/en/database/oracle/oracle-database/19/arpls/DBMS_PRIVILEGE_CAPTURE.html

131.1 DBMS_PRIVILEGE_CAPTURE Overview

Database privilege analysis enables you to create a policy that records the usage of system and object privileges that have been granted to users. You then can determine the privileges that your users are using and not using. From there, you can revoke any unused privileges, thereby reducing the number of excess privilege grants for users.

upvoted 1 times

  **auwia** 6 months, 2 weeks ago



F confirmed by this:

Usage Notes

When using role-based analysis for the CREATE_CAPTURE procedure, privilege use is analyzed even if the privilege is indirectly granted to the specified role.

For example, say role R2 contains role R1, and R1 contains privilege P1. If the privilege policy includes only role R2, any use of the P1 privilege is still analyzed, as privilege P1 is an indirect part of role R2.

upvoted 1 times

  **Darkseid1231** 7 months, 1 week ago

A C E is the correct answer



upvoted 2 times

  **Kuraudio** 1 year, 1 month ago

A says: " including administrative users". The definition of administrative users is here: <https://docs.oracle.com/database/121/ADMQS/GUID-CF1CD853-AF15-41EC-BC80-61918C73FDB5.htm#ADMQS12003>

So as SYS is excluded from the analysis, I can say A is wrong.


upvoted 1 times

  **XhostSI** 1 year, 11 months ago

A. An undo segment may be used by multiple transactions simultaneously - True
B. Undo segments can wrap around to the first extent when a transaction fills the last extent of the undo segment - True (undo segment act as a circular buffer)
C. Undo segments have a minimum of three extents - False (each segment has a minimum of two extents)
D. Undo segments can extend when a transaction fills the last extent of the undo segment - True (Undo tablespace are locally managed with auto extent allocation)
E. A single transaction may use multiple undo segments simultaneously - False
F. Undo segments must be stored in a BIGFILE tablespace - False
G. Undo segments must be stored in a SMALLFILE tablespace - False (DBCA automatically creates a smallfile undo tablespesce, but it can be configured as bigfile table space as well)

So correct answers are: A, B, D

upvoted 1 times

  **hadiwuu** 8 months, 1 week ago

HOLD UP!

upvoted 1 times

  **Guhborges** 2 years, 1 month ago

A - E - F CORRECT!!!!

A - Right - We can analyze the administrative user, the oracle just excludes the SYS user

B - Wrong, oracle does not delete administrative user, like question A, only SYS user

C - Wrong - We cannot analyze the permission of the scheme itself

D - Wrong - We cannot analyze the privilege of our schema

E - Right - we can analyze the permission that was given directly.

F - Right - we can analyze the permission that was given indirectly.

upvoted 4 times

  **Borislone** 2 years, 6 months ago

sorry BEF is definitely correct

upvoted 1 times

🗨️ **Borislone** 2 years, 6 months ago

the correct answer is AEF

upvoted 2 times

🗨️ **nobody347** 2 years, 7 months ago

Anybody have answer for this question after exam?

upvoted 1 times

🗨️ **ama** 3 years, 5 months ago

we can analyze all users except SYS, i dunno but answer B saying all users except administrative Users!

You cannot analyze the privileges of the SYS user.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/dbseg/performing-privilege-analysis-find-privilege-use.html#GUID-AD683DDB-D345-4702-B1F4-37B79F276B66>

upvoted 3 times

🗨️ **KtNow** 3 years ago

administrative users - SYS, SYSTEM.

DBMS_PRIVILEGE_CAPTURE Analyzes all privilege use, except privileges used by the SYS (ONLY SYS!!!!)

DEF sounds right

upvoted 3 times

🗨️ **Guhborges** 2 years, 1 month ago

You are wrong bout the user SYSTEM, we can't analyze only the user SYS

```
SQL> BEGIN
DBMS_PRIVILEGE_CAPTURE.CREATE_CAPTURE(
name =>
'validando_permissao',
description => 'Analyzes system
privilege use',
type =>
DBMS_PRIVILEGE_CAPTURE.G_CONTEXT,
condition =>
'SYS_CONTEXT("USERENV",
"SESSION_USER")="system");
END;
/ 2 3 4 5 6 7 8 9 10 11 12 13
```

PL/SQL procedure successfully completed.

upvoted 1 times

🗨️ **leozaan94** 2 years, 5 months ago

DEF sounds right also for me.

A and B: are false for administrative users.

C: is false cause this procedure doesn't analyze privileges not used.

D: is true cause you can set the type as G_DATABASE --> : "Captures all privilege use in the database, except privileges used by the SYS user." So including privileges in own schema.

E and F are true cause it analyzes privileges directly and indirectly

https://docs.oracle.com/en/database/oracle/oracle-database/21/arpls/DBMS_PRIVILEGE_CAPTURE.html#GUID-8E855D12-AF7B-4995-878B-B954FA92FF08

upvoted 2 times

🗨️ **ama** 3 years, 5 months ago

C, D are wrong!

Privileges in own Schema can not be analyzed!

upvoted 1 times

🗨️ **NowOrNever** 3 years, 6 months ago

DEF - i think

Database privilege analysis enables you to create a policy that records the usage of system and object privileges that have been granted to users. You then can determine the privileges that your users are using and not using. From there, you can revoke any unused privileges, thereby reducing the number of excess privilege grants for users.

C- you can analyze only used privileges, but then determine the unsued privileges

F - When using role-based analysis for the CREATE_CAPTURE procedure, privilege use is analyzed even if the privilege is indirectly granted to the specified role.

upvoted 2 times

🗨️ **Sharif1** 3 years, 6 months ago

The option B seems to be more close. The answer is given in the same link.

Analyzes all privilege use, except privileges used by the SYS user.

So the options which are correct seems to be B C F.

upvoted 2 times

Which three statements are true about undo segments and the use of undo by transactions in an Oracle database instance? (Choose three.)

- A. An undo segment may be used by multiple transactions simultaneously
- B. Undo segments can wrap around to the first extent when a transaction fills the last extent of the undo segment
- C. Undo segments have a minimum of three extents
- D. Undo segments can extend when a transaction fills the last extent of the undo segment
- E. A single transaction may use multiple undo segments simultaneously
- F. Undo segments must be stored in a BIGFILE tablespace
- G. Undo segments must be stored in a SMALLFILE tablespace

Correct Answer: ADG

Community vote distribution

ABD (75%)

ADE (25%)

  **adoptc94** Highly Voted 3 years, 5 months ago

F and G are both false!

While the default is smallfile you can also create a bigfile undo tablespace which means both smallfile and bigfile are possible for undo tablespaces

E is also false as each transaction is assigned to only one undo segment, but an undo segment can service more than one transaction at a time -> A is right



Which leaves us with B,C,D as possible answers.

B and D are true -> taken from the Oracle Administrator Workshop Studyguide: Transactions and Undo Data (14 - 5): If an extent fills up and more space is needed, the transaction acquires that space from the next extent in the segment. After all extents have been consumed, the transaction either wraps around back into the first extent or requests a new extent to be allocated to the undo segment (supporting answers B and D)

The right answers for this question are: A,B,D

I would also say that C is false as I am pretty sure that UNDO segments have a minimum of 2 extents (and not 3 extents) to support the circular use of undo segments and this can be achieved by a minimum of 2 extents

upvoted 17 times

  **tlcecertified** 3 years, 2 months ago

C is correct, in 19c, any simple transaction (e.g. updating only 1 field of 1 record) would generate 3 extents, you can observe the number of extent_id in dba_undo_extents view to justify this.


my test script

```
set transaction name 'testforanswer';
update dual set dummy = 'y';
```

```
select
r.name "RBS name",
t.used_ublk "Undo blocks",
t.used_urec "Undo recs"
from
v$session s,
v$transaction t,
v$rollname r
where
t.addr = s.taddr
and
r.usn = t.xidusn;
```

```
SELECT SEGMENT_NAME, EXTENT_ID FROM DBA_UNDO_EXTENTS WHERE SEGMENT_NAME='<what you found from above r.name column>'
rollback;
```

upvoted 2 times

  **ogdru** 2 years, 10 months ago

```
SQL> create undo tablespace undotbs02;
```

Tablespace created.

```
SQL> select distinct tablespace_name, extent_id from dba_undo_extents order by 1,2;
```


UNDOTBS02 0
UNDOTBS02 1

upvoted 4 times

 **avanand** Highly Voted 3 years, 5 months ago

I think it should be A D G

upvoted 5 times

 **nautil2** Most Recent 1 month, 4 weeks ago

Selected Answer: ABD

- A - true; query `SELECT addr,xidusn,start_time FROM v$transaction ORDER BY xidusn`; can return multiple ADDR (address of the transaction state object) for a single XIDUSN (undo segment number)
- B - true; when an extent (within a segment) is filled up, a transaction either wraps around back into the first extent, or requests a new extent to be allocated
- C - false; an undo segment can be formed of a single (one) extent. One extent is the total minimum.
- D - true; see answer B
- E - false; when a transaction starts, it is assigned to only one undo segment
- F - false; a command `CREATE SMALLFILE UNFO TABLESPACE tablespace_name`; is valid
- G - false; a command `CREATE BIGFILE UNFO TABLESPACE tablespace_name`; is valid

Sources:

F, G - Oracle Database 19c Documentation: SQL Language Reference > 15 SQL Statements: CREATE SEQUENCE to DROP CLUSTER > CREATE TABLESPACE

upvoted 1 times

 **musafir** 6 months ago

Selected Answer: ABD

- A. An undo segment may be used by multiple transactions simultaneously
Correct - "Multiple active transactions can write concurrently to the same undo segment or to different segments."
- B. Undo segments can wrap around to the first extent when a transaction fills the last extend of the undo segment
Correct - If the 1st extent is full then the next extent will be used and overwritten if it contains unactive undo.
- C. Undo segments have a minimum of three extents
Wrong - No documentation about minimum number of extents.
- D. Undo segments can extend when a transaction fills the last extent of the undo segment
Correct, AUTOEXTEND Option with Automatic UNDO management allows extension of UNDO.
- E. A single transaction may use multiple undo segments simultaneously
Wrong - Documentation says a transaction may use multiple extents simultaneously not segments.
<https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html#GUID-BB42D7D8-2C93-4F5B-B0C1-6361AEEF4B37>
- F. Undo segments must be stored in a BIGFILE tablespace
- G. Undo segments must be stored in a SMALLFILE tablespace
Wrong - Can be either Bigfile or Smallfile, smallfile is the default.

A B D are correct I think

upvoted 1 times

 **auwia** 6 months, 2 weeks ago

Selected Answer: ADE

- A. An undo segment may be used by multiple transactions simultaneously --> True
- B. Undo segments can wrap around to the first extent when a transaction fills the last extend of the undo segment. --> False, an undo that represents active transactions will not be overwritten by a wrap. Of course it will after the transition finished and undo retention timeout expired.
- C. Undo segments have a minimum of three extents. --> False, There is no link about that, probably is false. For people trying to demonstrate by SQL plus, I recommend to simulate a case when only 1 or 2 segment are free for extension... I'm sure oracle will extend the necessary and put data safe.
- D. Undo segments can extend when a transaction fills the last extent of the undo segment. --> True
- E. A single transaction may use multiple undo segments simultaneously. --> True
- F. Undo segments must be stored in a BIGFILE tablespace. --> clearly False
- G. Undo segments must be stored in a SMALLFILE tablespace. --> Clearly False


upvoted 1 times

 **Kuraudio** 1 year, 1 month ago

Selected Answer: ABD



ABD is correct



upvoted 1 times

 **XhostSI** 1 year, 11 months ago



- A. An undo segment may be used by multiple transactions simultaneously - True
- B. Undo segments can wrap around to the first extent when a transaction fills the last extent of the undo segment - True (undo segment act as a circular buffer)
- C. Undo segments have a minimum of three extents - False (each segment has a minimum of two extents)
- D. Undo segments can extend when a transaction fills the last extent of the undo segment - True (Undo tablespace are locally managed with auto



extent allocation)
E. A single transaction may use multiple undo segments simultaneously - False
F. Undo segments must be stored in a BIGFILE tablespace - False
G. Undo segments must be stored in a SMALLFILE tablespace - False (DBCA automatically creates a smallfile undo tablespesce, but it can be configured as bigfile table space as well)
So correct answers are: A, B, D
upvoted 2 times

  **Borislone** 2 years, 6 months ago
A is correct for parallel DML or DDL transaction multiple undo segments can be used simultaneously
upvoted 1 times



  **Schmeili** 3 years, 5 months ago
seems to be A, D, G

A because of:
Multiple active transactions can write concurrently to the same undo segment or to different segments. For example, transactions T1 and T2 can both write to undo segment U1, or T1 can write to U1 while T2 writes to undo segment U2.
upvoted 3 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times



  **NowOrNever** 3 years, 5 months ago
Correct Answer: CDG
Wrong Answer: ABEF
I could not complete verify C, but in my DB it seems to be the default size on creation, i have not tried to minimize them



UNDO segments are used in a circular fashion. If there are ten UNDO segments, the first transaction uses the first UNDO segment, the second transaction uses the second UNDO segment, and the eleventh transaction circles back to the first UNDO segment. Once a transaction starts in an UNDO segment, it must finish in that segment.
upvoted 1 times

  **ama** 3 years, 5 months ago
are you really sure that C is correct!!!!???

look below

```
SQL> select segment_name,count(EXTENT_ID) from dba_undo_extents group by segment_name;
SEGMENT_NAME COUNT(EXTENT_ID)
-----
_SYSSMU11_3443795589$ 24
_SYSSMU6_1437887270$ 22
_SYSSMU135_1421094332$ 2
_SYSSMU161_3597468454$ 2
upvoted 1 times
```

  **NowOrNever** 3 years, 4 months ago
you are right C is wrong
A is the correct one
<https://coggle.it/diagram/WuHx64IXs-uDQScg/t/managing-undo-data>
<https://www.coursehero.com/file/p5uvnu8/10-6-Transactions-and-Undo-Data-Each-transaction-is-assigned-to-only-one-undo/>
it's hidden in the course material
So i also go with ADG
upvoted 1 times

  **saad3577** 1 year, 3 months ago
G is wrong
you can ether Changing the Datafiles for an Undo Tablespace to a Fixed Size in 12c

<https://docs.oracle.com/database/121/ADMQS/GUID-D42E755D-4B5E-48E5-8D73-8356907DD86C.htm#ADMQS12465>
upvoted 1 times

Which two statements are true about Enterprise Manager Database Express? (Choose two.)

- A. It is available only when the database is open
- B. It can be used to perform database recovery
- C. The same port number can be used for Database Express configurations for databases on different hosts
- D. It can be used to switch a database into ARCHIVELOGMODE
- E. The same port number can be used for multiple Database Express configurations for multiple databases on the same host

Correct Answer: AC



Community vote distribution

AC (100%)

  **dbaclaudio** Highly Voted 3 years ago



Correct Answers:
A and C
A - Correct -> to startup the enterprise manager database express, the database have to be in open state
B - Incorrect -> to perform a full recovery, the database has to be in mount state, and EM express not work if database is not in open state
C - Correct -> since they are in different hosts, you can use the same port for then (default 5500)
D - Incorrect -> to switch a database to ARCHIVELOGMODE the database need to be in mount state (like B answer)
E - Incorrect -> the same port cannot be used in the same time for different services in the same host (here a basic of tcp/ip protocol)
If need more reference:
EM Express (Chapter 3.3)
<https://docs.oracle.com/en/database/oracle/oracle-database/19/admqsl/2-day-dba.pdf>

ARCHIVELOGMODE (Chapter 12.3.2)
<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/database-administrators-guide.pdf>
upvoted 12 times



  **avanand** Highly Voted 3 years, 5 months ago

Whats wrong with A and C.



And if A is correct how can D be correct?
upvoted 8 times

  **adoptc94** 3 years, 5 months ago



The correct answers are A and C
upvoted 9 times

  **Ekos** 3 years, 1 month ago

i agree
upvoted 1 times

  **auwia** Most Recent 6 months, 2 weeks ago

Selected Answer: AC
Provided answer are corrects.
upvoted 1 times



  **ama** 3 years, 6 months ago

A, B are correct!



EM Express is available only when the database is open. This means that EM Express cannot be used to start up the database. Other operations that require that the database change state, such as enable or disable ARCHIVELOG mode, are also not available in EM Express.
upvoted 1 times

  **ama** 3 years, 6 months ago

Sorry, A and C are corrects, because Database Recovery requires sometimes changing the DB state
upvoted 13 times

  **danito** 3 years, 7 months ago


A E <https://www.oracle.com/technetwork/database/manageability/emx-intro-1965965.html>
upvoted 1 times

  **elcaso34** 3 years, 7 months ago

OEM Database Express is designed to manage only one database, and is intended for database administrators. When you invoke Database Express, you will be prompted for a

username and password to connect to the database. You should provide a user account that has administrative privileges.

upvoted 1 times

 **ama** 3 years, 6 months ago

E is wrong...

Use the PL/SQL procedure DBMS_XDB_CONFIG.SETHTTPSPORT to set the HTTPS port for EM Express for the database to a port that is not in use. This will update the HTTPS port in the xdbconfig.xml file in the Oracle XML DB Repository. You must connect as SYS / AS SYSDBA to run the procedure.

so you Need to use a port that is not in use

upvoted 1 times

Table EMPLOYEES contains columns including EMPLOYEE_ID, JOB_ID and SALARY.
Only the EMPLOYEE_ID column is indexed.
Rows exist for employees 100 and 200.
Examine this statement:

```
UPDATE employees
  SET (job_id, salary) =
    (SELECT job_id, salary
     FROM employees
     WHERE employee_id = 200)
 WHERE employee_id = 100;
```



Which two statements are true? (Choose two.)



- A. Employee 100 will have SALARY set to the same value as the SALARY of employee 200
- B. Employee 200 will have JOB_ID set to the same value as the JOB_ID of employee 100
- C. Employee 200 will have SALARY set to the same value as the SALARY of employee 100
- D. Employee 100 will have JOB_ID set to the same value as the JOB_ID of employee 200
- E. Employees 100 and 200 will have the same JOB_ID as before the update command
- F. Employees 100 and 200 will have the same SALARY as before the update command

Correct Answer: AB



Community vote distribution



AD (100%)



  **danito** Highly Voted 3 years, 7 months ago
A D the correct options
upvoted 28 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times



  **you1234** Highly Voted 3 years, 6 months ago
A & D is correct
upvoted 13 times

  **auwia** Most Recent 6 months, 3 weeks ago
Selected Answer: AD
Provided answer are wrong.
The base! :-) Part IV
upvoted 1 times

  **fthusa** 8 months ago
AD are correct
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago
AD, Employee 100 will have Employee 200 data ...
upvoted 1 times

  **ryuah** 2 years ago
A,D is correct
upvoted 1 times

  **Guhborges** 2 years, 1 month ago
Selected Answer: AD
A D the correct!
upvoted 1 times

Which two are true about a SQL statement using SET operators such as UNION? (Choose two.)

- A. The data type group of each column returned by the second query must match the data type of the corresponding column returned by the first query.
- B. The number, but not names, of columns must be identical for all SELECT statements in the query.
- C. The data type of each column returned by the second query must exactly match the data type of the corresponding column returned by the first query.
- D. The names and number of columns must be identical for all SELECT statements in the query.
- E. The data type of each column returned by the second query must be implicitly convertible to the data type of the corresponding column returned by the first query.

Correct Answer: AC

Community vote distribution

AB (100%)

- danito

Highly Voted

3 years, 8 months ago

A B are the correct options

upvoted 13 times
- Ekos

3 years, 1 month ago

i agree

upvoted 1 times
- elbelgounetos

2 years, 10 months ago

reference <https://docs.oracle.com/en/database/oracle/oracle-database/21/sqlrf/The-UNION-ALL-INTERSECT-MINUS-Operators.html#GUID-B64FE747-586E-4513-945F-80CB197125EE>

upvoted 1 times
- Rivaldo11

Highly Voted

3 years, 6 months ago

I would also say A and B - "...Tables must have the same number of columns presented in the same order. In addition, the data type of the corresponding column must be in the same data type group such as number or character..."

upvoted 6 times
- Hell3n

Most Recent

4 months, 2 weeks ago

why not E?

upvoted 1 times
- Oracle2020

4 months, 4 weeks ago

I think B, C

upvoted 1 times
- mamadu

6 months, 2 weeks ago

I was confused with C and A since A. mentions "data type group".
i can confirm that a CHAR can be matched with VARCHAR2.
ans is A B

upvoted 2 times
- auwia

6 months, 3 weeks ago

Selected Answer: AB

A and B are correct

upvoted 2 times
- fthusa

8 months ago

AB correct

upvoted 1 times
- LeandroHPN

1 year, 6 months ago

Selected Answer: AB

A , B is correct.

upvoted 1 times
- ryuah

2 years ago


A,B is correct

upvoted 1 times

  **avanand** 3 years, 5 months ago

A and B are correct.
For SET operator, Data type group should be matched.
For example, a NUMBER column in first query can be UNIONed with INTERGER column in second query.

upvoted 3 times

  **you1234** 3 years, 6 months ago

C & E is correct answer
upvoted 1 times

Which three statements are true about time zones, date data types, and timestamp data types in an Oracle database? (Choose three.)

- A. The CURRENT_TIMESTAMP function returns data without time zone information
- B. A TIMESTAMP WITH LOCAL TIMEZONE data type column is stored in the database using the time zone of the session that inserted the row
- C. A TIMESTAMP data type column contains information about year, month, and day
- D. The DBTIMEZONE function can return an offset from Universal Coordinated Time (UTC)
- E. The SESSIONTIMEZONE function can return an offset from Universal Coordinated Time (UTC)

Correct Answer: ADE

Community vote distribution

BDE (100%)

 **SimoneF** Highly Voted 3 years ago


C D E:
"TIMESTAMP WITH LOCAL TIME ZONE is another variant of TIMESTAMP that is sensitive to time zone information. It differs from TIMESTAMP WITH TIME ZONE in that data stored in the database is normalized to the database time zone, and the time zone information is not stored as part of the column data. When a user retrieves the data, Oracle returns it in the user's local session time zone."


So, the data is stored with the DB Timezone, even if it's retrieved with the session timezone.
A is clearly wrong by example, so the rest are true. With D and E we can also have an easy sample, while for C, which is apparently the most ambiguous, I agree with whom was saying that the TIMESTAMP data type contains indeed informations about day, month and year, and in the answer is never stated that they are the only ones present. So as it's put, it's definitely right.


<https://docs.oracle.com/database/121/NLSPG/ch4datetime.htm#GUID-CD2954CE-45E2-4938-A599-CCB96879510F>
upvoted 9 times


 **auwia** Most Recent 6 months, 3 weeks ago

Selected Answer: BDE
Provided answers are wrong, by exclusion:
A. The CURRENT_TIMESTAMP function returns data without time zone information -->False, it is returned.
B. A TIMESTAMP WITH LOCAL TIMEZONE data type column is stored in the database using the time zone of the session that inserted the row
C. A TIMESTAMP data type column contains information about year, month, and day --> False, incomplete sentence.
D. The DBTIMEZONE function can return an offset from Universal Coordinated Time (UTC)
E. The SESSIONTIMEZONE function can return an offset from Universal Coordinated Time (UTC)
upvoted 1 times

 **yarsalan** 1 year, 12 months ago
C is wrong. The TIMESTAMP datatype is an extension of the DATE datatype. It stores a year, month, day, hour, minute, and second values. It also stores fractional seconds, which are not stored by the DATE datatype.
https://docs.oracle.com/cd/B19306_01/server.102/b14225/ch4datetime.htm#:~:text=The%20TIMESTAMP%20datatype%20is%20an,stored%20by%20the%20DATE%20datatype.
upvoted 1 times

 **MCzombie** 3 years ago
(CDE) IAM SURE!!
upvoted 3 times

 **heaaa** 3 years, 2 months ago
B IS wrong.
TIMESTAMP WITH LOCAL TIME ZONE Data Type
TIMESTAMP WITH LOCAL TIME ZONE is another variant of TIMESTAMP. It differs from TIMESTAMP WITH TIME ZONE as follows: data stored in the database is normalized to the database time zone, and the time zone offset is not stored as part of the column data.
so session time zone must transform .
upvoted 2 times

 **Chansi** 3 years, 5 months ago
B, is only right if you supply the offset . Here is the test . Still B ? any thoughts?

alter session set time_zone='-7:00';
SQL> desc test3
Name Null? Type

ID NUMBER
TIMES1 TIMESTAMP(6) WITH LOCAL TIME
ZONE

insert into test3 values (8, TIMESTAMP '2020-07-01 2:00:00 -8:00');

insert into test3 values (9,TIMESTAMP '2020-07-01 2:00:00')

店长微信：hj feng128



SQL> select id,times1 from test3 where id in (8,9);



ID

TIMES1



8
01-JUL-20 03.00.00.000000 AM

9
01-JUL-20 02.00.00.000000 AM
upvoted 2 times



  **NowOrNever** 3 years, 4 months ago
B - ... column is stored in the database ***using*** the time zone ...
in the first insert you are overwriting the offset
the second insert uses your session offset
upvoted 2 times



  **ama** 3 years, 5 months ago
i still believe B, D, E are correct..

C is wrong because The TIMESTAMP data type allows you to store date and time data including year, month, day, hour, minute and second.
upvoted 4 times

  **adoptc94** 3 years, 4 months ago
I don't see any issue with answer C!
While it does also store hour, minute and seconds the answer doesn't state that it ONLY stores year, month and day - it does contain it which is true.



I'd go with C,D,E
upvoted 8 times

  **auwia** 6 months, 3 weeks ago
You cannot exclude B, sentence C is incomplete, so it's false.
upvoted 1 times

  **ama** 3 years, 5 months ago
D is correct > SQL> select DBTIMEZONE from dual;

DBTIME

+01:00
upvoted 2 times

  **NowOrNever** 3 years, 6 months ago
Correct Answer: BCD
Wrong Answer: AE

<https://docs.oracle.com/database/121/NLSPG/ch4datetime.htm#NLSPG004>



CURRENT_TIMESTAMP
Returns the current date and time in the session time zone as a TIMESTAMP WITH TIME ZONE value



select CURRENT_TIMESTAMP from dual;
-- CURRENT_TIMESTAMP
-- 28-JUL-20 09.56.37.618148000 AM AMERICA/LOS_ANGELES

DBTIMEZONE
Returns the value of the database time zone. The value is a time zone offset or a time zone region name

SESSIONTIMEZONE
Returns the value of the current session's time zone
SELECT sessiontimezone FROM DUAL;
-- SESSIONTIMEZONE
-- America/Los_Angeles

B is driving me crazy, but AE are wrong, so ...
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago
also agree BCD
E returns the timezone name NOT the offset.
upvoted 1 times

  **ama** 3 years, 5 months ago

but in one of my database the sessiontimezone Returns this.

店长微信：hj feng128



SQL> select SESSIONTIMEZONE from dual;

SESSIONTIMEZONE

+02:00

SQL>



what you think?
upvoted 1 times



  **ama** 3 years, 6 months ago
A is wrong

SQL> select CURRENT_TIMESTAMP from dual;



CURRENT_TIMESTAMP



23.07.20 12:18:00,735965 +02:00
upvoted 4 times



  **you1234** 3 years, 6 months ago
B & C & E is correct answer
upvoted 3 times

  **ama** 3 years, 6 months ago
B, D, E are correct...

C is wrong see my Reply to danito
upvoted 4 times



  **ama** 3 years, 6 months ago
The TIMESTAMP data type allows you to store date and time data including year, month, day, hour, minute and second.
upvoted 1 times

  **danito** 3 years, 7 months ago
C D E i think
upvoted 4 times

  **ama** 3 years, 6 months ago
C is wrong ...
SQL> select systimestamp from dual;

SYSTIMESTAMP

01.07.20 11:49:00,748290 +02:00
upvoted 2 times

  **elbelgounetos** 2 years, 10 months ago
A is wrong It's local_timezone

B is wrong, it depend of the function used to insert
C is not "only" it can contain more information. => true

D true see documentation.
SESSIONTIMEZONE returns the time zone of the current session. The return type is a time zone offset (a character type in the format '[+|-]TZH:TZM') or a time zone region name, depending on how the user specified the session time zone value in the most recent ALTER SESSION statement.

E true see documentation
DBTIMEZONE returns the value of the database time zone. The return type is a time zone offset (a character type in the format '[+|-]TZH:TZM') or a time zone region name, depending on how the user specified the database time zone value in the most recent CREATE DATABASE or ALTER DATABASE statement.

SO C,D,E
upvoted 2 times

Which three statements are true about inner and outer joins? (Choose three.)

- A. A full outer join returns matched and unmatched rows
- B. An inner join returns matched rows
- C. Outer joins can only be used between two tables per query
- D. A full outer join must use Oracle syntax
- E. Outer joins can be used when there are multiple join conditions on two tables
- F. A left or right outer join returns only unmatched rows

Correct Answer: ABE

Community vote distribution

ABE (100%)

  **danito** Highly Voted 3 years, 7 months ago

I think A B E
upvoted 18 times

  **Ekos** Highly Voted 3 years, 2 months ago

i think A, B, E is the correct answer
E. Outer joins can be used when there are multiple join conditions on two tables
can be like following query :
- select * from table a left join table1 b on a.col1 = b.col1 and a.col2=b.col2;
upvoted 6 times

  **Lalala8** Most Recent 3 months, 3 weeks ago

Selected Answer: ABE

The three statements that are true about inner and outer joins are:

A. A full outer join returns matched and unmatched rows.
B. An inner join returns matched rows.
E. Outer joins can be used when there are multiple join conditions on two tables.
The other statements are false:



C. Outer joins can only be used between two tables per query. Outer joins can be used between any number of tables per query.
D. A full outer join must use Oracle syntax. Full outer joins can be implemented in any SQL database, not just Oracle.
F. A left or right outer join returns only unmatched rows. A left or right outer join returns all rows from the left or right table, respectively, even if there are no matching rows in the other table.
upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago

A, B , E is correct.
upvoted 1 times



  **ryuah** 2 years ago

A,B,E is correct
upvoted 2 times

  **krontar** 2 years, 4 months ago

A B E

F is wrong. To be true It should be "FULL outer join".
Left outer (only rows from the left table)
Right outer (only rows from the right table)
upvoted 1 times

  **krontar** 2 years, 4 months ago
F. A left or right outer join returns only unmatched rows
To be TRUE it should be
"F. A left or right outer join returns only unmatched rows FOR THE JOINED TABLE"
upvoted 1 times



  **SimoneF** 3 years ago

A B E, no doubts for me here.
upvoted 5 times

  **escoletsgo1** 3 years, 4 months ago

F is wrong guys. How can a left or right join 'ONLY" return unmatched rows?
ABE is the answer.

upvoted 4 times

  **ama** 3 years, 5 months ago

A,B, F

upvoted 1 times

  **NiciMilo** 3 years, 5 months ago

Why F? i think E is correct

upvoted 1 times

  **ama** 3 years, 5 months ago

because F : LEFT JOIN returns only unmatched rows from the left table. RIGHT JOIN returns only unmatched rows from the right table.

<https://mode.com/sql-tutorial/sql-outer-joins/>

upvoted 2 times

  **adoptc94** 3 years, 5 months ago

F is wrong!! A left join returns all rows from the left table and rows from the right table where the join condition is met! Left or Right Joins DO NOT return only unmatched rows!
So if we have table1 (left table) and table2 (right table): The left join would return all records from table1 and only those records from table2 that intersect with table1.

I would go with answers: A,B,E

upvoted 6 times

  **Yema** 2 years, 11 months ago

ama is correct. I read his reference.

There are a few types of outer joins:

LEFT JOIN returns only unmatched rows from the left table.
RIGHT JOIN returns only unmatched rows from the right table.
FULL OUTER JOIN returns unmatched rows from both tables.

upvoted 1 times

  **saad3577** 1 year, 3 months ago

F is wrong
it's not "ONLY"
- Right Outer Join returns all the rows from the table on the right and columns of the table on the left is null padded.

- Left Outer Join returns all the rows from the table on the left and columns of the table on the right is null padded.

<https://www.geeksforgeeks.org/difference-between-left-right-and-full-outer-join/>

upvoted 1 times

  **you1234** 3 years, 6 months ago

A & B & F is correct answer

upvoted 4 times

  **saad3577** 1 year, 3 months ago

F is wrong
it's not "ONLY"
- Right Outer Join returns all the rows from the table on the right and columns of the table on the left is null padded.

- Left Outer Join returns all the rows from the table on the left and columns of the table on the right is null padded.

<https://www.geeksforgeeks.org/difference-between-left-right-and-full-outer-join/>

upvoted 1 times

  **ama** 3 years, 6 months ago

F?

The left outer join returns a resultset table with the matched data from the two tables and then the remaining rows of the left table and null from the right table's columns.

upvoted 1 times



Examine this description of the TRANSACTIONS table:

Name	Null?	Type
-----	-----	-----
TRANSACTION_ID	NOT NULL	VARCHAR2 (6)
TRANSACTION_DATE		DATE
AMOUNT		NUMBER (10, 2)
CUSTOMER_ID		VARCHAR2 (6)



Which two SQL statements execute successfully? (Choose two.)

- A. SELECT customer_id AS "CUSTOMER-ID", transaction_date AS DATE, amount + 100 "DUES" FROM transactions;
- B. SELECT customer_id AS "CUSTOMER-ID", transaction_date AS "DATE", amount + 100 DUES FROM transactions;
- C. SELECT customer_id AS CUSTOMER-ID, transaction_date AS TRANS_DATE, amount + 100 "DUES AMOUNT" FROM transactions;
- D. SELECT customer_id CUSTID, transaction_date TRANS_DATE, amount + 100 DUES FROM transactions;
- E. SELECT customer_id AS 'CUSTOMER-ID', transaction_date AS DATE, amount + 100 'DUES AMOUNT' FROM transactions;



Correct Answer: BD

-   **zouve** 5 months, 2 weeks ago

A. SELECT customer_id AS "CUSTOMER-ID", transaction_date AS DATE, amount + 100 "DUES" FROM transactions;
B. SELECT customer_id AS "CUSTOMER-ID", transaction_date AS "DATE", amount + 100 DUES FROM transactions;
C. SELECT customer_id AS CUSTOMER-ID, transaction_date AS TRANS_DATE, amount + 100 "DUES AMOUNT" FROM transactions;
D. SELECT customer_id CUSTID, transaction_date TRANS_DATE, amount + 100 DUES FROM transactions;
E. SELECT customer_id AS 'CUSTOMER-ID', transaction_date AS DATE, amount + 100 'DUES AMOUNT' FROM transactions;

upvoted 1 times
-   **Darkseid1231** 6 months, 2 weeks ago

Yes BD is okay

upvoted 1 times
-   **Givin** 1 year, 2 months ago

BD - Correct

upvoted 3 times

The CUSTOMERS table has a CUST_CREDIT_LIMIT column of data type NUMBER.
Which two queries execute successfully? (Choose two.)

- A. SELECT NVL(cust_credit_limit * .15, 'Not Available') FROM customers;
- B. SELECT NVL2(cust_credit_limit * .15, 'Not Available') FROM customers;
- C. SELECT NVL(TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;
- D. SELECT TO_CHAR(NVL(cust_credit_limit * .15, 'Not Available')) FROM customers;
- E. SELECT NVL2(cust_credit_limit, TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;

Correct Answer: CE



Community vote distribution



CE (100%)

  **you1234** Highly Voted 3 years, 6 months ago

- A) SELECT NVL(cust_credit_limit * .15, 'Not Available') FROM customers;
- B) SELECT NVL2(cust_credit_limit * .15, 'Not Available') FROM customers;
- C) SELECT NVL(TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;
- D) SELECT TO_CHAR(NVL(cust_credit_limit * .15, 'Not Available')) FROM customers;
- E) SELECT NVL2(cust_credit_limit, TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;

Correct C & E.
upvoted 16 times

  **ama** 3 years, 5 months ago
YES C & E
upvoted 3 times

  **Ekos** 3 years, 2 months ago
correct, C & E
upvoted 3 times



  **Lalala8** Most Recent 3 months, 3 weeks ago



Selected Answer: CE
C. SELECT NVL(TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;



This query will return the value of the cust_credit_limit column multiplied by 0.15, converted to a character value, if the cust_credit_limit column is not null. Otherwise, it will return the string 'Not Available'.



E. SELECT NVL2(cust_credit_limit, TO_CHAR(cust_credit_limit * .15), 'Not Available') FROM customers;

This query will return the value of the cust_credit_limit column if it is not null. Otherwise, it will return the value of the expression TO_CHAR(cust_credit_limit * .15), converted to a character value. If the expression TO_CHAR(cust_credit_limit * .15) is null, the NVL2() function will return the string 'Not Available'.
upvoted 1 times

  **j_tw** 9 months, 1 week ago
I don't know why not A ?
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Implicit cast is never easy to understand on fly, so oracle (and many others) recommend to try to avoid them.
upvoted 1 times

  **XhostSI** 1 year, 11 months ago
Adopted to HR Schema:
SELECT NVL(TO_CHAR(commission_pct * .15), 'Not Available') FROM employees;
SELECT NVL2(commission_pct, TO_CHAR(commission_pct * .15), 'Not Available') FROM employees;
Correct answers: C, E
upvoted 1 times

  **ama** 3 years, 6 months ago
C , D are corrects!
upvoted 1 times

  **amorimleandro** 3 years, 6 months ago

Nope, C, E are correct. You cannot convert all the function NLE in L, because its result is an error. But, in E, we have a similar situation as in C.
upvoted 2 times

  **LrnsTgh** 2 years, 5 months ago

D has no different with A. they aren't correct.
upvoted 1 times

店长微信：hifeng128

Which statement is true about aggregate functions?



- A. Aggregate functions can be nested to any number of levels
- B. The AVG function implicitly converts NULLS to zero
- C. Aggregate functions can be used in any clause of a SELECT statement
- D. The MAX and MIN functions can be used on columns with character data types



Correct Answer: A

Community vote distribution

D: (100%)

  **you1234** Highly Voted 3 years, 6 months ago
D is correct
upvoted 19 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times

  **Lalala8** Most Recent 3 months, 3 weeks ago
Selected Answer: D



The only statement that is true about aggregate functions is:



D. The MAX and MIN functions can be used on columns with character data types



The other statements are false:



A. Aggregate functions can be nested to any number of levels - Aggregate functions can be nested to a maximum depth of 15 levels.
B. The AVG function implicitly converts NULLS to zero - The AVG function ignores NULLS when calculating the average.
C. Aggregate functions can be used in any clause of a SELECT statement - Aggregate functions can only be used in the SELECT clause and the GROUP BY clause of a SELECT statement.
Aggregate functions are functions that perform calculations on groups of rows. Some common aggregate functions include COUNT(), SUM(), AVG(), MAX(), and MIN().



Aggregate functions can be used to calculate a variety of statistics, such as the total number of rows in a table, the average value of a column, or the highest value in a column.
upvoted 1 times

  **Oracle2020** 5 months, 1 week ago
The correct answer is D It is true that MAX and MIN can be used with character data types.
B is incorrect because AVG ignores NULL values
C is incorrect because aggregate functions cannot be used in any clause of the select. That is, they are used with SELECT statements that select many rows, scan a set of rows, and return a single response for all of them.
upvoted 1 times

  **RaNik69** 8 months, 1 week ago
Selected Answer: D
d is correct
upvoted 1 times

  **Lingy** 1 year, 2 months ago
A is wrong.
Aggregates are not allowed on an outer query.
D is correct
MAX is allowed only on expressions that evaluate to built-in data types (including CHAR, VARCHAR, DATE, TIME, CHAR FOR BIT DATA, etc.).
[https://docs.oracle.com/javadb/10.6.2.1/ref/rrefsqlj27781.html#:~:text=MAX%20is%20an%20aggregate%20function,BIT%20DATA%2C%20etc.\).](https://docs.oracle.com/javadb/10.6.2.1/ref/rrefsqlj27781.html#:~:text=MAX%20is%20an%20aggregate%20function,BIT%20DATA%2C%20etc.).)
<https://docs.oracle.com/javadb/10.6.2.1/ref/rrefsqlj33923.html#rrefsqlj33923>
upvoted 2 times

  **LeandroHPN** 1 year, 6 months ago
D is correct. Max level aggregate functions is 2.
upvoted 1 times

  **danito** 3 years, 7 months ago

A is wrong. I think the correct is D
upvoted 2 times

店长微信：hj feng128



Your database instance was shut down normally and then started in NOMOUNT state. You then execute this command:
ALTER DATABASE MOUNT;
Which two actions are performed? (Choose two.)



- A. The online redo logs are opened
- B. The online data files are opened
- C. The alert log records the execution details
- D. The Oracle background processes are started
- E. The initialization parameter file is read
- F. The control file is read



Correct Answer: DF



Community vote distribution



CF (100%)

  **baeji** Highly Voted 3 years, 1 month ago
A. The online redo logs are opened - OPEN
B. The online data files are opened - OPEN
C. The alert log records the execution details - MOUNT
D. The Oracle background processes are started - NOMOUNT
E. The initialization parameter file is read - NOMOUNT
F. The control file is read - MOUNT
upvoted 18 times

  **Rivaldo11** Highly Voted 3 years, 6 months ago
E and D are during NOMOUNT stage, A and B are during OPEN stage, so correct answer should be C and F.
upvoted 15 times

  **NowOrNever** 3 years, 5 months ago
i agree with CF
<http://facedba.blogspot.com/2014/07/oracle-database-startup-stages-and.html>
upvoted 7 times



  **Ekos** 3 years, 1 month ago
i agree with CF
upvoted 2 times



  **Lalala8** Most Recent 3 months, 3 weeks ago
Selected Answer: CF

When you execute the ALTER DATABASE MOUNT command, the following two actions are performed:

The control file is read. The control file contains information about the database, including the names and locations of the data files and redo log files.
The online data files are opened. The online data files contain the data for the database.
The other actions are not performed when you execute the ALTER DATABASE MOUNT command:

The online redo logs are opened. The online redo logs are used to record changes to the database. They are not opened until you execute the ALTER DATABASE OPEN command.
The Oracle background processes are started. The Oracle background processes are responsible for managing the database. They are started when you execute the STARTUP command.
The initialization parameter file is read. The initialization parameter file contains configuration parameters for the database. It is read when you start the Oracle instance.
upvoted 1 times



  **Oracle2020** 5 months, 1 week ago
I agree C and F
upvoted 1 times

  **Kuraudio** 1 year, 1 month ago
Selected Answer: CF
CF is correct
upvoted 1 times

  **[Removed]** 2 years, 4 months ago

Initialization parameters are read and Oracle background processes are started in NOMOUNT mode.
The control file is read , alert log records the execution details along with trace files starts in MOUNT mode because the control file is read at this stage.
Redo logs and datafiles are opened after OPEN mode.

E,D 100%
upvoted 1 times

  **[Removed]** 2 years, 4 months ago
I mean CF lol
upvoted 1 times

Which two are true about shrinking a segment online? (Choose two.)

- A. It is not possible to shrink either indexes or Index Organized Tables (IOTs)
- B. It always eliminates all migrated rows if any exist in the table
- C. To shrink a table it must have a PRIMARY KEY constraint
- D. To shrink a table it must have a UNIQUE KEY constraint
- E. To shrink a table it must have row movement enabled
- F. It must be in a tablespace that uses Automatic Segment Space Management (ASSM)

Correct Answer: CE

Reference:


https://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_3001.htm

Community vote distribution

EF (100%)

  **you1234** Highly Voted 3 years, 6 months ago



E & F is correct one
upvoted 20 times

  **ama** 3 years, 6 months ago

i agree
upvoted 5 times

  **leozanon94** 2 years, 7 months ago

i agree, but during exam what did you chose?
upvoted 2 times

  **Ekos** 3 years, 1 month ago

i agree E & F is the correct answer
upvoted 4 times

  **Lalala8** Most Recent 3 months, 3 weeks ago



Selected Answer: EF

The two statements that are true about shrinking a segment online are:

- E. To shrink a table it must have row movement enabled
- F. It must be in a tablespace that uses Automatic Segment Space Management (ASSM)

The other statements are false:

- A. It is not possible to shrink either indexes or Index Organized Tables (IOTs) - It is possible to shrink indexes and IOTs, but it is more complex than shrinking a table.
 - B. It always eliminates all migrated rows if any exist in the table - Shrinking a segment does not always eliminate all migrated rows. Only migrated rows that are no longer referenced by any other rows will be eliminated.
 - C. To shrink a table it must have a PRIMARY KEY constraint - It is not necessary for a table to have a PRIMARY KEY constraint to be shrunk.
 - D. To shrink a table it must have a UNIQUE KEY constraint - It is not necessary for a table to have a UNIQUE KEY constraint to be shrunk.
- upvoted 1 times

  **Oracle2020** 5 months, 1 week ago

I agree E,F
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: EF

E and F correct for me.
upvoted 1 times

  **Kuraudio** 1 year, 1 month ago

Selected Answer: EF

EF is correct
upvoted 1 times

Examine these commands:

```
[oracle@host01 ~]$ sqlplus u1/oracle
SQL> SELECT * FROM emp;
ENO          ENAME          DN
-----
1            Alan            2
2            Ben             2

SQL> exit
[oracle@host01 ~]$ cat emp.dat

3,Curl,4
4,Bob,4
[oracle@host01 ~]$ sqlldr u1/oracle TABLE=emp
```

Which two statements are true about the sqlldr execution? (Choose two.)

- A. It overwrites data in EMP with data in EMP.DAT
- B. It uses the database buffer cache to load data
- C. It generates a log that contains control file entries, which can be used with normal SQL*Loader operations
- D. It generates a sql script that it uses to load data from EMP.DAT to EMP
- E. It appends data from EMP.DAT to EMP

Correct Answer: BD



Community vote distribution

CE (100%)

  **Rivaldo11** Highly Voted 3 years, 6 months ago

C & E are correct - ..."SQL*Loader express mode generates a log file that includes a SQL*Loader control file. The log file also contains SQL scripts for creating the external table and performing the load using a SQL INSERT AS SELECT statement. Neither the control file nor the SQL scripts are used by SQL*Loader express mode. They are made available to you in case you want to use them as a starting point to perform operations using regular SQL*Loader or standalone external tables; the control file is for use with SQL*Loader, whereas the SQL scripts are for use with standalone external tables operations."

upvoted 28 times

  **Ekos** 3 years, 1 month ago

i agree

upvoted 3 times

  **Oracle2020** Most Recent 5 months, 1 week ago

I agree C,E
why examtopic answers wrong?

upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: CE

C and E are correct, wrong are:

A. SQL*Loader doesn't overwrite existing data in a table unless specified with the REPLACE option in the control file.

B. SQL*Loader uses direct path load method that bypasses the buffer cache and writes data blocks directly into the data files.

D. SQL*Loader doesn't generate a SQL script that it uses to load data. The loading operation is controlled by a control file, which is a text file that contains data loading instructions.

upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: CE

C & E are correct

upvoted 1 times

  **Borislone** 2 years, 6 months ago

CD is correct
upvoted 2 times



店长微信：hj feng128

  **Datajimm** 1 year, 10 months ago

D is not correct.
It is true, that is generates a script, but "Neither the control file nor the SQL scripts are used by SQL*Loader express mode." while answer D states "...that it uses to load data..."

Source: <https://docs.oracle.com/en/database/oracle/oracle-database/19/sutil/oracle-sql-loader-express-mode.html#GUID-0F35B551-861B-450D-8BF3-2312893A67D7>

upvoted 1 times

  **ama** 3 years, 6 months ago

This is tricky Question, as the Oracle documentation states:
A conventional Path insert takes the data in the source file and constructs Basic insert Statements using the individual source data as bind variables, These Statements are then then run against the table and data is inserted, the data is read into the Buffer Cache.

upvoted 1 times

  **you1234** 3 years, 6 months ago

D & E is correct answer
upvoted 2 times

A script abc.sql must be executed to perform a job.

A database user HR, who is defined in this database, executes this command:


\$ sqlplus hr/hr@orcl @abc.sql

What will happen upon execution?

- A. The command succeeds and HR will be connected to the orcl and abc.sql databases
- B. The command succeeds and HR will be connected to the orcl database instance, and the abc.sql script will be executed
- C. The command fails because the script must refer to the full path name
- D. The command fails and reports an error because @ is used twice

Correct Answer: B

Community vote distribution



you1234 Highly Voted 3 years, 6 months ago
B looks OK the answer?
upvoted 16 times

Yifan_ 4 days, 1 hour ago
B is correct
upvoted 1 times

Ekos 3 years, 1 month ago
yes, i agree
upvoted 2 times

ama 3 years, 6 months ago
yes B is correct
upvoted 7 times

Lalala8 Most Recent 3 months, 3 weeks ago
Selected Answer: B

The correct answer is B. The command succeeds and HR will be connected to the orcl database instance, and the abc.sql script will be executed.

The sqlplus command is used to connect to an Oracle database and execute SQL statements. The @ symbol is used to specify a script file to be executed.

In the command \$ sqlplus hr/hr@orcl @abc.sql, the user HR is connecting to the orcl database instance and specifying the abc.sql script to be executed.

The command will succeed and HR will be connected to the orcl database instance and the abc.sql script will be executed.
upvoted 2 times

doyinbare 4 months, 2 weeks ago
B is correct. I tried it in my personal environment and it worked. What happens is that the hr user connects to the orcl instance and then executes the sql script successfully.
upvoted 1 times

Oracle2020 5 months, 1 week ago
The correct answer is B, this time I agree with the one indicated.
upvoted 1 times

Oracle2020 5 months, 1 week ago
la respuesta correcta es la B, esta vez coincido con la indicada.
upvoted 1 times

auwia 6 months, 3 weeks ago
Selected Answer: B
Provided answer is correct.
upvoted 1 times

RaNik69 8 months, 1 week ago
Selected Answer: B
B is correct



upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago

Selected Answer: B

B is correct

upvoted 1 times

  **ryuah** 1 year, 7 months ago

Selected Answer: C



hey C is correct because orcl is alias not database instance. and full path need.

upvoted 1 times

  **auwia** 6 months, 3 weeks ago



It's false, as the example is exposed you can also omit the extension of the file (.sql)

upvoted 1 times

  **Guhborges** 9 months, 4 weeks ago


Sorry but you're wrong. if you're in linux path that have the script, full path not needed.

upvoted 2 times

  **danito** 3 years, 7 months ago

C is correct. orcl is a host not database instance

upvoted 1 times

  **ama** 3 years, 6 months ago

You are wrong!!!!

upvoted 2 times

  **Guhborges** 2 years, 1 month ago

Yeah, on this case you're wrong


upvoted 1 times

Which three statements are true regarding indexes? (Choose three.)

- A. A UNIQUE index can be altered to be non-unique
- B. A SELECT statement can access one or more indices without accessing any tables
- C. A table belonging to one user can have an index that belongs to a different user
- D. An update to a table can result in updates to any or all of the table's indexes
- E. When a table is dropped and is moved to the RECYCLE BIN, all indexes built on that table are permanently dropped
- F. An update to a table can result in no updates to any of the table's indexes

Correct Answer: ADE

Community vote distribution



ama Highly Voted 3 years, 5 months ago
B, C, F
upvoted 11 times

Phat 3 years ago
I would select BCF
upvoted 2 times

ama Highly Voted 3 years, 6 months ago
so in my opinion , B & C & D are correct!
upvoted 7 times

MCzombie 3 years ago
agreed!
upvoted 2 times

zouve Most Recent 5 months, 2 weeks ago
Selected Answer: BDF
regarding C the owner of the table must grant the INDEX privilege on the table to another user allowing that user to create an index on the table. In general, it is best practice to keep tables and indexes owned by the same user. This will help to ensure the security of your data.
upvoted 1 times

auwia 6 months, 3 weeks ago
Selected Answer: BDF
A. A UNIQUE index can be altered to be non-unique --> False
B. A SELECT statement can access one or more indices without accessing any tables. True
C. A table belonging to one user can have an index that belongs to a different user --> False, schema is still the owner.
D. An update to a table can result in updates to any or all of the table's indexes. True, it depends if the update goes on column with index or without.
E. When a table is dropped and is moved to the RECYCLE BIN, all indexes built on that table are permanently dropped. --> False, indexes will follow the table.
F. An update to a table can result in no updates to any of the table's indexes. True, you do an update on a single column that has no index.
upvoted 2 times

jonsnoow 4 months, 1 week ago
C is OK. Here is the test.
As sys user grant:
GRANT CREATE ANY INDEX TO HR;

As hr user:
CREATE INDEX cust_first_name_idx ON oe.customers (cust_first_name);
Index CUST_FIRST_NAME_IDX created.



As user sys:



SELECT D.OWNER, D.table_owner, D.table_name FROM dba_indexes D WHERE D.index_name = 'CUST_FIRST_NAME_IDX';



"OWNER" "TABLE_OWNER" "TABLE_NAME"
"HR" "OE" "CUSTOMERS"
upvoted 1 times



piontk 7 months, 1 week ago

B is right, when you use a Fast Full Index Scan.
https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/indexes-and-index-organized-tables.html#GUID-2C8A7262-DEBE-4932-8EFB-64E61AD041C1
upvoted 1 times

  **Marcello86CT** 1 year, 7 months ago
For Me is B , E and F
A. is wrong ,because i have to drop the index and then create it again
B. is correct since Index are phisically and logically indipendent from data to which they refer
C. is Wrong , Index are schema object so is not possible that a table on a schema use an index created on another schema.
D is wrong , it follow B. Indexes are independent from data they just speed up the retrieve process.
E. is Correct , when we drop a table then its indexes are dropped as well
F. is correct due to its logical and physical independent nature then an update can result in a no update on any indexes of a table.
upvoted 2 times



  **yarsalan** 1 year, 12 months ago
Selected Answer: BCF
B, C, F are correct
upvoted 1 times



  **Borislone** 2 years, 4 months ago
The Answer should be BDE.
C is missing the "Not" making it incorrect.
upvoted 1 times



  **SimoneF** 3 years ago
I think the biggest confusion here is given by option D.
While A is false (can't change the uniqueness of an index) as E (Indexes of a dropped table are moved to the recycled bin if enabled)



B is true if you select more indexed columns on different indexes and no unindexed columns.
C is true for sure cause it's possible, given the right privileges, to create an index as schema X on a table owned by schema Y.
F Is true cause no index may be updated, if you update a non-indexed column and the row address is unchanged

D also sounds right, I think the key might be in the phrasing here: if it's intended that the update of any or all indexes are the only possible outcome of a table update, then it would in fact be false. Otherwise, as someone else noted, I also suspect that more answers may be right, even though you only have to select three.
upvoted 4 times



  **gbab1792** 3 years, 1 month ago
C D E is correct
upvoted 1 times



  **MartinY** 3 years, 2 months ago
BCF is correct
B: index-only scan
C: global temporary table
F: changing the non-index column
upvoted 2 times



  **rosiiiee** 3 years, 4 months ago
So what is the EXACT ANSWER? I think CDE... I saw this question when i study 1z0-071 test, but after i see these discusses, i cannot have confidence about my answer...
upvoted 1 times

  **danito** 3 years, 4 months ago
Be careful with the exam. I had this question but with a small detail:
one of answer was: very similar to C but "cannot" instead of "can"
C A table belonging to one user can have an index that belongs to a different use
(my exam) A table belonging to one user cannot have an index that belongs to a different use

So you must study a lot. To study this dump is not enough
upvoted 3 times

  **NowOrNever** 3 years, 4 months ago
Thx
so BDF is right
and C is maybe missing the "not" which would make it to a wrong answer
upvoted 3 times

  **Phat** 3 years ago
if this is changed in the question, so I guess the final answer is BDF.
upvoted 2 times

  **BananaSlug** 3 years, 5 months ago
A C E

A - True: To view all indexes, you can simply type: SELECT * FROM all_indexes;

B - False: "when rows are inserted or deleted, all indexes on the table must be updated as well. Also, when a column is updated, all indexes that contain the column must be updated." - https://docs.oracle.com/cd/B10501_01/server.920/a96521/indexes.htm
upvoted 1 times

 **BananaSlug** 3 years, 5 months ago

C - True: "To create an index in another schema, you must have CREATE ANY INDEX system privilege. Also, the owner of the schema to contain the index must have either space quota on the tablespaces to contain the index or index partitions, or UNLIMITED TABLESPACE system privilege." - https://docs.oracle.com/cd/A87860_01/doc/server.817/a76956/indexes.htm
upvoted 1 times


 **BananaSlug** 3 years, 5 months ago

D - False: All answers on Stack Overflow say the only way to do this is to first drop the unique index and recreate it as non-unique. - <https://stackoverflow.com/questions/11979152/oracle-database-converting-unique-index-to-non-unique-one>

E - True: (same as B) "when rows are inserted or deleted, all indexes on the table must be updated as well. Also, when a column is updated, all indexes that contain the column must be updated." - https://docs.oracle.com/cd/B10501_01/server.920/a96521/indexes.htm
upvoted 1 times

 **BananaSlug** 3 years, 5 months ago

F - False: "If you dropped a table, and recovered it back from the recycle bin, what happens to its indexes? Are all indexes recovered back automatically? The answer is that all indexes will be recovered, if you recover a dropped table from the recycle bin. However, the indexes' names will not be the original names." - <http://dba.fyicenter.com/faq/oracle/Recovered-Table-with-Index.html>
upvoted 1 times


 **NowOrNever** 3 years, 5 months ago

AE - wrong that i agree

D - correct i agree too

F - if i update a non indexed column and the row does not change the position in the datafile (because the updated value uses the same amount of space) the index should not be touched right?

upvoted 1 times

 **ama** 3 years, 5 months ago

You are correct that updating an non-indexed column will not cause changes to the Indexes, probably F should be picked instead D

what you think?

upvoted 1 times

 **ama** 3 years, 5 months ago

again here is the evidence that F correct and D is wrong


When rows are inserted or deleted, all indexes on the table must be updated. When a column is updated, all indexes on the column must be updated.

https://docs.oracle.com/cd/B12037_01/appdev.101/b10795/adfns_in.htm

so final answers should be B, C , F

Thx

upvoted 1 times

 **NowOrNever** 3 years, 5 months ago

E - is not entirely true, you can't recover it, but it still remains in the recycle bin and you have to drop it to recreate an index with the same name on the column again - so what means "permanently dropped"?

B - https://docs.oracle.com/cd/A87860_01/doc/server.817/a76992/ch6_acce.htm

If all columns selected by a query are in a composite index, then Oracle can return these values from the index without accessing the table.

Using Bitmap Access Plans on Regular B*-tree Indexes

Note: This statement is executed by accessing indexes only, so no table access is necessary.

My tendency goes to BCD

I thing EF may also be possible, F more likely than E

upvoted 1 times

 **adoptc94** 3 years, 5 months ago

I would also suggest: BCD

Especially for D Oracle says following:



The database automatically maintains and uses indexes after they are created. The database also automatically reflects changes to data, such as adding, updating, and deleting rows, in all relevant indexes with no additional actions required by users. Retrieval performance of indexed data remains almost constant, even as rows are inserted. However, the presence of many indexes on a table degrades DML performance because the database must also update the indexes.

Source: https://docs.oracle.com/cd/E11882_01/server.112/e40540/indexiot.htm#CNCPT1895



But F seems also like it's true, just like you mentioned if someone updates only non indexed columns in the table, there shouldn't be any changes done to the index.

I am confused as to what the right answers are for sure - is it possible that out of the 6 choices 4 are right, but you only have to choose 3 of them?

upvoted 3 times

  **ama** 3 years, 6 months ago
E is wrong >

What Is the Recycle Bin?
The recycle bin is actually a data dictionary table containing information about dropped objects. Dropped tables and any associated objects such as indexes, constraints, nested tables, and the likes are not removed and still occupy space. They continue to count against user space quotas, until specifically purged from the recycle bin or the unlikely situation where they must be purged by the database because of tablespace space constraints.
upvoted 1 times

  **ama** 3 years, 6 months ago
C is correct >

To create an index in another schema, you must have CREATE ANY INDEX system privilege. Also, the owner of the schema to contain the index must have either the UNLIMITED TABLESPACE system privilege or space quota on the tablespaces to contain the index or index partitions.
upvoted 2 times

Which three statements are true about Oracle synonyms? (Choose three.)

- A. A synonym cannot be created for a PL/SQL package
- B. A synonym can be available to all users
- C. A SEQUENCE can have a synonym
- D. A synonym created by one user can refer to an object belonging to another user
- E. Any user can drop a PUBLIC synonym



Correct Answer: BCD



Reference:



https://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_7001.htm



Community vote distribution



BCD (100%)



  **NiciMilo** Highly Voted 3 years, 5 months ago
B, C and D
upvoted 19 times



  **Ekos** 3 years, 1 month ago
i agree, B, C, and D is the correct answer
upvoted 2 times

  **Kuraudio** Most Recent 1 year, 1 month ago
Selected Answer: BCD
BCD is correct
upvoted 1 times



  **LeandroHPN** 1 year, 6 months ago
B , C , D is correct
upvoted 1 times

  **Mandar79** 3 years, 4 months ago
https://docs.oracle.com/cd/B19306_01/server.102/b14200/statements_7001.htm
Public synonyms are accessible to all users.
upvoted 1 times

  **Mandar79** 3 years, 4 months ago
https://docs.oracle.com/cd/B28359_01/server.111/b28286/statements_9002.htm#SQLRF01805
To drop a private synonym, either the synonym must be in your own schema or you must have the DROP ANY SYNONYM system privilege.
To drop a PUBLIC synonym, you must have the DROP PUBLIC SYNONYM system privilege.
upvoted 2 times

  **ama** 3 years, 6 months ago
E is wrong!!!!!!!!!!!!

To drop a PUBLIC synonym, you must have the DROP PUBLIC SYNONYM system privilege
upvoted 3 times

  **you1234** 3 years, 6 months ago
B & D & E is correct answer
upvoted 2 times

View the Exhibit and examine the structure of the PRODUCTS table

Which two tasks require subqueries? (Choose two.)

Table PRODUCTS		
Name	Null?	Type
PROD_ID	NOT NULL	NUMBER(6)
PROD_NAME	NOT NULL	VARCHAR2(50)
PROD_DESC	NOT NULL	VARCHAR2(4000)
PROD_CATEGORY	NOT NULL	VARCHAR2(50)
PROD_CATEGORY_ID	NOT NULL	NUMBER
PROD_UNIT_OF_MEASURE		VARCHAR2(20)
SUPPLIER_ID	NOT NULL	NUMBER(6)
PROD_STATUS	NOT NULL	VARCHAR2(20)
PROD_LIST_PRICE	NOT NULL	NUMBER(8,2)
PROD_MIN_PRICE	NOT NULL	NUMBER(8,2)

- A. Display the number of products whose PROD_LIST_PRICE is more than the average PROD_LIST_PRICE
- B. Display suppliers whose PROD_LIST_PRICE is less than 1000
- C. Display products whose PROD_MIN_PRICE is more than the average PROD_LIST_PRICE of all products, and whose status is orderable
- D. Display the total number of products supplied by supplier 102 which have a product status of obsolete
- E. Display the minimum PROD_LIST_PRICE for each product status

Correct Answer: AC

Community vote distribution

AC (100%)

- Ekos

Highly Voted

3 years, 1 month ago

i agree, A & C is the correct answer

upvoted 11 times
- mamadu

Most Recent

6 months, 2 weeks ago

True A & C

A. select * from products where PROD_LIST_PRICE > (select AVG(PROD_LIST_PRICE) from products);

B. select SUPPLIER_ID from products where PROD_LIST_PRICE < 1000;

C. select * from products where PROD_LIST_PRICE > (select AVG(PROD_LIST_PRICE) from products) and PROD_STATUS = 'orderable';

D. select * from products where SUPPLIER_ID = 102 and PROD_STATUS = 'obsolete';

E. select MIN(PROD_LIST_PRICE) from products group by PROD_STATUS;

upvoted 1 times
- LeandroHPN

1 year, 6 months ago

Selected Answer: AC

A e C is correct.

upvoted 1 times
- algerianphoenix

1 year, 8 months ago

AC, correct answer by testing.

upvoted 1 times
- iamjavacompiler

2 years, 8 months ago

Why is D wrong?

upvoted 1 times
- leozanon94

2 years, 7 months ago

Select * from products where supplier_id = 102 And prod_status = 'obsolete';

upvoted 3 times

The SALES_Q1 and USERS tablespaces exist in one of your databases and TEMP is a temporary tablespace.

Segment creation is not deferred.

You execute this command:

```
CREATE USER sales
  IDENTIFIED BY sales_123
  DEFAULT TABLESPACE sales_q1
  TEMPORARY TABLESPACE temp
  QUOTA 5M ON users;
```

Which three statements must be true so that the SALES user can create tables in SALES_Q1? (Choose three.)

- A. The sales user must have a quota on the TEMP tablespace
- B. The sales user must have a quota on the SALES_Q1 tablespace to hold the initial extents of all tables they plan to create in their schema
- C. The sales user must have been granted the CREATE SESSION privilege
- D. The sales user must have their quota on the users tablespace removed
- E. The sales user must have a quota on the SALES_Q1 tablespace to hold all the rows to be inserted into any table in their schema
- F. The sales user must have been granted the CREATE TABLE privilege

Correct Answer: BDF


Community vote distribution

BCF (100%)

  **Rivaldo11** Highly Voted 3 years, 6 months ago

B, C and F

upvoted 24 times

  **NowOrNever** 3 years, 5 months ago

BDF

<https://blogs.oracle.com/sql/how-to-create-users-grant-them-privileges-and-remove-them-in-oracle-database>

<https://docs.oracle.com/database/121/DBSEG/users.htm#DBSEG002>

You can assign each user a tablespace quota for any tablespace, except a temporary tablespace.

upvoted 1 times

  **adoptc94** 3 years, 5 months ago

Why would D be right?

The database checks the users quota in the tablespace that he wants to create the object in. Having quota in another tablespace isn't any issue so why would you need to remove the quota from tablespace users to be able to create objects in tablespace sales_q1?

The right answers are:

B -> The user need quota in the tablespace he wants to create the object (sales_q1)

C -> User needs to be able to create a session (login in) to the database to create objects in the first place

F -> User needs to have the right privilege to create the object (in this case a table)

D is complete nonsense given the question and won't help in any way to create objects in sales_q1 tablespace

upvoted 4 times

  **NowOrNever** 3 years, 4 months ago

Sorry

BCF of course



upvoted 2 times

  **RaNik69** Most Recent 8 months, 1 week ago

Selected Answer: BCF

BCF are correct

upvoted 1 times

  **Rcravi** 1 year, 10 months ago

Correct - CEF

A - you cant assign quota on TEMP tablespace

B - you do not need any quota for creating table in any tablespace

D - is wrong as you are not creating table in users, so this is not concerning to the question

Correct -

C,F - these are the basic privileges required.

E - For holding the data, quota is required.

upvoted 1 times

🗨️ 👤 **Rcravi** 1 year, 10 months ago

Update/correction

E - is correct when asked for inserting the data. So E is wrong for this question

B- is correct when deferred_segment_creation is set to FALSE, you need quota on sales_Q1 tablespace.

correct answers are B C and F.

upvoted 3 times

🗨️ 👤 **Ekos** 3 years, 1 month ago

C,E,F is the correct answer.

b is wrong. we dont need any initial extents when we create table. the segments are needed when we start insert the data.

upvoted 2 times

🗨️ 👤 **SimoneF** 3 years ago

Read the question again: it clearly states "Segment creation is not deferred." exactly to prevent this objection.

And E is false anyway, cause the question also asks for the requirements for the creation of the tables, not for their population.

upvoted 1 times

🗨️ 👤 **SimoneF** 3 years ago

Sorry one more thing: E is also false cause SALES_Q1 is not the only tablespace where sales has quota, so we don't know if all the tables in that schema will reside there or on USERS. So, SALES_Q1 doesn't necessarily have to be as big as all the data that the schema will have.

upvoted 1 times

🗨️ 👤 **Mandar79** 3 years, 4 months ago

BCF is correct

<https://docs.oracle.com/database/121/DBSEG/users.htm#DBSEG99778>

When you create a new user account, you should enable this user to access the database

upvoted 3 times

🗨️ 👤 **Chansi** 3 years, 5 months ago

Isn't it CEF ? What if the user is executing Create as Select so just B won't work rt ?

upvoted 2 times

🗨️ 👤 **SimoneF** 3 years ago

There are two problems with E:

first and most notably, user sales can (and not randomly, does) have quota on other tablespaces, so not all of his tables have to reside on SALES_Q1.

Second, the question asks what's necessary for the user to be able to create tables. To create a select, meaning creating and filling them, is not a necessary requirement. Even if the space wouldn't be necessary to insert all the data, to be enough for hosting the initial extents would satisfy the minimum requirement, which is to create the tables (hence, the structure) to begin with.

upvoted 1 times

Which three statements are true about table data storage in an Oracle Database? (Choose three.)

- A. Data block headers contain their own Data Block Address (DBA)
- B. A table row piece can be chained across several database blocks
- C. Multiple row pieces from the same row may be stored in different database blocks
- D. Multiple row pieces from the same row may be stored in the same block
- E. Data block free space is always contiguous in the middle of the block
- F. Index block free space is always contiguous in the middle of the block

Correct Answer: *BCE*

Reference:

https://books.google.com.pk/books?id=INtVjrmbzckC&pg=PA267&lpg=PA267&dq=Multiple+row+pieces+from+the+same+row+may+be+stored+in+different+database+blocks&source=bl&ots=IElhFe2Q52&sig=ACfU3U1TETJPY194pARxfIJ5R0aQdySCMQ&hl=en&sa=X&ved=2ahUKEwivveLa97_mAhWOzaQKHahYAVEQ6AEwAnoECAoQAQ#v=onepage&q=Multiple%20row%20pieces%20from%20the%20same%20row%20may%20be%20stored%20in%20different%20database%20blocks&f=false

Community vote distribution



ama Highly Voted 3 years, 6 months ago

I would say A, B, C are correct!
upvoted 12 times

baeji 3 years, 1 month ago
agreed A, B, C are correct
upvoted 2 times

zouve Most Recent 5 months ago

Selected Answer: ABC
After reviewing this <https://antognini.ch/2016/10/whats-the-difference-between-row-migration-and-row-chaining/>
upvoted 1 times

musafir 6 months ago

Selected Answer: ACD
A. Data block headers contain their own Data Block Address (DBA) - True
B. A table row piece can be chained across several database blocks - Wrong, a row may be stored in multiple row pieces which are then chained across several database blocks. The row is chained not the row piece.
C. Multiple row pieces from the same row may be stored in different database blocks - True
D. Multiple row pieces from the same row may be stored in the same block - True
E. Data block free space is always contiguous in the middle of the block - Wrong
F. Index block free space is always contiguous in the middle of the block - Wrong
"As the database fills a data block from the bottom up, the amount of free space between the row data and the block header decreases." No mention of middle of the block anywhere in the Oracle Docs.

<https://docs.oracle.com/en/database/oracle/oracle-database/23/cncpt/tables-and-table-clusters.html#GUID-37546C88-24EE-47BE-8662-A9CED99BB90F>
upvoted 2 times

auwia 6 months, 2 weeks ago

Selected Answer: ADE
A. Data block headers contain their own Data Block Address (DBA) --> True
B. A table row piece can be chained across several database blocks. --> False, it's not "database block", it's "data block"
C. Multiple row pieces from the same row may be stored in different database blocks. --> False, it's not "database block", it's "data block"
D. Multiple row pieces from the same row may be stored in the same block --> True
E. Data block free space is always contiguous in the middle of the block --> True, in the oracle official guide (online) there many pictures showing the "data block" and the free space is always included between header and data stored.
F. Index block free space is always contiguous in the middle of the block. --> False, "index block" does not exists, it is "data block with type table or index, or cluster".

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: ACD



A, C, and D sounds better to me.

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Please ignore this comment.

upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: ABC

D is false in 19c because there is no more limitation about the number of "intra-block" columns. This is an example for 400 columns

```
declare
l_txt long;
begin
l_txt := 'create table t (c0 varchar2(20)';
for i in 1 .. 400
loop
l_txt := l_txt || ', c' || i || ' varchar2(20)';
end loop;
l_txt := l_txt || ') ' ;
execute immediate l_txt;
end;
/
```

```
select count(dbms_rowid.ROWID_BLOCK_NUMBER(rowid)) BLOCK_NUMBER from t;
```

```
BLOCK_NUMBER
-----
0
```

```
declare
v_txt varchar2(4000);
begin
v_txt := 'insert into t values ( 0';
for i in 1 .. 400
loop
v_txt := v_txt || ', ' || i ;
end loop;
v_txt := v_txt || ') ' ;
execute immediate v_txt;
commit;
end;
/
```

```
select count(dbms_rowid.ROWID_BLOCK_NUMBER(rowid)) BLOCK_NUMBER from t;
```

```
BLOCK_NUMBER
-----
1
```

so I think the right answers are a (for sure), b, c (for sure)

upvoted 1 times

  **calibre_04** 1 year, 2 months ago

- A. Data block headers contain their own Data Block Address (DBA) ==> Yes
- B. A table row piece can be chained across several database blocks ==> No, ROWS are chained not ROW PIECE
- C. Multiple row pieces from the same row may be stored in different database blocks ==> Yes (Row chaining)
- D. Multiple row pieces from the same row may be stored in the same block ==> Yes (Intra-block chaining e.g. table with more than 255 columns)
- E. Data block free space is always contiguous in the middle of the block ==> NO
- F. Index block free space is always contiguous in the middle of the block ==> NO

upvoted 2 times

  **emburria** 2 years ago

A, C, D. B is not correct. A row can be chained, but not a row piece...if so, then is another piece

upvoted 3 times

  **flaviogcmelo** 2 years, 9 months ago

A,B and C.
<https://antognini.ch/2016/10/whats-the-difference-between-row-migration-and-row-chaining/>

upvoted 1 times

  **yukclam9** 2 years, 9 months ago

I dont think B is correct - the most granular unit of row is a single row piece. row piece can not be chained. a row can be chained by multiple row pieces + pointers.

upvoted 2 times

店长微信：hj feng128

  **saad3577** 1 year, 3 months ago

B is correct :
Simply put, chained rows happen when a row is too big to fit into a single database block. Chained rows usually result from an insert. For example, if the blocksize for the database is 4 kilobytes and there is an 8 kilobyte row to be inserted, Oracle will break the data into pieces and store it in 3 different blocks that are chained together. There is forwarding information available to allow the database to collect all of the bits of chained row information from multiple blocks.

http://www.dba-oracle.com/t_identify_chained_rows.htm

upvoted 1 times

  **Phat** 3 years ago

so what is the final answers for this?

upvoted 3 times

  **Mandar79** 3 years, 4 months ago

The Oracle block header contains the following:
[http://www.dba-oracle.com/t_data_block_format.htm#:~:text=The%20header%20of%20a%20data,and%20a%20%22variable%22%20area.&text=The%20data%20block%20address,displacement%20into%20the%20data%20file\).](http://www.dba-oracle.com/t_data_block_format.htm#:~:text=The%20header%20of%20a%20data,and%20a%20%22variable%22%20area.&text=The%20data%20block%20address,displacement%20into%20the%20data%20file).)
A fixed-size block header.

A block directory.

The interested transaction list (ITL).

Space management information.

The data block address.

upvoted 2 times

  **Mandar79** 3 years, 4 months ago

The header contains general block information, such as the block address and the type of segment (for example, data or index).
[https://docs.oracle.com/cd/B10501_01/server.920/a96524/c03block.htm#:~:text=The%20header%20contains%20general%20block,example%2C%20data%20or%20index\).](https://docs.oracle.com/cd/B10501_01/server.920/a96524/c03block.htm#:~:text=The%20header%20contains%20general%20block,example%2C%20data%20or%20index).)

upvoted 1 times

  **adoptc94** 3 years, 5 months ago

Right answers: A,B,C

E and F are false because the free space of a block gets fragmented over time as data is added and removed from it -> this is called fragmentation (meaning the free space is not contiguous anymore). If the free space of a block is fragmented to a point where no new rows can be inserted, the free space is coalesced by the oracle server so that it's contiguous again and can be used for inserts or updates.

Coalescing free blocks doesn't change the actual size of free space in the block.

For why answer D is wrong, I am not completely sure but I guess that you wouldn't store a row as multiple row pieces in one block. If you can fit the whole row into a single block than it would be stored there as a single row piece and not multiple - but that's just my guess

upvoted 2 times

  **KtNow** 3 years ago


about D.
--ora.doc--
Oracle Database can only store 255 columns in a row piece. if you insert a row into a table that has 1000 columns, then the database creates 4 row pieces, typically chained over multiple blocks. --
typically chained, but possible situation that it can be in one block.
for example oracle block 4K(8K). row 1000 columns=4 row pieces, if column is small then at least 2 row pieces possible insert in one block. why not? for my point of view answer D possible also.

upvoted 4 times

  **KtNow** 3 years ago

so i would say A-B-C-D

upvoted 1 times

  **ama** 3 years, 5 months ago



C is correct > <https://antognini.ch/2016/10/whats-the-difference-between-row-migration-and-row-chaining/>

upvoted 2 times

  **SimoneF** 3 years ago

Interestingly, this also says that D is true.
I was actually in doubt on the B answer, which is then false cause it is a row that is chained in different row pieces, while the right ones are A-C-D

upvoted 6 times

  **ama** 3 years, 6 months ago

E is wrong ...

Free space can be also Noncontiguous also called fragmented space.

https://docs.oracle.com/cd/E25054_01/server.1111/e25789/logical.htm

upvoted 4 times



ama 3 years, 6 months ago

A is correct!!

The header of a data block has two areas a "common" area and a "variable" area. The Oracle block header contains the following:

A fixed-size block header.

A block directory.

The interested transaction list (ITL).

Space management information.

The data block address. This "DBA" (48 bytes), consisting of the file number (v\$datafile.file#) and the relative block number (the displacement into the data file). Note that the relative block address is not the same as the absolute block address.

upvoted 1 times

Examine the description of the BOOKS table:

Name	Null?	Type
-----	-----	-----
TRANSACTION_ID	NOT NULL	VARCHAR2 (6)
TRANSACTION_DATE		DATE
AMOUNT		NUMBER (10, 2)
CUSTOMER_ID		VARCHAR2 (6)

The table has 100 rows.

Examine this sequence of statements issued in a new session:

INSERT INTO books VALUES ('ADV112', 'Adventures of Tom Sawyer', NULL, NULL);

SAVEPOINT a;

DELETE FROM books;

ROLLBACK TO SAVEPOINT a;

ROLLBACK;

Which two statements are true? (Choose two.)

- A. The second ROLLBACK command does nothing
- B. The second ROLLBACK command replays the delete
- C. The first ROLLBACK command restores the 101 rows that were deleted, leaving the inserted row still to be committed
- D. The second ROLLBACK command undoes the insert
- E. The first ROLLBACK command restores the 101 rows that were deleted and commits the inserted row

Correct Answer: DE

Community vote distribution



you1234

Highly Voted

3 years, 7 months ago

C & D is correct one

upvoted 24 times

Guhborges

2 years, 2 months ago

I agree, C & D

upvoted 1 times

Ekos

3 years, 1 month ago

i agree

upvoted 2 times

yukclam9

Highly Voted

2 years, 9 months ago

to me the insert does not make sense. why we are inserting characters to a date column?

upvoted 6 times

Reiwardy

1 year, 2 months ago

That's why they are trying to rollback the stuff.

upvoted 2 times

LrnsTgh

2 years, 5 months ago

LOL. very good eyes

upvoted 2 times

feixiang

Most Recent

5 months, 1 week ago

the SQL INSERT is wrong

upvoted 1 times



musafir

5 months, 3 weeks ago

Selected Answer: CD



C D is correct

upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

Selected Answer: A

There is only one correct answer here A.
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

the first rollback finishes the DELETE transaction, so no transaction open for second rollback.
upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago

Selected Answer: CD

C , D is correct.
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago

Even though it is asked for "two" correct answers, I only found D as correct, the others make no sense to me, especially the SQL INSERT itself is wrong.
upvoted 1 times

Which three statements are true about external tables in Oracle 18c and later releases? (Choose three.)

- A. External table files can be used for other external tables in a different database
- B. The ORACLE_LOADER access driver can be used to unload data from a database into an external table
- C. The ORACLE_DATAPUMP access driver can be used to unload data from a database into an external table
- D. They cannot be partitioned
- E. The ORACLE_DATAPUMP access driver can be used to load data into a database from an external table
- F. They support UPDATES but not INSERTs and DELETES



Correct Answer: ABF

Community vote distribution



ACE (100%)

  **GraceYu** Highly Voted 3 years, 7 months ago

ACE. B is wrong. ORACLE_LOADER loads data not unload data. ORACLE_DATAPUMP is able to load and unload data.
upvoted 21 times

  **danito** 3 years, 7 months ago

I agree
upvoted 2 times

  **Ekos** 3 years, 1 month ago

i agree
upvoted 1 times



  **Mandar79** Highly Voted 3 years, 4 months ago

ACE - Correct
https://docs.oracle.com/cd/B19306_01/server.102/b14215/et_concepts.htm
B - Incorrect
The ORACLE_LOADER access driver is the default. It can perform only data loads, and the data must come from text datafiles
C and E - Correct
The ORACLE_DATAPUMP access driver can perform both loads and unloads.
F - Incorrect
<https://docs.oracle.com/en/database/oracle/oracle-database/18/admin/managing-tables.html#GUID-697B86CC-875C-4F68-AF80-49B41F33AE45>
You can, for example, select, join, or sort external table data. You can also create views and synonyms for external tables. However, no DML operations (UPDATE, INSERT, or DELETE) are possible, and no indexes can be created, on external tables.
D - Incorrect
Partitioned external tables were introduced in Oracle Database 12c Release 2 (12.2), allowing external tables to benefit from partition pruning and partition-wise.
<https://oracle-base.com/articles/12c/partitioned-external-tables-12cr2#:~:text=Partitioned%20external%20tables%20were%20introduced,are%20supported%20with%20some%20restrictions.>
upvoted 7 times

  **Kuraudio** Most Recent 1 year, 1 month ago

Selected Answer: ACE

ACE is correct
upvoted 1 times

  **danito** 3 years, 7 months ago

F is wrong. The correct answers are ABC
<https://docs.oracle.com/en/database/oracle/oracle-database/18/admin/managing-tables.html#GUID-F6948F0E-0557-4C42-9145-1897DE974CC3>

Another access driver, ORACLE_DATAPUMP, lets you unload data....
No DML operations (UPDATE, INSERT, or DELETE) are possibl...
upvoted 1 times



Which three statements are true about the Oracle join and ANSI join syntax? (Choose three.)

- A. The Oracle join syntax supports creation of a Cartesian product of two tables
- B. The Oracle join syntax performs better than the SQL:1999 compliant ANSI join syntax
- C. The SQL:1999 compliant ANSI join syntax supports natural joins
- D. The SQL:1999 compliant ANSI join syntax supports creation of a Cartesian product of two tables
- E. The Oracle join syntax only supports right outer joins
- F. The Oracle join syntax supports natural joins
- G. The Oracle join syntax performs less well than the SQL:1999 compliant ANSI join syntax



Correct Answer: CDF



Community vote distribution



ACD (100%)



  **ama** Highly Voted 3 years, 5 months ago
A,C, D



<https://oracle-base.com/articles/9i/ansi-iso-sql-support>
upvoted 13 times



  **Ekos** 3 years, 1 month ago
i agree
upvoted 2 times

  **saad3577** Most Recent 1 year, 3 months ago
I think it's question with 4 choices
my choice is (A,C,D and F)
upvoted 2 times

  **mamadu** 6 months, 2 weeks ago
F is false
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago
since you do not specify a join condition, you just use the NATURAL JOIN keyword.
upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago
Selected Answer: ACD
A, C , D is correct.
upvoted 1 times

  **yarsalan** 1 year, 12 months ago
A and C and D
upvoted 1 times

Which two tasks can you perform using DBCA for databases? (Choose two.)

- A. Configure a nonstandard block size for a new database
- B. Register a new database with an available Enterprise Manager Management server
- C. Change the standard block size of an existing database
- D. Configure incremental backups for a new database
- E. Enable flashback database for an existing database

Correct Answer: AB

Reference:

https://docs.oracle.com/cd/B16254_01/doc/server.102/b14196/install003.htm

🗲️ 👤 **Rivaldo11** Highly Voted 👍 3 years, 6 months ago

A and B.

upvoted 13 times

🗲️ 👤 **Ekos** 3 years, 1 month ago

i agree

upvoted 1 times

🗲️ 👤 **yarsalan** Most Recent ⌚ 1 year, 12 months ago

A and B

upvoted 1 times

🗲️ 👤 **ofranco** 2 years, 4 months ago

For me is A,B

upvoted 1 times

🗲️ 👤 **Borislone** 2 years, 6 months ago

the correct answer is BD

upvoted 3 times

🗲️ 👤 **Borislone** 2 years, 6 months ago

D is correct. When creating new database

upvoted 1 times

🗲️ 👤 **Guhborges** 2 years, 2 months ago

You can't create incremental backups with dbca.

upvoted 1 times

🗲️ 👤 **shotcom** 4 months, 1 week ago

https://www.dba-oracle.com/t_adv_techniques_oem_grid_control_35.htm

according to that link its now possible to create incremental backups with dbca

upvoted 1 times

🗲️ 👤 **pr0glinx** 3 years, 2 months ago

A and B

You can change by DBCA: But for new databases only!

upvoted 4 times

You execute this query:

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd `Monday for` fmMonth rrrr')
```

What is the result?

- A. It executes successfully but does not return any result
- B. It returns the date for the first Monday of the next month
- C. It generates an error
- D. It returns the date for the last Monday of the current month

Correct Answer: B

Community vote distribution



zouve 5 months, 2 weeks ago

Selected Answer: B

the correct query is

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd "Monday for" fmMonth rrrr') FROM DUAL;
```

I had this question also on the 1Z0-071 exam.

upvoted 1 times

musafir 5 months, 3 weeks ago

Selected Answer: C

C is correct since FROM DUAL is missing and ('Monday for') should actually be ("Monday for"). Could be an issue where the website didn't post the question correctly. I've seen this question in 071 exam which has the correct format but different answers.

<https://www.examttopics.com/discussions/oracle/view/32798-exam-1z0-071-topic-2-question-17-discussion/>

upvoted 1 times

BeomK 9 months ago

B is correct.

```
SQL> SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd "Monday for" fmMonth rrrr') from dual;
```

```
TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE),'MON'),'DD"MONDAYFOR"FMMONTHRRRR')
```

```
-----
01 Monday for May 2023
```

upvoted 1 times

Guhborges 9 months, 4 weeks ago

The correc query:

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd "Monday for" fmMonth rrrr') from dual;
```

B is right.

upvoted 1 times

ivanadj 10 months, 2 weeks ago

Selected Answer: C

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd `Monday for` fmMonth rrrr') from dual;
```

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd `Monday for` fmMonth rrrr') from dual
```

*

ERROR at line 1:

ORA-01821: date format not recognized

upvoted 1 times

stefan_a 11 months, 3 weeks ago

Selected Answer: C

C - it generates an error, even if FROM dual is added.

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd `Monday for` fmMonth rrrr')
```

returns ORA-00923: FROM keyword not found where expected

```
SELECT TO_CHAR(NEXT_DAY(LAST_DAY(SYSDATE), 'MON'), 'dd `Monday for` fmMonth rrrr') FROM dual;
```



returns ORA-01821: date format not recognized

upvoted 1 times



  **Guhborges** 10 months, 2 weeks ago

店长微信：hj feng128

It was probably a mistake when moving to the website, some characters are wrong. For me, the righ answers is B.
upvoted 1 times

  **Wh0au** 1 year, 3 months ago

FROM DUAL; must be added after the query statement
so answer is B
upvoted 3 times

  **steves** 1 year, 4 months ago

I think answer should be C. There is no from clause!
upvoted 3 times

Examine this command:

```
CREATE UNDO TABLESPACE undotbs01
DATAFILE 'undotbs_01.dbf'
```

SIZE 100M -

AUTOEXTEND ON;

Which two actions must you take to ensure UNDOTBS01 is used as the default UNDO tablespace? (Choose two.)

- A. Add the SEGMENT SPACE MANAGEMENT AUTO clause
- B. Set UNDO_TABLESPACE to UNDOTBS01
- C. Add the NOLOGGING clause
- D. Make certain that the database operates in automatic undo management mode
- E. Add the ONLINE clause

Correct Answer: BD

  **Asmodeus** Highly Voted 2 years, 11 months ago

Also vote for B & D.
SQL> alter system set undo_management=manual scope=spfile;
Bounce instance.

```
SQL> show parameter undo
NAME TYPE VALUE
-----
temp_undo_enabled boolean FALSE
undo_management string MANUAL
undo_retention integer 900
undo_tablespace string UNDOTBS1
```

```
SQL> alter system set undo_tablespace='UNDOTBS2';
alter system set undo_tablespace='UNDOTBS2'
*
ERROR at line 1:
ORA-02097: parameter cannot be modified because specified value is invalid
ORA-30014: operation only supported in Automatic Undo Management mode
upvoted 8 times
```

  **doyinbare** Most Recent 4 months, 2 weeks ago



B & D. Mind you, the parameter UNDO_TABLESPACE is only valid in Automatic undo management mode, hence both go hand in hand.
upvoted 1 times

  **[Removed]** 2 years, 4 months ago

B and D 100%
upvoted 2 times

  **Mandar79** 3 years, 4 months ago

https://docs.oracle.com/cd/B19306_01/server.102/b14231/undo.htm
You can switch from using one undo tablespace to another. Because the UNDO_TABLESPACE initialization parameter is a dynamic parameter, the ALTER SYSTEM SET statement can be used to assign a new undo tablespace.
ALTER SYSTEM SET UNDO_TABLESPACE = undotbs_02;
upvoted 1 times

  **NowOrNever** 3 years, 4 months ago



BD - correct
ALTER SYSTEM SET UNDO_TABLESPACE = UNDOTBS01;
show parameter undo

```
NAME TYPE VALUE
-----
temp_undo_enabled Boolean FALSE
undo_management string AUTO
undo_retention integer 900
undo_tablespace string UNDOTBS01
upvoted 4 times
```

  **Mandar79** 3 years, 4 months ago

BD seems to be correct can anyone confirm
upvoted 4 times

店长微信：hj feng128

  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times

Which three statements are true about views in an Oracle database? (Choose three.)

- A. Views can be updated without the need to re-grant privileges on the view
- B. Tables in the defining query of a view must always exist in order to create the view
- C. The WITH CHECK clause prevents certain rows from being displayed when querying the view
- D. Data Manipulation Language (DML) can always be used on views
- E. Inserting one or more rows using a view whose defining query contains a GROUP BY clause will cause an error
- F. Deleting one or more rows using a view whose defining query contains a GROUP BY clause will cause an error
- G. The WITH CHECK clause prevents certain rows from being updated or inserted

Correct Answer: AEF

Community vote distribution

AFG (60%)

EFG (40%)

 **Rivaldo11** Highly Voted 3 years, 6 months ago

G is tricky answer - "The WITH CHECK clause prevents certain rows from being updated or inserted". Techically, it is correct, but I would rather say - "The WITH CHECK clause allow certain rows to be updated or inserted", since only rows which are inline with WITH CHECK constraint can be affected. Speaking about A - "...Specify OR REPLACE to re-create the view if it already exists. You can use this clause to change the definition of an existing view without dropping, re-creating, and regranting object privileges previously granted on it..." So, it should be A, E and F.

upvoted 8 times

 **noobasty** 2 years, 6 months ago

Answer is EFG

A is wrong because it should be views can be altered*** without the need to re-grant privileges on the view. Not update

G is correct, certain rows cannot be updated or inserted . for example. you try to update the department_id to 100 from 20 but the view was created with a where department id_ = 20

or insert a deparment_id other than 20 . it will cause a ORA-01402 error

upvoted 8 times

 **shotcom** 4 months ago

answered G was not finished, certain rows cannot be updated or inserted, where?

the correct form of the answer is "The WITH CHECK clause prevents certain rows from being updated or inserted in the underlying table through the view."

which makes G wrong

upvoted 1 times

 **noobasty** 2 years, 6 months ago


Answer is EFG

A is wrong because it should be views can be replaced*** without the need to re-grant privileges on the view. Not update

G is correct, certain rows cannot be updated or inserted . for example. you try to update the department_id to 100 from 20 but the view was created with a where department id_ = 20

or insert a deparment_id other than 20 . it will cause a ORA-01402 error

upvoted 2 times

 **Ekos** 3 years, 1 month ago

hmm... make sense

upvoted 1 times

 **NowOrNever** Highly Voted 3 years, 5 months ago

Correct Answer: AEF

Wrong Answer: BCD

Uncertain: G

https://docs.oracle.com/cd/B28359_01/server.111/b28286/statements_8004.htm#SQLRF01504

for me G could also be a valid answer

upvoted 5 times

 **Ajinkya_Tambe** Most Recent 2 months, 2 weeks ago

I think A,E,F,G all are correct

upvoted 1 times

 **zouve** 5 months, 2 weeks ago

Selected Answer: AFG

I had this question also on the 1Z0-071 exam.
AG was the right answers in there but the question is a little bit different
upvoted 1 times

🗨️ **auwia** 6 months, 2 weeks ago

Selected Answer: EFG

For E, F, and G ; I cannot find any point against them, so probably they are corrects in my opinion.
A, because the "updated" word I would say it's false, because we go in the case of DML op. when they are allowed and not on the view.

The view must not contain any of the following constructs:

A set operator
a DISTINCT operator
An aggregate or analytic function
A GROUP BY, ORDER BY, MODEL, CONNECT BY, or START WITH clause
A collection expression in a SELECT list
A subquery in a SELECT list
A subquery designated WITH READ ONLY
Joins, with some exceptions, as documented in Oracle Database Administrator's Guide.
upvoted 1 times

🗨️ **auwia** 6 months, 2 weeks ago

For A, it should be ALTERED without needs of re-grant, not UPDATED.
upvoted 1 times

🗨️ **piontk** 6 months, 3 weeks ago

Selected Answer: AFG

E -> not always, in this case is a wrong statement.
Test case:

```
CREATE TABLE employees (
id NUMBER,
name VARCHAR2(100),
salary NUMBER,
department_id NUMBER
);
```

```
INSERT INTO employees VALUES(1, 'John Doe', 50000, 1);
INSERT INTO employees VALUES(2, 'Jane Doe', 60000, 1);
INSERT INTO employees VALUES(3, 'Robert Smith', 70000, 2);
```

```
CREATE VIEW dept1_employees5 AS
SELECT id, count(*) cnt
FROM employees
group by id;
```

```
INSERT INTO dept1_employees5 (id)
VALUES (4);
```

This will not cause an error. Otherwise DELETE statement will:

```
SQL> delete from dept1_employees5 where id=4;
delete from dept1_employees5 where id=4
*
```

ERROR at line 1:
ORA-01732: data manipulation operation not legal on this view

In my opinion correct answer is AFG
upvoted 1 times

🗨️ **CMjer** 8 months, 3 weeks ago

Selected Answer: EFG

EFG is correct.
- E, F: insert/update/delete is prevented when you have GROUP BY in defining query of a view.
- G: true, that's what WITH CHECK clause does :)
upvoted 1 times

🗨️ **Kuraudio** 1 year, 1 month ago

A is correct:

Specify OR REPLACE to re-create the view if it already exists. You can use this clause to change the definition of an existing view without dropping, re-creating, and regranting object privileges previously granted on it.

Remember that removing an object's privileges has a cascade effect, so with the "OR REPLACE" you avoid assigning the privileges again.
upvoted 1 times

🗨️ **Kuraudio** 1 year, 1 month ago

A is wrong: Source: <https://www.oracletutorial.com/oracle-view/oracle-with-check-option/>

upvoted 1 times

🗨️ 👤 **Kuraudio** 1 year, 1 month ago

Sorry, A is correct

upvoted 1 times

🗨️ 👤 **yarsalan** 1 year, 12 months ago

Selected Answer: AFG

I think, A,F,G are correct

upvoted 1 times

🗨️ 👤 **Borislone** 2 years, 6 months ago

D is correct

upvoted 1 times

🗨️ 👤 **Mandar79** 3 years, 4 months ago

AEF

https://docs.oracle.com/cd/B28359_01/server.111/b28310/views001.htm#ADMIN11780

B - Incorrect

For example, when a view is created that refers to a nonexistent table or an invalid column of an existing table ...CREATE FORCE VIEW AS

E & F - Correct

If a view is defined by a query that contains SET or DISTINCT operators, a GROUP BY clause, or a group function, then rows cannot be inserted into, updated in, or deleted from the base tables using the view.

upvoted 3 times

🗨️ 👤 **adoptc94** 3 years, 4 months ago

This question is a nightmare... there are 4 possible answers

A -> Specify OR REPLACE to re-create the view if it already exists. You can use this clause to change the definition of an existing view without dropping, re-creating, and regranting object privileges previously granted on it.

Source: https://docs.oracle.com/cd/B28359_01/server.111/b28286/statements_8004.htm#SQLRF01504

E & F-> If a view is defined by a query that contains SET or DISTINCT operators, a GROUP BY clause, or a group function, then rows cannot be inserted into, updated in, or deleted from the base tables using the view.

Source: <https://docs.oracle.com/database/121/ADMIN/views.htm#ADMIN11781>

G -> Specify WITH CHECK OPTION to indicate that Oracle Database prohibits any changes to the table or view that would produce rows that are not included in the subquery. When used in the subquery of a DML statement, you can specify this clause in a subquery in the FROM clause but not in subquery in the WHERE clause.

Source: https://docs.oracle.com/cd/B28359_01/server.111/b28286/statements_8004.htm#SQLRF01504

I'd go with answer A, E, F and hope for the best!

upvoted 5 times

🗨️ 👤 **ama** 3 years, 6 months ago

A, F, G

Source > <https://juejin.im/post/5e43caf4f265da57213ef2e3>

Q138. Which three statements are true about views in an Oracle database?

A) The WITH CHECK clause prevents certain rows from being displayed when querying the view.

B) The WITH CHECK clause prevents certain rows from being updated or inserted.

C) Tables in the defining query of a view must always exist in order to create the view.

D) Data Manipulation Language (DML) can always be used on views

E) Deleting one or more rows using a view whose defining query contains a GROUP BY clause will cause an error.

F) Views can be updated without the need to re-grant privileges on the view.

G) Inserting one or more rows using a view whose defining query contains a GROUP BY clause will cause an error

Answer: BEF

upvoted 2 times

🗨️ 👤 **Rivaldo11** 3 years, 6 months ago

G should also be correct.

upvoted 1 times

🗨️ 👤 **you1234** 3 years, 6 months ago

A & E & F is correct answer

upvoted 3 times

🗨️ 👤 **danito** 3 years, 7 months ago

B is not correct, you can create a broken view with force option. I think DEG

upvoted 1 times

You execute this command:

```
CREATE SMALLFILE TABLESPACE sales  
DATAFILE '/u01/app/oracle/sales01.dbf
```

SIZE 5G -

SEGMENT SPACE MANAGEMENT AUTO;

Which two statements are true about the SALES tablespace? (Choose two.)

- A. It must be smaller than the smallest BIGFILE tablespace
- B. Free space is managed using freelists
- C. Any data files added to the tablespace must have a size of 5 gigabytes
- D. It uses the database default blocksize
- E. It is a locally managed tablespace

Correct Answer: *DE*

  **Ekos**  3 years, 1 month ago

D, E is the correct answer
upvoted 8 times

  **Mandar79**  3 years, 4 months ago

Incorrect segment management clause It lets you specify whether Oracle Database should track the used and free space in the segments in the tablespace using free lists or bitmaps. This clause is not valid for a temporary tablespace.

AUTO :Specify AUTO if you want the database to manage the free space of segments in the tablespace using a bitmap. MANUAL :Specify MANUAL if you want the database to manage the free space of segments in the tablespace using free lists

E - Correct

https://docs.oracle.com/cd/B28359_01/server.111/b28310/tspaces002.htm#ADMIN11360

Create a locally managed tablespace by specifying LOCAL in the EXTENT MANAGEMENT clause of the CREATE TABLESPACE statement. This is the default for new permanent tablespaces,

D - Correct

https://docs.oracle.com/cd/B28359_01/server.111/b28310/create005.htm#ADMIN11105

The most commonly used block size should be picked as the standard block size. In many cases, this is the only block size that you need to specify. Typically, DB_BLOCK_SIZE is set to either 4K or 8K. If you do not set a value for this parameter, the default data block size is operating system specific.

upvoted 5 times

  **[Removed]**  2 years, 4 months ago

D, E both are by default lol.
upvoted 1 times

  **Mandar79** 3 years, 4 months ago

B - Incorrect segment management clause It lets you specify whether Oracle Database should track the used and free space in the segments in the tablespace using free lists or bitmaps. This clause is not valid for a temporary tablespace.

AUTO :Specify AUTO if you want the database to manage the free space of segments in the tablespace using a bitmap. MANUAL :Specify MANUAL if you want the database to manage the free space of segments in the tablespace using free lists

upvoted 2 times

In the SALES database, DEFERRED_SEGMENT_CREATION is TRUE.
Examine this command:
SQL> CREATE TABLE T1(c1 INT PRIMARY KEY, c2 CLOB);
Which segment or segments, if any, are created as a result of executing the command?



- A. T1, an index segment for the primary key, a LOB segment, and a lobindex segment
- B. no segments are created
- C. T1 only
- D. T1 and an index segment created for the primary key only
- E. T1, an index segment for the primary key, and a LOB segment only

Correct Answer: C



Community vote distribution



B (100%)

  **Rivaldo11** Highly Voted 3 years, 6 months ago
Answer is B.
upvoted 12 times

  **ama** 3 years, 6 months ago
I agree

DEFERRED_SEGMENT_CREATION specifies the semantics of deferred segment creation. If set to true, then segments for tables and their dependent objects (LOBs, indexes) will not be created until the first row is inserted into the table.
Before creating a set of tables, if it is known that a significant number of them will not be populated, then consider setting this parameter to true. This saves disk space and minimizes install time.
upvoted 1 times

  **Ekos** 3 years, 1 month ago
i agree
upvoted 1 times

  **Chansi** Highly Voted 3 years, 5 months ago
SQL> alter session set deferred_segment_creation=TRUE;

Session altered.



SQL> CREATE TABLE T1(c1 INT PRIMARY KEY, c2 CLOB);



Table created.



SQL> select segment_name,segment_type,bytes from user_segments;

no rows selected

Answer is B
upvoted 7 times

  **PKSOA** Most Recent 4 months, 2 weeks ago
Selected Answer: B
B is correct answer
<https://oracle-base.com/articles/11g/segment-creation-on-demand-11gr2>
upvoted 1 times



  **Kuraudio** 1 year, 1 month ago
Selected Answer: B
B is correct
upvoted 1 times

  **[Removed]** 2 years, 4 months ago
B (until rows are inserted in non- partitioned table, it doesn't create any segments) when DEFERRED_SEGMENT_CREATION is true
upvoted 1 times



  **Viki5099** 3 years, 1 month ago

ALTER SESSION SET DEFERRED_SEGMENT_CREATION=TRUE
SHOW PARAMETERS DEF
CREATE TABLE TEMP_DEF_SEQ_CHK (C1 INT PRIMARY KEY, LOB1 CLOB)
SELECT * FROM USER_SEGMENTS;
no rows selected
SELECT SEGMENT_CREATED FROM USER_INDEXES;
SEG



NO
NO
upvoted 1 times

  **Mandar79** 3 years, 4 months ago
DEFERRED_SEGMENT_CREATION specifies the semantics of deferred segment creation. If set to true, then segments for tables and their dependent objects (LOBs, indexes) will not be created until the first row is inserted into the table.



<https://oracle-base.com/articles/11g/segment-creation-on-demand-11gr2>
upvoted 2 times



  **khalilshahin01** 3 years, 6 months ago
SQL> alter session set DEFERRED_SEGMENT_CREATION=true;
SQL> CREATE TABLE T1(c1 INT PRIMARY KEY, c2 CLOB);
SQL> select * from user_segments where segment_name = upper('T1');
no row selected

Answer is B.
upvoted 2 times

  **khalilshahin01** 3 years, 6 months ago
sorry Answer is A.

You can also query the SEGMENT_CREATED column of the USER_TABLES, USER_INDEXES, or USER_LOBS views. For non-partitioned tables, indexes, and LOBs, this column shows YES if the segment is created.
upvoted 1 times

  **ama** 3 years, 6 months ago
DEFERRED_SEGMENT_CREATION specifies the semantics of deferred segment creation. If set to true, then segments for tables and their dependent objects (LOBs, indexes) will not be created until the first row is inserted into the table.
upvoted 2 times

  **SimoneF** 3 years ago
Exactly, khalilshahin01 you got it right the first time: as you saw from the user_segments, if nothing is there then no segment has been created.
upvoted 1 times

  **you1234** 3 years, 6 months ago
12c.

SQL> show parameter def

NAME TYPE VALUE

deferred_segment_creation boolean TRUE

SQL> CREATE TABLE T1(c1 INT PRIMARY KEY, c2 CLOB);

Table created.

SQL> select CONSTRAINT_NAME,CONSTRAINT_TYPE,TABLE_NAME from user_constraints where table_name='T1';

CONSTRAINT_NAME C TABLE_NAME

SYS_C0081022 P T1



SQL> select COLUMN_NAME,SEGMENT_NAME,INDEX_NAME from user_lobs where table_name='T1';

COLUMN_NAME SEGMENT_NAME INDEX_NAME

C2 SYS_LOB0000686210C00002\$\$ SYS_IL0000686210C00002\$\$

SQL>
not done any DML operations.

A is correct.
upvoted 1 times

  **Muhab** 3 years, 6 months ago

店长微信：hj feng128

```
no B is correct.  
SQL> select COLUMN_NAME,SEGMENT_NAME,INDEX_NAME,SEGMENT_CREATED from user_lobs where table_name='T1';  
COLUMN_NAME SEGMENT_NAME INDEX_NAME SEGMENT_CREATED  
-----  
C2 SYS_LOB0000686210C00002$$ SYS_IL0000686210C00002$$ NO
```

see SEGMENT_CREATED is no that mean no segment created.
upvoted 5 times

Which three activities are recorded in the database alert log? (Choose three.)

- A. Data Definition Language (DDL) statements
- B. non-default database parameters
- C. block corruption errors
- D. deadlock errors
- E. session logins and logouts

Correct Answer: BCD

Reference:

https://docs.oracle.com/cd/B28359_01/server.111/b28310/monitoring001.htm#ADMIN11247

Community vote distribution

BCD (100%)

  **Mandar79** Highly Voted 3 years, 4 months ago



BCD

https://docs.oracle.com/cd/B28359_01/server.111/b28310/monitoring001.htm#ADMIN11247

All internal errors (ORA-600), block corruption errors (ORA-1578), and deadlock errors (ORA-60) that occur

The values of all initialization parameters that had nondefault values at the time the database and instance start

upvoted 13 times

  **Ekos** 3 years, 1 month ago

i agree

upvoted 2 times

  **auwia** Most Recent 6 months, 2 weeks ago

Selected Answer: BCD

Provided answer are corrects:

The alert log is a chronological log of messages and errors, and includes the following items:

All internal errors (ORA-600), block corruption errors (ORA-1578), and deadlock errors (ORA-60) that occur

Administrative operations, such as CREATE, ALTER, and DROP statements and STARTUP, SHUTDOWN, and ARCHIVELOG statements

Messages and errors relating to the functions of shared server and dispatcher processes

Errors occurring during the automatic refresh of a materialized view

The values of all initialization parameters that had nondefault values at the time the database and instance start.

upvoted 1 times

  **Frozencc** 7 months ago

Selected Answer: BCD

BCD are correct

upvoted 1 times

  **flaviogcmelo** 2 years, 9 months ago

In item 8.1.1.1 (About Monitoring Errors with Trace Files and the Alert Log) in the version 19c documentation, the following information is included:

The alert log is a chronological log of messages and errors, and includes the following items:

All internal errors (ORA-00600), block corruption errors (ORA-01578), and deadlock errors (ORA-00060) that occur

Administrative operations, such as some CREATE, ALTER, and DROP statements and STARTUP, SHUTDOWN, and ARCHIVELOG statements

Messages and errors relating to the functions of shared server and dispatcher processes

Errors occurring during the automatic refresh of a materialized view

The values of all initialization parameters that had nondefault values at the time the database and instance start

Therefore, I understand that there are four correct options in this item: ABCD

upvoted 1 times

  **auwia** 6 months, 2 weeks ago

A is wrong because only CREATE. ALTER and DROP are tracked, RENAME and TRUNCATE not!.

DDL = CREATE, ALTER, DROP, RENAME AND TRUNCATE.

upvoted 2 times

🗨️ 👤 **s_sani** 3 years, 4 months ago

CDE https://docs.oracle.com/cd/B28359_01/server.111/b28310/monitoring001.htm#ADMIN11242
upvoted 1 times

🗨️ 👤 **kemsalim** 2 months, 2 weeks ago

E is trully wrong because when you open that alert log u can't find it in alert log for sure
upvoted 1 times

🗨️ 👤 **NowOrNever** 3 years, 4 months ago

E is wrong
upvoted 2 times

🗨️ 👤 **Mandar79** 3 years, 4 months ago

https://docs.oracle.com/cd/B28359_01/server.111/b28310/monitoring001.htm#ADMIN11247
C and D are correct
All internal errors (ORA-600), block corruption errors (ORA-1578), and deadlock errors (ORA-60) that occur
The values of all initialization parameters that had nondefault values at the time the database and instance start
A - Incorrect
The DDL log is stored in the log/ddl subdirectory of the ADR home
11g : DDL go both log.xml & alert.log
12c : DDL only go to log.xml
upvoted 2 times

🗨️ 👤 **jonsnoow** 4 months, 1 week ago

"The DDL log is created only for the RDBMS component and only if the ENABLE_DDL_LOGGING initialization parameter is set to true. When this parameter is set to false, DDL statements are not included in the DDL log, but they are included in the alert log."


https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/ENABLE_DDL_LOGGING.html#GUID-6FBA6147-D545-4E7D-94F0-A97EC1C721AE
upvoted 1 times

Which two statements are true about Enterprise Manager (EM) Express?

- A. You can use a single instance of EM Express to manage multiple database running on the same server.
- B. EM Express uses a separate repository database to store target database metadata.
- C. By default, EM express is available for a database after database creation using DBCA.
- D. You can shut down a database instance using EM Express.
- E. You cannot start up a database instance using EM Express.

Correct Answer: CE

Community vote distribution



Mandar79 Highly Voted 3 years, 4 months ago
CE seems correct
upvoted 9 times

Guhborges 2 years, 2 months ago
You can't shutdown with EM on 18c and higher!
upvoted 1 times

Ekos 3 years, 1 month ago
yes, agree
upvoted 1 times

kemsalim Most Recent 2 months, 2 weeks ago
Selected Answer: CD
i chosee C and D,

Because you can shutdown the database using Enterprise Manager,
upvoted 1 times

Lalala8 3 months, 3 weeks ago
Selected Answer: CE
C and E correct answers.
upvoted 1 times

[Removed] 2 years, 4 months ago
Answer 1: C (after database is created using DBCA, thats exactly when EM is available)
Answer 2: E (You can never start database instance using an EM)

Why is A wrong?
=> On a same server, 1 (instance) of EM can manage only 1 database not multiple databases.
Many instances of EM can manage many databases.
upvoted 1 times

pepelu69 3 years ago
AE are the correct answer
upvoted 4 times

Phat 3 years ago
I agree
upvoted 1 times

Fan 2 years, 10 months ago
that's right. you can find this question in 062.
upvoted 1 times

Fan 2 years, 10 months ago
type. It's C E.
upvoted 1 times

Fan 2 years, 10 months ago
TYPO. C E are correct.
upvoted 1 times

🗨️ 👤 **Mandar79** 3 years, 4 months ago

店长微信 : hj feng128

EM Express is available only when the database is open. This means that EM Express cannot be used to start up the database
<https://docs.oracle.com/en/database/oracle/oracle-database/19/admqslg/getting-started-with-database-administration.html#GUID-BA75AD46-D22E-4914-A31E-C395CD6A2BBA>

<https://www.oracle.com/technetwork/database/manageability/emx-intro-1965965.html>
When dbca completes, it indicates the port on which EM Express is configured
upvoted 4 times

Examine this SQL statement:

SELECT cust_id, cust_last_name `Last Name`

FROM customers -

WHERE country_id = 10 -

UNION -

SELECT cust_id CUST_NO, cust_last_name

FROM customers -

WHERE country_id = 30 -



Identify three ORDER BY clauses, any one of which can complete the query successfully. (Choose three.)



- A. ORDER BY `Last Name`
- B. ORDER BY 2, 1
- C. ORDER BY 2, cust_id
- D. ORDER BY CUST_NO
- E. ORDER BY `CUST_NO`



Correct Answer: ABC



Community vote distribution



ABC (100%)



  **escoletsgo1** Highly Voted 3 years, 4 months ago
ABC
You cant sort by CUST_NO, it is not in the first query.
upvoted 16 times



  **Ekos** 3 years, 1 month ago
yes, agree
upvoted 2 times



  **Mandar79** Highly Voted 3 years, 4 months ago
ABC is correct, you cannot sort if cust_no is not present in the first select query
upvoted 6 times

  **musafir** Most Recent 6 months ago
Selected Answer: ABC
ABC correct
upvoted 1 times

  **Wh0au** 1 year, 3 months ago
corret answer is BCD!! why choice the ABC??
upvoted 1 times

  **saad3577** 1 year, 3 months ago
A is correct regarding oracle documentation :
The alias can be used in the order_by_clause but not other clauses in the query.
upvoted 1 times

  **LeandroHPN** 1 year, 6 months ago
Selected Answer: ABC
A, B, C is correct.
upvoted 1 times

  **algerianphoenix** 1 year, 8 months ago
BCD, correct answer by test
upvoted 2 times

  **[Removed]** 2 years, 4 months ago

店长微信：hj feng128

A, B, C (In union the first select statement must have the columns specified to make the usage of the order by clause)
upvoted 1 times

  **BananaSlug** 3 years, 5 months ago

B C D is correct
upvoted 3 times

Which two statements are true about the configuration and use of UNDO_RETENTION with GURANTEED RETENTION? (Choose two.)

- A. UNDO_RETENTION specifies for how long Oracle attempts to keep expired and unexpired UNDO.
- B. UNDO_RETENTION specifies how long all types of UNDO are retained.
- C. Unexpired UNDO is always retained.
- D. Active UNDO is always retained.
- E. UNDO_RETENTION specifies for how long Oracle attempts to keep unexpired UNDO.

Correct Answer: CE

Community vote distribution

CD (100%)

- Rivaldo11** Highly Voted

3 years, 6 months ago

"...If retention guarantee is enabled, the specified minimum undo retention is guaranteed; the database never overwrites unexpired undo data even if it means that transactions fail due to lack of space in the undo tablespace. If retention guarantee is not enabled, the database can overwrite unexpired undo when space is low, thus lowering the undo retention for the system..." Acoording to this, correct answers should be C and E.

upvoted 9 times
- _gio_**

2 years, 4 months ago

E is wrong because Oracle doesn't "attempt to keep". It's mandatory to keep the unexpired undo. So C, D

upvoted 4 times
- SimoneF** Highly Voted

3 years ago

By definition, undo extents that still have to overlive the retention period are unexpired. So, when undo_retention is set and retention is guaranteed, unexpired undo is always retained.

As opposed to the case in which retention is noguarantee, where also unexpired extents can be overwritten IF and when needed.

In this particular case, so, I would say C-D are the correct answers.

upvoted 7 times
- Phat**

3 years ago

After reading the oracle docs, section 16.2.2.3 Retention Guarantee. I can say the correct answer shall be C,E. retention guarantee & undo_retention are about unexpired undo information.

<https://docs.oracle.com/database/121/ADMIN/undo.htm#ADMIN11466>

upvoted 2 times
- elbelgounetos**

2 years, 10 months ago

I Agree.

Undo retention is the threshold after which UNDO go from unexpired to expired. So it's pointless to talk about undo retention that keep some kind of undo. UNDO RETENTION defines what kind of undo it is.

So C and D

upvoted 1 times
- RaNik69** Most Recent

8 months, 1 week ago

Selected Answer: CD

C D are correct

upvoted 1 times
- dianiss1050**

10 months, 2 weeks ago

C and D are correct

upvoted 1 times
- saad3577**

1 year, 3 months ago

Unexpired undo can be reused if not enough space in undo tbs ,
The UNDO_RETENTION parameter can only be honored if the current undo tablespace has enough space. If an active transaction requires undo space and the undo tablespace does not have available space, then the system starts reusing unexpired undo space. This action can potentially cause some queries to fail with a "snapshot too old" message.

I go with D,E

upvoted 2 times
- [Removed]**

2 years, 4 months ago

Answer 1. C (always retains only unexpired undo not expired one)
Answer 2. E vs. D

Analysis:
D: Active undo is definitely retained by undo, yes, but it is not the use/configuration of UNDO_RETENTION. Ref: <https://blog.oracle48.nl/oracle-database-undo-space-explained/>

E: Undo retention does specify for how long Oracle must keep unexpired undo, yes, there is (attempt) in the sentence but attempt does mean a task for a computer, unlike human a computer attempts are accurate they don't pause and say I don't want to attempt this like us when we get afraid.

So,
As the question demands configuration and use of UNDO_RETENTION with GUARANTEED RETENTION,

The true answers are C and E
upvoted 1 times

  **[Removed]** 2 years, 4 months ago

Answer 1. C (always retains only unexpired undo not expired one)
Answer 2. E vs. D

Analysis:
D: Active undo is definitely retained by undo, yes, but it is not the use of UNDO_RETENTION.



E: Undo retention does specify for how long Oracle must keep unexpired undo, yes, there is (attempt) in the sentence but attempt does mean a task for a computer, unlike human a computer attempts are accurate they don't pause and say i don't want to attempt this like us when we get afraid.

So,
As the question demands configuration and use of UNDO_RETENTION with GUARANTEED RETENTION,



The true answers are C and E
upvoted 1 times

  **Borislone** 2 years, 6 months ago


BC is correct since undo retention is guarantee, all type of undo has to be retained event if the new undo transaction has to fail
upvoted 1 times

  **Ekos** 3 years, 1 month ago


i choose D and E for the correct answer
upvoted 5 times

  **Chansi** 3 years, 5 months ago

the question is undo retention with guaranteed retention
so its C and D as unexpired undo in this case is always retained
upvoted 4 times

  **NowOrNever** 3 years, 5 months ago



Correct Answer: CE
Wrong Answer: AB
A - keep only unexpired UNDO
B - not all
D - true, but not described by UNDO_RETENTION, only the confirmed UNDO
upvoted 3 times

  **ogdru** 2 years, 10 months ago

The question is ask about the configuration and use of UNDO_RETENTION with GURANTEED so I sure D is correct answer.
upvoted 1 times

  **NiciMilo** 3 years, 5 months ago

I think should be DE
upvoted 3 times

  **ama** 3 years, 6 months ago

I bet on A, C

When automatic undo management is enabled, there is always a current undo retention period, which is the minimum amount of time that Oracle Database attempts to retain old undo information before overwriting it. Old (committed) undo information that is older than the current undo retention period is said to be expired. Old undo information with an age that is less than the current undo retention period is said to be unexpired.
upvoted 2 times

  **you1234** 3 years, 6 months ago

this is real question asked in the exam

the one more question is

which two statments are true about UNDO and UNDO tablespace?

- A) There can be only one UNDO tablespace created in database.
 B) UNDO segments are owned by SYSBACKUP.
 C) UNDO Segments are owned by SYSTEM
 D) An instance will crash if the active undo tablespace is lost.
 E) An undo tablespace may be owned by only one instance.


correct answer is A & D. any suggestion for this answer? this question is asked in real exam

upvoted 1 times

  **Rivaldo11** 3 years, 6 months ago

D and E are correct.

upvoted 7 times

  **ama** 3 years, 6 months ago

E is correct

UNDO_RETENTION

(seconds)

Once a transaction commits the undo data for that transaction stays in the undo tablespace until space is required in which case it will be overwritten.



committed undo information (unexpired) :

also known as unexpired undo, required to support undo_retention interval

after undo_retention period or undo tablespace space pressure unless guaranteed option is set (see below)

http://datadisk.org.uk/html_docs/oracle/undo.htm

upvoted 2 times

  **ama** 3 years, 6 months ago

C is correct

Retention Guarantee

Oracle Database 10g lets you guarantee undo retention. When you enable this option, the database never overwrites unexpired undo data i.e. undo data whose age is less than the undo retention period. This option is disabled by default, which means that the database can overwrite the unexpired undo data in order to avoid failure of DML operations if there is not enough free space left in the undo tablespace.

satya-dba.blogspot.com/2009/09/undo-tablespace-undo-management.html

upvoted 2 times

  **adoptc94** 3 years, 5 months ago

This question is just tricky and not clear at all...

C is correct but only if you're using guarantee retention.

Normally the statement 'Unexpired UNDO is always retained' is not true! If you're not using guarantee retention unexpired undo can be overwritten, which means it is not always retained

For me two statements that are 100% true in all cases are D and E

- active undo is always retained no matter what

- undo_retention is the parameter that specifies how long to retain unexpired undo

But answer C would be true in this case too, so I wouldn't know what two answers to pick... considering the question is about using undo with guarantee retention C&E could be true but then again D is also true no matter how you look at it

upvoted 2 times

  **Phat** 3 years ago

the undo guarantee and retention are about unexpired undo, not talk about active undo. And the question is asking clearly about the use of these 2 parameters.

upvoted 1 times

  **elcaso34** 3 years, 6 months ago

I think that:

Answer A is correct

UNDO_RETENTION specifies (in seconds) how long the confirmed UNDO information will be retained:

UNCONFIRMED UNDO - ACTIVE.

CONFIRMED UNDO - NOT EXPIRED (WITHIN THE RETENTION INTERVAL)

EXPIRED (OUTSIDE THE RETENTION INTERVAL)

Answer C is correct, as it has the guarantee of retention, the UNDO information of the confirmed but not expired transactions will be kept until it reaches the UNDO_RETENTION value.

upvoted 1 times

Table ORDER_ITEMS contains columns ORDER_ID, UNIT_PRICE and QUANTITY, of data type NUMBER.

Examine these SQL statements:

Statement 1:

```
SELECT MAX(unit_price * quantity) `Maximum Order`  
FROM order_items;
```

Statement 2:

```
SELECT MAX(unit_price * quantity) `Maximum Order`  
  
FROM order_items -  
GROUP BY order_id;  
Which two statements are true?
```

- A. Statement 1 returns only one row of output.
- B. Statement 2 returns only one row of output.
- C. Both statements will return NULL if either UNIT_PRICE or QUANTITY contains NULL.
- D. Both the statements give the same output.
- E. Statement 2 may return multiple rows of output.

Correct Answer: AE

- danito** Highly Voted

3 years, 8 months ago

A E because <https://docs.oracle.com/javadb/10.8.3.0/ref/rrefsqlj27781.html>
The Expression can contain multiple column references or expressions, but it cannot contain another aggregate or subquery. It must evaluate to a built-in data type. You can therefore call methods that evaluate to built-in data types. (For example, a method that returns a java.lang.Integer or int evaluates to an INTEGER.) If an expression evaluates to NULL, the aggregate skips that value.
upvoted 15 times
- Ekos**

3 years, 1 month ago

i agree
upvoted 2 times
- [Removed]** Most Recent

2 years, 4 months ago

A and E (max function gives only one output in st.1 / when group by order_id is used, max function is used for every order_id
upvoted 1 times
- flaviogcmelo**

2 years, 9 months ago

A , E are the correct statements. I've tested on this same sample schema.
upvoted 2 times
- guestart**

2 years, 10 months ago

Yes, A & E are correct answers.

The followding demo is what I have tested on LiveSQL.

CREATE TABLE order_items (order_id number, unit_price number, quantity number);

INSERT INTO order_items VALUES (1, 234.35, 300);
INSERT INTO order_items VALUES (2, 95.10, 20);
INSERT INTO order_items VALUES (3, 150.75, 400);
INSERT INTO order_items VALUES (4, 50.65, NULL);
INSERT INTO order_items VALUES (5, 20.15, 150);
INSERT INTO order_items VALUES (6, 80.50, 70);

COMMIT;

SELECT MAX(unit_price * quantity) "Maximum Order"
FROM order_items;



Maximum Order



70305



SELECT order_id, MAX(unit_price * quantity) "Maximum Order"
FROM order_items



ORDER_ID Maximum Order

6 5635
1 70305
2 1902
4 -
5 3022.5
3 60300
upvoted 4 times

  **Mandar79** 3 years, 4 months ago
A - Correct in all scenarios
E - Correct if there is more than 1 row
C - Correct if there is only one row which has either UNIT_PRICE or QUANTITY contains NULL
upvoted 3 times



  **spike_ge** 2 years, 5 months ago
C is wrong. Statement 2 may result in multiple rows and some of them may not be NULL
upvoted 2 times



  **Yanal98** 1 year, 1 month ago
C is incorrect. Statement one will not give null although if there are null values within either the "unit_price" or the "quantity".
upvoted 1 times

  **Yanal98** 1 year, 1 month ago
You can try this using this statement on HR Schema:

select max(commission_pct*employee_id) from employees group by first_name ;

And this manipulates Statement 2 execution in C.
upvoted 1 times

  **ama** 3 years, 6 months ago
A, E seems ok
upvoted 4 times

  **you1234** 3 years, 7 months ago
any body pass this exam ? Please suggest
upvoted 3 times

Which three files are used by conventional path SQL*Loader when the TABLE option is not specified? (Choose three.)

- A. dump files
- B. control files
- C. password files
- D. bad files
- E. input files

Correct Answer: *BDE*



Community vote distribution

BDE (100%)

  **Mandar79** Highly Voted 3 years, 4 months ago

BDE seems to be correct, can someone confirm.
SQLldr CONTROL=sample.ctl, LOG=sample.log, BAD=baz.bad, DATA=etc.dat
USERID=scott/tiger, ERRORS=999, LOAD=2000, DISCARD=toss.dsc,
DISCARDMAX=5

https://docs.oracle.com/cd/B19306_01/server.102/b14215/ldr_params.htm
upvoted 10 times

  **NowOrNever** 3 years, 4 months ago
BDE - I agree
https://docs.oracle.com/cd/B28359_01/server.111/b28319/ldr_modes.htm#i1011474
upvoted 1 times

  **auwia** Most Recent 6 months, 2 weeks ago

Selected Answer: BDE
Provided answer are corrects.
upvoted 1 times

  **aediansyahsari** 1 year, 8 months ago

Selected Answer: BDE
correct
upvoted 1 times

  **[Removed]** 2 years, 4 months ago

BDE 100%
upvoted 1 times

Which three statements are true about the naming methods and their features supported by Oracle database used to resolve connection information? (Choose three.)

- A. Local Naming requires setting the TNS_ADMIN environment variable on the client side.
- B. A client can connect to an Oracle database instance even if no client side network admin has been configured.
- C. Directory Naming can be used if Connect-Time Failover is required.
- D. Easy Connect supports TCP/IP and SSL.
- E. Local naming can be used if Connect-Time Failover is required.
- F. Directory Naming requires setting the TNS_ADMIN environment variable on the client side.

Correct Answer: ADF

Community vote distribution

BCE (100%)

- ismoil

3 months ago

A,B,C i have checked because local naming relies on tnsnames.ora file

upvoted 1 times
- guretto

3 months ago

MOS note "Does EZCONNECT work with SSL / TCPS? ORA-12547 is received with 18c and below. (Doc ID 2502638.1)" states that from 19c and above, we now have an option for TCPS (TLS) with Easy Connect Plus.

upvoted 1 times
- Darkseid1231

6 months, 3 weeks ago

I will go for ABE

upvoted 3 times
- RaNik69

8 months, 1 week ago

Selected Answer: BCE

BCE correct

upvoted 1 times
- jonsnoow

8 months, 3 weeks ago

I think BCE.
B - Easy Connect requires no client-side configuration.
C - Local Naming supports the advanced connection option Connect-time failover.
E - Directory Naming supports Connect-time failover.

upvoted 2 times
- piontk

10 months, 2 weeks ago

ABF I think

upvoted 1 times
- dianiss1050

10 months, 2 weeks ago

D is wrong because Easy Connect only support TCP/IP

upvoted 2 times
- guretto

2 months, 1 week ago

before 19c

upvoted 1 times

Which two statements are true about User Authentication in an Oracle Database? (Choose two.)

- A. Password authentication must be used for system-privileged administrative users.
- B. Password File authentication must be used for system-privileged administrative users.
- C. Operating System authentication may be used for system-privileged administrative users.
- D. Password File authentication is supported for any type of database user.
- E. REMOTE_LOGIN_PASSWORDFILE must be set to exclusive to permit password changes for system-privileged administrative users.

Correct Answer: AC

Community vote distribution

CD (67%)

AC (33%)

🗲️ 👤 **jonsnoow** 4 months, 1 week ago

I think C and E.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/getting-started-with-database-administration.html#GUID-662ACF2E-546A-4EF0-89EA-35D4633A0996>

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/REMOTE_LOGIN_PASSWORDFILE.html#GUID-6619299E-95E8-4821-B123-3B5899F046C7

upvoted 1 times

🗲️ 👤 **zouve** 5 months, 2 weeks ago

Selected Answer: CD

Statements A and B are incorrect. Password authentication and Password File authentication are both supported for system-privileged administrative users, but they are not required. Operating System authentication is also supported for system-privileged administrative users.

Statement E is incorrect. REMOTE_LOGIN_PASSWORDFILE must be set to exclusive to permit password changes for all users, not just system-privileged administrative users.

upvoted 1 times

🗲️ 👤 **musafir** 6 months ago

Selected Answer: AC

A. Password authentication must be used for system-privileged administrative users.

Correct, can't find any reference where password authentication can be disabled. Some type of authentication has to be enabled.

B. Password File authentication must be used for system-privileged administrative users. Wrong - Can use different methods.

C. Operating System authentication may be used for system-privileged administrative users. Correct

D. Password File authentication is supported for any type of database user. Wrong - "Password files are used to authenticate database administrators."

E. REMOTE_LOGIN_PASSWORDFILE must be set to exclusive to permit password changes for system-privileged administrative users.

Wrong - Can also be set to NONE which will allow password changes.

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/REMOTE_LOGIN_PASSWORDFILE.html#GUID-6619299E-95E8-4821-B123-3B5899F046C7

<https://docs.oracle.com/en/database/oracle/oracle-database/19/dbseg/configuring-authentication.html#GUID-30E91ECB-1DA1-4C9E-B630-3C9E541CD915>

upvoted 1 times

🗲️ 👤 **mamadu** 6 months, 2 weeks ago

Selected Answer: CD



I believe.

upvoted 1 times

Which is the default column or columns for sorting output from compound queries using SET operators such as INTERSECT in a SQL statement?

- A. the first NUMBER column in the first SELECT of the compound query
- B. the first NUMBER or VARCHAR2 column in the last SELECT of the compound query
- C. the first column in the last SELECT of the compound query
- D. the first column in the first SELECT of the compound query
- E. the first VARCHAR2 column in the first SELECT of the compound query

Correct Answer: *D*

  **Guhborges** 9 months, 4 weeks ago

D. the first column in the first SELECT of the compound query.

When using SET operators such as INTERSECT in a SQL statement, the default column or columns for sorting output is based on the column order in the first SELECT statement of the compound query. Specifically, the first column in the first SELECT statement is used to sort the results. If you want to change the sort order or use a different column for sorting, you can use an ORDER BY clause at the end of the compound query.

upvoted 1 times

Examine the description of the PRODUCT_STATUS table:

Name	Null?	Type
-----	-----	-----
PROD_ID	NOT NULL	NUMBER(2)
STATUS	NOT NULL	VARCHAR2(15)

The STATUS column contains the values 'IN STOCK' or 'OUT OF STOCK' for each row.

Which two queries will execute successfully? (Choose two.)

- A.

```
SELECT prod_id "CURRENT AVAILABILITY" || q('s not available)' FROM
product_status WHERE status = 'OUT OF STOCK';
```
- B.

```
SELECT prod_id q's not available" FROM product_status WHERE status = 'OUT OF
STOCK';
```
- C.

```
SELECT prod_id || q('s not available)' 'CURRENT AVAILABILITY' FROM
product_status WHERE status = 'OUT OF STOCK';
```
- D.

```
SELECT prod_id || q('s not available)' FROM product_status WHERE status = 'OUT
OF STOCK';
```
- E.



```
SELECT prod_id || q"'s not available" FROM product_status WHERE status = 'OUT OF
STOCK';
```
- F.



```
SELECT prod_id || q('s not available)' "CURRENT AVAILABILITY" FROM
product_status WHERE status = 'OUT OF STOCK';
```



Correct Answer: BC

Community vote distribution

BC (100%)

  **canijho** Highly Voted 10 months, 2 weeks ago
A - Bad alias, ORA-00923: FROM keyword not found where expected
B - Bad syntax, ORA-01756: quoted string not properly terminated
C - Bad alias, ORA-00923: FROM keyword not found where expected
D --OK
E - Bad syntax, ORA-00904: "Q": invalid identifier
F --OK
upvoted 5 times

  **auwia** Most Recent 6 months, 3 weeks ago
Selected Answer: DF
D and F are correct.
upvoted 1 times

  **BeomK** 9 months ago
D,F is correct
upvoted 3 times

Which two statements are true about UNDO and UNDO tablespaces? (Choose two.)

- A. There can be only one UNDO tablespace created in a database.
- B. UNDO segments are owned by SYSTEM.
- C. UNDO segments are owned by SYSBACKUP.
- D. An UNDO tablespace may be owned by only one instance.
- E. An instance will crash if the active undo tablespace is lost.

Correct Answer: *BD*



Community vote distribution

DE (100%)

  **auwia** 6 months, 3 weeks ago


Selected Answer: DE

Agree with D and E.
upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: DE



D (RAC) E are correct
upvoted 1 times

  **markwu01** 9 months, 1 week ago

When an undo tablespace is created, the corresponding undo segments are created as well. The ownership of the undo segments is then given to the user who creates the undo tablespace. Therefore, the ownership of the undo segments belongs to the user, not to the SYS or SYSTEM user. so the correct answer are D and E
upvoted 2 times

  **j_tw** 9 months, 1 week ago

I know that Undo segment is owned by SYS user, not SYSTEM. B - false
upvoted 2 times

  **piontk** 10 months, 2 weeks ago

D and E are the correct answers
upvoted 2 times

You must create a tablespace of non-standard block size in a new file system and plan to use this command:

```
CREATE TABLESPACE ns_tbs
  DATAFILE '/u02/oracle/data/nstbs_f01.dbf'
  SIZE 100G
  BLOCKSIZE 32K;
```

The standard block size is 8k but other non-standard block sizes will also be used.



Which two are requirements for this command to succeed? (Choose two.)

- A. DB_32K_CACHE_SIZE must be less than DB_CACHE_SIZE.
- B. DB_32K_CACHE_SIZE must be set to a value that can be accommodated in the SGA.
- C. DB_32K_CACHE_SIZE should be set to a value greater than DB_CACHE_SIZE.
- D. The /u02 file system must have at least 100g space for the datafile.
- E. The operating system must use a 32k block size.
- F. DB_CACHE_SIZE must be set to a size that is smaller than DB_32K_CACHE_SIZE.

Correct Answer: BD

Community vote distribution

BD (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: BD

Provided answer are corrects.

upvoted 1 times

Examine the description of the PRODUCTS table:

Name	Null?	Type
-----	-----	-----
PROD_ID	NOT NULL	NUMBER
PROD_NAME		VARCHAR2 (40)
COST		NUMBER (8,2)
RELEASE_DATE		DATE

Which query is valid?

- A. SELECT prod_id, release_date, SUM(cost) FROM products GROUP BY prod_id;
- B. SELECT prod_id, release_date, SUM(cost) FROM products GROUP BY prod_id, release_date;
- C. SELECT prod_id, AVG(MAX(cost)) FROM products GROUP BY prod_id;
- D. SELECT prod_id, MAX(AVG(cost)) FROM products GROUP BY prod_id;

Correct Answer: B

Community vote distribution

B (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: B

Provided answer is correct.
upvoted 1 times

Which two statements are true about the Oracle Data Dictionary? (Choose two.)

- A. All data dictionary view join base tables to dynamic performance views.
- B. Data dictionary base tables can be queried directly.
- C. It is owned by the SYS user.
- D. It is owned by the SYSTEM user.
- E. Data dictionary views are always created with queries that join two or more base tables.

Correct Answer: *BD*

Community vote distribution

BC (100%)

  **ivanadj** Highly Voted 11 months, 1 week ago

Selected Answer: BC

https://docs.oracle.com/cd/B10501_01/server.920/a96524/c05dicti.htm
upvoted 5 times

  **j_tw** 9 months, 1 week ago

I agree that. Data dictionary is owned by SYS user
upvoted 1 times



  **hezzzy** Most Recent 3 weeks ago

BC
https://docs.oracle.com/cd/B10500_01/server.920/a96524/c05dicti.htm
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: BC

B and C are corrects.
upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: BC

b c are correct
upvoted 1 times



Which two Oracle database space management features require the use of locally managed tablespaces? (Choose two.)



- A. Automatic data file extension (AUTOEXTEND)
- B. Server-generated tablespace space alerts
- C. Online segment shrink
- D. Free space management with bitmaps
- E. Oracle Managed Files (OMF)



Correct Answer: CD



Community vote distribution



CD (100%)



  **Darkseid1231** 7 months ago
D E is the correct answer
upvoted 1 times

  **hotink57** 7 months, 2 weeks ago
Selected Answer: CD
A - false, as it can be used with both types
B - false
C - true, as it requires locally managed tablespaces with automatic segment space management (ASSM)
D - true, as the bitmap structure is stored in the data files of the tablespace
E - false
upvoted 2 times

  **Darkseid1231** 6 months, 3 weeks ago
This is correct
upvoted 1 times

  **piontk** 10 months, 2 weeks ago
AF is correct
upvoted 2 times

  **piontk** 5 months, 3 weeks ago
CD are the correct answers, after investigation
upvoted 1 times

  **piontk** 10 months, 2 weeks ago
AE* is the correct
upvoted 2 times

Which three statements are true about connection strings and service names used to connect to an Oracle database instance? (Choose three.)

- A. A connection string must include the SID of a database instance.
- B. A connection string including a service name must be defined in the tnsnames.ora file.
- C. A single database instance can support connections for multiple service names.
- D. A single connection string can refer to multiple database instances.
- E. A service name is created by a listener.
- F. Different connection strings in the same tnsnames.ora file can contain the same service name, host and port parameters.

Correct Answer: ABE

Community vote distribution



CDF (100%)

  bertchiang Highly Voted 11 months ago

I think the correct answer is CDF
upvoted 5 times

  j_tw 9 months, 2 weeks ago



Why not B?
upvoted 1 times

  RaNik69 8 months, 1 week ago

B is wrong because you can connect with easy connect without tnsnames
upvoted 2 times

  auwia Most Recent 6 months, 3 weeks ago

Selected Answer: CDF
C,D,F in my opinion.
upvoted 2 times

  RaNik69 8 months, 1 week ago

C D F are correct. B is wrong
upvoted 3 times

Which three statements are true about using SQL*Plus? (Choose three.)

- A. It has both command-line and graphical user interfaces (GUI).
- B. It must be downloaded from the Oracle Technology Network (OTN).
- C. It can run Recovery Manager (RMAN) commands.
- D. It has its own commands that are separate from any SQL statements.
- E. It can run scripts entered at the SQL prompt.
- F. It can run scripts passed to it by a shell script.

Correct Answer: ADE



Community vote distribution

DEF (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: DEF

Provided answer are wrong.
upvoted 2 times

  **Juv07** 7 months, 4 weeks ago

D E F is a correct answers
upvoted 2 times

18. The orders table has a column ORDER_DATE of data type DATE.

The default display format for a date is DD-MON-RR.



Which two WHERE conditions demonstrate the correct usage of conversion functions? (Choose two.)



- A. WHERE order_date > TO_DATE(ADD_MONTHS(SYSDATE, 6), 'MON DD YYYY')
- B. WHERE order dare > TO DATE('JUL 10 2018', 'MON DD YYYY')
- C. WHERE TO_CHAR(order date, 'MON DD YYYY') = 'JAN 20 2019'
- D. WHERE order_date > TO_CHAR(ADD_MONTHS(SYSDATE, 6), 'MON DD YYYY')
- E. WHERE order_date IN (TO_DATE('Oct 21 2018', 'Mon DD YYYY'), TO_CHAR('Nov 21 2018', 'Mon DD YYYY'))



Correct Answer: AC

Community vote distribution



BC (100%)

  **auwia** 6 months, 3 weeks ago
All are wrong! Someone should fix the text :-)
upvoted 1 times

  **RaNik69** 8 months, 1 week ago
Selected Answer: BC
b c are correct if ORDER DARE >> ORDER_DATE
upvoted 1 times

  **BeomK** 9 months ago
B,C is correct.

A. false - ORA-01843: not a valid month
B. true - TO DATE >> TO_DATA : wrong entry?? - Normal operation after modification
C. true - Normal operation after modification
D. false - ORA-01858: a non-numeric character was found where a numeric was expected
E. false - ORA-01722: invalid number
upvoted 2 times

  **j_tw** 9 months, 2 weeks ago
I think C & E are correct.

A(false) - ADD_MONTHS(SYSDATE, 6) is already DATE type, so it will be an error to use TO_DATE function.



B(false) - same as A

C(true)

D(false) - It can't be possible to compare diffrent data type. (date > string)

E(maybe true)

Am I wrong?
upvoted 1 times

  **j_tw** 9 months ago
NVM this comment. I perfectly misunderstood about TO_DATE function. BC may be correct
upvoted 1 times

Which three statements are true about the Automatic Diagnostic Repository (ADR)? (Choose three.)

- A. It can be used for problem diagnosis of a database when that database's instance is down.
- B. The ADR base is specified in the DIAGNOSTIC_DEST database parameter.
- C. It is held inside an Oracle database schema.
- D. It is a file-based repository held outside any database.
- E. It is only used for Oracle Database diagnostic information.

Correct Answer: ABE

Community vote distribution

ABD (100%)


  j_tw Highly Voted 9 months, 1 week ago

Selected Answer: ABD

ABD



A - The Automatic Diagnostic Repository (ADR) is a directory structure that is stored outside of the database. It is therefore available for problem diagnosis when the database is down.
B - They are replaced by the initialization parameter DIAGNOSTIC_DEST, which identifies the location of the ADR.
D - The data is then stored in the Automatic Diagnostic Repository (ADR)—a file-based repository outside the database—where it can later be retrieved by incident number and analyzed.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/diagnosing-and-resolving-problems.html#GUID-43A0C671-CE6F-492F-9032-A010E96DC257>
upvoted 5 times

  auwia Most Recent 6 months, 2 weeks ago

Selected Answer: ABD

Provided answer are wrong:
The Automatic Diagnostic Repository (ADR) is a directory structure that is stored outside of the database. It is therefore available for problem diagnosis when the database is down.
The ADR root directory is known as ADR base. Its location is set by the DIAGNOSTIC_DEST initialization parameter.
Ref:
<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/diagnosing-and-resolving-problems.html#GUID-EA97EEF6-9207-4536-B808-3B91DACA7AD6>
upvoted 1 times

  RaNik69 8 months, 1 week ago

Selected Answer: ABD

A B D are correct
upvoted 1 times

Which two statements are true about a self join? (Choose two.)

- A. It can be an inner join.
- B. It can be a left outer join.
- C. It must be an equijoin.
- D. The join key column must have an index.
- E. It must be a full outer join.

Correct Answer: AB



Community vote distribution

AB (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: AB

Provided answer are correct, self join can be any kind of join! ;-)
upvoted 1 times

  **mamadu** 6 months, 2 weeks ago

I think this is true because we can use other join conditions and not necessarily an equality join (=).
You can also use any outer join to join the same table. (most probably using a non-equijoin operator)
upvoted 1 times

  **Darkseid1231** 6 months, 3 weeks ago

A C seems okay
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

"must" is not welcomed in this case! :-)
upvoted 1 times

  **Darkseid1231** 6 months, 2 weeks ago

in a self join the table is join with itself. Equijoin is the most common join used.
upvoted 1 times

Which statement is true about database links?

- A. Private database link creation requires the same user to exist in both the local and the remote databases.
- B. A database link created in a database allows a connection from that database's instance to the target database's instance, but not vice versa.
- C. A public database link can be used by a user connected to the local database instance to connect to any schema in the remote database instance.
- D. A public database link can be created only by SYS.
- E. A database link can be created only between two Oracle databases.

Correct Answer: C

Community vote distribution



B (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: B

Which statement is true about database links?

- A. Private database link creation requires the same user to exist in both the local and the remote databases. --> False.
 - B. A database link created in a database allows a connection from that database's instance to the target database's instance, but not vice versa. --> True
 - C. A public database link can be used by a user connected to the local database instance to connect to any schema in the remote database instance. --> False, you don't see any schema.
 - D. A public database link can be created only by SYS. --> False
 - E. A database link can be created only between two Oracle databases. --> False, you can create a db link to itself! :-)
- upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: B

b is correct

upvoted 2 times

  **j_tw** 9 months, 1 week ago

B

PUBLIC does not means a user can use any schema. It means other users can use db link what you made it.

When you create a private or public database link, you can determine which schema on the remote database the link will establish connections to by creating fixed user, current user, and connected user database links.

<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-a-distributed-database.html#GUID-B5950167-35F8-4C88-B063-382E0290CB53>

upvoted 2 times

Which two queries execute successfully? (Choose two.)

- A. SELECT NULLIF(NULL, 100) FROM DUAL;
- B. SELECT NULLIF(100, 'A') FROM DUAL;
- C. SELECT COALESCE(100, 'A') FROM DUAL;
- D. SELECT COALESCE(100, NULL, 200) FROM DUAL;
- E. SELECT NULLIF(100, 100) FROM DUAL;

Correct Answer: AC



Community vote distribution

DE (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: DE

D and E are correct.
upvoted 2 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: DE

D E is correct. Easy to test
upvoted 2 times

  **BeomK** 9 months ago

D,E is correct

- A. ORA-00932: inconsistent datatypes: expected - got CHAR
- B. ORA-00932: inconsistent datatypes: expected NUMBER got CHAR
- C. ORA-00932: inconsistent datatypes: expected NUMBER got CHAR
- D. SQL> SELECT COALESCE(100, NULL, 200) FROM DUAL;



COALESCE(100,NULL,200)

100

E. SQL> SELECT NULLIF(100, 100) FROM DUAL;



NULLIF(100,100)

upvoted 2 times

  **Bolillo** 11 months, 1 week ago

Selected Answer: DE

Tested both
upvoted 3 times

  **ivanadj** 11 months, 2 weeks ago

D and E are correct
SQL> SELECT COALESCE(100, NULL, 200) FROM DUAL;

COALESCE(100,NULL,200)

100

SQL> SELECT NULLIF(100, 100) FROM DUAL;

NULLIF(100,100)


upvoted 3 times

Which three statements are true about data block storage in an Oracle Database? (Choose three.)

- A. An index block can contain row data.
- B. A data block header is of a fixed length.
- C. Row data is stored starting at the end of the block.
- D. A table block must always contain row data.
- E. A block header contains a row directory pointing to all rows in the block.

Correct Answer: BCE

Community vote distribution



☐ **zouve** 5 months ago

Header (Common and Variable)

The header contains general block information, such as the block address and the type of segment; for example, data, index, or rollback. While some block overhead is fixed in size, the total block overhead size is variable. On average, the fixed and variable portions of data block overhead total 84 to 107 bytes.

upvoted 2 times

☐ **auwia** 6 months, 2 weeks ago

Selected Answer: ACE

A. An index block can contain row data. --> True.

ref: https://docs.oracle.com/cd/A57673_01/DOC/server/doc/SCN73/ch3.htm

B. A data block header is of a fixed length. --> False.

ref: https://www.dba-oracle.com/t_data_block_format.htm

C. Row data is stored starting at the end of the block. --> True.

ref: <https://docs.oracle.com/en/database/oracle/oracle-database/19/cncpt/logical-storage-structures.html>

D. A table block must always contain row data. --> False.

ref: <https://richardfoote.wordpress.com/category/block-dumps/>

E. A block header contains a row directory pointing to all rows in the block. --> True.

ref: <https://medium.com/geekculture/explain-about-oracle-database-data-block-interview-a-be137c52fe4d>

upvoted 4 times

☐ **auwia** 6 months, 2 weeks ago

A. An index block can contain row data. --> True, "Row Data" portion : This portion of the block contains table or index data.

B. A data block header is of a fixed length. --> False, the header of a data block has two areas a "common" area and a "variable" area. In the common area you have a fixed-size block header. ==> But it's different to say, I guess.

C. Row data is stored starting at the end of the block. --> True, "Space Management in Data Blocks" -> As the database fills a data block from the bottom up, the amount of free space between the row data and the block header decreases.

D. A table block must always contain row data. --> False, it also contains info about tab and col: If we look at a partial block dump of one of the table blocks.

E. A block header contains a row directory pointing to all rows in the block. --> True, a block header includes the information about block type information, whether it is a row block or an index block, and the information of tables that have the rows in it, and the row directory, means the addresses of each row in this block.

upvoted 1 times

☐ **auwia** 6 months, 3 weeks ago

Selected Answer: BCE

Provided answer are corrects!

upvoted 3 times

☐ **auwia** 6 months, 2 weeks ago

I'm sorry I was wrong. please remove this full comment with vote.

upvoted 1 times

Which two are benefits of external tables? (Choose two.)

- A. They support DELETES which transparently deletes records in the file system as if they were table rows.
- B. They can be queried, transformed, and joined with other tables without having to load the data first.
- C. They support UPDATES which transparently updates records in the file system as if they were table rows.
- D. The results of a complex join or aggregating function or both can be unloaded to a file for transportation to another database.
- E. They can be queried while the database is in the MOUNT state like dynamic performance views.

Correct Answer: *BD*

Community vote distribution

BD (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: BD

Provided answer are corrects.
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Only select is allowed on external tables + DB has to be in "open" state.
upvoted 1 times

Examine the description of the MEMBERS table:

Name	Null?	Type
-----	-----	-----
MEMBER_ID	NOT NULL	VARCHAR2 (6)
FIRST_NAME		VARCHAR2 (50)
LAST_NAME	NOT NULL	VARCHAR2 (50)
ADDRESS		VARCHAR2 (50)
CITY		VARCHAR2 (25)

Examine the partial query:

SELECT city, last_name LNAME FROM members ...;

You want to display all cities that contain the string an. The cities must be returned in ascending order, with the last names further sorted in descending order.

Which two clauses must you add to the query? (Choose two.)

- A. ORDER BY 1, 2
- B. WHERE city = '%AN%'
- C. WHERE city LIKE '%AN%'
- D. ORDER BY last_name DESC, city ASC
- E. WHERE city IN ('%AN%')
- F. ORDER BY 1, LNAME DESC

Correct Answer: CD


Community vote distribution

CF (100%)

 **auwia** 6 months, 3 weeks ago

Selected Answer: CF

The base! Part V :-)
upvoted 3 times

 **RaNik69** 8 months, 1 week ago

Selected Answer: CF

c f are correct
upvoted 2 times

 **Juv07** 11 months, 1 week ago

The good answers is C,F
upvoted 3 times

 **PrasenC** 11 months, 1 week ago

If the order is not specified, ASC is the default.
upvoted 1 times

 **PrasenC** 11 months, 1 week ago

Selected Answer: CF

Refer:
<https://docs.oracle.com/javadb/10.6.2.1/ref/rrefsqlj13658.html#:~:text=If%20the%20order%20is%20not%20specified%2C%20ASC%20is%20the%20default.&text=Specifies%20that%20the%20results%20should%20be%20returned%20in%20descending%20order.&text=Specifies%20that%20NULL%20values%20should%20be%20returned%20before%20non%2DNULL%20values.&text=Specifies%20that%20NULL%20values%20should%20be%20returned%20after%20non%2DNULL%20values>
If the order is not specified, ASC is the default.
upvoted 3 times

Which three statements are true about sequences in a single instance Oracle database? (Choose three.)

- A. Sequences can always have gaps.
- B. Two or more tables cannot have keys generated from the same sequence.
- C. A sequence's unallocated cached values are lost if the instance shuts down.
- D. A sequence can issue duplicate values.
- E. A sequence can only be dropped by a DBA.
- F. A sequence number that was allocated can be rolled back if a transaction fails.

Correct Answer: BCF



Community vote distribution

ACD (100%)

  **auwia** 6 months, 3 weeks ago


Selected Answer: ACD

Agree with A, C and D.
upvoted 2 times

  **RaNik69** 8 months, 1 week ago



Selected Answer: ACD

ACD are correct
upvoted 2 times



  **piontk** 10 months, 2 weeks ago

Selected Answer: ACD

Answer is ACD
upvoted 2 times

  **piontk** 10 months, 2 weeks ago

B is wrong -> "Sequence numbers are generated independently of tables, so the same sequence can be used for one or for multiple tables"
<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-SEQUENCE.html>
upvoted 1 times

  **ivanadj** 11 months, 1 week ago

Selected Answer: ACD

<https://www.examttopics.com/discussions/oracle/view/20160-exam-1z0-071-topic-2-question-27-discussion/>
upvoted 2 times

Your database instance is started with an SPFILE.

A PFILE is also available.

You execute this command:

```
ALTER SYSTEM SET DB_CACHE_SIZE=100K;
```

Where is the value changed?

- A. in the SPFILE, PFILE, and memory
- B. in the SPFILE and in memory
- C. only in the SPFILE
- D. only in memory
- E. in the SPFILE and PFILE

Correct Answer: B



Community vote distribution

B (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: B

Correct answer is B.
upvoted 1 times

  **Darkseid1231** 7 months ago

I will go for D as the correct answer.
Scope=spfile is not specify.
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

The scope is intrinsic because you started the instance by the spfile, so the scope is already set.
upvoted 1 times

  **RaNik69** 9 months ago

Selected Answer: B

When you issue the ALTER SYSTEM statement while connected to a non-CDB or a CDB root, the scope depends on whether you started up the database using a traditional plain-text parameter file (pfile) or server parameter file (spfile).

MEMORY indicates that the change is made in memory, takes effect immediately, and persists until the database is shut down. If you started up the database using a parameter file (pfile), then this is the only scope you can specify.

SPFILE indicates that the change is made in the server parameter file. The new setting takes effect when the database is next shut down and started up again. You must specify SPFILE when changing the value of a static parameter that is described as not modifiable in Oracle Database Reference.

BOTH indicates that the change is made in memory and in the server parameter file. The new setting takes effect immediately and persists after the database is shut down and started up again.

If a server parameter file was used to start up the database, then BOTH is the default. If a parameter file was used to start up the database, then MEMORY is the default, as well as the only scope you can specify.
upvoted 1 times

You want to apply the principle of Least Privilege in all your live databases.

One of your requirements is to revoke unnecessary privileges from all users who have them using Privilege Analysis.

Which two are types of analyses can be done using the DBMS_PRIVILEGE_CAPTURE package? (Choose two.)

- A. analysis of privileges that a user has on other schema’s objects
- B. analysis of privileges that a user has on their own schema objects
- C. analysis of all privileges used by the sys user.
- D. analysis of privileges granted indirectly to a role that are then used by a user who has been granted that role
- E. analysis of privileges granted directly to a role that are then used by a user who has been granted that role

Correct Answer: AE

Community vote distribution



zouve 5 months, 1 week ago

Selected Answer: DE

DE for me
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: AD

A. analysis of privileges that a user has on other schema’s objects. --> True, it makes sense.
B. analysis of privileges that a user has on their own schema objects. --> False, it makes no sense. The owner is the owner why do you want to check his objects? :-)
C. analysis of all privileges used by the sys user. --> False, SYS is the only user excluded by three tool.
D. analysis of privileges granted indirectly to a role that are then used by a user who has been granted that role. --> True, it makes sense.
E. analysis of privileges granted directly to a role that are then used by a user who has been granted that role. --> False, it makes no sense.
upvoted 1 times

RaNik69 8 months, 1 week ago

Selected Answer: DE

D E are correct
upvoted 1 times

j_tw 9 months ago

I think A also can be correct...
upvoted 1 times

ivanadj 11 months, 1 week ago

Selected Answer: DE

<https://www.examttopics.com/discussions/oracle/view/24044-exam-1z0-082-topic-1-question-63-discussion/>
upvoted 1 times

Examine the description of the SALES table:

Name	Null?	Type
-----	-----	-----
PRODUCT_ID	NOT NULL	NUMBER(10)
CUSTOMER_ID	NOT NULL	NUMBER(10)
TIME_ID	NOT NULL	DATE
CHANNEL_ID	NOT NULL	NUMBER(5)
PROMO_ID	NOT NULL	NUMBER(5)
QUANTITY_SOLD	NOT NULL	NUMBER(10,2)
PRICE		NUMBER(10,2)
AMOUNT_SOLD	NOT NULL	NUMBER(10,2)

The SALES table has 55,000 rows.



Examine this statement:



```
CREATE TABLE sales1 (prod_id, cust_id, quantity_sold, price)
AS
SELECT product_id, customer_id, quantity_sold, price
FROM sales
WHERE 1 = 1;
```

Which two statements are true? (Choose two.)

- A. SALES1 is created with 1 row.
- B. SALES1 is created with no rows.
- C. SALES1 has primary KEY and UNIQUE constraints on any selected columns which had those constraints in the SALES table.
- D. SALES1 has NOT NULL constraints on any selected columns which had those constraints in the SALES table.
- E. SALES1 is created with 55,000 rows.

Correct Answer: DE

  **Darkseid1231** 7 months ago
D E is correct
upvoted 3 times

  **j_tw** 9 months, 2 weeks ago
DE
//
D - At using subquery to make table, only 'Not Null' constraint will be copied
upvoted 3 times


Which three statements are true about performing Data Manipulation Language (DML) operations on a view with no INSTEAD OF triggers defined? (Choose three.)

- A. Views cannot be used to query rows from an underlying table if the table has a PRIMARY KEY and the PRIMARY KEY columns are not referenced in the defining query of the view.
- B. The WITH CHECK clause has no effect when deleting rows from the underlying table through the view.
- C. Delete statements can always be done on a table through a view.
- D. Views cannot be used to add or modify rows in an underlying table if the defining query of the view contains the DISTINCT keyword.
- E. Views cannot be used to add rows to an underlying table if the table has columns with NOT NULL constraints lacking default values which are not referenced in the defining query of the view.
- F. Insert statements can always be done on a table through a view.

Correct Answer: CEF

Community vote distribution

BDE (100%)

  **ivanadj** 11 months, 1 week ago

Selected Answer: BDE


<https://www.examttopics.com/discussions/oracle/view/20172-exam-1z0-071-topic-2-question-37-discussion/>
upvoted 1 times

Which three statements are true about indexes and their administration in an Oracle database? (Choose three.)

- A. An index can be created as part of a CREATE TABLE statement.
- B. If a query filters on an indexed column, then it will always be used during execution of the query.
- C. A descending index is a type of function-based index.
- D. A DROP INDEX statement always prevents updates to the table during the drop operation.
- E. A UNIQUE and non-unique index can be created on the same table column.
- F. An INVISIBLE INDEX is not maintained when Data Manipulation Language (DML) is performed on its underlying table.

Correct Answer: ABE

Community vote distribution



☐ **Sajib_Arafat** 3 months, 1 week ago

SQL> create index a1 on x1(b1);

Index created.

SQL> create unique index a2 on x1(b1);

create unique index a2 on x1(b1)

*

ERROR at line 1:

ORA-01408: such column list already indexed

SQL>

upvoted 1 times

☐ **zouve** 5 months, 2 weeks ago

Selected Answer: ACE

ACE for me

upvoted 1 times

☐ **piontk** 5 months, 2 weeks ago

Selected Answer: ACE

A-

CREATE TABLE Books(
ID NUMBER PRIMARY KEY USING INDEX (CREATE INDEX idx_custom_books_id ON books (ID)),
Title VARCHAR2(100) NOT NULL,
Author VARCHAR2(100) NOT NULL,
PublishYear NUMBER NOT NULL
);

C- "Oracle Database treats descending indexes as if they were function-based indexes"

https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-INDEX.html#GUID-1F89BBC0-825F-4215-AF71-7588E31D8BFE_GUID-92F4F0FB-499A-4ED7-8630-B219F8A50B90

E- "When you have multiple indexes on the same set of columns, only one of these indexes can be visible at a time, and any other indexes must be invisible."

"You can create multiple indexes on the same set of columns when at least one of the following index characteristics is different:

The indexes have different uniqueness properties.

You can create both a unique and a non-unique index on the same set of columns."

<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-indexes.html#GUID-B4B30CF4-2B95-44D6-8596-EC2A378251EF>

Correct is ACE, according the above evidences.

upvoted 2 times

☐ **shotcom** 4 months, 1 week ago

you can create both unique and non unique on the same set of columns not on one column

option E is saying on the same table column making it wrong

upvoted 1 times

☐ **auwia** 6 months, 3 weeks ago

Selected Answer: ABC

Correct answer in my opinion.

upvoted 1 times

  **auwia** 6 months, 3 weeks ago



```
E is false:
-- Create table
CREATE TABLE test_table (
id NUMBER,
value VARCHAR2(255)
);

-- Populate the table
INSERT INTO test_table (id, value) VALUES (1, 'One');
INSERT INTO test_table (id, value) VALUES (2, 'Two');
COMMIT;

-- Create a unique index
CREATE UNIQUE INDEX test_table_uniq_idx ON test_table(id);

-- Create a non-unique index
CREATE INDEX test_table_idx ON test_table(id);
```

ORA-01408: such column list already indexed
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Correct answer for me are: A, B, and C; conclusion got by exclusion and priority: D, E, and F are wrong with high priority for me! :-)
upvoted 1 times

  **ivanadj** 11 months, 1 week ago

Selected Answer: ACE

<https://www.examttopics.com/discussions/oracle/view/9439-exam-1z0-071-topic-1-question-300-discussion/>
upvoted 2 times

Which three actions are ways to apply the principle of least privilege? (Choose three.)

- A. setting the O7_DICTIONARY_ACCESSIBILITY parameter to true
- B. setting the REMOTE_OS_AUTHENT parameter to true
- C. revoking execute privilege on UTL_SMTP, UTL_TCP, UTL_HTTP, and UTL_FILE from the SYSTEM user
- D. enabling Unified Auditing
- E. using Access Control Lists (ACLs)
- F. revoking execute privilege on UTL_SMTP, UTL_TCP, UTL_HTTP, and UTL_FILE from the public user

Correct Answer: CDE

Community vote distribution

DEF (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: DEF

Very good link, thanks dianiss1050. Yes correct answer are DEF.



False answer:

A - Protect the data dictionary. When parameter O7_DICTIONARY_ACCESSIBILITY is FALSE, people who has ANY TABLE privileges are restricted from accessing data dictionary-based tables. The default value of the parameter is FALSE and should not be changed.

B - Restrict remote database authentication. REMOTE_OS_AUTHENT specifies whether remote clients will be authenticated with the value of the OS_AUTHENT_PREFIX parameter. The default value is FALSE and should be like this. If set to TRUE, users created as "CREATE USER... IDENTIFIED EXTERNALLY" will be able to log in to the system.

C - The user granted by default for those packages is PUBLIC and NOT SYSTEM.

upvoted 2 times

  **dianiss1050** 9 months ago

The answer is DEF

<https://dbtut.com/index.php/2020/01/09/principle-of-least-privilege-in-oracle-databases/>

upvoted 3 times

The STORES table has a column START_DATE of data type DATE, containing the date the row was inserted.

You only want to display details of rows where START_DATE is within the last 25 months.

Which WHERE clause can be used?

- A. WHERE ADD_MONTHS(start_date, 25) <= SYSDATE
- B. WHERE MONTHS_BETWEEN(start_date, SYSDATE) <= 25
- C. WHERE MONTHS_BETWEEN(SYSDATE, start_date) <=25
- D. WHERE TO_NUMBER(start date - SYSDATE) <= 25

Correct Answer: A

Community vote distribution

C (100%)

  **Jinas** 5 months ago

C is correct
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: C



Provided answer is wrong. C is correct.

```
-- Creating the table
CREATE TABLE STORES (
ID NUMBER,
NAME VARCHAR2(50),
START_DATE DATE
);


-- Inserting data into the table
INSERT INTO STORES VALUES (1, 'Store 1', SYSDATE);
INSERT INTO STORES VALUES (2, 'Store 2', ADD_MONTHS(SYSDATE, -10));
INSERT INTO STORES VALUES (3, 'Store 3', ADD_MONTHS(SYSDATE, -20));
INSERT INTO STORES VALUES (4, 'Store 4', ADD_MONTHS(SYSDATE, -30));

-- Commit to save data
COMMIT;
```



-- Selecting rows where START_DATE is within the last 25 months
SELECT * FROM STORES
WHERE MONTHS_BETWEEN(SYSDATE, START_DATE) <= 25;
upvoted 1 times

  **fdelasie** 7 months, 3 weeks ago



A is correct, in C id start_date > sysdate MONTHS_BETWEEN is negative and <= 25
upvoted 1 times

  **Dragon67** 8 months, 1 week ago

when i execute C, i get no row selected, but when i execute B, i get a value, i dont understand why, so i have no choice than to select B as answer.
upvoted 1 times

  **ivanadj** 11 months, 1 week ago

<https://www.examttopics.com/exams/oracle/1z0-071/view/2/>
upvoted 1 times

  **ivanadj** 11 months, 2 weeks ago

Selected Answer: C

The same question as question 10 in 1z0-071. The answer is C.
upvoted 1 times

Examine the description of the BOOKS_TRANSACTIONS table:

Name	Null?	Type
-----	-----	-----
TRANSACTION_ID	NOT NULL	VARCHAR2 (6)
TRANSACTION_TYPE		VARCHAR2 (3)
BORROWED_DATE		DATE
BOOK_ID		VARCHAR2 (6)
MEMBER_ID		VARCHAR2 (6)

Examine this partial SQL statement:

SELECT * FROM books_transactions



Which two WHERE conditions give the same result? (Choose two.)



- A. WHERE borrowed_date = SYSDATE AND (transaction_type = 'RM' AND member_id = 'A101' OR member_id = 'A102');
- B. WHERE borrowed_date = SYSDATE AND transaction_type = 'RM' OR member_id IN ('A101', 'A102');
- C. WHERE borrowed_date = SYSDATE AND (transaction_type = 'RM' AND (member_id = 'A101' OR member_id = 'A102'));
- D. WHERE borrowed_date = SYSDATE AND (transaction_type = 'RM' OR member_id IN ('A101', 'A102'));
- E. WHERE (borrowed_date = SYSDATE AND transaction_type = 'RM') OR member_id IN ('A101', 'A102');



Correct Answer: AC



Community vote distribution

BE (100%)

  **psycrotic5** 1 month, 3 weeks ago
but in E, there are mistake VA102, should be A102, this mistake was done on purpose? or its misstype?
upvoted 1 times

  **nautil2** 8 months, 3 weeks ago
Selected Answer: BE
The AND operator is evaluated earlier, it has higher priority than the OR operator. B and E examples differ only in parentheses, and those parentheses does not change the priority, because the priority of the AND is higher than the priority of the OR operator.
upvoted 2 times

  **ivanadj** 10 months, 2 weeks ago
Selected Answer: BE
<https://www.examttopics.com/discussions/oracle/view/8434-exam-1z0-071-topic-1-question-274-discussion/>
upvoted 1 times

  **piontk** 10 months, 2 weeks ago
Selected Answer: BE
BE is the correct answer
upvoted 2 times


Which two statements are true about views? (Choose two.)

- A. The WITH CHECK clause prevents certain rows from being updated or inserted in the underlying table through the view.
- B. The WITH CHECK clause prevents certain rows from being displayed when querying the view.
- C. Views can be updated without the need to re-grant privileges on the view.
- D. A view must only refer to tables in its defining query.
- E. Views can be indexed.

Correct Answer: AC


Community vote distribution

AC (100%)

 **auwia** 6 months, 3 weeks ago

Selected Answer: AC

Provided answer are corrects.
upvoted 1 times

 **nautil2** 8 months, 3 weeks ago

Selected Answer: AC

A) TRUE; see <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-VIEW.html#GUID-61D2D2B4-DACC-4C7C-89EB-7E50D9594D30> "Notes on Updatable Views". WITH CHECK OPTION prevents INSERT, UPDATE and DELETE in join views
B) FALSE; same URL as in A), "subquery_restriction_clause, WITH CHECK OPTION" ... Oracle Database prohibits any changes to the table or view . This statements talks only about changes, not about reading the rows of a view.
C) TRUE; <https://www.oracletutorial.com/oracle-view/oracle-updatable-view/> no extra grants are mentioned
D) FALSE; a view can be based on another view, see this example:
SQL> CREATE TABLE a (b int);
SQL> INSERT INTO a VALUES (0);
SQL> CREATE VIEW v_a AS SELECT * FROM a;
SQL> CREATE VIEW vv_a AS SELECT * FROM v_a;
SQL> SELECT * FROM vv_a;
B
-
0
E) FALSE; <https://asktom.oracle.com/pls/apex/asktom.search?tag=indexes-on-view>
upvoted 1 times



Which compression method is recommended for Direct-Path Insert operations?



- A. ROW STORE COMPRESS BASIC
- B. ROW STORE COMPRESS ADVANCED
- C. COLUMN STORE COMPRESS ADVANCED
- D. COLUMN STORE COMPRESS BASIC



Correct Answer: B



Community vote distribution



A (100%)



  **guretto** 3 months, 1 week ago
why A? it is not good for OLTP.
B is good for both DWH and OLTP
upvoted 2 times

  **shotcom** 4 months, 2 weeks ago
correct answer is A
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Selected Answer: A
Correct answer is A.
upvoted 2 times

  **Darkseid1231** 7 months ago
You sure it is A ? I will go for B instead
upvoted 2 times

  **RaNik69** 8 months, 1 week ago
Selected Answer: A
<https://blogs.oracle.com/dbstorage/post/when-to-use-and-not-use-the-different-types-of-oracle-database-19c-data-compression>
upvoted 2 times

  **dianiss1050** 9 months, 1 week ago
The correct answer is A
upvoted 2 times

Which three statements are true about Deferred Segment Creation in Oracle databases? (Choose three.)

- A. It is supported for Index Organized Tables (IOTs) contained in locally managed tablespaces.
- B. It is the default behavior for tables and indexes.
- C. Sessions may dynamically switch back and forth from DEFERRED to IMMEDIATE segment creation.
- D. Indexes inherit the DEFERRED to IMMEDIATE segment creation attribute from their parent table
- E. It is supported for SYS-owned tables contained in locally managed tablespaces.

Correct Answer: ADE

Community vote distribution

BCD (86%)

14%

  **piontk** 5 months, 2 weeks ago

Selected Answer: ABD

B and D are correct:

B - DEFERRED_SEGMENT_CREATION default value is true

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/DEFERRED_SEGMENT_CREATION.html#GUID-29066C9A-1CC7-48E4-A278-BB9678EE3362

D- "Index segment creation is deferred when the associated table defers segment creation. This is because index segment creation reflects the behavior of the table with which it is associated."

<https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-indexes.html#GUID-EF7459D8-6396-4275-A73F-75BBF2A7EC2C>

upvoted 1 times

  **piontk** 5 months, 2 weeks ago

There's no documentation supporting C.

A is correct, tested as follows:

-- as sys

CONN sys/senha@orclpdb as sysdba

```
CREATE TABLESPACE tbs_local DATAFILE SIZE 50M
EXTENT MANAGEMENT LOCAL
SEGMENT SPACE MANAGEMENT AUTO;
```

```
alter user hr quota 10m on tbs_local;
```

-- as hr

conn hr/hr@orclpdb

```
CREATE TABLE countries_iot(
country_id CHAR(2) NOT NULL,
country_name VARCHAR2(40),
currency_name VARCHAR2(25),
currency_symbol VARCHAR2(3),
region VARCHAR2(15),
CONSTRAINT country_iot_id_pk_demo PRIMARY KEY (country_id )
)
SEGMENT CREATION DEFERRED
ORGANIZATION INDEX
TABLESPACE tbs_local;
```

CONN sys/senha@orclpdb as sysdba

```
SELECT TABLESPACE_NAME, EXTENT_MANAGEMENT, SEGMENT_SPACE_MANAGEMENT FROM DBA_TABLESPACES WHERE
TABLESPACE_NAME='TBS_LOCAL';
```

COL OWNER FOR A5

COL TABLE_NAME FOR A20

```
SELECT OWNER, TABLE_NAME, IOT_TYPE, SEGMENT_CREATED FROM DBA_TABLES WHERE TABLE_NAME='COUNTRIES_IOT';
```

```
select segment_name, bytes/1024/1024 MB from dba_segments where segment_name='COUNTRIES_IOT';
```

upvoted 1 times

  **jonsnoow** 3 months ago

店长微信：hitong128

https://docs.oracle.com/en/database/oracle/oracle-database/19/ferfn/DEFERRED_SEGMENT_CREATION.html#GUID-29066C9A-1CC7-48E4-A278-BB9678EE3362



upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: BCD

Agree with B, C, and D.


upvoted 2 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: BCD

BCD are correct <http://blog.itpub.net/70018393/viewspace-2902513/>

upvoted 3 times

  **piontk** 10 months, 2 weeks ago

Selected Answer: BCD

BCD is correct.

SYS and SYSTEM users tables will always have a segment.

upvoted 1 times

  **dianiss1050** 10 months, 2 weeks ago

The answer B is correct

upvoted 1 times

Examine these statements executed in a single Oracle session:

```
CREATE TABLE product (pcode NUMBER(2), pname VARCHAR2(20));

INSERT INTO product VALUES (1, 'pen');

INSERT INTO product VALUES (2, 'pencil');

INSERT INTO product VALUES (3, 'fountain pen');

SAVEPOINT a;

UPDATE product SET pcode = 10 WHERE pcode = 1;

COMMIT;

DELETE FROM product WHERE pcode = 2;

SAVEPOINT b;

UPDATE product SET pcode = 30 WHERE pcode = 3;

SAVEPOINT c;

DELETE FROM product WHERE pcode = 10;

ROLLBACK TO SAVEPOINT b;

COMMIT;
```

Which three statements are true? (Choose three.)

- A. The code for pen is 10.
- B. There is no row containing fountain pen.
- C. There is no row containing pen.
- D. The code for fountain pen is 3.
- E. The code for pen is 1.
- F. There is no row containing pencil.

Correct Answer: BCF



Community vote distribution

ADF (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: ADF



Provided answer are wrong! The base! :-) part VI
upvoted 1 times

  **nautil2** 8 months, 1 week ago

Selected Answer: ADF

Correct answer are: A,D,F.
Before SAVEPOINT b; command was issued, content of the table is:
SQL> SELECT * FROM product;
PCODE PNAME



10 pen
3 fountain pen
The command "ROLLBACK TO SAVEPOINT b;" returns the table to the state created by the SAVEPOINT b;
upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: ADF


a d f are correct

upvoted 1 times

  **piontk** 10 months, 2 weeks ago

Selected Answer: ADF

ADF is right
upvoted 2 times

  **ivanadj** 11 months, 1 week ago

Selected Answer: ADF

<https://www.examttopics.com/discussions/oracle/view/8225-exam-1z0-071-topic-1-question-243-discussion/>
upvoted 3 times

Question #129

Topic 1

Examine the description of the SALES1 table:

Name	Null	Type
-----	-----	-----
SALES_ID	NOT NULL	NUMBER
STORE_ID	NOT NULL	NUMBER
ITEMS_ID		NUMBER
QUANTITY		NUMBER
SALES_DATE		DATE

SALES2 is a table with the same description as SALES1.

Some sales data is contained erroneously in both tables.

You must display rows from SALES1 and SALES2 and wish to see the duplicates too.



Which set operator generates the required output?

- A. UNION ALL
- B. SUBTRACT
- C. MINUS
- D. UNION
- E. INTERSECT

Correct Answer: A

Community vote distribution

A (100%)

  **nautil2** 8 months, 1 week ago

Selected Answer: A

The UNION ALL operator does not eliminate duplicate selected rows. <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/The-UNION-ALL-INTERSECT-MINUS-Operators.html#GUID-B64FE747-586E-4513-945F-80CB197125EE>
upvoted 1 times

In the promotions table, the PROMO_BEGIN_DATE column is of data type DATE and the default date format is DD-MON-RR.

Which two statements are true about expressions using PROMO_BEGIN_DATE contained in a query? (Choose two.)

- A. TO_NUMBER(PROMO_BEGIN_DATE) - 5 will return a number.
- B. PROMO_BEGIN_DATE - SYSDATE will return an error.
- C. PROMO_BEGIN_DATE - SYSDATE will return a number.
- D. PROMO_BEGIN_DATE - 5 will return a date.
- E. TO_DATE(PROMO_BEGIN_DATE * 5) will return a date.

Correct Answer: *BD*

Community vote distribution

CD (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: CD

C and D are corrects.
upvoted 2 times

  **nautil2** 8 months, 1 week ago



Selected Answer: CD

CD are correct answers
upvoted 1 times

  **RaNik69** 8 months, 1 week ago

Selected Answer: CD

C D are correct
upvoted 1 times

  **ivanadj** 11 months, 1 week ago

Selected Answer: CD

<https://www.examttopics.com/discussions/oracle/view/10240-exam-1z0-071-topic-1-question-269-discussion/>
upvoted 3 times



Which two statements are true regarding indexes? (Choose two.)

- A. A non-unique index can be altered to be unique.
- B. An update to a table can result in no updates to any of the table's indexes.
- C. The RECYCLE BIN never contains indexes.
- D. An update to a table can result in updates to any or all of the table's indexes.
- E. A table belonging to one user cannot have an index that belongs to a different user.

Correct Answer: BD

Community vote distribution

BD (100%)

  **musafir** 6 months, 1 week ago

Selected Answer: BD

B and D are correct
<https://www.examttopics.com/discussions/oracle/view/20178-exam-1z0-071-topic-2-question-42-discussion/>
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: BD

Provided answer are corrects.
upvoted 1 times

  **j_tw** 9 months ago


Maybe BE?
upvoted 2 times

  **ahmadkrishan** 10 months, 1 week ago

sry A + D correct
upvoted 1 times

  **ahmadkrishan** 10 months, 2 weeks ago

E is true
upvoted 1 times

  **shotcom** 4 months, 1 week ago

To create an index in another schema, all of the following conditions must be true:
You have CREATE ANY INDEX system privilege.
The owner of the other schema has a quota for the tablespaces to contain the index or index partitions, or UNLIMITED TABLESPACE system privilege.
so E is not correct
upvoted 1 times

Which three statements are true regarding single row subqueries? (Choose three.)

- A. A SQL statement may have multiple single row subquery blocks.
- B. They can be used in the WHERE clause.
- C. They must be placed on the left side of the comparison operator or condition.
- D. They can be used in the HAVING clause.
- E. They must be placed on the right side of the comparison operator or condition.
- F. They must return a row to prevent errors in the SQL statement.

Correct Answer: ABD



Community vote distribution

ABD (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: ABD

Provided answer are corrects.
upvoted 1 times

  **Darkseid1231** 7 months ago

I will go for A B F
upvoted 1 times

  **nautil2** 8 months ago

Selected Answer: ABD

ABD seems to be correct.
A,B: TRUE
SELECT column1,
(SELECT MAX(column2) FROM table1) AS max_value,
(SELECT COUNT(*) FROM table2) AS row_count
FROM table3
WHERE column4 = (SELECT MIN(column4) FROM table4);
C: FALSE - column name is placed on the left side of the condition/operator, while subquery is placed on the right side of it.
D: TRUE
SELECT deptno, MIN(sal) FROM emp GROUP BY deptno HAVING MIN(sal) > (SELECT MIN(sal) FROM emp WHERE deptno = 30);
E: FALSE, a subquery does not need to be surrounded by a condition or operator, see example for A+B answer
F: FALSE, single-row subquery can return no row without raising an error
SELECT ename, init, bdate FROM employees WHERE bdate > (SELECT bdate FROM employees WHERE empno = 99999);
upvoted 1 times


Which two statements are true about Oracle synonyms? (Choose two.)

- A. Any user can create a public synonym.
- B. A synonym has an object number.
- C. A synonym can be created on an object in a package.
- D. All private synonym names must be unique in the database.
- E. A synonym can have a synonym.

Correct Answer: AD

Community vote distribution



  **zouve** 5 months, 2 weeks ago

Selected Answer: CE

CE same question in 1z0-071
upvoted 1 times

  **zouve** 5 months ago

select object_id, object_name, OBJECT_TYPE from all_objects;

i can see object_id to every synonym
upvoted 1 times

  **auwia** 6 months, 2 weeks ago

Selected Answer: CE


Provided answer are wrong, C and E are corrects in my opinion:
A. Any user can create a public synonym. --> False, you need grants.
B. A synonym has an object number. --> False, OBJECT ID
C. A synonym can be created on an object in a package. --> True. you can grant to object type like function, store procedure, etc.
D. All private synonym names must be unique in the database. --> False, must be unique in the SCHEMA.
E. A synonym can have a synonym. --> True.
upvoted 1 times

  **shotcom** 4 months, 1 week ago

https://asktom.oracle.com/pls/apex/f%3Fp%3D100:11:0::::P11_QUESTION_ID:4743759113231 according to that link you cant create a synonym for an object in a package. yes you can create a synonym for the package itself but not for the objects inside the package so C is wrong
upvoted 1 times

  **auwia** 6 months, 2 weeks ago



<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-SYNONYM.html#GUID-A806C82F-1171-478E-A910-F9C6C42739B2>
upvoted 1 times

  **jonsnoow** 3 months, 2 weeks ago



From the same link:

"Restriction on the FOR Clause
The schema object cannot be contained in a package."

From the same link
upvoted 1 times

  **auwia** 6 months, 2 weeks ago

C. A synonym can be created on an object in a package. --> True. you can link* to object type like function, store procedure, etc.
upvoted 1 times

  **Darkseid1231** 7 months ago

B D is correct
upvoted 1 times



  **nautil2** 8 months ago

Selected Answer: BE

A: FALSE, a user needs privilege CREATE PUBLIC SYNONYM to do it
B: TRUE
SQL> SELECT owner, object_type, object_name, object_id FROM dba_objects WHERE object_id=20770;

GSMADMIN_INTERNAL SYNONYM DDL_REQUESTS_PWD 20770
C: FALSE;
D: FALSE
SQL> SELECT object_name, object_type, object_id, owner FROM dba_objects WHERE object_type='SYNONYM' AND owner LIKE 'KOTELNIK%';
OBJECT_NAME OBJECT_TYPE OBJECT_ID OWNER

EMPLOYEES_S SYNONYM 76119 KOTELNIK1
EMPLOYEES_S SYNONYM 76128 KOTELNIK2E:
E: TRUE, <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-SYNONYM.html#GUID-A806C82F-1171-478E-A910-F9C6C42739B2> Use the CREATE SYNONYM statement to create a synonym, which is an alternative name for ... another synonym
upvoted 2 times

  **ivanadj** 11 months, 2 weeks ago

Selected Answer: BE

The same question as in 1z0-071
Question #: 270
upvoted 2 times

Examine this command:


SQL> ALTER TABLE ORDERS SHRINK SPACE COMPACT

Which two statements are true? (Choose two.)

- A. The SHRINK operation causes rows to be moved to empty space starting from the beginning of the ORDERS segment.
- B. Queries and DML statements are allowed on ORDERS while the SHRINK is executing.
- C. Dependent indexes become UNUSABLE.
- D. Only queries are allowed on ORDERS while the shrink is executing.
- E. The high-water mark (HWM) of ORDERS is adjusted.
- F. The SHRINK operation causes rows to be moved to empty space starting toward the end of the ORDERS segment.

Correct Answer: CE

Community vote distribution



psycrotic5 1 month, 2 weeks ago

FALSE:
F.Contrary to the claim, the SHRINK operation moves rows to empty space starting from the beginning of the segment, not toward the end.
E: If you specify COMPACT, then Oracle Database only defragments the segment space and compacts the table rows for subsequent release. The database does not readjust the high water mark and does not release the space immediately. You must issue another ALTER TABLE ... SHRINK SPACE statement later to complete the operation. This clause is useful if you want to accomplish the shrink operation in two shorter steps rather than one longer step.
i go for AB also.
upvoted 1 times

zouve 5 months ago

Selected Answer: BF
ALTER TABLE SHRINK SPACE
The ALTER TABLE SHRINK SPACE algorithm starts from the bottom of the segment and starts moving rows to the beginning of the segment. Compacting the segment requires row movement. Therefore, you must enable row movement for the object you want to shrink before specifying this clause.
By default, Oracle Database compacts the segment, adjusts the High Water Mark, and releases the recuperated space immediately.
B. ALTER TABLE MY_TABLE SHRINK SPACE COMPACT
Compact rows, without adjusting HWM.

<http://aleckaplandba.blogspot.com/2015/07/shrink-table-by-example.html>
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: BF
Correct answers are B and F, because rows are moved up and free space goes down. For COMPACT clause, watermarks are not updated. DML and queries are both supported in this "phase".
upvoted 2 times



nautil2 8 months ago

Selected Answer: AB
A: TRUE; It wouldn't make sense if the data were filled from the end of segment, because it could remain some empty space after shrinking. Data are moved to empty space to use maximum of empty space. I could not find any proof in the documentation. If anyone find it, feel free and post it here.
B: TRUE, <https://docs.oracle.com/en/database/oracle/oracle-database/19/admin/managing-space-for-schema-objects.html#GUID-37DF35E0-116B-4AAB-BFB3-DDBDBAB29BEB> DML operations and queries can be issued during the data movement phase of segment shrink.
C: FALSE, same URL as B, Indexes are maintained during the shrink operation and remain usable after the operation is complete.
D: FALSE, also DML is allowed, see B
E: FALSE; <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/ALTER-TABLE.html#GUID-552E7373-BF93-477D-9DA3-B2C9386F2877>
F: FALSE; opposite to A
upvoted 1 times

RaNik69 9 months ago

Selected Answer: AB
<https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/ALTER-TABLE.html#GUID-552E7373-BF93-477D-9DA3-B2C9386F2877>

upvoted 2 times

  **piontk** 10 months, 2 weeks ago

Selected Answer: AB

AB are correct
upvoted 3 times

  **ahmadkrishan** 10 months, 1 week ago

can u explain
upvoted 1 times

Question #135



Topic 1

The sales table has columns PROD_ID and QUANTITY_SOLD of data type number.

Which two queries execute successfully? (Choose two.)

- A. SELECT COUNT(prod_id) FROM sales WHERE quantity_sold > 55000 GROUP BY prod_id;
- B. SELECT COUNT(prod_id) FROM sales GROUP BY prod_id WHERE quantity_sold > 55000;
- C. SELECT prod_id FROM sales WHERE quantity_sold > 55000 AND COUNT(*) > 10 GROUP BY COUNT(*) > 10;
- D. SELECT prod_id FROM sales WHERE quantity_sold > 55000 GROUP BY prod_id HAVING COUNT(*) > 10;
- E. SELECT prcd_id FROM sales WHERE quantity_sold > 55000 AND COUNT (*) > 10 GROUP BY prod_id HAVING COUNT(*) > 10;

Correct Answer: AD

  **Darkseid1231** 7 months ago

A D is correct
upvoted 1 times

Examine the description of the PRODUCT_DETAILS table:

NAME	NULL	TYPE
-----	-----	-----
PRODUCT_ID	NOT NULL	NUMBER(2)
PRODUCT_NAME	NOT NULL	VARCHAR2(25)
PRODUCT_PRICE		NUMBER(8,2)
EXPIRY_DATE		DATE

Which two statements are true? (Choose two.)

- A. PRODUCT_ID can be assigned the PRIMARY KEY constraint.
- B. EXPIRY_DATE cannot be used in arithmetic expressions.
- C. PRODUCT_NAME cannot contain duplicate values.
- D. EXPIRY_DATE contains the SYSDATE by default if no date is assigned to it.
- E. PRODUCT_PRICE can be used in an arithmetic expression even if it has no value stored in it.
- F. PRODUCT_PRICE contains the value zero by default if no value is assigned to it.

Correct Answer: AC

Community vote distribution

AE (100%)



  **auwia** 6 months, 3 weeks ago

Selected Answer: AE

Provided answer are wrong. The base! :-) part VII
upvoted 1 times

  **Juv07** 11 months ago

A E is the good answers
upvoted 2 times

  **ivanadj** 11 months, 2 weeks ago

Selected Answer: AE

The same question as 1z0-071
Question #: 229
upvoted 2 times

Which two statements are true about INTERVAL data types? (Choose two.)

- A. INTERVAL DAY TO SECOND columns support fractions of seconds.
- B. The YEAR field in an INTERVAL YEAR TO MONTH column must be a positive value.
- C. The value in an INTERVAL DAY TO SECOND column can be copied into an INTERVAL YEAR TO MONTH column.
- D. INTERVAL YEAR TO MONTH columns only support monthly intervals within a range of years.
- E. INTERVAL YEAR TO MONTH columns only support monthly intervals within a single year.
- F. INTERVAL YEAR TO MONTH support yearly intervals.

Correct Answer: AD

Community vote distribution



nautil2 5 months ago

Selected Answer: AD

True statements about intervals in Oracle Database 19c:
A. True - INTERVAL DAY TO SECOND columns support fractions of seconds.
D. True - INTERVAL YEAR TO MONTH columns only support monthly intervals within a range of years.
False statements:
B. False - The YEAR field in an INTERVAL YEAR TO MONTH column can be a negative value.
C. False - The value in an INTERVAL DAY TO SECOND column cannot be copied into an INTERVAL YEAR TO MONTH column.
E. False - INTERVAL YEAR TO MONTH columns can include monthly intervals spanning multiple years.
F. False - INTERVAL YEAR TO MONTH columns do not support yearly intervals.
upvoted 2 times

auwia 6 months, 3 weeks ago

Selected Answer: AF

Provided answer are wrong.
upvoted 2 times

ivanadj 11 months, 2 weeks ago

Selected Answer: AF

1z0-071
Question #: 292
upvoted 1 times

Examine the description of the EMPLOYEES table:

Name	Null?	Type
-----	-----	-----
EMP_ID	NOT NULL	NUMBER
EMP_NAME		VARCHAR2 (10)
DEPT_ID		NUMBER (2)
SALARY		NUMBER (8, 2)
JOIN_DATE		DATE

NLS_DATE_FORMAT is set to DD-MON-YY.

Which query requires explicit data type conversion?

- A. SELECT join_date FROM employees WHERE join_date > '10-02-2018';
- B. SELECT join_date + '20' FROM employees;
- C. SELECT salary + '120.50' FROM employees;
- D. SELECT SUBSTR(join_date, 1, 2) - 10 FROM employees;
- E. SELECT join_date || ' ' || salary FROM employees;

Correct Answer: C

Community vote distribution



psycrotic5 1 month, 2 weeks ago

what about
Option D?: SELECT SUBSTR(join_date, 1, 2) - 10 FROM employees; This query is attempting to perform a substring operation on a DATE type, which is not valid without first converting the date to a string. Oracle does not support taking substrings directly from DATE data types. The date must be converted to a string using TO_CHAR before SUBSTR can be used. Furthermore, subtracting 10 from a substring of a date does not make sense unless the substring represents a number, which is not the case here since join_date is a date and not a string of numbers. Therefore, explicit conversion is necessary, and the query itself is somewhat nonsensical without additional context or correction. This makes Option D the most clearly incorrect and the one that certainly requires explicit data type conversion to even make sense as a SQL statement.

upvoted 1 times

nautil2 5 months ago

Selected Answer: C

AC is correct.
A returns:
SELECT join_date FROM employees WHERE join_date > '10-02-2018';
SELECT join_date FROM employees WHERE join_date > '10-02-2018'
*
ERROR at line 1:
ORA-01843: not a valid month

C returns:
SQL> SELECT salary+'120.50' FROM employees;
SELECT salary+'120.50' FROM employees
*
ERROR at line 1:
ORA-01722: invalid number
upvoted 1 times

nautil2 5 months ago

Let me correct my suggestion, please. Only C is correct.
Datatype conversion in A does not make the query valid.
SQL> SELECT join_date FROM employees WHERE TO_DATE(join_date,'DD-MM-YYYY') > '10-02-2018';
SELECT join_date FROM employees WHERE TO_DATE(join_date,'DD-MM-YYYY') > '10-02-2018'
*
ERROR at line 1:
ORA-01843: not a valid month

Only altering session variable NLS_DATE_FORMAT with
ALTER SESSION SET nls_date_format='DD-MM-YY';
makes the A query to return valid response. So there's no datatype conversion, just altering session variable, therefore A is false.



upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: A

Provided answer is wrong! The base! :-) part VIII

upvoted 1 times

  **Pivoine** 6 months, 3 weeks ago

C

tested

upvoted 1 times

  **ivanadj** 11 months, 1 week ago

Selected Answer: A

<https://www.examttopics.com/discussions/oracle/view/13453-exam-1z0-071-topic-1-question-296-discussion/>

upvoted 3 times

Question #139

Topic 1

You start your database instance in NOMOUNT state.

Which two actions are performed? (Choose two.)

- A. SYS can access the database.
- B. The consistency of the database is checked.
- C. The control files are opened.
- D. All required background processes are started.
- E. Memory is allocated for the SGA.

Correct Answer: BC

Community vote distribution

DE (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: DE

Agree with D and E.

upvoted 1 times

  **PrasenC** 11 months, 1 week ago

Selected Answer: DE

During mount control file is read or open.

upvoted 4 times

Which two statements are true about the PMON background process? (Choose two.)

- A. It registers database services with all local and remote listeners known to the database instance.
- B. It frees unused temporary segments.
- C. It rolls back transactions when a process fails.
- D. It records checkpoint information in the control file.
- E. It frees resources held by abnormally terminated processes.

Correct Answer: AE

Community vote distribution

CE (100%)

  **psycrotic5** 1 month, 2 weeks ago

A incorrect cuz:
In releases before Oracle Database 12c, PMON was responsible for listener registration, but this task is now the responsibility of the Listener Registration Process (LREG)
upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: CE

To don't confuse with same question, but different answer at page 7.
upvoted 1 times

  **Dragon67** 9 months, 3 weeks ago

I THINK ITS CE
upvoted 2 times

  **j_tw** 9 months, 2 weeks ago

I agree that.
A - LREG
B - SMON
C - PMON
D - CKPT
E - PMON
upvoted 3 times


Which two statements are true about the Oracle join and ANSI join syntax? (Choose two.)

- A. The Oracle join syntax performs less well than the SQL:1999 compliant ANSI join syntax.
- B. The Oracle join syntax lacks the ability to do outer joins.
- C. The SQL:1999 compliant ANSI join syntax supports creation of a Cartesian product of two tables.
- D. The Oracle join syntax supports creation of a Cartesian product of two tables.
- E. The Oracle join syntax performs better than the SQL:1999 compliant ANSI join syntax.

Correct Answer: *CD*

Community vote distribution

CD (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: CD

Provided answer is correct
upvoted 1 times

  **Darkseid1231** 7 months ago

C D seems to be correct
upvoted 2 times

Which two statements are true about single row functions? (Choose two.)

- A. CEIL : can be used for positive and negative numbers.
- B. MOD : returns the remainder of a division operation.
- C. CONCAT : can be used to combine any number of values.
- D. FLOOR : returns the smallest integer greater than or equal to a specified number.
- E. TRUNC : can be used only with NUMBER data types.

Correct Answer: AE

Community vote distribution

AB (100%)

- 👤

psycrotic5

1 month, 2 weeks ago

agree with nautil2,
E. TRUNC: can be used only with NUMBER data types. This statement is incorrect. TRUNC can be used with both NUMBER and DATE data types. It truncates to the specified precision for numbers, and to the specified date component for dates

upvoted 1 times
- 👤

nautil2

4 months, 3 weeks ago

Selected Answer: AB

C – false, max 2 arguments
D – false, returns biggest integer smaller than specified number
E – false, can be also used with DATE datatype

upvoted 2 times
- 👤

auwia

6 months, 3 weeks ago

Selected Answer: AB

A and B are correct in my opinion.

upvoted 1 times
- 👤

TECHNO_GUY

9 months, 2 weeks ago

AB are the correct answers

upvoted 3 times
- 👤

Dragon67

9 months, 2 weeks ago

when are you preparing to take the exams

upvoted 1 times
- 👤

j_tw

9 months, 2 weeks ago

agree that.
E - can be used number and also date data type.

upvoted 2 times
- 👤

Dragon67

9 months, 3 weeks ago

THERE ONLY 1 ANSWER WHICH IS A

upvoted 1 times
- 👤

Dragon67

9 months, 2 weeks ago

*error it should be AB

upvoted 1 times



What is true about non-equijoin statement performance?



- A. The BETWEEN condition used with a non-equijoin always performs better than when using the >= and <= conditions.
- B. The BETWEEN condition used with a non-equijoin sometimes performs better than using the >= and <= conditions.
- C. The Oracle join syntax performs better than the SQL:1999 compliant ANSI join syntax.
- D. The Oracle join syntax performs less well than the SQL:1999 compliant ANSI join syntax.
- E. The join syntax used makes no difference to performance.

Correct Answer: E

Community vote distribution

E (100%)

  **Darkseid1231** 6 months, 2 weeks ago
I will go for C
upvoted 1 times

  **auwia** 6 months, 3 weeks ago
Selected Answer: E
Provided answer is correct.
upvoted 1 times

Examine the description of the CUSTOMERS table:

Name	Null?	Type
-----	-----	-----
CUST_ID	NOT NULL	VARCHAR2 (2)
CUST_LAST_NAME		VARCHAR2 (30)
CITY		VARCHAR2 (10)
CUST_CREDIT_LIMIT		NUMBER (6, 2)

You need to display last names and credit limits of all customers whose last name starts with A or B in lower or upper case, and whose credit limit is below 1000.

Examine this partial query:

```
SELECT cust_last_name, cust_credit_limit FROM customers
```

Which two WHERE conditions give the required result? (Choose two.)

- A. WHERE (UPPER(cust_last_name) LIKE 'A%' OR UPPER(cust_last_name) LIKE 'B%') AND ROUND(cust_credit_limit) < 1000;
- B. WHERE (INITCAP(cust_last_name) LIKE 'A%' OR INITCAP(cust_last_name) LIKE 'B%') AND cust_credit_limit < 1000;
- C. Where UPPER(cust_last_name) IN ('A%', 'B%') AND cust_credit_limit < 1000;
- D. WHERE (UPPER(cust_last_name) LIKE INITCAP('A') OR UPPER(cust_last_name) LIKE INITCAP('B')) AND ROUND(cust_credit_limit) < ROUND(1000);
- E. WHERE UPPER(cust_last_name) BETWEEN UPPER('A%' AND 'B%') AND ROUND(cust_credit_limit) < 1000;

Correct Answer: BC

Community vote distribution



psycrotic5 1 month, 2 weeks ago
AB CORRECT - TESTED
upvoted 1 times

nautil2 4 months, 3 weeks ago
Selected Answer: B
Only B is correct. Answer A is almost correct, however function ROUND negatively affects condition cust_credit limit < 1000 for customers with credit ranging from 999.50 to 999.99. Their credit is rounded to 1000 and condition < 1000 removes these customers from query results. So these customers are not shown in final results, thought their credit is below 1000. Initial instruction requesting to show customers with credit below 1000 is not fulfilled. Answer A is false. Other answers C,D,E do not return correct results.
upvoted 2 times

auwia 6 months, 3 weeks ago
Selected Answer: AB
Provided answer are wrong! The base! Part IX! :-)
upvoted 1 times

Juv07 7 months, 4 weeks ago
Correct Answer AB
upvoted 1 times

j_tw 9 months, 2 weeks ago
Selected Answer: AB

Why is not AB?
I know that '%' is only used in LIKE operator.
Doesn't it?
upvoted 4 times

店长微信：hj feng128

Question #145

Topic 1

Which two statements are true about substitution variables? (Choose two.)

- A. A substitution variable can be used with any clause in a SELECT statement.
- B. A substitution variable used to prompt for a column name must be enclosed in double quotation marks.
- C. A substitution variable used to prompt for a column name must be enclosed in single quotation marks.
- D. A substitution variable prefixed with & always prompts only once for a value in a session.
- E. A substitution variable can be used only in a SELECT statement.
- F. A substitution variable prefixed with && prompts only once for a value in a session unless it is set to undefined in the session.

Correct Answer: AF



Community vote distribution

AF (100%)

  **auwia** 6 months, 3 weeks ago

Selected Answer: AF

Provided answer are correct
upvoted 2 times

  **dianiss1050** 9 months ago

AF is correct
upvoted 1 times

The CUSTOMERS table has a CUST_LAST_NAME column of data type VARCHAR2.

The table has two rows whose CUST_LAST_NAME values are Anderson and Ausson.

Which query produces output for CUST_LAST_NAME containing Oder for the first row and Aus for the second?

- A. SELECT REPLACE(SUBSTR(cust_last_name, -3), 'An', 'O') FROM customers;
- B. SELECT REPLACE(TRIM(TRAILING 'son' FROM cust_last_name), 'An', 'O') FROM customers;
- C. SELECT INITCAP(REPLACE(TRIM('son' FROM cust_last_name), 'An', 'O')) FROM customers;
- D. SELECT REPLACE(REPLACE(cust_last_name, 'son', ''), 'An', 'O') FROM customers;

Correct Answer: D

Community vote distribution



musafir 6 months, 1 week ago

Selected Answer: D

D is correct
upvoted 2 times

Darkseid1231 6 months, 2 weeks ago

D is correct
upvoted 1 times

auwia 6 months, 3 weeks ago

Selected Answer: B

Provided answer is correct
upvoted 1 times

auwia 6 months, 3 weeks ago

D I'm sorry
upvoted 2 times

jonsnoow 8 months, 2 weeks ago

D is OK. B is not the wright answer. You will got an error:

ORA-30001: trim set should have only one character
30001. 00000 - "trim set should have only one character"
*Cause: Trim set contains more or less than 1 character. This is not allowed in TRIM function.
upvoted 3 times

blu 8 months, 2 weeks ago

Selected Answer: D

D is correct
upvoted 1 times

j_tw 9 months ago

I think B also can be the answer. TRIM(TRAILINIG 'son' from CUST_LAST_NAME) command will delete 'son' from right side. Than Anderson will be Ander, and Ausson will be Aus.
upvoted 1 times

auwia 6 months, 3 weeks ago

Trim/trailing works with 1 character only. D is correct.
upvoted 3 times



Which two statements are true about GLOBAL TEMPORARY TABLES? (Choose two.)

- A. GLOBAL TEMPORARY TABLE space allocation occurs at session start.
- B. GLOBAL TEMPORARY TABLE rows inserted by a session are available to any other session whose user has been granted select on the table.
- C. A TRUNCATE command issued in a session causes all news in a GLOBAL TEMPORARY TABLE for the issuing session to be deleted.
- D. A GLOBAL TEMPORARY TABLE 's definition is available to multiple sessions.
- E. A DELETE command on a GLOBAL TEMPORARY TABLE cannot be rolled back.

Correct Answer: CD


Community vote distribution

CD (100%)

  **nautil2** 4 months, 3 weeks ago


Selected Answer: CD

A – false; space is allocated when it is requested/needed, see <https://asktom.oracle.com/pls/apex/asktom.search?tag=about-temporary-tables>
 B – false; table is visible to other sessions but each session can see only its own data, see <https://www.oracletutorial.com/oracle-basics/oracle-global-temporary-table/>
 C – true
 D – true
 E – false; commit and rollback are possible in Global Temporary Tables
 upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: CD

Provided answer are corrects.
 upvoted 1 times

  **Darkseid1231** 7 months ago

D C are correct
 upvoted 1 times


Which three statements are true about Resumable Space Allocation in Oracle databases? (Choose three.)

- A. A user's session may be suspended and resumed multiple times.
- B. The AFTER SUSPEND event trigger can itself be suspended due to space conditions.
- C. All sessions must have the same timeout value when waiting for resumable space allocations.
- D. Resumable space allocation is only possible with locally managed tablespaces.
- E. A user's session may be suspended even if the user has the UNLTMTTED TABLESPACE system privilege.
- F. Resumable space allocation may be enabled for some sessions and not others.

Correct Answer: ADF

Community vote distribution



  **nautil2** 4 months, 3 weeks ago



Selected Answer: AEF

A - true, "A resumable statement can be suspended and resumed multiple times during execution."
 B - false, "SQL statements executed within a AFTER SUSPEND trigger are always non-resumable and are always autonomous."
 C - false, ALTER SESSION SET RESUMABLE_TIMEOUT sets specific timeout for the session it is executed in, it does not affect other sessions
 D - false, "There are two ways to specify alert thresholds for both locally managed and dictionary managed tablespaces:"
 E - true, there is no notice about UNLIMITED TABLESPACE in documentation claiming this system privilege prevents raising Resumable Space Allocation event
 F - true, ALTER SESSION ENABLE RESUMABLE can be executed in some sessions and it is not necessary to run it in all sessions
 upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: ABF



Correct answer in my opinion.
 upvoted 1 times

  **jonsnoow** 8 months, 2 weeks ago

I also think that AEF a correct, but I find this statement in a book...

"If you create an AFTER SUSPEND ON DATABASE trigger that attempts to fix problems, remember that it might hit a space problem itself."

Watson, John; Watson, John. OCA Oracle Database 12c Installation and Administration Exam Guide (Exam 1Z0-062) (Oracle Press) . McGraw Hill LLC. Kindle Edition.
 upvoted 1 times

  **Humber31** 8 months, 3 weeks ago

Selected Answer: AEF

Resumable Space Allocation is possible also with Dictionary Managed Tablespaces; with some limitations, but possible, so D is false.

There is nothing that indicates that UNLIMITED TABLESPACE privilege can avoid the Resumable Space Allocation, so E is true.

Also, A and F are correct.
 upvoted 2 times

Examine the description of the PRODUCT_INFORMATION table:

Name	Null?	Type
-----	-----	-----
PROD_ID	NOT NULL	NUMBER(2)
PROD_NAME		VARCHAR2(10)
LIST_PRICE		NUMBER(6,2)


Which query retrieves the number of products with a null list price?

- A. SELECT COUNT(list_price) FROM product_information WHERE list_price IS NULL;
- B. SELECT COUNT(NVL(list_price, 0)) FROM product_information WHERE list_price IS NULL;
- C. SELECT COUNT(DISTINCT list price) FROM product_information WHERE list_price IS NULL;
- D. SELECT COUNT(list_price) FROM product_information WHERE list_price = NULL;

Correct Answer: A


Community vote distribution

B (100%)

 **nautil2** 4 months, 3 weeks ago


Selected Answer: B

A – false; function COUNT does not cover NULL values in its output
B – true; NULLs become zeros and COUNT counts up all zero values, which is the same as the number of NULL values
C – false; DISTINCT NULL is still NULL, see A
D – false, NULL value is examined with WHERE IS NOT NULL and also see feature mentioned in A
upvoted 1 times


 **Ichdb** 6 months, 1 week ago

Selected Answer: B

B ! :)
upvoted 1 times

 **Juv07** 7 months, 4 weeks ago

Correct Answer is B. Count function can't calculate NULL
upvoted 1 times

 **j_tw** 9 months ago

Selected Answer: B

B is correct.
COUNT function can calculate NULL only it use asterik.
COUNT(*) : can count NULL
COUNT(Column Name) : cannot count NULL
upvoted 2 times

Examine the description of the CUSTOMERS table:

Name	Null?	Type
-----	-----	-----
CUSTINO	NOT NULL	NUMBER (3)
CUSTNAME	NOT NULL	VARCHAR2 (25)
CUSTADDRESS		VARCHAR2 (35)
CUST_CREDIT_LIMIT		NUMBER (5)

CUSTINO is the PRIMARY KEY.

You must determine if any customers' details have been entered more than once using a different CUSTINO, by listing all duplicate names.


Which two methods can you use to get the required result? (Choose two.)

- A. subquery
- B. self join
- C. LEFT OUTER JOIN with self join
- D. RIGHT OUTER JOIN With self join
- E. FULL OUTER JOIN with self join

Correct Answer: AB

Community vote distribution

AB (100%)

 **nautil2** 4 months, 2 weeks ago

Selected Answer: AB

A, B are correct answers. See example:
A query:
SELECT custino,custname,custaddress,cust_credit_limit FROM customers
WHERE custname IN (
SELECT custname FROM customers
GROUP BY custname HAVING COUNT(custname)>1
)
AND custaddress IN (
SELECT custaddress FROM customers
GROUP BY custaddress HAVING COUNT(custaddress)>1
)
AND cust_credit_limit IN (
SELECT cust_credit_limit FROM customers
GROUP BY cust_credit_limit HAVING COUNT(cust_credit_limit)>1
)
ORDER BY custname;
B query:
SELECT DISTINCT(c.custino),c.custname,c.custaddress,c.cust_credit_limit FROM customers c, customers u WHERE c.custname = u.custname AND
c.custaddress = u.custaddress AND c.cust_credit_limit = u.cust_credit_limit AND c.custino <> u.custino ORDER BY c.custname, c.custino;
upvoted 1 times

 **auwia** 6 months, 3 weeks ago

Selected Answer: AB

Provided answer are correct.
upvoted 2 times

 **Darkseid1231** 6 months, 2 weeks ago

You right

Question #151


Topic 1

Which two statements are true about undo segments and the use of undo by transactions in an Oracle database instance? (Choose two.)

- A. Undo segments can be stored in the SYSAUX tablespace.
- B. Undo segments can wrap around to the first extent when a transaction fills the last extend of the undo segment.
- C. A single transaction may use multiple undo segments simultaneously.
- D. Undo segments can extend when a transaction fills the last extent of the undo segment.
- E. Undo segments can be stored in the SYSTEM tablespace.

Correct Answer: BD

Community vote distribution



nautil2 4 months, 2 weeks ago

Selected Answer: DE

A – false
B – false; Undo wrap doesn't happen if a transaction needs an extra space within Undo. If you don't have space, you would see a message like can't extend Undo tablespace to XXX extents and the statement given would be aborted. Wrap, as Hemant has mentioned, would be for the inactive undo segments which are not in use now and can be re-used by the on-going transactions. <https://forums.oracle.com/ords/apexds/post/why-does-undo-segment-wrap-around-4006>
C – false; each transaction is assigned to only one undo segment, but an undo segment can service more than one transaction at a time; <https://www.examttopics.com/discussions/oracle/view/26057-exam-1z0-082-topic-1-question-64-discussion/>
D – true
E – true; When the database instance starts, the database automatically selects the first available undo tablespace. If no undo tablespace is available, then the instance starts without an undo tablespace and stores undo records in the SYSTEM tablespace.
upvoted 1 times

zouve 5 months, 1 week ago

maybe BDE ?
BD from the 64 and E because
"When the database instance starts, the database automatically selects the first available undo tablespace. If no undo tablespace is available, then the instance starts without an undo tablespace and stores undo records in the SYSTEM tablespace. This is not recommended, and an alert message is written to the alert log file to warn that the system is running without an undo tablespace."
upvoted 2 times

Melih1991tr 6 months ago

Selected Answer: DE

If there is no undo tablespace, undo segments can be stored in SYSTEM. And in question 64, people said "B wrong" in comments
upvoted 1 times

heukwoo 6 months, 2 weeks ago

Selected Answer: BD

1z0-082-question-64
upvoted 2 times

Darkseid1231 6 months, 2 weeks ago

D E perhaps !
upvoted 1 times

DATADIR1 and DATATDIR2 are database directory objects.

Examine this command:

```
[oracle@host01 ~] expdp system/oracle \  
> FULL=y \  
> DUMPFILE=datadir1:full%U.dmp, datadir2:full2%U.dmp \  
> FILESIZE=2G \  
> PARALLEL=4 \  
> LOGFILE=datadir1:/expfull.log \  
> JOB_NAME=expfull
```

Which two statements are true about the expdp operation? (Choose two.)

- A. It fails if the total size of the dump file is more than 2GB.
- B. It starts only when four worker processes are available.
- C. It creates a master table to store details of the export operation.
- D. It creates dump files for the entire database.
- E. It creates dump files only for objects in the SYSTEM schema.

Correct Answer: BE

Community vote distribution

BC (50%)



CD (50%)

  **piontk** 5 months, 2 weeks ago

Selected Answer: BC

Correct is BC



upvoted 1 times

  **piontk** 5 months, 2 weeks ago

D - Is wrong because "Several system schemas cannot be exported, because they are not user schemas; they contain Oracle-managed data and metadata. Examples of schemas that are not exported include SYS, ORDSYS, and MDSYS."

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sutil/oracle-data-pump-export-utility.html#GUID-FCB1FA15-5798-48C0-85ED-08E9DAC5E687>

upvoted 1 times

  **piontk** 5 months, 2 weeks ago

B - Oracle uses the MAX_DATAPUMP_PARALLEL_PER_JOB parameter to set the maximum number of parallel processes allowed per Oracle Data Pump job.

https://docs.oracle.com/en/database/oracle/oracle-database/19/refrn/MAX_DATAPUMP_PARALLEL_PER_JOB.html#GUID-33B1F962-B8C3-4DCE-BE68-66FC5D34ECA3

How to test it:

alter session set container=cdb\$root;

alter system set max_datapump_parallel_per_job=2 scope=both;

If you try to start a job with PARALLEL=4 Oracle will return an error:

ORA-31626: job does not exist

ORA-31637: cannot create job EXPFULL for user SYSTEM

upvoted 1 times

  **piontk** 5 months, 2 weeks ago

C - <https://docs.oracle.com/en/database/oracle/oracle-database/19/sutil/oracle-data-pump-overview.html#GUID-B9A676E5-B3CB-417B-BC97-65FE33448C2F>

upvoted 1 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: CD

Agree with C and D.

upvoted 1 times

Darkseid1231

7 months ago

C D is correct

upvoted 1 times

Gian1988

7 months, 3 weeks ago

Correct answer: C - D

upvoted 2 times

Question #153

Topic 1

Which two statements are true about Enterprise Manager Database Express? (Choose two.)

- A. It is available only when the database is open
- B. It can be used to perform database recovery when the database is opened
- C. The same port number can be used for Database Express configurations for databases on different hosts
- D. It can be used to switch a database into ARCHIVELOGMODE
- E. The same port number can be used for multiple Database Express configurations for multiple databases on the same host
- F. It can be used to perform database recovery when the database is mounted

Correct Answer: AC

Community vote distribution

AC (100%)

Darkseid1231

6 months, 2 weeks ago

a c is correct

upvoted 1 times

auwia

6 months, 3 weeks ago

Selected Answer: AC

Provided answer are correct.

upvoted 1 times

Examine these SQL statements which execute successfully:

```
CREATE TABLE emp
  (emp_no NUMBER(2) CONSTRAINT emp_emp_no_pk PRIMARY KEY,
   ename VARCHAR2(15),
   salary NUMBER(8,2),
   mgr_no NUMBER(2));

ALTER TABLE emp ADD CONSTRAINT emp_mgr_fk
  FOREIGN KEY(mgr_no)
  REFERENCES emp(emp_no)
  ON DELETE SET NULL;

ALTER TABLE emp
  DISABLE CONSTRAINT emp_emp_no_pk
  CASCADE;

ALTER TABLE emp
  ENABLE CONSTRAINT emp_emp_no_pk;
```


Which two statements are true after execution? (Choose two.)

- A. The primary key constraint will be enabled and IMMEDIATE.
- B. The primary key constraint will be enabled and DEFERRED.
- C. The foreign key constraint will be enabled and DEFERRED.
- D. The foreign key constraint will be disabled.
- E. The foreign key constraint will be enabled and IMMEDIATE.

Correct Answer: AD

Community vote distribution

AD (100%)

 **nautil2** 4 months, 2 weeks ago

Selected Answer: AD

A, D - true
Experimental – after executing provided commands, the state of constraints is following:

```
SQL> SELECT constraint_name,constraint_type,table_name,status,deferred FROM user_constraints WHERE table_name='EMP';

CONSTRAINT_NAME CONSTRAINT_TYPE TABLE_NAME STATUS DEFERRED
-----
EMP_EMP_NO_PK P EMP ENABLED IMMEDIATE
EMP_MGR_FK R EMP DISABLED IMMEDIATE
upvoted 1 times
```

 **auwia** 6 months, 3 weeks ago

Selected Answer: AD

Provided answer are correct.
upvoted 1 times

 **Darkseid1231** 7 months ago

Ummm A D seems to be okay.
Any suggestion ??
upvoted 1 times

Which two statements are true about the ORDER BY clause? (Choose two.)

- A. Only columns that are specified in the SELECT list can be used in the ORDER BY clause.
- B. Numeric values are displayed in descending order if they have decimal positions.
- C. NULLS are not included in the sort operation.
- D. Column aliases can be used in the ORDER BY clause.
- E. In a character sort, the values are case-sensitive.

Correct Answer: CE

Community vote distribution

DE (100%)

  **nautil2** 4 months, 2 weeks ago

Selected Answer: DE



A – false; following queries are executed correctly:
SELECT * FROM employees ORDER BY 1;
SELECT emp_name, dept_id FROM employees ORDER BY emp_id;
B – false; no mention in Oracle documentation about numeric values order
C – false; Quoting Oracle documentation - options NULLS FIRST | NULLS LAST Specify whether returned rows containing null values should appear first or last in the ordering sequence. NULLS LAST is the default for ascending order, and NULLS FIRST is the default for descending order.
D – true; example: SELECT emp_name AS en, dept_id FROM employees ORDER BY en;
E – true; example:
SQL> SELECT emp_id,emp_name FROM employees WHERE UPPER(emp_name)='JONES' ORDER BY emp_name ASC;
EMP_ID EMP_NAME



1 JONES
7 JONES
10 JOnes
9 Jones
upvoted 2 times

  **auwia** 6 months, 3 weeks ago

Selected Answer: DE

Provided answer are wrong! The base! :-) part X
upvoted 2 times

  **Darkseid1231** 6 months, 2 weeks ago
Hi, have you already attend the exam ? @auwia
upvoted 1 times

  **thanhnx1979** 7 months, 2 weeks ago
D, E is correct
upvoted 3 times